



SCALA

for the Impatient

Second Edition

Cay S. Horstmann

Foreword by Martin Odersky



Scala for the Impatient

Second Edition

This page intentionally left blank

Scala for the Impatient

Second Edition

Cay S. Horstmann

 Addison-Wesley

Boston • Columbus • Indianapolis • New York • San Francisco • Amsterdam • Cape Town
Dubai • London • Madrid • Milan • Munich • Paris • Montreal • Toronto • Delhi • Mexico City
São Paulo • Sydney • Hong Kong • Seoul • Singapore • Taipei • Tokyo

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and the publisher was aware of a trademark claim, the designations have been printed with initial capital letters or in all capitals.

The author and publisher have taken care in the preparation of this book, but make no expressed or implied warranty of any kind and assume no responsibility for errors or omissions. No liability is assumed for incidental or consequential damages in connection with or arising out of the use of the information or programs contained herein.

For information about buying this title in bulk quantities, or for special sales opportunities (which may include electronic versions; custom cover designs; and content particular to your business, training goals, marketing focus, or branding interests), please contact our corporate sales department at corpsales@pearsoned.com or (800) 382-3419.

For government sales inquiries, please contact governmentsales@pearsoned.com.

For questions about sales outside the United States, please contact intlcs@pearson.com.

Visit us on the Web: informit.com/aw

Library of Congress Control Number: 2016954825

Copyright © 2017 Pearson Education, Inc.

All rights reserved. Printed in the United States of America. This publication is protected by copyright, and permission must be obtained from the publisher prior to any prohibited reproduction, storage in a retrieval system, or transmission in any form or by any means, electronic, mechanical, photocopying, recording, or likewise. For information regarding permissions, request forms and the appropriate contacts within the Pearson Education Global Rights & Permissions Department, please visit www.pearsoned.com/permissions/.

ISBN-13: 978-0-13-454056-6

ISBN-10: 0-13-454056-5

Text printed in the United States of America.

1 16

*To my wife, who made writing this book possible,
and to my children, who made it necessary.*

This page intentionally left blank

Contents

Foreword to the First Edition xvii

Preface xix

About the Author xxi

1 THE BASICS **A1** 1

1.1	The Scala Interpreter	1
1.2	Declaring Values and Variables	4
1.3	Commonly Used Types	5
1.4	Arithmetic and Operator Overloading	6
1.5	More about Calling Methods	8
1.6	The <code>apply</code> Method	9
1.7	Scaladoc	10
	Exercises	15

2 CONTROL STRUCTURES AND FUNCTIONS **A1** 17

2.1	Conditional Expressions	18
2.2	Statement Termination	19
2.3	Block Expressions and Assignments	20

2.4	Input and Output	21
2.5	Loops	22
2.6	Advanced for Loops	24
2.7	Functions	25
2.8	Default and Named Arguments L1	26
2.9	Variable Arguments L1	26
2.10	Procedures	28
2.11	Lazy Values L1	28
2.12	Exceptions	29
	Exercises	31

3 WORKING WITH ARRAYS **A1** 35

3.1	Fixed-Length Arrays	35
3.2	Variable-Length Arrays: Array Buffers	36
3.3	Traversing Arrays and Array Buffers	37
3.4	Transforming Arrays	38
3.5	Common Algorithms	40
3.6	Deciphering Scaladoc	41
3.7	Multidimensional Arrays	42
3.8	Interoperating with Java	43
	Exercises	44

4 MAPS AND TUPLES **A1** 47

4.1	Constructing a Map	48
4.2	Accessing Map Values	48
4.3	Updating Map Values	49
4.4	Iterating over Maps	50
4.5	Sorted Maps	50
4.6	Interoperating with Java	50
4.7	Tuples	51
4.8	Zipping	52
	Exercises	52

5 CLASSES **A1** 55

5.1	Simple Classes and Parameterless Methods	55
5.2	Properties with Getters and Setters	56
5.3	Properties with Only Getters	59
5.4	Object-Private Fields	60
5.5	Bean Properties L1	61
5.6	Auxiliary Constructors	62
5.7	The Primary Constructor	63
5.8	Nested Classes L1	66
	Exercises	68

6 OBJECTS **A1** 71

6.1	Singlets	71
6.2	Companion Objects	72
6.3	Objects Extending a Class or Trait	73
6.4	The <code>apply</code> Method	73
6.5	Application Objects	74
6.6	Enumerations	75
	Exercises	77

7 PACKAGES AND IMPORTS **A1** 79

7.1	Packages	80
7.2	Scope Rules	81
7.3	Chained Package Clauses	83
7.4	Top-of-File Notation	83
7.5	Package Objects	83
7.6	Package Visibility	84
7.7	Imports	85
7.8	Imports Can Be Anywhere	85
7.9	Renaming and Hiding Members	86
7.10	Implicit Imports	86
	Exercises	87

8	INHERITANCE	A1	91
8.1	Extending a Class		91
8.2	Overriding Methods		92
8.3	Type Checks and Casts		93
8.4	Protected Fields and Methods		94
8.5	Superclass Construction		94
8.6	Overriding Fields		95
8.7	Anonymous Subclasses		97
8.8	Abstract Classes		97
8.9	Abstract Fields		97
8.10	Construction Order and Early Definitions	L3	98
8.11	The Scala Inheritance Hierarchy		100
8.12	Object Equality	L1	102
8.13	Value Classes	L2	103
	Exercises		105
9	FILES AND REGULAR EXPRESSIONS	A1	109
9.1	Reading Lines		109
9.2	Reading Characters		110
9.3	Reading Tokens and Numbers		111
9.4	Reading from URLs and Other Sources		111
9.5	Reading Binary Files		112
9.6	Writing Text Files		112
9.7	Visiting Directories		112
9.8	Serialization		113
9.9	Process Control	A2	114
9.10	Regular Expressions		116
9.11	Regular Expression Groups		117
	Exercises		118
10	TRAITS	L1	121
10.1	Why No Multiple Inheritance?		121
10.2	Traits as Interfaces		123
10.3	Traits with Concrete Implementations		124

10.4	Objects with Traits	125
10.5	Layered Traits	125
10.6	Overriding Abstract Methods in Traits	127
10.7	Traits for Rich Interfaces	127
10.8	Concrete Fields in Traits	128
10.9	Abstract Fields in Traits	130
10.10	Trait Construction Order	130
10.11	Initializing Trait Fields	132
10.12	Traits Extending Classes	133
10.13	SelfTypes L2	134
10.14	What Happens under the Hood	135
	Exercises	137

11 OPERATORS **L1** 141

11.1	Identifiers	142
11.2	Infix Operators	143
11.3	Unary Operators	143
11.4	Assignment Operators	144
11.5	Precedence	144
11.6	Associativity	145
11.7	The <code>apply</code> and <code>update</code> Methods	146
11.8	Extractors L2	147
11.9	Extractors with One or No Arguments L2	149
11.10	The <code>unapplySeq</code> Method L2	149
11.11	Dynamic Invocation L2	150
	Exercises	153

12 HIGHER-ORDER FUNCTIONS **L1** 157

12.1	Functions as Values	157
12.2	Anonymous Functions	159
12.3	Functions with Function Parameters	160
12.4	Parameter Inference	160
12.5	Useful Higher-Order Functions	161
12.6	Closures	162

12.7	SAM Conversions	163
12.8	Currying	164
12.9	Control Abstractions	166
12.10	The return Expression	167
	Exercises	168

13 COLLECTIONS **A2** 171

13.1	The Main Collections Traits	172
13.2	Mutable and Immutable Collections	173
13.3	Sequences	174
13.4	Lists	175
13.5	Sets	177
13.6	Operators for Adding or Removing Elements	178
13.7	Common Methods	180
13.8	Mapping a Function	182
13.9	Reducing, Folding, and Scanning A3	184
13.10	Zipping	187
13.11	Iterators	188
13.12	Streams A3	189
13.13	LazyViews A3	190
13.14	Interoperability with Java Collections	191
13.15	Parallel Collections	193
	Exercises	194

14 PATTERN MATCHING AND CASE CLASSES **A2** 197

14.1	A Better Switch	198
14.2	Guards	199
14.3	Variables in Patterns	199
14.4	Type Patterns	200
14.5	Matching Arrays, Lists, and Tuples	201
14.6	Extractors	202
14.7	Patterns in Variable Declarations	203
14.8	Patterns in for Expressions	204
14.9	Case Classes	205

14.10	The copy Method and Named Parameters	205
14.11	Infix Notation in case Clauses	206
14.12	Matching Nested Structures	207
14.13	Are Case Classes Evil?	208
14.14	Sealed Classes	209
14.15	Simulating Enumerations	209
14.16	The Option Type	210
14.17	Partial Functions L2	211
	Exercises	212

15 ANNOTATIONS **A2** 215

15.1	What Are Annotations?	216
15.2	What Can Be Annotated?	216
15.3	Annotation Arguments	217
15.4	Annotation Implementations	218
15.5	Annotations for Java Features	219
	15.5.1 Java Modifiers	219
	15.5.2 Marker Interfaces	220
	15.5.3 Checked Exceptions	220
	15.5.4 Variable Arguments	221
	15.5.5 JavaBeans	221
15.6	Annotations for Optimizations	222
	15.6.1 Tail Recursion	222
	15.6.2 Jump Table Generation and Inlining	223
	15.6.3 Eliding Methods	224
	15.6.4 Specialization for Primitive Types	225
15.7	Annotations for Errors and Warnings	226
	Exercises	227

16 XML PROCESSING **A2** 229

16.1	XML Literals	230
16.2	XML Nodes	230
16.3	Element Attributes	232
16.4	Embedded Expressions	233

16.5	Expressions in Attributes	234
16.6	Uncommon Node Types	235
16.7	XPath-like Expressions	235
16.8	Pattern Matching	237
16.9	Modifying Elements and Attributes	238
16.10	Transforming XML	239
16.11	Loading and Saving	239
16.12	Namespaces	242
	Exercises	243

17 FUTURES **A2** 247

17.1	Running Tasks in the Future	248
17.2	Waiting for Results	250
17.3	The Try Class	251
17.4	Callbacks	251
17.5	Composing Future Tasks	252
17.6	Other Future Transformations	255
17.7	Methods in the Future Object	256
17.8	Promises	258
17.9	Execution Contexts	260
	Exercises	260

18 TYPE PARAMETERS **L2** 265

18.1	Generic Classes	266
18.2	Generic Functions	266
18.3	Bounds for Type Variables	266
18.4	View Bounds	268
18.5	Context Bounds	268
18.6	The <code>ClassTag</code> Context Bound	269
18.7	Multiple Bounds	269
18.8	Type Constraints L3	269
18.9	Variance	271
18.10	Co- and Contravariant Positions	272

18.11 Objects Can't Be Generic	274
18.12 Wildcards	275
Exercises	275

19 ADVANCED TYPES **L2** 279

19.1 Singleton Types	280
19.2 Type Projections	281
19.3 Paths	282
19.4 Type Aliases	283
19.5 Structural Types	283
19.6 Compound Types	284
19.7 Infix Types	285
19.8 Existential Types	286
19.9 The Scala Type System	287
19.10 Self Types	288
19.11 Dependency Injection	289
19.12 Abstract Types L3	291
19.13 Family Polymorphism L3	293
19.14 Higher-Kinded Types L3	296
Exercises	299

20 PARSING **A3** 303

20.1 Grammars	304
20.2 Combining Parser Operations	305
20.3 Transforming Parser Results	307
20.4 Discarding Tokens	308
20.5 Generating Parse Trees	309
20.6 Avoiding Left Recursion	310
20.7 More Combinators	311
20.8 Avoiding Backtracking	314
20.9 Packrat Parsers	314
20.10 What Exactly Are Parsers?	315
20.11 Regex Parsers	316

20.12 Token-Based Parsers 317

20.13 Error Handling 319

Exercises 320

21 IMPLICITS L3 323

21.1 Implicit Conversions 324

21.2 Using Implicits for Enriching Existing Classes 324

21.3 Importing Implicits 325

21.4 Rules for Implicit Conversions 326

21.5 Implicit Parameters 328

21.6 Implicit Conversions with Implicit Parameters 329

21.7 Context Bounds 329

21.8 Type Classes 331

21.9 Evidence 333

21.10 The `@implicitNotFound` Annotation 334

21.11 `CanBuildFrom` Demystified 334

Exercises 336

Index 338

Foreword to the First Edition

When I met Cay Horstmann some years ago he told me that Scala needed a better introductory book. My own book had come out a little bit earlier, so of course I had to ask him what he thought was wrong with it. He responded that it was great but too long; his students would not have the patience to read through the eight hundred pages of *Programming in Scala*. I conceded that he had a point. And he set out to correct the situation by writing *Scala for the Impatient*.

I am very happy that his book has finally arrived because it really delivers on what the title says. It gives an eminently practical introduction to Scala, explains what's particular about it, how it differs from Java, how to overcome some common hurdles to learning it, and how to write good Scala code.

Scala is a highly expressive and flexible language. It lets library writers use highly sophisticated abstractions, so that library users can express themselves simply and intuitively. Therefore, depending on what kind of code you look at, it might seem very simple or very complex.

A year ago, I tried to provide some clarification by defining a set of levels for Scala and its standard library. There were three levels each for application programmers and for library designers. The junior levels could be learned quickly and would be sufficient to program productively. Intermediate levels would make programs more concise and more functional and would make libraries more flexible to use. The highest levels were for experts solving specialized tasks. At the time I wrote:

I hope this will help newcomers to the language decide in what order to pick subjects to learn, and that it will give some advice to teachers and book authors in what order to present the material.

Cay's book is the first to have systematically applied this idea. Every chapter is tagged with a level that tells you how easy or hard it is and whether it's oriented towards library writers or application programmers.

As you would expect, the first chapters give a fast-paced introduction to the basic Scala capabilities. But the book does not stop there. It also covers many of the more "senior" concepts and finally progresses to very advanced material which is not commonly covered in a language introduction, such as how to write parser combinators or make use of delimited continuations. The level tags serve as a guideline for what to pick up when. And Cay manages admirably to make even the most advanced concepts simple to understand.

I liked the concept of *Scala for the Impatient* so much that I asked Cay and his editor, Greg Doench, whether we could get the first part of the book as a free download on the Typesafe web site. They have graciously agreed to my request, and I would like to thank them for that. That way, everybody can quickly access what I believe is currently the best compact introduction to Scala.

Martin Odersky

January 2012

Preface

The evolution of traditional languages such as Java, C#, and C++ has slowed down considerably, and programmers who are eager to use more modern language features are looking elsewhere. Scala is an attractive choice; in fact, I think it is by far the most attractive choice for programmers who want to improve their productivity. Scala has a concise syntax that is refreshing after the Java boilerplate. It runs on the Java virtual machine, providing access to a huge set of libraries and tools. And Scala doesn't just target the Java virtual machine. The ScalaJS project emits JavaScript code, enabling you to write both the server-side and client-side parts of a web application in a language that isn't JavaScript. Scala embraces the functional programming style without abandoning object orientation, giving you an incremental learning path to a new paradigm. The Scala interpreter lets you run quick experiments, which makes learning Scala very enjoyable. Last but not least, Scala is statically typed, enabling the compiler to find errors, so that you don't waste time finding them—or not—later in the running program.

I wrote this book for *impatient* readers who want to start programming in Scala right away. I assume you know Java, C#, or C++, and I don't bore you with explaining variables, loops, or classes. I don't exhaustively list all the features of the language, I don't lecture you about the superiority of one paradigm over another, and I don't make you suffer through long and contrived examples. Instead, you will get the information that you need in compact chunks that you can read and review as needed.

Scala is a big language, but you can use it effectively without knowing all of its details intimately. Martin Odersky, the creator of Scala, has identified levels of expertise for application programmers and library designers—as shown in the following table.

Application Programmer	Library Designer	Overall Scala Level
Beginning A1		Beginning
Intermediate A2	Junior L1	Intermediate
Expert A3	Senior L2	Advanced
	Expert L3	Expert

For each chapter (and occasionally for individual sections), I indicate the experience level required. The chapters progress through levels **A1**, **L1**, **A2**, **L2**, **A3**, **L3**. Even if you don't want to design your own libraries, knowing about the tools that Scala provides for library designers can make you a more effective library user.

This is the second edition of this book, and I updated it thoroughly for Scala 2.12. I added coverage of recent Scala features such as string interpolation, dynamic invocation, implicit classes, and futures, and updated all chapters to reflect current Scala usage.

I hope you enjoy learning Scala with this book. If you find errors or have suggestions for improvement, please visit <http://horstmann.com/scala> and leave a comment. On that page, you will also find a link to an archive file containing all code examples from the book.

I am very grateful to Dmitry Kirsanov and Alina Kirsanova who turned my manuscript from XHTML into a beautiful book, allowing me to concentrate on the content instead of fussing with the format. Every author should have it so good!

Reviewers include Adrian Cumiskey, Mike Davis, Rob Dickens, Steve Haines, Susan Potter, Daniel Sobral, Craig Tataryn, David Walend, and William Wheeler. Thanks so much for your comments and suggestions!

Finally, as always, my gratitude goes to my editor, Greg Doench, for encouraging me to write this book, and for his insights during the development process.

*Cay Horstmann
San Francisco, 2016*

About the Author

Cay S. Horstmann is author of *Core Java™, Volumes I & II, Tenth Edition* (Prentice Hall, 2016), as well as a dozen other books for professional programmers and computer science students. He is a professor of computer science at San Jose State University and a Java Champion.

The Basics

Topics in This Chapter **A1**

- 1.1 The Scala Interpreter — page 1
- 1.2 Declaring Values and Variables — page 4
- 1.3 Commonly Used Types — page 5
- 1.4 Arithmetic and Operator Overloading — page 6
- 1.5 More about Calling Methods — page 8
- 1.6 The `apply` Method — page 9
- 1.7 Scaladoc — page 10
- Exercises — page 15

Chapter

1

In this chapter, you will learn how to use Scala as an industrial-strength pocket calculator, working interactively with numbers and arithmetic operations. We introduce a number of important Scala concepts and idioms along the way. You will also learn how to browse the Scaladoc documentation at a beginner's level.

Highlights of this introduction are:

- Using the Scala interpreter
- Defining variables with `var` and `val`
- Numeric types
- Using operators and functions
- Navigating Scaladoc

1.1 The Scala Interpreter

To start the Scala interpreter:

1. Install Scala.
2. Make sure that the `scala/bin` directory is on the PATH.
3. Open a command shell in your operating system.
4. Type `scala` followed by the Enter key.

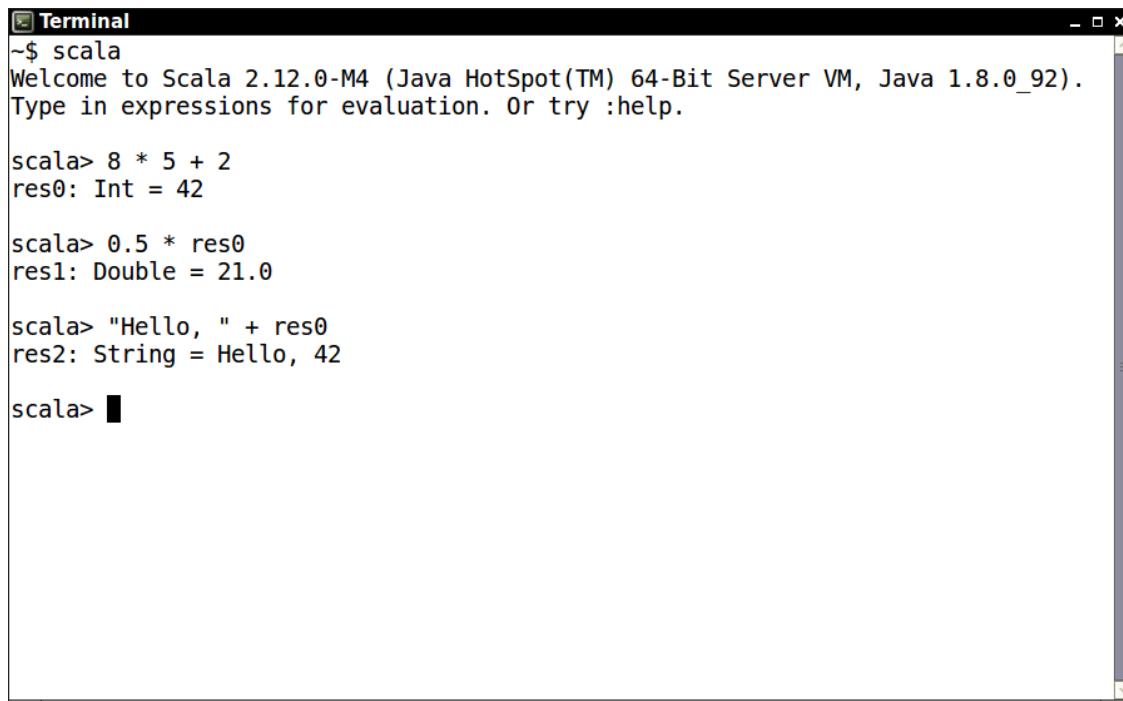
Now type commands followed by Enter. Each time, the interpreter displays the answer, as shown in Figure 1–1. For example, if you type `8 * 5 + 2` (as shown in boldface below), you get 42.

```
scala> 8 * 5 + 2
res0: Int = 42
```

The answer is given the name `res0`. You can use that name in subsequent computations:

```
scala> 0.5 * res0
res1: Double = 21.0
scala> "Hello, " + res0
res2: java.lang.String = Hello, 42
```

As you can see, the interpreter also displays the type of the result—in our examples, `Int`, `Double`, and `java.lang.String`.



The screenshot shows a terminal window titled "Terminal". The window contains the following Scala interpreter session:

```
~$ scala
Welcome to Scala 2.12.0-M4 (Java HotSpot(TM) 64-Bit Server VM, Java 1.8.0_92).
Type in expressions for evaluation. Or try :help.

scala> 8 * 5 + 2
res0: Int = 42

scala> 0.5 * res0
res1: Double = 21.0

scala> "Hello, " + res0
res2: String = Hello, 42

scala> █
```

Figure 1–1 The Scala Interpreter



TIP: Don't like the command shell? Several integrated development environments that support Scala have a “worksheet” feature for entering expressions and displaying their result whenever the sheet is saved. Figure 1–2 shows a worksheet in the Eclipse-based Scala IDE.

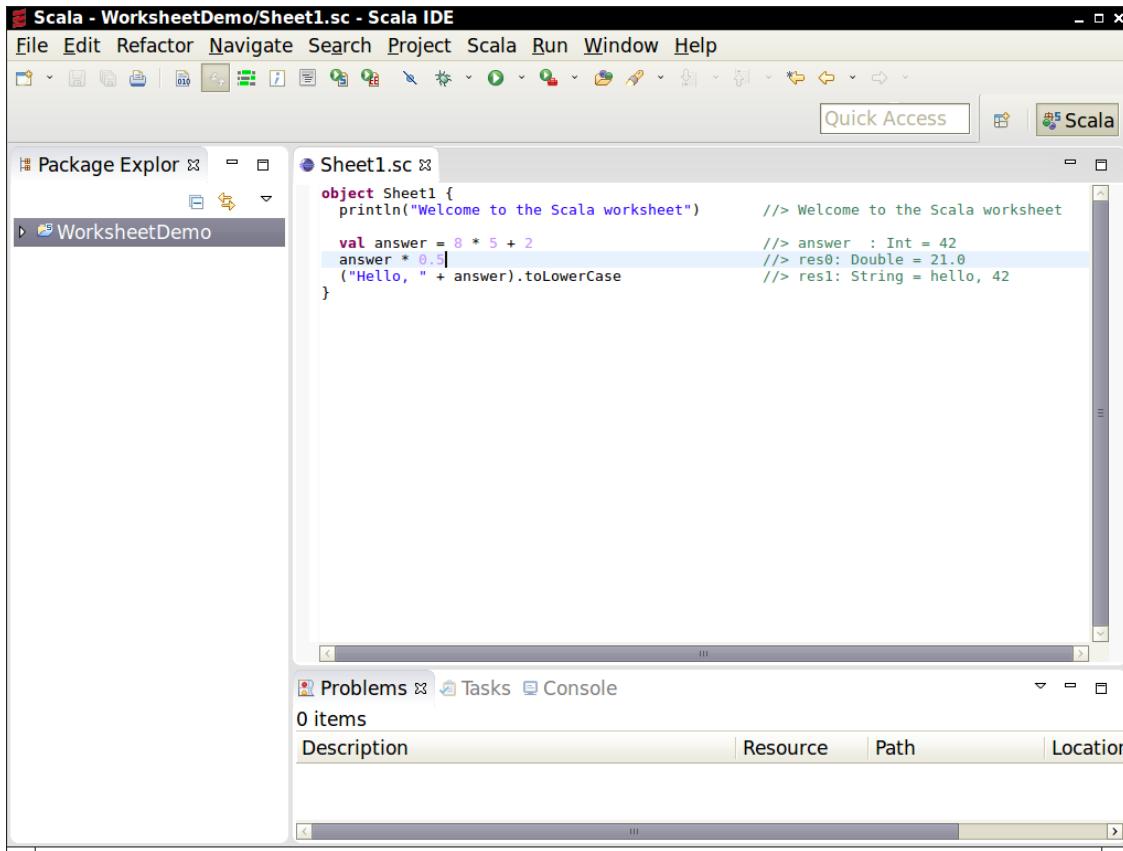


Figure 1–2 A Scala Worksheet

When calling methods, try using *tab completion* for method names. Type `res2.to` and then hit the Tab key. If the interpreter offers choices such as

`toCharArray` `toLowerCase` `toString` `toUpperCase`

this means tab completion works in your environment. Type a U and hit the Tab key again. You now get a single completion:

`res2.toUpperCase`

Hit the Enter key, and the answer is displayed. (If you can't use tab completion in your environment, you'll have to type the complete method name yourself.)

Also try hitting the ↑ and ↓ arrow keys. In most implementations, you will see the previously issued commands, and you can edit them. Use the ←, →, and Del keys to change the last command to

`res2.toLowerCase`

As you can see, the Scala interpreter reads an expression, evaluates it, prints it, and reads the next expression. This is called the *read-eval-print loop*, or REPL.

Technically speaking, the `scala` program is *not* an interpreter. Behind the scenes, your input is quickly compiled into bytecode, and the bytecode is executed by the Java virtual machine. For that reason, most Scala programmers prefer to call it “the REPL”.



TIP: The REPL is your friend. Instant feedback encourages experimenting, and you will feel good whenever something works.

It is a good idea to keep an editor window open at the same time, so you can copy and paste successful code snippets for later use. Also, as you try more complex examples, you may want to compose them in the editor and then paste them into the REPL.



TIP: In the REPL, type `:help` to see a list of useful commands. All commands start with a colon. For example, the `:warnings` command gives more detailed information about the most recent compiler warning. You only have to enter the unique prefix of each command. For example, `:w` is the same as `:warnings`—at least for now, since there isn’t currently another command starting with `w`.

1.2 Declaring Values and Variables

Instead of using `res0`, `res1`, and so on, you can define your own names:

```
scala> val answer = 8 * 5 + 2
answer: Int = 42
```

You can use these names in subsequent expressions:

```
scala> 0.5 * answer
res3: Double = 21.0
```

A value declared with `val` is actually a constant—you can’t change its contents:

```
scala> answer = 0
<console>:6: error: reassignment to val
```

To declare a variable whose contents can vary, use a `var`:

```
var counter = 0
counter = 1 // OK, can change a var
```

In Scala, you are encouraged to use a `val` unless you really need to change the contents. Perhaps surprisingly for Java or C++ programmers, most programs don’t need many `var` variables.

Note that you need not specify the type of a value or variable. It is inferred from the type of the expression with which you initialize it. (It is an error to declare a value or variable without initializing it.)

However, you can specify the type if necessary. For example,

```
val greeting: String = null
val greeting: Any = "Hello"
```



NOTE: In Scala, the type of a variable or function is written *after* the name of the variable or function. This makes it easier to read declarations with complex types.

As I move back and forth between Scala and Java, I find that my fingers write Java declarations such as `String greeting` on autopilot, so I have to rewrite them as `greeting: String`. This is a bit annoying, but when I work with complex Scala programs, I really appreciate that I don't have to decrypt C-style type declarations.



NOTE: You may have noticed that there were no semicolons after variable declarations or assignments. In Scala, semicolons are only required if you have multiple statements on the same line.

You can declare multiple values or variables together:

```
val xmax, ymax = 100 // Sets xmax and ymax to 100
var greeting, message: String = null
// greeting and message are both strings, initialized with null
```

1.3 Commonly Used Types

You have already seen some of the data types of Scala, such as `Int` and `Double`. Like Java, Scala has seven numeric types: `Byte`, `Char`, `Short`, `Int`, `Long`, `Float`, and `Double`, and a `Boolean` type. However, unlike Java, these types are *classes*. There is no distinction between primitive types and class types in Scala. You can invoke methods on numbers, for example:

```
1.toInt() // Yields the string "1"
```

or, more excitingly,

```
1.to(10) // Yields Range(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

(We will discuss the `Range` class in Chapter 13. For now, just view it as a collection of numbers.)

In Scala, there is no need for wrapper types. It is the job of the Scala compiler to convert between primitive types and wrappers. For example, if you make an array of Int, you get an int[] array in the virtual machine.

As you saw in Section 1.1, “The Scala Interpreter,” on page 1, Scala relies on the underlying java.lang.String class for strings. However, it augments that class with well over a hundred operations in the StringOps class. For example, the intersect method yields the characters that are common to two strings:

```
"Hello".intersect("World") // Yields "lo"
```

In this expression, the java.lang.String object "Hello" is implicitly converted to a StringOps object, and then the intersect method of the StringOps class is applied. So, remember to look into the StringOps class when you use the Scala documentation (see Section 1.7, “Scaladoc,” on page 10).

Similarly, there are classes RichInt, RichDouble, RichChar, and so on. Each of them has a small set of convenience methods for acting on their poor cousins—Int, Double, or Char. The to method that you saw above is actually a method of the RichInt class. In the expression

```
1.to(10)
```

the Int value 1 is first converted to a RichInt, and the to method is applied to that value.

Finally, there are classes BigInt and BigDecimal for computations with an arbitrary (but finite) number of digits. These are backed by the java.math.BigInteger and java.math.BigDecimal classes, but, as you will see in the next section, they are much more convenient because you can use them with the usual mathematical operators.



NOTE: In Scala, you use methods, not casts, to convert between numeric types. For example, 99.44.toInt is 99, and 99.toChar is 'c'. Of course, as in Java, the toString method converts any object to a string.

To convert a string containing a number into the number, use toInt or toDouble. For example, "99.44".toDouble is 99.44.

1.4 Arithmetic and Operator Overloading

Arithmetic operators in Scala work just as you would expect in Java or C++:

```
val answer = 8 * 5 + 2
```

The + - * / % operators do their usual job, as do the bit operators & | ^ >> <<. There is just one surprising aspect: These operators are actually methods. For example,

`a + b`

is a shorthand for

`a.+(b)`

Here, `+` is the name of the method. Scala has no silly prejudice against non-alphanumeric characters in method names. You can define methods with just about any symbols for names. For example, the `BigInt` class defines a method called `%` that returns a pair containing the quotient and remainder of a division.

In general, you can write

`a method b`

as a shorthand for

`a.method(b)`

where `method` is a method with two parameters (one implicit, one explicit). For example, instead of

`1.to(10)`

you can write

`1 to 10`

Use whatever you think is easier to read. Beginning Scala programmers tend to stick to the Java syntax, and that's fine. Of course, even the most hardened Java programmers seem to prefer `a + b` over `a.+(b)`.

There is one notable difference between Scala and Java or C++. Scala does not have `++` or `--` operators. Instead, simply use `+=1` or `-=1`:

`counter+=1 // Increments counter—Scala has no ++`

Some people wonder if there is any deep reason for Scala's refusal to provide a `++` operator. (Note that you can't simply implement a method called `++`. Since the `Int` class is immutable, such a method cannot change an integer value.) The Scala designers decided it wasn't worth having yet another special rule just to save one keystroke.

You can use the usual mathematical operators with `BigInt` and `BigDecimal` objects:

```
val x: BigInt = 1234567890
x * x * x // Yields 1881676371789154860897069000
```

That's much better than Java, where you would have had to call `x.multiply(x).multiply(x)`.



NOTE: In Java, you cannot overload operators, and the Java designers claimed this is a good thing because it stops you from inventing crazy operators like !@\$&* that would make your program impossible to read. Of course, that's silly; you can make your programs just as unreadable by using crazy method names like qxywz. Scala allows you to define operators, leaving it up to you to use this feature with restraint and good taste.

1.5 More about Calling Methods

You have already seen how to call methods on objects, such as

```
"Hello".intersect("World")
```

If the method has no parameters, you don't have to use parentheses. For example, the API of the `StringOps` class shows a method `sorted`, without `()`, which yields a new string with the letters in sorted order. Call it as

```
"Bonjour".sorted // Yields the string "Bjnooru"
```

The rule of thumb is that a parameterless method that doesn't modify the object has no parentheses. We discuss this further in Chapter 5.

In Java, mathematical methods such as `sqrt` are defined as static methods of the `Math` class. In Scala, you define such methods in *singleton objects*, which we will discuss in detail in Chapter 6. A package can have a *package object*. In that case, you can import the package and use the methods of the package object without any prefix:

```
import scala.math._ // In Scala, the _ character is a "wildcard," like * in Java
sqrt(2) // Yields 1.4142135623730951
pow(2, 4) // Yields 16.0
min(3, Pi) // Yields 3.0
```

If you don't import the `scala.math` package, add the package name:

```
scala.math.sqrt(2)
```



NOTE: If a package that starts with `scala.`, you can omit the `scala` prefix. For example, `import math._` is equivalent to `import scala.math._`, and `math.sqrt(2)` is the same as `scala.math.sqrt(2)`. However, in this book, we always use the `scala` prefix for clarity.

You can find more information about the `import` statement in Chapter 7. For now, just use `import packageName._` whenever you need to import a particular package.

Often, a class has a *companion object* whose methods act just like static methods do in Java. For example, the `BigInt` companion object to the `scala.math.BigInt` class

has a method `probablePrime` that generates a random prime number with a given number of bits:

```
BigInt.probablePrime(100, scala.util.Random)
```

Here, `Random` is a singleton random number generator object, defined in the `scala.util` package. Try this in the REPL; you'll get a number such as `1039447980491200275486540240713`.

1.6 The apply Method

In Scala, it is common to use a syntax that looks like a function call. For example, if `s` is a string, then `s(i)` is the *i*th character of the string. (In C++, you would write `s[i]`; in Java, `s.charAt(i)`.) Try it out in the REPL:

```
val s = "Hello"
s(4) // Yields 'o'
```

You can think of this as an overloaded form of the `()` operator. It is implemented as a method with the name `apply`. For example, in the documentation of the `StringOps` class, you will find a method

```
def apply(n: Int): Char
```

That is, `s(4)` is a shortcut for

```
s.apply(4)
```

Why not use the `[]` operator? You can think of a sequence `s` of element type `T` as a function from $\{0, 1, \dots, n - 1\}$ to `T` that maps i to $s(i)$, the *i*th element of the sequence.

This argument is even more convincing for maps. As you will see in Chapter 4, you look up a map value for a given key as `map(key)`. Conceptually, a map is a function from keys to values, and it makes sense to use the function notation.



CAUTION: Occasionally, the `()` notation conflicts with another Scala feature: implicit parameters. For example, the expression

```
"Bonjour".sorted(3)
```

yields an error because the `sorted` method can optionally be called with an ordering, but 3 is not a valid ordering. You can use parentheses:

```
("Bonjour".sorted)(3)
```

or call `apply` explicitly:

```
"Bonjour".sorted.apply(3)
```

When you look at the documentation for the `BigInt` companion object, you will see `apply` methods that let you convert strings or numbers to `BigInt` objects. For example, the call

```
BigInt("1234567890")
```

is a shortcut for

```
BigInt.apply("1234567890")
```

It yields a new `BigInt` object, *without having to use new*. For example:

```
BigInt("1234567890") * BigInt("112358111321")
```

Using the `apply` method of a companion object is a common Scala idiom for constructing objects. For example, `Array(1, 4, 9, 16)` returns an array, thanks to the `apply` method of the `Array` companion object.



NOTE: All through this chapter, we have assumed that Scala code is executed on the Java virtual machine. That is in fact true for the standard Scala distribution. However, the Scala.js project (www.scala-js.org) provides tools to translate Scala to JavaScript. If you take advantage of that project, you can write both the client-side and the server-side code of web applications in Scala.

1.7 Scaladoc

Java programmers use Javadoc to navigate the Java API. Scala has its own variant, called Scaladoc (see Figure 1–3).

Using Scaladoc can be a bit more overwhelming than Javadoc. Scala classes tend to have many more convenience methods than Java classes. Some methods use advanced features that are more meaningful to library implementors than to library users.

Here are some tips for navigating Scaladoc as a newcomer to the language.

You can browse Scaladoc online at www.scala-lang.org/api, but it is a good idea to download a copy from <http://scala-lang.org/download/all.html> and install it locally.

Unlike Javadoc, which presents an alphabetical listing of classes, Scaladoc is organized by packages. If you know a class or method name, don't bother navigating to the package. Simply use the search bar on the top of the entry page (see Figure 1–4).

The screenshot shows a Mozilla Firefox browser window with the title "Scala Standard Library 2.12.0-M4 - Mozilla Firefox". The address bar contains the URL "www.scala-lang.org/api/2.12.0-M4/". The page content is the Scaladoc documentation for the root package of the Scala standard library. The left sidebar features a large blue circular icon with a white letter 'p' and the text "root package". The main content area starts with a heading "Package structure" and describes the core types available in the `scala` package. It lists several notable packages, each with a brief description and a bulleted list of sub-packages or components. The "Notable packages include:" section lists packages like `scala.collection`, `scala.concurrent`, `scala.io`, `scala.math`, `scala.sys`, and `scala.util.matching`. The "Other packages exist. See the complete list on the left." section points to the sidebar. The "Automatic imports" section notes that identifiers in the `scala` package and the `scala.Predef` object are always in scope by default. The "Copyright" notice at the bottom states: "Scala programming documentation. Copyright (c) 2003-2016 EPFL, with contributions from Lightbend."

Figure 1–3 The entry page for Scaladoc

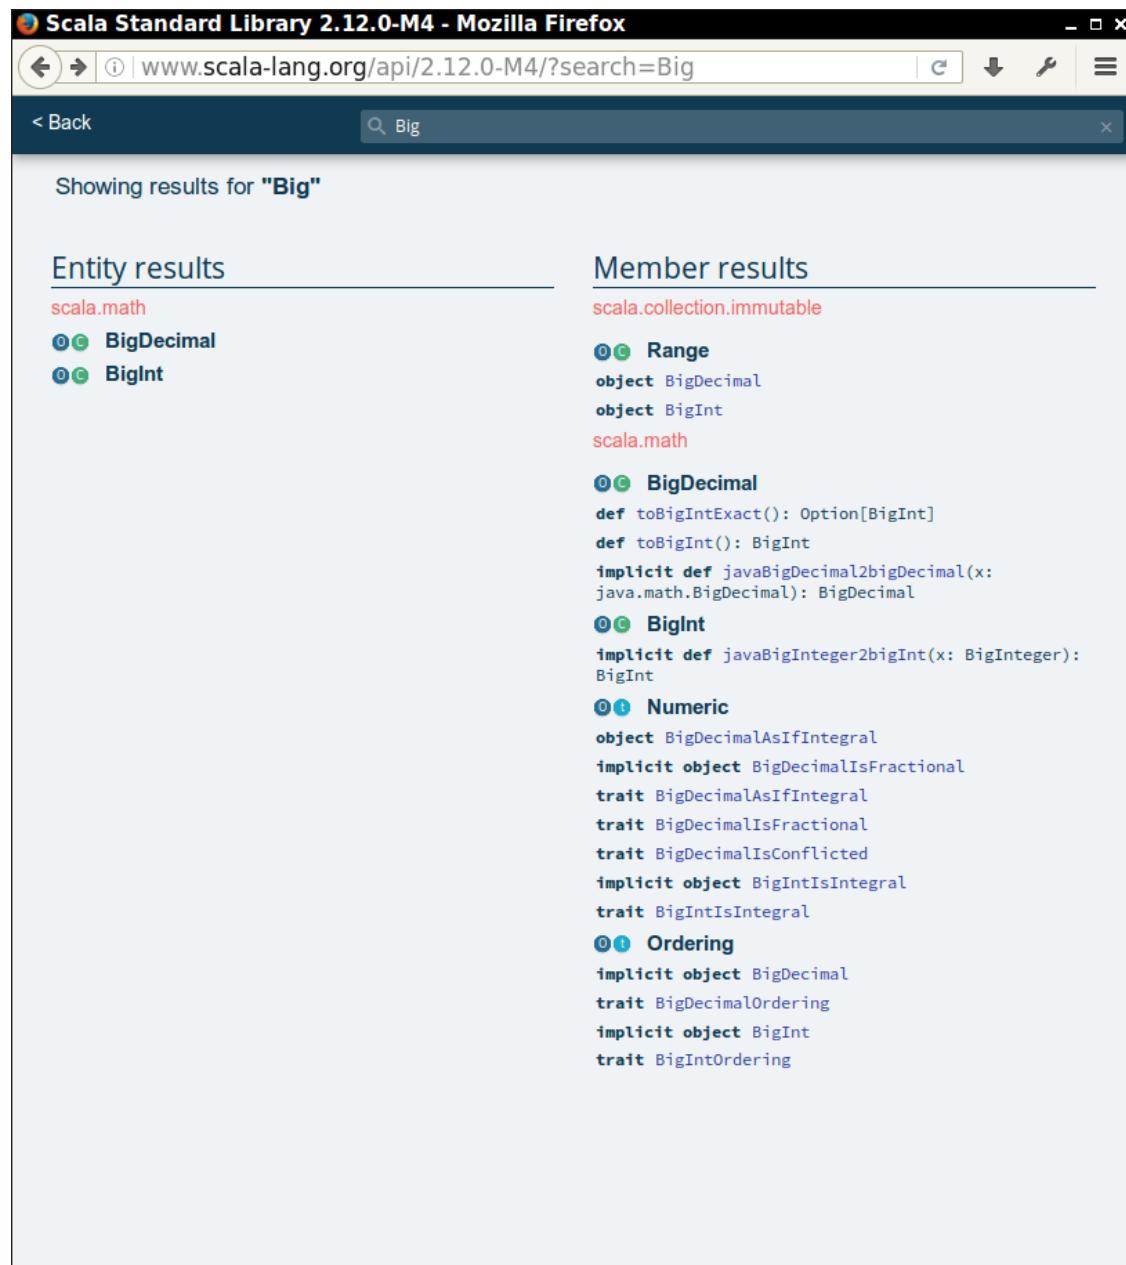


Figure 1–4 The search bar in Scaladoc

Click on the X symbol to clear the search pattern, or click on a matching class or method (Figure 1–5).

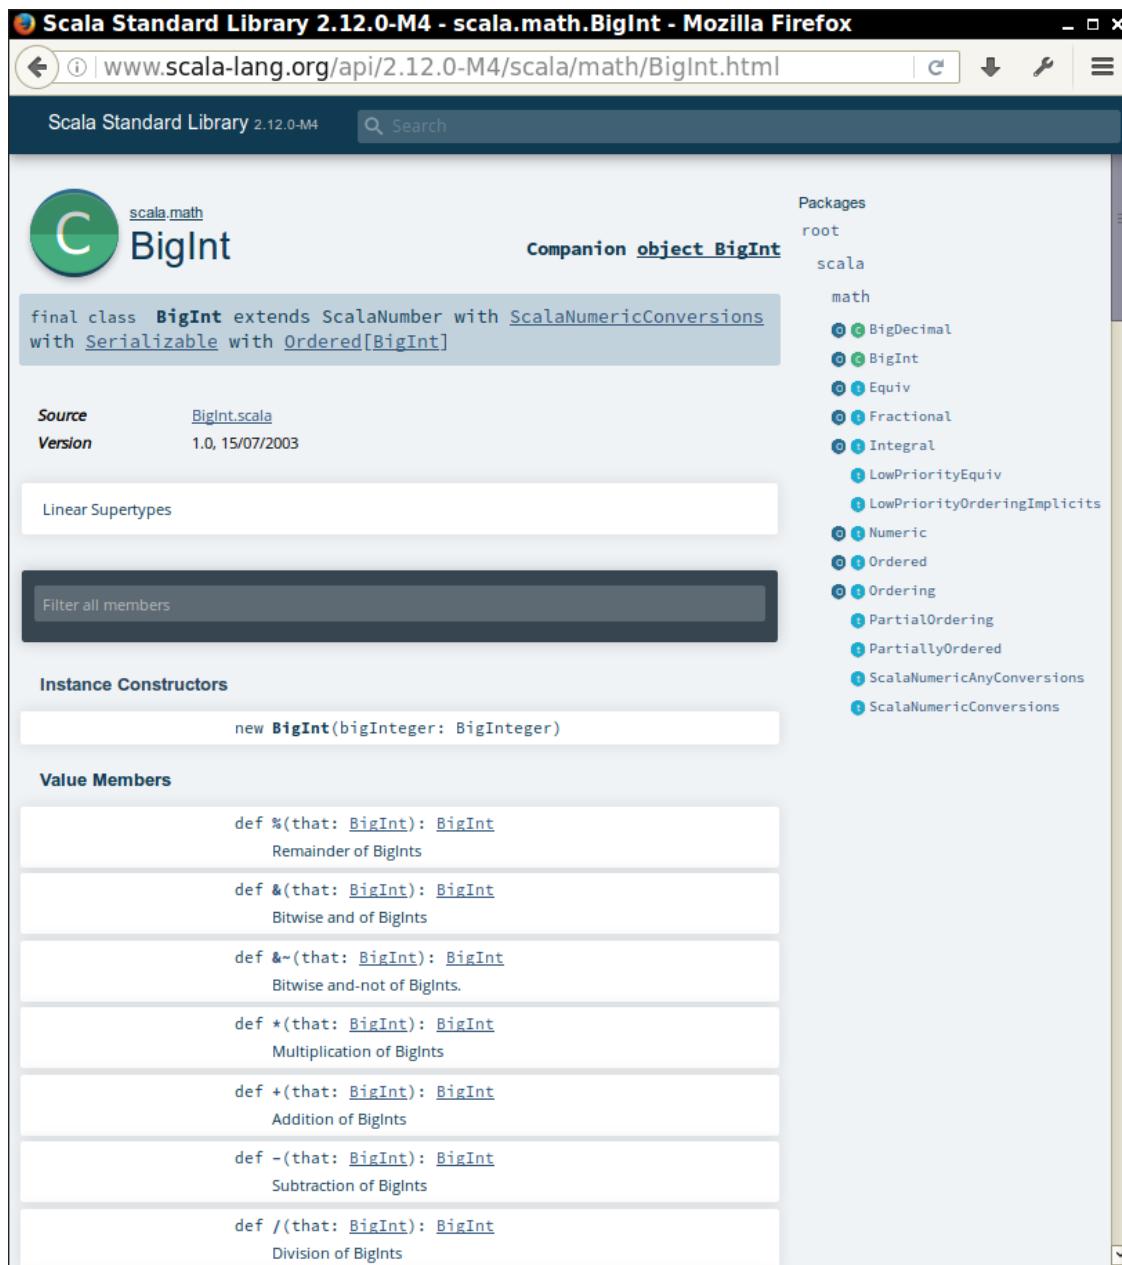


Figure 1–5 Class documentation in Scaladoc

Note the C and O symbols next to the class name. They let you navigate to the class (C) or the companion object (O). For traits (which are similar to Java interfaces), you see t and O symbols instead.

Keep these tips in mind:

- Remember to look into `RichInt`, `RichDouble`, and so on, if you want to know how to work with numeric types. Similarly, to work with strings, look into `StringOps`.
- The mathematical functions are in the package `scala.math`, not in any class.
- Sometimes, you'll see functions with funny names. For example, `BigInt` has a method `unary_-`. As you will see in Chapter 11, this is how you define the prefix negation operator `-x`.
- Methods can have functions as parameters. For example, the `count` method in `StringOps` requires a function that returns true or false for a `Char`, specifying which characters should be counted:

```
def count(p: (Char) => Boolean) : Int
```

You supply a function, often in a very compact notation, when you call the method. As an example, the call `s.count(_.isUpper)` counts the number of uppercase characters. We will discuss this style of programming in much more detail in Chapter 12.

- You'll occasionally run into classes such as `Range` or `Seq[Char]`. They mean what your intuition tells you—a range of numbers, a sequence of characters. You will learn all about these classes as you delve more deeply into Scala.
- In Scala, you use square brackets for type parameters. A `Seq[Char]` is a sequence of elements of type `Char`, and `Seq[A]` is a sequence of elements of some type `A`.
- There are many slightly different types for sequences such as `GenSeq`, `GenIterable`, `GenTraversableOnce`, and so on. The differences between them are rarely important. When you see such a construct, just think “sequence.” For example, the `StringOps` class defines a method

```
def containsSlice[B](that: GenSeq[B]): Boolean
```

This method tests whether the string contains with a given sequence. If you like, you can pass a `Range`:

```
"Bierstube".containsSlice('r'.to('u'))  
// Yields true since the string contains Range('r', 's', 't', 'u')
```

- Don't get discouraged that there are so many methods. It's the Scala way to provide lots of methods for every conceivable use case. When you need to solve a particular problem, just look for a method that is useful. More often than not, there is one that addresses your task, which means you don't have to write so much code yourself.
- Some methods have an “implicit” parameter. For example, the `sorted` method of `StringOps` is declared as

```
def sorted[B >: Char](implicit ord: math.Ordering[B]): String
```

That means that an ordering is supplied “implicitly,” using a mechanism that we will discuss in detail in Chapter 21. You can ignore these implicit parameters for now.

- Finally, don’t worry if you run into the occasional indecipherable incantation, such as the `[B >: Char]` in the declaration of `sorted`. The expression `B >: Char` means “any supertype of `Char`,” but for now, ignore that generality.
- Whenever you are confused what a method does, just try it out in the REPL:

```
"Scala".sorted // Yields "Saacl"
```

Now you can clearly see that the method returns a new string that consists of the characters in sorted order.

Exercises

1. In the Scala REPL, type `3.` followed by the Tab key. What methods can be applied?
2. In the Scala REPL, compute the square root of 3, and then square that value. By how much does the result differ from 3? (Hint: The `res` variables are your friend.)
3. Are the `res` variables `val` or `var`?
4. Scala lets you multiply a string with a number—try out `"crazy" * 3` in the REPL. What does this operation do? Where can you find it in Scaladoc?
5. What does `10 max 2` mean? In which class is the `max` method defined?
6. Using `BigInt`, compute 2^{1024} .
7. What do you need to import so that you can get a random prime as `probablePrime(100, Random)`, without any qualifiers before `probablePrime` and `Random`?
8. One way to create random file or directory names is to produce a random `BigInt` and convert it to base 36, yielding a string such as `"qsnvbewtomcj38o06kul"`. Poke around Scaladoc to find a way of doing this in Scala.
9. How do you get the first character of a string in Scala? The last character?
10. What do the `take`, `drop`, `takeRight`, and `dropRight` string functions do? What advantage or disadvantage do they have over using `substring`?

Control Structures and Functions

Topics in This Chapter **A1**

- 2.1 Conditional Expressions — page 18
- 2.2 Statement Termination — page 19
- 2.3 Block Expressions and Assignments — page 20
- 2.4 Input and Output — page 21
- 2.5 Loops — page 22
- 2.6 Advanced for Loops — page 24
- 2.7 Functions — page 25
- 2.8 Default and Named Arguments **L1** — page 26
- 2.9 Variable Arguments **L1** — page 26
- 2.10 Procedures — page 28
- 2.11 Lazy Values **L1** — page 28
- 2.12 Exceptions — page 29
- Exercises — page 31

Chapter

2

In this chapter, you will learn how to implement conditions, loops, and functions in Scala. You will encounter a fundamental difference between Scala and other programming languages. In Java or C++, we differentiate between *expressions* (such as `3 + 4`) and *statements* (for example, an `if` statement). An expression has a value; a statement carries out an action. In Scala, almost all constructs have values. This feature can make programs more concise and easier to read.

Here are the highlights of this chapter:

- An `if` expression has a value.
- A block has a value—the value of its last expression.
- The Scala `for` loop is like an “enhanced” Java `for` loop.
- Semicolons are (mostly) optional.
- The `void` type is `Unit`.
- Avoid using `return` in a function.
- Beware of missing `=` in a function definition.
- Exceptions work just like in Java or C++, but you use a “pattern matching” syntax for `catch`.
- Scala has no checked exceptions.

2.1 Conditional Expressions

Scala has an `if/else` construct with the same syntax as in Java or C++. However, in Scala, an `if/else` has a value, namely the value of the expression that follows the `if` or `else`. For example,

```
if (x > 0) 1 else -1
```

has a value of 1 or -1, depending on the value of `x`. You can put that value in a variable:

```
val s = if (x > 0) 1 else -1
```

This has the same effect as

```
if (x > 0) s = 1 else s = -1
```

However, the first form is better because it can be used to initialize a `val`. In the second form, `s` needs to be a `var`.

(As already mentioned, semicolons are mostly optional in Scala—see Section 2.2, “Statement Termination,” on page 19.)

Java and C++ have a `?:` operator for this purpose. The expression

```
x > 0 ? 1 : -1 // Java or C++
```

is equivalent to the Scala expression `if (x > 0) 1 else -1`. However, you can't put statements inside a `?:` expression. The Scala `if/else` combines the `if/else` and `?:` constructs that are separate in Java and C++.

In Scala, every expression has a type. For example, the expression `if (x > 0) 1 else -1` has the type `Int` because both branches have the type `Int`. The type of a mixed-type expression, such as

```
if (x > 0) "positive" else -1
```

is the common supertype of both branches. In this example, one branch is a `java.lang.String`, and the other an `Int`. Their common supertype is called `Any`. (See Section 8.11, “The Scala Inheritance Hierarchy,” on page 100 for details.)

If the `else` part is omitted, for example in

```
if (x > 0) 1
```

then it is possible that the `if` statement yields no value. However, in Scala, every expression is supposed to have *some* value. This is finessed by introducing a class `Unit` that has one value, written as `()`. The `if` statement without an `else` is equivalent to

```
if (x > 0) 1 else ()
```

Think of `()` as a placeholder for “no useful value,” and of `Unit` as an analog of `void` in Java or C++.

(Technically speaking, `void` has no value whereas `Unit` has one value that signifies “no value.” If you are so inclined, you can ponder the difference between an empty wallet and a wallet with a bill labeled “no dollars.”)



NOTE: Scala has no switch statement, but it has a much more powerful pattern matching mechanism that we will discuss in Chapter 14. For now, just use a sequence of `if` statements.



CAUTION: The REPL is more nearsighted than the compiler—it only sees one line of code at a time. For example, when you type

```
if (x > 0) 1
else if (x == 0) 0 else -1
```

the REPL executes `if (x > 0) 1` and shows the answer. Then it gets confused about the `else` keyword.

If you want to break the line before the `else`, use braces:

```
if (x > 0) { 1
} else if (x == 0) 0 else -1
```

This is only a concern in the REPL. In a compiled program, the parser will find the `else` on the next line.



TIP: If you want to paste a block of code into the REPL without worrying about its nearsightedness, use *paste mode*. Type

```
:paste
```

Then paste in the code block and type `Ctrl+D`. The REPL will then analyze the block in its entirety.

2.2 Statement Termination

In Java and C++, every statement ends with a semicolon. In Scala—like in JavaScript and other scripting languages—a semicolon is never required if it falls just before the end of the line. A semicolon is also optional before an `,` an `else`, and similar locations where it is clear from context that the end of a statement has been reached.

However, if you want to have more than one statement on a single line, you need to separate them with semicolons. For example,

```
if (n > 0) { r = r * n; n -= 1 }
```

A semicolon is needed to separate `r = r * n` and `n -= 1`. Because of the `}`, no semicolon is needed after the second statement.

If you want to continue a long statement over two lines, make sure that the first line ends in a symbol that *cannot be* the end of a statement. An operator is often a good choice:

```
s = s0 + (v - v0) * t + // The + tells the parser that this is not the end
  0.5 * (a - a0) * t * t
```

In practice, long expressions usually involve function or method calls, and then you don't need to worry much—after an opening `(`, the compiler won't infer the end of a statement until it has seen the matching `)`.

In the same spirit, Scala programmers favor the Kernighan & Ritchie brace style:

```
if (n > 0) {
  r = r * n
  n -= 1
}
```

The line ending with a `{` sends a clear signal that there is more to come.

Many programmers coming from Java or C++ are initially uncomfortable about omitting semicolons. If you prefer to put them in, feel free to—they do no harm.

2.3 Block Expressions and Assignments

In Java or C++, a block statement is a sequence of statements enclosed in `{ }` . You use a block statement whenever you need to put multiple actions in the body of a branch or loop statement.

In Scala, a `{ }` block contains a sequence of *expressions*, and the result is also an expression. The value of the block is the value of the last expression.

This feature can be useful if the initialization of a `val` takes more than one step. For example,

```
val distance = { val dx = x - x0; val dy = y - y0; sqrt(dx * dx + dy * dy) }
```

The value of the `{ }` block is the last expression, shown here in bold. The variables `dx` and `dy`, which were only needed as intermediate values in the computation, are neatly hidden from the rest of the program.

In Scala, assignments have no value—or, strictly speaking, they have a value of type `Unit`. Recall that the `Unit` type is the equivalent of the `void` type in Java and C++, with a single value written as `()`.

A block that ends with an assignment, such as

```
{ r = r * n; n -= 1 }
```

has a `Unit` value. This is not a problem, just something to be aware of when defining functions—see Section 2.7, “Functions,” on page 25.

Since assignments have `Unit` value, don’t chain them together.

```
x = y = 1 // No
```

The value of `y = 1` is `()`, and it’s highly unlikely that you wanted to assign a `Unit` to `x`. (In contrast, in Java and C++, the value of an assignment is the value that is being assigned. In those languages, chained assignments are useful.)

2.4 Input and Output

To print a value, use the `print` or `println` function. The latter adds a newline character after the printout. For example,

```
print("Answer: ")  
println(42)
```

yields the same output as

```
println("Answer: " + 42)
```

There is also a `printf` function with a C-style format string:

```
printf("Hello, %s! You are %d years old.%n", name, age)
```

Or better, use *string interpolation*

```
print(f"Hello, ${name}! In six months, you'll be ${age + 0.5}%7.2f years old.%n")
```

A formatted string is prefixed with the letter `f`. It contains expressions that are prefixed with `$` and optionally followed by C-style format strings. The expression `$name` is replaced with the value of the variable `name`. The expression `${age + 0.5}%7.2f` is replaced with the value of `age + 0.5`, formatted as a floating-point number of width 7 and precision 2. You need `${...}` around expressions that are not variable names.

Using the `f` interpolator is better than using the `printf` method because it is type-safe. If you accidentally use `%f` with an expression that isn’t a number, the compiler reports an error.



NOTE: Formatted strings are one of three predefined string interpolators in the Scala library. With a prefix of `s`, strings can contain expressions but not format directives. With a prefix of `raw`, escape sequences in a string are not evaluated. For example, `raw"\n` is a “newline” that starts with a backslash and the letter `n`, not a newline character.

To include a `$` sign in an interpolated string, double it. For example, `s"$$price"` yields a dollar sign followed by the value of `price`.

You can also define your own interpolators—see Exercise 11 on page 32. However, interpolators that produce compile-time errors (such as the `f` interpolator) need to be implemented as “macros,” an experimental Scala feature that is beyond the scope of this book.

You can read a line of input from the console with the `readLine` method of the `scala.io.StdIn` class. To read a numeric, Boolean, or character value, use `readInt`, `readDouble`, `readByte`, `readShort`, `readLong`, `readFloat`, `readBoolean`, or `readChar`. The `readLine` method, but not the other ones, takes a prompt string:

```
import scala.io
val name = StdIn.readLine("Your name: ")
print("Your age: ")
val age = StdIn.readInt()
println(s"Hello, ${name}! Next year, you will be ${age + 1}.")
```

2.5 Loops

Scala has the same `while` and `do` loops as Java and C++. For example,

```
while (n > 0) {
    r = r * n
    n -= 1
}
```

Scala has no direct analog of the `for (initialize; test; update)` loop. If you need such a loop, you have two choices. You can use a `while` loop. Or, you can use a `for` statement like this:

```
for (i <- 1 to n)
    r = r * i
```

You saw the `to` method of the `RichInt` class in Chapter 1. The call `1 to n` returns a Range of the numbers from 1 to `n` (inclusive).

The construct

```
for (i <- expr)
```

makes the variable `i` traverse all values of the expression to the right of the `<-`. Exactly how that traversal works depends on the type of the expression. For a Scala collection, such as a `Range`, the loop makes `i` assume each value in turn.



NOTE: There is no `val` or `var` before the variable in the `for` loop. The type of the variable is the element type of the collection. The scope of the loop variable extends until the end of the loop.

When traversing a string, you can loop over the index values:

```
val s = "Hello"
var sum = 0
for (i <- 0 to s.length - 1)
    sum += s(i)
```

In this example, there is actually no need to use indexes. You can directly loop over the characters:

```
var sum = 0
for (ch <- "Hello") sum += ch
```

In Scala, loops are not used as often as in other languages. As you will see in Chapter 12, you can often process the values in a sequence by applying a function to all of them, which can be done with a single method call.



NOTE: Scala has no `break` or `continue` statements to break out of a loop. What to do if you need a break? Here are a few options:

- Use a Boolean control variable.
- Use nested functions—you can return from the middle of a function.
- Use the `break` method in the `Breaks` object:

```
import scala.util.control.Breaks._
breakable {
    for (...) {
        if (...) break; // Exits the breakable block
        ...
    }
}
```

Here, the control transfer is done by throwing and catching an exception, so you should avoid this mechanism when time is of essence.



NOTE: In Java, you cannot have two local variables with the same name and overlapping scope. In Scala, there is no such prohibition, and the normal shadowing rule applies. For example, the following is perfectly legal:

```
val n = ...
for (n <- 1 to 10) {
    // Here n refers to the loop variable
}
```

2.6 Advanced for Loops

In the preceding section, you saw the basic form of the `for` loop. However, this construct is much richer in Scala than in Java or C++. This section covers the advanced features.

You can have multiple *generators* of the form *variable* \leftarrow *expression*. Separate them by semicolons. For example,

```
for (i <- 1 to 3; j <- 1 to 3) print(f"${10 * i + j}%3d")
// Prints 11 12 13 21 22 23 31 32 33
```

Each generator can have a *guard*, a Boolean condition preceded by `if`:

```
for (i <- 1 to 3; j <- 1 to 3 if i != j) print(f"${10 * i + j}%3d")
// Prints 12 13 21 23 31 32
```

Note that there is no semicolon before the `if`.

You can have any number of *definitions*, introducing variables that can be used inside the loop:

```
for (i <- 1 to 3; from = 4 - i; j <- from to 3) print(f"${10 * i + j}%3d")
// Prints 13 22 23 31 32 33
```

When the body of the `for` loop starts with `yield`, the loop constructs a collection of values, one for each iteration:

```
for (i <- 1 to 10) yield i % 3
// Yields Vector(1, 2, 0, 1, 2, 0, 1, 2, 0, 1)
```

This type of loop is called a `for comprehension`.

The generated collection is compatible with the first generator.

```
for (c <- "Hello"; i <- 0 to 1) yield (c + i).toChar
// Yields "HIfelmlop"
for (i <- 0 to 1; c <- "Hello") yield (c + i).toChar
// Yields Vector('H', 'e', 'l', 'l', 'o', 'I', 'f', 'm', 'm', 'p')
```



NOTE: If you prefer, you can enclose the generators, guards, and definitions of a for loop in braces, and you can use newlines instead of semicolons to separate them:

```
for { i <- 1 to 3
      from = 4 - i
      j <- from to 3 }
```

2.7 Functions

Scala has functions in addition to methods. A method operates on an object, but a function doesn't. C++ has functions as well, but in Java, you have to imitate them with static methods.

To define a function, specify the function's name, parameters, and body like this:

```
def abs(x: Double) = if (x >= 0) x else -x
```

You must specify the types of all parameters. However, as long as the function is not recursive, you need not specify the return type. The Scala compiler determines the return type from the type of the expression to the right of the = symbol.

If the body of the function requires more than one expression, use a block. The last expression of the block becomes the value that the function returns. For example, the following function returns the value of r after the for loop.

```
def fac(n : Int) = {
  var r = 1
  for (i <- 1 to n) r = r * i
  r
}
```

There is no need for the return keyword in this example. It is possible to use return as in Java or C++, to exit a function immediately, but that is not commonly done in Scala.



TIP: While there is nothing wrong with using return in a named function (except the waste of seven keystrokes), it is a good idea to get used to life without return. Pretty soon, you will be using lots of *anonymous functions*, and there, return doesn't return a value to the caller but breaks out to the enclosing named function. Think of return as a kind of break statement for functions, and only use it when you want that breakout functionality.

With a recursive function, you must specify the return type. For example,

```
def fac(n: Int): Int = if (n <= 0) 1 else n * fac(n - 1)
```

Without the return type, the Scala compiler couldn't verify that the type of $n * \text{fac}(n - 1)$ is an Int.



NOTE: Some programming languages (such as ML and Haskell) *can* infer the type of a recursive function, using the Hindley-Milner algorithm. However, this doesn't work well in an object-oriented language. Extending the Hindley-Milner algorithm so it can handle subtypes is still a research problem.

2.8 Default and Named Arguments L1

You can provide default arguments for functions that are used when you don't specify explicit values. For example,

```
def decorate(str: String, left: String = "[", right: String = "]") =
    left + str + right
```

This function has two parameters, `left` and `right`, with default arguments "`[`" and "`]`".

If you call `decorate("Hello")`, you get "`[Hello]`". If you don't like the defaults, supply your own: `decorate("Hello", "<<<", ">>>")`.

If you supply fewer arguments than there are parameters, the defaults are applied from the end. For example, `decorate("Hello", ">>[")` uses the default value of the `right` parameter, yielding "`>>[Hello]`".

You can also specify the parameter names when you supply the arguments. For example,

```
decorate(left = "<<<", str = "Hello", right = ">>>")
```

The result is "`<<<Hello>>>`". Note that the named arguments need not be in the same order as the parameters.

Named arguments can make a function call more readable. They are also useful if a function has many default parameters.

You can mix unnamed and named arguments, provided the unnamed ones come first:

```
decorate("Hello", right = "]<<<") // Calls decorate("Hello", "[", "]<<<")
```

2.9 Variable Arguments L1

Sometimes, it is convenient to implement a function that can take a variable number of arguments. The following example shows the syntax:

```
def sum(args: Int*) = {
    var result = 0
    for (arg <- args) result += arg
    result
}
```

You can call this function with as many arguments as you like.

```
val s = sum(1, 4, 9, 16, 25)
```

The function receives a single parameter of type `Seq`, which we will discuss in Chapter 13. For now, all you need to know is that you can use a `for` loop to visit each element.

If you already have a sequence of values, you cannot pass it directly to such a function. For example, the following is not correct:

```
val s = sum(1 to 5) // Error
```

If the `sum` function is called with one argument, that must be a single integer, not a range of integers. The remedy is to tell the compiler that you want the parameter to be considered an argument sequence. Append `: _*`, like this:

```
val s = sum(1 to 5: _) // Consider 1 to 5 as an argument sequence
```

This call syntax is needed in a recursive definition:

```
def recursiveSum(args: Int*) : Int = {
    if (args.length == 0) 0
    else args.head + recursiveSum(args.tail : _*)
}
```

Here, the `head` of a sequence is its initial element, and `tail` is a sequence of all other elements. That's again a `Seq`, and we have to use `: _*` to convert it to an argument sequence.



CAUTION: When you call a Java method with variable arguments of type `Object`, such as `PrintStream.printf` or `MessageFormat.format`, you need to convert any primitive types by hand. For example,

```
val str = MessageFormat.format("The answer to {0} is {1}",
    "everything", 42.asInstanceOf[AnyRef])
```

This is the case for any `Object` parameter, but I mention it here because it is most common with `varargs` methods.

2.10 Procedures

Scala has a special notation for a function that returns no value. If the function body is enclosed in braces *without a preceding = symbol*, then the return type is `Unit`. Such a function is called a *procedure*. A procedure returns no value, and you only call it for its side effect. For example, the following procedure prints a string inside a box, like

```
-----
|Hello|
-----
```

Since the procedure doesn't return any value, we omit the = symbol.

```
def box(s : String) { // Look carefully: no =
    val border = "-" * (s.length + 2)
    print(f"$border%n|$s|%n$border%n")
}
```

Some people (not me) dislike this concise syntax for procedures and suggest that you always use an explicit return type of `Unit`:

```
def box(s : String): Unit = {
    ...
}
```



CAUTION: The concise procedure syntax can be a surprise for Java and C++ programmers. It is a common error to accidentally omit the = in a function definition. You then get an error message at the point where the function is called: You are told that `Unit` is not acceptable at that location.

2.11 Lazy Values L1

When a `val` is declared as `lazy`, its initialization is deferred until it is accessed for the first time. For example,

```
lazy val words = scala.io.Source.fromFile("/usr/share/dict/words").mkString
```

(We will discuss file operations in Chapter 9. For now, just take it for granted that this call reads all characters from a file into a string.)

If the program never accesses `words`, the file is never opened. To verify this, try it out in the REPL, but misspell the file name. There will be no error when the initialization statement is executed. However, if you access `words`, you will get an error message that the file is not found.

Lazy values are useful to delay costly initialization statements. They can also deal with other initialization issues, such as circular dependencies. Moreover, they are essential for developing lazy data structures—see Section 13.12, “Streams,” on page 189.

You can think of lazy values as halfway between val and def. Compare

```
val words = scala.io.Source.fromFile("/usr/share/dict/words").mkString
  // Evaluated as soon as words is defined
lazy val words = scala.io.Source.fromFile("/usr/share/dict/words").mkString
  // Evaluated the first time words is used
def words = scala.io.Source.fromFile("/usr/share/dict/words").mkString
  // Evaluated every time words is used
```



NOTE: Laziness is not cost-free. Every time a lazy value is accessed, a method is called that checks, in a threadsafe manner, whether the value has already been initialized.

2.12 Exceptions

Scala exceptions work the same way as in Java or C++. When you throw an exception, for example

```
throw new IllegalArgumentException("x should not be negative")
```

the current computation is aborted, and the runtime system looks for an exception handler that can accept an `IllegalArgumentException`. Control resumes with the innermost such handler. If no such handler exists, the program terminates.

As in Java, the objects that you throw need to belong to a subclass of `java.lang.Throwable`. However, unlike Java, Scala has no “checked” exceptions—you never have to declare that a function or method might throw an exception.



NOTE: In Java, “checked” exceptions are checked at compile time. If your method might throw an `IOException`, you must declare it. This forces programmers to think where those exceptions should be handled, which is a laudable goal. Unfortunately, it can also give rise to monstrous method signatures such as `void doSomething() throws IOException, InterruptedException, ClassNotFoundException`. Many Java programmers detest this feature and end up defeating it by either catching exceptions too early or using excessively general exception classes. The Scala designers decided against checked exceptions, recognizing that thorough compile-time checking isn’t always a good thing.

A throw expression has the special type `Nothing`. That is useful in `if/else` expressions. If one branch has type `Nothing`, the type of the `if/else` expression is the type of the other branch. For example, consider

```
if (x >= 0) { sqrt(x)
} else throw new IllegalArgumentException("x should not be negative")
```

The first branch has type `Double`, the second has type `Nothing`. Therefore, the `if/else` expression also has type `Double`.

The syntax for catching exceptions is modeled after the pattern matching syntax (see Chapter 14).

```
val url = new URL("http://horstmann.com/fred-tiny.gif")
try {
    process(url)
} catch {
    case _: MalformedURLException => println(s"Bad URL: $url")
    case ex: IOException => ex.printStackTrace()
}
```

As in Java or C++, the more general exception types should come after the more specific ones.

Note that you can use `_` for the variable name if you don't need it.

The `try/finally` statement lets you dispose of a resource whether or not an exception has occurred. For example:

```
val in = new URL("http://horstmann.com/fred.gif").openStream()
try {
    process(in)
} finally {
    in.close()
}
```

The `finally` clause is executed whether or not the `process` function throws an exception. The reader is always closed.

This code is a bit subtle, and it raises several issues.

- What if the `URL` constructor or the `openStream` method throws an exception? Then the `try` block is never entered, and neither is the `finally` clause. That's just as well—`in` was never initialized, so it makes no sense to invoke `close` on it.
- Why isn't `val in = new URL(...).openStream()` inside the `try` block? Then the scope of `in` would not extend to the `finally` clause.
- What if `in.close()` throws an exception? Then that exception is thrown out of the statement, superseding any earlier one. (This is just like in Java, and it isn't very nice. Ideally, the old exception would stay attached to the new one.)

Note that try/catch and try/finally have complementary goals. The try/catch statement handles exceptions, and the try/finally statement takes some action (usually cleanup) when an exception is not handled. You can combine them into a single try/catch/finally statement:

```
try { ... } catch { ... } finally { ... }
```

This is the same as

```
try { try { ... } catch { ... } } finally { ... }
```

However, that combination is rarely useful.



NOTE: Scala does not have an analog to the Java try-with-resources statement. Consider using the scala-ARM library (<http://jsuereth.com/scala-arm>). Then you can write

```
import resource._
import java.nio.file._
for (in <- resource(Files.newBufferedReader(inPath));
     out <- resource(Files.newBufferedWriter(outPath))) {
  ...
}
```



NOTE: The Try class is designed to work with computations that may fail with exceptions. We will look at it more closely in Chapter 17. Here is a simple example:

```
import scala.io._
val result =
  for (a <- Try { StdIn.readLine("a: ").toInt };
       b <- Try { StdIn.readLine("b: ").toInt })
    yield a / b
```

If an exception occurs in either of the calls to `toInt`, or because of division by zero, then `result` is a `Failure` object, containing the exception that caused the computation to fail. Otherwise, `result` is a `Success` object holding the result of the computation.

Exercises

1. The *signum* of a number is 1 if the number is positive, -1 if it is negative, and 0 if it is zero. Write a function that computes this value.
2. What is the value of an empty block expression {}? What is its type?

3. Come up with one situation where the assignment $x = y = 1$ is valid in Scala.
(Hint: Pick a suitable type for x .)
4. Write a Scala equivalent for the Java loop


```
for (int i = 10; i >= 0; i--) System.out.println(i);
```
5. Write a procedure `countdown(n: Int)` that prints the numbers from n to 0.
6. Write a `for` loop for computing the product of the Unicode codes of all letters in a string. For example, the product of the characters in "Hello" is 9415087488L.
7. Solve the preceding exercise without writing a loop. (Hint: Look at the `StringOps` Scaladoc.)
8. Write a function `product(s : String)` that computes the product, as described in the preceding exercises.
9. Make the function of the preceding exercise a recursive function.
10. Write a function that computes x^n , where n is an integer. Use the following recursive definition:
 - $x^n = y \cdot y$ if n is even and positive, where $y = x^{n/2}$.
 - $x^n = x \cdot x^{n-1}$ if n is odd and positive.
 - $x^0 = 1$.
 - $x^n = 1 / x^{-n}$ if n is negative.

Don't use a return statement.

11. Define a string interpolator `date` so that you can define a `java.time.LocalDate` as `date"$year-$month-$day"`. You need to define an "implicit" class with a `date` method, like this:

```
implicit class DateInterpolator(val sc: StringContext) extends AnyVal {
  def date(args: Any*): LocalDate = . . .
}
```

`args(i)` is the value of the i th expression. Convert each to a string and then to an integer, and pass them to the `LocalDate.of` method. If you already know some Scala, add error handling. Throw an exception if there aren't three arguments, or if they aren't integers, or if they aren't separated by dashes. (You get the strings in between the expressions as `sc.parts`.)

This page intentionally left blank

Working with Arrays

Topics in This Chapter **A1**

- 3.1 Fixed-Length Arrays — page 35
- 3.2 Variable-Length Arrays: Array Buffers — page 36
- 3.3 Traversing Arrays and Array Buffers — page 37
- 3.4 Transforming Arrays — page 38
- 3.5 Common Algorithms — page 40
- 3.6 Deciphering Scaladoc — page 41
- 3.7 Multidimensional Arrays — page 42
- 3.8 Interoperating with Java — page 43
- Exercises — page 44

Chapter

3

In this chapter, you will learn how to work with arrays in Scala. Java and C++ programmers usually choose an array or its close relation (such as array lists or vectors) when they need to collect a bunch of elements. In Scala, there are other choices (see Chapter 13), but for now, I'll assume you are impatient and just want to get going with arrays.

Key points of this chapter:

- Use an `Array` if the length is fixed, and an `ArrayBuffer` if the length can vary.
- Don't use `new` when supplying initial values.
- Use `()` to access elements.
- Use `for (elem <- arr)` to traverse the elements.
- Use `for (elem <- arr if ...) ... yield ...` to transform into a new array.
- Scala and Java arrays are interoperable; with `ArrayBuffer`, use `scala.collection.JavaConversions`.

3.1 Fixed-Length Arrays

If you need an array whose length doesn't change, use the `Array` type in Scala. For example,

```

val nums = new Array[Int](10)
    // An array of ten integers, all initialized with zero
val a = new Array[String](10)
    // A string array with ten elements, all initialized with null
val s = Array("Hello", "World")
    // An Array[String] of length 2—the type is inferred
    // Note: No new when you supply initial values
s(0) = "Goodbye"
    // Array("Goodbye", "World")
    // Use () instead of [] to access elements

```

Inside the JVM, a Scala `Array` is implemented as a Java array. The arrays in the preceding example have the type `java.lang.String[]` inside the JVM. An array of `Int`, `Double`, or another equivalent of the Java primitive types is a primitive type array. For example, `Array(2,3,5,7,11)` is an `int[]` in the JVM.

3.2 Variable-Length Arrays: Array Buffers

Java has `ArrayList` and C++ has `vector` for arrays that grow and shrink on demand. The equivalent in Scala is the `ArrayBuffer`.

```

import scala.collection.mutable.ArrayBuffer
val b = ArrayBuffer[Int]()
    // Or new ArrayBuffer[Int]
    // An empty array buffer, ready to hold integers
b += 1
    // ArrayBuffer(1)
    // Add an element at the end with +=
b += (1, 2, 3, 5)
    // ArrayBuffer(1, 1, 2, 3, 5)
    // Add multiple elements at the end by enclosing them in parentheses
b ++= Array(8, 13, 21)
    // ArrayBuffer(1, 1, 2, 3, 5, 8, 13, 21)
    // You can append any collection with the ++= operator
b.trimEnd(5)
    // ArrayBuffer(1, 1, 2)
    // Removes the last five elements

```

Adding or removing elements at the end of an array buffer is an efficient (“amortized constant time”) operation.

You can also insert and remove elements at an arbitrary location, but those operations are not as efficient—all elements after that location must be shifted. For example:

```
b.insert(2, 6)
// ArrayBuffer(1, 1, 6, 2)
// Insert before index 2
b.insert(2, 7, 8, 9)
// ArrayBuffer(1, 1, 7, 8, 9, 6, 2)
// You can insert as many elements as you like
b.remove(2)
// ArrayBuffer(1, 1, 8, 9, 6, 2)
b.remove(2, 3)
// ArrayBuffer(1, 1, 2)
// The second parameter tells how many elements to remove
```

Sometimes, you want to build up an Array, but you don't yet know how many elements you will need. In that case, first make an array buffer, then call

```
b.toArray
// Array(1, 1, 2)
```

Conversely, call `a.toBuffer` to convert the array `a` to an array buffer.

3.3 Traversing Arrays and Array Buffers

In Java and C++, there are several syntactical differences between arrays and array lists/vectors. Scala is much more uniform. Most of the time, you can use the same code for both.

Here is how you traverse an array or array buffer with a `for` loop:

```
for (i <- 0 until a.length)
  println(s"$i: ${a(i)}")
```

The `until` method is similar to the `to` method, except that it excludes the last value. Therefore, the variable `i` goes from `0` to `a.length - 1`.

In general, the construct

```
for (i <- range)
```

makes the variable `i` traverse all values of the range. In our case, the loop variable `i` assumes the values `0, 1, and so on until (but not including) a.length`.

To visit every second element, let `i` traverse

```
0 until a.length by 2
// Range(0, 2, 4, ...)
```

To visit the elements starting from the end of the array, traverse

```
0 until a.length by -1
// Range(..., 2, 1, 0)
```



TIP: Instead of `0 until a.length` or `0 until a.length by -1`, you can use `a.indices` or `a.indices.reverse`.

If you don't need the array index in the loop body, visit the array elements directly:

```
for (elem <- a)
  println(elem)
```

This is similar to the “enhanced” `for` loop in Java or the “range-based” `for` loop in C++. The variable `elem` is set to `a(0)`, then `a(1)`, and so on.

3.4 Transforming Arrays

In the preceding sections, you saw how to work with arrays just like you would in Java or C++. But in Scala, you can go further. It is easy to take an array (or array buffer) and transform it in some way. Such transformations don't modify the original array but yield a new one.

Use a `for` comprehension like this:

```
val a = Array(2, 3, 5, 7, 11)
val result = for (elem <- a) yield 2 * elem
// result is Array(4, 6, 10, 14, 22)
```

The `for/yield` loop creates a new collection of the same type as the original collection. If you started with an array, you get another array. If you started with an array buffer, that's what you get from `for/yield`.

The result contains the expressions after the `yield`, one for each iteration of the loop.

Oftentimes, when you traverse a collection, you only want to process the elements that match a particular condition. This is achieved with a *guard*: an `if` inside the `for`. Here we double every even element, dropping the odd ones:

```
for (elem <- a if elem % 2 == 0) yield 2 * elem
```

Keep in mind that the result is a new collection—the original collection is not affected.



NOTE: Alternatively, you could write

```
a.filter(_ % 2 == 0).map(2 * _)
```

or even

```
a filter { _ % 2 == 0 } map { 2 * _ }
```

Some programmers with experience in functional programming prefer `filter` and `map` to guards and `yield`. That's just a matter of style—the `for/yield` loop does exactly the same work. Use whichever you find easier.

Suppose we want to remove all negative elements from an array buffer of integers. A traditional sequential solution might traverse the array buffer and remove unwanted elements as they are encountered.

```
var n = a.length
var i = 0
while (i < n) {
  if (a(i) >= 0) i += 1
  else { a.remove(i); n -= 1 }
}
```

That's a bit fussy—you have to remember *not* to increment `i` when you remove the element, and to decrement `n` instead. It is also not efficient to remove elements from the middle of the array buffer. This loop unnecessarily moves elements that will later be removed.

In Scala, the obvious solution is to use a `for/yield` loop and keep all non-negative elements:

```
val result = for (elem <- a if elem >= 0) yield elem
```

The result is a new array buffer. Suppose that we want to modify the original array buffer instead, removing the unwanted elements. Then we can collect their positions:

```
val positionsToRemove = for (i <- a.indices if a(i) < 0) yield i
```

Now remove the elements at those positions, starting from the back:

```
for (i <- positionsToRemove.reverse) a.remove(i)
```

Or better, remember the positions to keep, copy them over, and then shorten the buffer:

```
val positionsToKeep = for (i <- a.indices if a(i) >= 0) yield i
for (j <- positionsToKeep.indices) a(j) = a(positionsToKeep(j))
a.trimEnd(a.length - positionsToKeep.length)
```

The key observation is that it is better to have *all index values together* instead of seeing them one by one.

3.5 Common Algorithms

It is often said that a large percentage of business computations are nothing but computing sums and sorting. Fortunately, Scala has built-in functions for these tasks.

```
Array(1, 7, 2, 9).sum
// 19
// Works for ArrayBuffer too
```

In order to use the `sum` method, the element type must be a numeric type: either an integral or floating-point type or `BigInteger`/`BigDecimal`.

Similarly, the `min` and `max` methods yield the smallest and largest element in an array or array buffer.

```
ArrayBuffer("Mary", "had", "a", "little", "lamb").max
// "little"
```

The `sorted` method sorts an array or array buffer and *returns* the sorted array or array buffer, without modifying the original:

```
val b = ArrayBuffer(1, 7, 2, 9)
val bSorted = b.sorted
// b is unchanged; bSorted is ArrayBuffer(1, 2, 7, 9)
```

You can also supply a comparison function, but then you should use the `sortWith` method:

```
val bDescending = b.sortWith(_ > _) // ArrayBuffer(9, 7, 2, 1)
```

See Chapter 12 for the function syntax.

You can sort an array, but not an array buffer, in place:

```
val a = Array(1, 7, 2, 9)
scala.util.Sorting.quickSort(a)
// a is now Array(1, 2, 7, 9)
```

For the `min`, `max`, and `quickSort` methods, the element type must have a comparison operation. This is the case for numbers, strings, and other types with the `Ordered` trait.

Finally, if you want to display the contents of an array or array buffer, the `mkString` method lets you specify the separator between elements. A second variant has parameters for the prefix and suffix. For example,

```
a.mkString(" and ")
// "1 and 2 and 7 and 9"
a.mkString("<", ", ", ">")
// "<1,2,7,9>"
```

Contrast with `toString`:

```
a.toString
// "[I@b73e5"
// This is the useless toString method from Java
b.toString
// "ArrayBuffer(1, 7, 2, 9)"
// The toString method reports the type, which is useful for debugging
```

3.6 Deciphering Scaladoc

There are lots of useful methods on arrays and array buffers, and it is a good idea to browse the Scala documentation to get an idea of what's there.

Scala has a richer type system than Java, so you may encounter some strange-looking syntax as you browse the Scala documentation. Fortunately, you don't have to understand all nuances of the type system to do useful work. Use Table 3–1 as a “decoder ring.”

Table 3–1 Scaladoc Decoder Ring

Scaladoc	Explanation
<code>def count(p: (A) => Boolean): Int</code>	This method takes a <i>predicate</i> , a function from A to Boolean. It counts for how many elements the function is true. For example, <code>a.count(_ > 0)</code> counts how many elements of <code>a</code> are positive.
<code>def append(elems: A*): Unit</code>	This method takes <i>zero or more</i> arguments of type A. For example, <code>b.append(1, 7, 2, 9)</code> appends four elements to <code>b</code> .
<code>def appendAll(xs: TraversableOnce[A]): Unit</code>	The <code>xs</code> parameter can be any collection with the <code>TraversableOnce</code> trait, a trait in the Scala collections hierarchy. Other common traits that you may encounter in Scaladoc are <code>Traversable</code> and <code>Iterable</code> . All Scala collections implement these traits, and the difference between them is academic for library users. Simply think “any collection” when you see one of these.

(Continues)

Table 3–1 Scaladoc Decoder Ring (*Continued*)

Scaladoc	Explanation
<code>def containsSlice[B](that: GenSeq[B]): Boolean</code>	A <code>GenSeq</code> or <code>Seq</code> is a collection whose elements are arranged in sequential order. Think “array, list, or string.”
<code>def += (elem: A): ArrayBuffer.this.type</code>	This method returns <code>this</code> , which allows you to chain calls, for example: <code>b += 4 -- 5</code> . When you work with an <code>ArrayBuffer[A]</code> , you can just think of the method as <code>def += (elem: A) : ArrayBuffer[A]</code> . If someone forms a subclass of <code>ArrayBuffer</code> , then the return type of <code>+=</code> is that subclass.
<code>def copyToArray[B >: A] (xs: Array[B]): Unit</code>	Note that the function copies an <code>ArrayBuffer[A]</code> into an <code>Array[B]</code> . Here, <code>B</code> is allowed to be a <i>supertype</i> of <code>A</code> . For example, you can copy from an <code>ArrayBuffer[Int]</code> to an <code>Array[Any]</code> . At first reading, just ignore the <code>[B >: A]</code> and replace <code>B</code> with <code>A</code> .
<code>def sorted[B >: A] (implicit cmp: Ordering[B]): ArrayBuffer[A]</code>	The element type <code>A</code> must have a supertype <code>B</code> for which an “ <i>implicit</i> ” object of type <code>Ordering[B]</code> exists. Such an ordering exists for numbers, strings, and other types with the <code>Ordered</code> trait, as well as for classes that implement the Java <code>Comparable</code> interface.
<code>def ++:[B >: A, That](that: collection.Traversable[B])(implicit bf: CanBuildFrom[ArrayBuffer[A], B, That]): That</code>	This declaration happens when the method creates a new collection. Most of the time, Scaladoc hides such complex descriptions and shows a simpler one tagged with “[use case]”. In this case, it doesn’t. You can simplify the declaration in your mind to a “happy day scenario,” like this: <code>def ++:(that: ArrayBuffer[A]) : ArrayBuffer[A]</code> What if that is an <code>ArrayBuffer[B]</code> or some other collection? Try it out in the REPL: <code>ArrayBuffer('a', 'b') ++: "cd" // Yields a string "abcd"</code>

3.7 Multidimensional Arrays

Like in Java, multidimensional arrays are implemented as arrays of arrays. For example, a two-dimensional array of `Double` values has the type `Array[Array[Double]]`. To construct such an array, use the `ofDim` method:

```
val matrix = Array.ofDim[Double](3, 4) // Three rows, four columns
```

To access an element, use two pairs of parentheses:

```
matrix(row)(column) = 42
```

You can make ragged arrays, with varying row lengths:

```
val triangle = new Array[Array[Int]](10)
for (i <- triangle.indices)
  triangle(i) = new Array[Int](i + 1)
```

3.8 Interoperating with Java

Since Scala arrays are implemented as Java arrays, you can pass them back and forth between Java and Scala.

This works in almost all cases, except if the array element type isn't an exact match. In Java, an array of a given type is automatically converted to an array of a supertype. For example, a Java `String[]` array can be passed to a method that expects a Java `Object[]` array. Scala does not permit this automatic conversion because it is unsafe. (See Chapter 18 for a detailed explanation.)

Suppose you want to invoke a Java method with an `Object[]` parameter, such as `java.util.Arrays.binarySearch(Object[] a, Object key)`:

```
val a = Array("Mary", "a", "had", "lamb", "little")
java.util.Arrays.binarySearch(a, "beef") // Does not work
```

This does not work because Scala will not convert an `Array[String]` into an `Array[Object]`. You can force the conversion like this:

```
java.util.Arrays.binarySearch(a.asInstanceOf[Array[Object]], "beef")
```



NOTE: This is just an example to show how to overcome element type differences. If you want to carry out binary search in Scala, do it like this:

```
import scala.collection.Searching._
val result = a.search("beef")
```

The result is `Found(n)` if the element was found at position `n` or `InsertionPoint(n)` if the element was not found but should be inserted before position `n`.

If you call a Java method that receives or returns a `java.util.List`, you could, of course, use a Java `ArrayList` in your Scala code—but that is unattractive. Instead, import the implicit conversion methods in `scala.collection.JavaConversions`. Then you can use Scala buffers in your code, and they automatically get wrapped into Java lists when calling a Java method.

For example, the `java.lang.ProcessBuilder` class has a constructor with a `List<String>` parameter. Here is how you can call it from Scala:

```
import scala.collection.JavaConversions.bufferAsJavaList
import scala.collection.mutable.ArrayBuffer
val command = ArrayBuffer("ls", "-al", "/home/cay")
val pb = new ProcessBuilder(command) // Scala to Java
```

The Scala buffer is wrapped into an object of a Java class that implements the `java.util.List` interface.

Conversely, when a Java method returns a `java.util.List`, you can have it automatically converted into a Buffer:

```
import scala.collection.JavaConversions.asScalaBuffer
import scala.collection.mutable.Buffer
val cmd : Buffer[String] = pb.command() // Java to Scala
// You can't use ArrayBuffer—the wrapped object is only guaranteed to be a Buffer
```

If the Java method returns a wrapped Scala buffer, then the implicit conversion unwraps the original object. In our example, `cmd == command`.

Exercises

1. Write a code snippet that sets `a` to an array of `n` random integers between 0 (inclusive) and `n` (exclusive).
2. Write a loop that swaps adjacent elements of an array of integers. For example, `Array(1, 2, 3, 4, 5)` becomes `Array(2, 1, 4, 3, 5)`.
3. Repeat the preceding assignment, but produce a new array with the swapped values. Use `for/yield`.
4. Given an array of integers, produce a new array that contains all positive values of the original array, in their original order, followed by all values that are zero or negative, in their original order.
5. How do you compute the average of an `Array[Double]`?
6. How do you rearrange the elements of an `Array[Int]` so that they appear in reverse sorted order? How do you do the same with an `ArrayBuffer[Int]`?
7. Write a code snippet that produces all values from an array with duplicates removed. (Hint: Look at Scaladoc.)
8. Suppose you are given an array buffer of integers and want to remove all but the first negative number. Here is a sequential solution that sets a flag when the first negative number is called, then removes all elements beyond.

```
var first = true
var n = a.length
var i = 0
while (i < n) {
    if (a(i) >= 0) i += 1
    else {
        if (first) { first = false; i += 1 }
        else { a.remove(i); n -= 1 }
    }
}
```

This is a complex and inefficient solution. Rewrite it in Scala by collecting positions of the negative elements, dropping the first element, reversing the sequence, and calling `a.remove(i)` for each index.

9. Improve the solution of the preceding exercise by collecting the positions that should be moved and their target positions. Make those moves and truncate the buffer. Don't copy any elements before the first unwanted element.
10. Make a collection of all time zones returned by `java.util.TimeZone.getAvailableIDs` that are in America. Strip off the "America/" prefix and sort the result.
11. Import `java.awt.datatransfer._` and make an object of type `SystemFlavorMap` with the call

```
val flavors = SystemFlavorMap.getDefaultFlavorMap().asInstanceOf[SystemFlavorMap]
```

Then call the `getNativesForFlavor` method with parameter `DataFlavor.imageFlavor` and get the return value as a Scala buffer. (Why this obscure class? It's hard to find uses of `java.util.List` in the standard Java library.)

Maps and Tuples

Topics in This Chapter **A1**

- 4.1 Constructing a Map — page 48
- 4.2 Accessing Map Values — page 48
- 4.3 Updating Map Values — page 49
- 4.4 Iterating over Maps — page 50
- 4.5 Sorted Maps — page 50
- 4.6 Interoperating with Java — page 50
- 4.7 Tuples — page 51
- 4.8 Zipping — page 52
- Exercises — page 52

Chapter 4

A classic programmer's saying is, "If you can only have one data structure, make it a hash table." Hash tables—or, more generally, maps—are among the most versatile data structures. As you will see in this chapter, Scala makes it particularly easy to use them.

Maps are collections of key/value pairs. Scala has a general notion of tuples—aggregates of n objects, not necessarily of the same type. A pair is simply a tuple with $n = 2$. Tuples are useful whenever you need to aggregate two or more values together, and we briefly discuss the syntax at the end of this chapter.

Highlights of the chapter are:

- Scala has a pleasant syntax for creating, querying, and traversing maps.
- You need to choose between mutable and immutable maps.
- By default, you get a hash map, but you can also get a tree map.
- You can easily convert between Scala and Java maps.
- Tuples are useful for aggregating values.

4.1 Constructing a Map

You can construct a map as

```
val scores = Map("Alice" -> 10, "Bob" -> 3, "Cindy" -> 8)
```

This constructs an immutable `Map[String, Int]` whose contents can't be changed. If you want a mutable map, use

```
val scores = scala.collection.mutable.Map("Alice" -> 10, "Bob" -> 3, "Cindy" -> 8)
```

If you want to start out with a blank map, you have to supply type parameters:

```
val scores = scala.collection.mutable.Map[String, Int]()
```

In Scala, a map is a collection of *pairs*. A pair is simply a grouping of two values, not necessarily of the same type, such as ("Alice", 10).

The `->` operator makes a pair. The value of

"Alice" -> 10

is

("Alice", 10)

You could have equally well defined the map as

```
val scores = Map(("Alice", 10), ("Bob", 3), ("Cindy", 8))
```

The `->` operator is just a little easier on the eyes than the parentheses. It also supports the intuition that a map data structure is a kind of function that maps keys to values. The difference is that a function computes values, and a map just looks them up.

4.2 Accessing Map Values

In Scala, the analogy between functions and maps is particularly close because you use the `()` notation to look up key values.

```
val bobsScore = scores("Bob") // Like scores.get("Bob") in Java
```

If the map doesn't contain a value for the requested key, an exception is thrown.

To check whether there is a key with the given value, call the `contains` method:

```
val bobsScore = if (scores.contains("Bob")) scores("Bob") else 0
```

Since this call combination is so common, there is a shortcut:

```
val bobsScore = scores.getOrElse("Bob", 0)
// If the map contains the key "Bob", return the value; otherwise, return 0.
```

Finally, the call `map.get(key)` returns an `Option` object that is either `Some(value for key)` or `None`. We discuss the `Option` class in Chapter 14.



NOTE: Given an immutable map, you can turn it into a map with a fixed default value for keys that are not present, or a function to compute such values.

```
val scores1 = scores.withDefaultValue(0)
val zeldasScore1 = scores1.get("Zelda")
// Yields 0 since "Zelda" is not present
val scores2 = scores.withDefault(_.length)
val zeldasScore2 = scores2.get("Zelda")
// Yields 5, applying the length function to the key that is not present
```

4.3 Updating Map Values

In a mutable map, you can update a map value, or add a new one, with a () to the left of an = sign:

```
scores("Bob") = 10
// Updates the existing value for the key "Bob" (assuming scores is mutable)
scores("Fred") = 7
// Adds a new key/value pair to scores (assuming it is mutable)
```

Alternatively, you can use the += operation to add multiple associations:

```
scores += ("Bob" -> 10, "Fred" -> 7)
```

To remove a key and its associated value, use the -= operator:

```
scores -= "Alice"
```

You can't update an immutable map, but you can do something that's just as useful—obtain a new map that has the desired update:

```
val newScores = scores + ("Bob" -> 10, "Fred" -> 7) // New map with update
```

The newScores map contains the same associations as scores, except that "Bob" has been updated and "Fred" added.

Instead of saving the result as a new value, you can update a var:

```
var scores = ...
scores = scores + ("Bob" -> 10, "Fred" -> 7)
```

You can even use the += operator:

```
scores += ("Bob" -> 10, "Fred" -> 7)
```

Similarly, to remove a key from an immutable map, use the - operator to obtain a new map without the key:

```
scores = scores - "Alice"
```

or

```
scores -= "Alice"
```

You might think that it is inefficient to keep constructing new maps, but that is not the case. The old and new maps share most of their structure. (This is possible because they are immutable.)

4.4 Iterating over Maps

The following amazingly simple loop iterates over all key/value pairs of a map:

```
for ((k, v) <- map) process k and v
```

The magic here is that you can use pattern matching in a Scala for loop. (Chapter 14 has all the details.) That way, you get the key and value of each pair in the map without tedious method calls.

If for some reason you want to visit only the keys or values, use the keySet and values methods, as you would in Java. The values method returns an Iterable that you can use in a for loop.

```
scores.keySet // A set such as Set("Bob", "Cindy", "Fred", "Alice")
for (v <- scores.values) println(v) // Prints 10 8 7 10 or some permutation thereof
```

To reverse a map—that is, switch keys and values—use

```
for ((k, v) <- map) yield (v, k)
```

4.5 Sorted Maps

There are two common implementation strategies for maps: hash tables and balanced trees. Hash tables use the hash codes of the keys to scramble entries, so iterating over the elements yields them in unpredictable order. By default, Scala gives you a map based on a hash table because it is usually more efficient. If you need to visit the keys in sorted order, use a `SortedMap` instead.

```
val scores = scala.collection.mutable.SortedMap("Alice" -> 10,
    "Fred" -> 7, "Bob" -> 3, "Cindy" -> 8)
```



TIP: If you want to visit the keys in insertion order, use a `LinkedHashMap`. For example,

```
val months = scala.collection.mutable.LinkedHashMap("January" -> 1,
    "February" -> 2, "March" -> 3, "April" -> 4, "May" -> 5, ...)
```

4.6 Interoperating with Java

If you get a Java map from calling a Java method, you may want to convert it to a Scala map so that you can use the pleasant Scala map API. This is also useful if you want to work with a mutable tree map, which Scala doesn't provide.

Simply add an `import` statement:

```
import scala.collection.JavaConversions.mapAsScalaMap
```

Then trigger the conversion by specifying the Scala map type:

```
val scores: scala.collection.mutable.Map[String, Int] =
  new java.util.TreeMap[String, Int]
```

In addition, you can get a conversion from `java.util.Properties` to a `Map[String, String]`:

```
import scala.collection.JavaConversions.propertiesAsScalaMap
val props: scala.collection.Map[String, String] = System.getProperties()
```

Conversely, to pass a Scala map to a method that expects a Java map, provide the opposite implicit conversion. For example:

```
import scala.collection.JavaConversions.mapAsJavaMap
import java.awt.font.TextAttribute._ // Import keys for map below
val attrs = Map(FAMILY -> "Serif", SIZE -> 12) // A Scala map
val font = new java.awt.Font(attrs) // Expects a Java map
```

4.7 Tuples

Maps are collections of key/value pairs. Pairs are the simplest case of *tuples*—aggregates of values of different types.

A tuple value is formed by enclosing individual values in parentheses. For example,

```
(1, 3.14, "Fred")
```

is a tuple of type

```
Tuple3[Int, Double, java.lang.String]
```

which is also written as

```
(Int, Double, java.lang.String)
```

If you have a tuple, say,

```
val t = (1, 3.14, "Fred")
```

then you can access its components with the methods `_1`, `_2`, `_3`, for example:

```
val second = t._2 // Sets second to 3.14
```

Unlike array or string positions, the component positions of a tuple start with 1, not 0.



NOTE: You can write `t._2` as `t _2` (with a space instead of a period), but not `t_2`.

Usually, it is better to use pattern matching to get at the components of a tuple, for example

```
val (first, second, third) = t // Sets first to 1, second to 3.14, third to "Fred"
```

You can use `_` if you don't need all components:

```
val (first, second, _) = t
```

Tuples are useful for functions that return more than one value. For example, the `partition` method of the `StringOps` class returns a pair of strings, containing the characters that fulfill a condition and those that don't:

```
"New York".partition(_.isUpper) // Yields the pair ("NY", "ew ork")
```

4.8 Zipping

One reason for using tuples is to bundle together values so that they can be processed together. This is commonly done with the `zip` method. For example, the code

```
val symbols = Array("<", "-", ">")
val counts = Array(2, 10, 2)
val pairs = symbols.zip(counts)
```

yields an array of pairs

```
Array(("<", 2), ("-", 10), (">", 2))
```

The pairs can then be processed together:

```
for ((s, n) <- pairs) print(s * n) // Prints <<----->>
```



TIP: The `toMap` method turns a collection of pairs into a map.

If you have a collection of keys and a parallel collection of values, zip them up and turn them into a map like this:

```
keys.zip(values).toMap
```

Exercises

1. Set up a map of prices for a number of gizmos that you covet. Then produce a second map with the same keys and the prices at a 10 percent discount.

2. Write a program that reads words from a file. Use a mutable map to count how often each word appears. To read the words, simply use a `java.util.Scanner`:

```
val in = new java.util.Scanner(new java.io.File("myfile.txt"))
while (in.hasNext()) process in.next()
```

Or look at Chapter 9 for a Scalaesque way.

At the end, print out all words and their counts.

3. Repeat the preceding exercise with an immutable map.
4. Repeat the preceding exercise with a sorted map, so that the words are printed in sorted order.
5. Repeat the preceding exercise with a `java.util.TreeMap` that you adapt to the Scala API.
6. Define a linked hash map that maps "Monday" to `java.util.Calendar.MONDAY`, and similarly for the other weekdays. Demonstrate that the elements are visited in insertion order.
7. Print a table of all Java properties reported by the `getProperties` method of the `java.lang.System` class, like this:

<code>java.runtime.name</code>	Java(TM) SE Runtime Environment
<code>sun.boot.library.path</code>	/home/apps/jdk1.6.0_21/jre/lib/i386
<code>java.vm.version</code>	17.0-b16
<code>java.vm.vendor</code>	Sun Microsystems Inc.
<code>java.vendor.url</code>	http://java.sun.com/
<code>path.separator</code>	:
<code>java.vm.name</code>	Java HotSpot(TM) Server VM

You need to find the length of the longest key before you can print the table.

8. Write a function `minmax(values: Array[Int])` that returns a pair containing the smallest and the largest values in the array.
9. Write a function `lteqgt(values: Array[Int], v: Int)` that returns a triple containing the counts of values less than `v`, equal to `v`, and greater than `v`.
10. What happens when you zip together two strings, such as `"Hello".zip("World")`? Come up with a plausible use case.

Classes

Topics in This Chapter A1

- 5.1 Simple Classes and Parameterless Methods — page 55
- 5.2 Properties with Getters and Setters — page 56
- 5.3 Properties with Only Getters — page 59
- 5.4 Object-Private Fields — page 60
- 5.5 Bean Properties L1 — page 61
- 5.6 Auxiliary Constructors — page 62
- 5.7 The Primary Constructor — page 63
- 5.8 Nested Classes L1 — page 66
- Exercises — page 68

Chapter

5

In this chapter, you will learn how to implement classes in Scala. If you know classes in Java or C++, you won't find this difficult, and you will enjoy the much more concise notation of Scala.

The key points of this chapter are:

- Fields in classes automatically come with getters and setters.
- You can replace a field with a custom getter/setter without changing the client of a class—that is the “uniform access principle.”
- Use the `@BeanProperty` annotation to generate the JavaBeans `getXxx`/`setXxx` methods.
- Every class has a primary constructor that is “interwoven” with the class definition. Its parameters turn into the fields of the class. The primary constructor executes all statements in the body of the class.
- Auxiliary constructors are optional. They are called `this`.

5.1 Simple Classes and Parameterless Methods

In its simplest form, a Scala class looks very much like its equivalent in Java or C++:

```
class Counter {
    private var value = 0 // You must initialize the field
    def increment() { value += 1 } // Methods are public by default
    def current() = value
}
```

In Scala, a class is not declared as `public`. A Scala source file can contain multiple classes, and all of them have public visibility.

To use this class, you construct objects and invoke methods in the usual way:

```
val myCounter = new Counter // Or new Counter()
myCounter.increment()
println(myCounter.current)
```

You can call a parameterless method (such as `current`) with or without parentheses:

```
myCounter.current // OK
myCounter.current() // Also OK
```

Which form should you use? It is considered good style to use `()` for a *mutator* method (a method that changes the object state), and to drop the `()` for an *accessor* method (a method that does not change the object state).

That's what we did in our example:

```
myCounter.increment() // Use () with mutator
println(myCounter.current) // Don't use () with accessor
```

You can enforce this style by declaring `current` without `()`:

```
class Counter {
    ...
    def current = value // No () in definition
}
```

Now class users must use `myCounter.current`, without parentheses.

5.2 Properties with Getters and Setters

When writing a Java class, we don't like to use public fields:

```
public class Person { // This is Java
    public int age; // Frowned upon in Java
}
```

With a public field, anyone could write to `fred.age`, making Fred younger or older. That's why we prefer to use getter and setter methods:

```
public class Person { // This is Java
    private int age;
    public int getAge() { return age; }
    public void setAge(int age) { this.age = age; }
}
```

A getter/setter pair such as this one is often called a *property*. We say that the class Person has an age property.

Why is this any better? By itself, it isn't. Anyone can call fred.setAge(21), keeping him forever twenty-one.

But if that becomes a problem, we can guard against it:

```
public void setAge(int newValue) { if (newValue > age) age = newValue; }
// Can't get younger
```

Getters and setters are better than public fields because they let you start with simple get/set semantics and evolve them as needed.



NOTE: Just because getters and setters are better than public fields doesn't mean they are always good. Often, it is plainly bad if every client can get or set bits and pieces of an object's state. In this section, I show you how to implement properties in Scala. It is up to you to choose wisely when a gettable/settable property is an appropriate design.

Scala provides getter and setter methods for every field. Here, we define a public field:

```
class Person {
    var age = 0
}
```

Scala generates a class for the JVM with a *private* age field and getter and setter methods. These methods are public because we did not declare age as private. (For a private field, the getter and setter methods are private.)

In Scala, the getter and setter methods are called age and age_=. For example,

```
println(fred.age) // Calls the method fred.age()
fred.age = 21 // Calls fred.age_=(21)
```

In Scala, the getters and setters are not named getXxx and setXxx, but they fulfill the same purpose. Section 5.5, “Bean Properties,” on page 61 shows how to generate Java-style getXxx and setXxx methods, so that your Scala classes can interoperate with Java tools.



NOTE: To see these methods with your own eyes, compile the Person class and then look at the bytecode with javap:

```
$ scalac Person.scala
$ javap -private Person
Compiled from "Person.scala"
public class Person extends java.lang.Object implements scala.ScalaObject{
    private int age;
    public int age(); public void age_$eq(int);
    public Person();
}
```

As you can see, the compiler created methods age and age_\$eq. (The = symbol is translated to \$eq because the JVM does not allow an = in a method name.)



TIP: You can run the javap command inside the REPL as
:javap -private Person

At any time, you can redefine the getter and setter methods yourself. For example,

```
class Person {
    private var privateAge = 0 // Make private and rename

    def age = privateAge
    def age_=(newValue: Int) {
        if (newValue > privateAge) privateAge = newValue; // Can't get younger
    }
}
```

The user of your class still accesses fred.age, but now Fred can't get younger:

```
val fred = new Person
fred.age = 30
fred.age = 21
println(fred.age) // 30
```



NOTE: Bertrand Meyer, the inventor of the influential Eiffel language, formulated the *Uniform Access Principle* that states: “All services offered by a module should be available through a uniform notation, which does not betray whether they are implemented through storage or through computation.” In Scala, the caller of fred.age doesn’t know whether age is implemented through a field or a method. (Of course, in the JVM, the service is *always* implemented through a method, either synthesized or programmer-supplied.)



TIP: It may sound scary that Scala generates getter and setter methods for every field. But you have some control over this process.

- If the field is private, the getter and setter are private.
- If the field is a `val`, only a getter is generated.
- If you don't want any getter or setter, declare the field as `private[this]` (see Section 5.4, “Object-Private Fields,” on page 60).

5.3 Properties with Only Getters

Sometimes you want a *read-only property* with a getter but no setter. If the value of the property never changes after the object has been constructed, use a `val` field:

```
class Message {
    val timeStamp = java.time.Instant.now
    ...
}
```

The Scala compiler produces a Java class with a private final field and a public getter method, but no setter.

Sometimes, however, you want a property that a client can't set at will, but that is mutated in some other way. The `Counter` class from Section 5.1, “Simple Classes and Parameterless Methods,” on page 55 is a good example. Conceptually, the counter has a `current` property that is updated when the `increment` method is called, but there is no setter for the property.

You can't implement such a property with a `val`—a `val` never changes. Instead, provide a private field and a property getter, like this:

```
class Counter {
    private var value = 0
    def increment() { value += 1 }
    def current = value // No () in declaration
}
```

Note that there are no `()` in the definition of the getter method. Therefore, you *must* call the method without parentheses:

```
val n = myCounter.current // Calling myCounter.current() is a syntax error
```

To summarize, you have four choices for implementing properties:

1. `var foo`: Scala synthesizes a getter and a setter.
2. `val foo`: Scala synthesizes a getter.

3. You define methods `foo` and `foo_=`.
4. You define a method `foo`.



NOTE: In Scala, you cannot have a write-only property (that is, a property with a setter and no getter).



TIP: When you see a field in a Scala class, remember that it is not the same as a field in Java or C++. It is a private field *together with* a getter (for a `val` field) or a getter and a setter (for a `var` field).

5.4 Object-Private Fields

In Scala (as well as in Java or C++), a method can access the private fields of *all* objects of its class. For example,

```
class Counter {
    private var value = 0
    def increment() { value += 1 }

    def isLess(other : Counter) = value < other.value
        // Can access private field of other object
}
```

Accessing `other.value` is legal because `other` is also a `Counter` object.

Scala allows an even more severe access restriction with the `private[this]` qualifier:

```
private[this] var value = 0 // Accessing someObject.value is not allowed
```

Now, the methods of the `Counter` class can only access the `value` field of the current object, not of other objects of type `Counter`. This access is sometimes called *object-private*, and it is common in some OO languages such as SmallTalk.

With a class-private field, Scala generates private getter and setter methods. However, for an object-private field, no getters and setters are generated at all.



NOTE: Scala allows you to grant access rights to specific classes. The `private[ClassName]` qualifier states that only methods of the given class can access the given field. Here, the `ClassName` must be the name of the class being defined or an enclosing class. (See Section 5.8, “Nested Classes,” on page 66 for a discussion of inner classes.)

In this case, the implementation will generate auxiliary getter and setter methods that allow the enclosing class to access the field. These methods will be public because the JVM does not have a fine-grained access control system, and they will have implementation-dependent names.

5.5 Bean Properties L1

As you saw in the preceding sections, Scala provides getter and setter methods for the fields that you define. However, the names of these methods are not what Java tools expect. The JavaBeans specification (www.oracle.com/technetwork/articles/javase/spec-136004.html) defines a Java property as a pair of `getFoo`/`setFoo` methods (or just a `getFoo` method for a read-only property). Many Java tools rely on this naming convention.

When you annotate a Scala field with `@BeanProperty`, then such methods are automatically generated. For example,

```
import scala.beans.BeanProperty

class Person {
    @BeanProperty var name: String = _
}
```

generates *four* methods:

1. `name: String`
2. `name_=(newValue: String): Unit`
3. `getName(): String`
4. `setName(newValue: String): Unit`

Table 5–1 shows which methods are generated in all cases.

Table 5–1 Generated Methods for Fields

Scala Field	Generated Methods	When to Use
val/var name	public name name_= (var only)	To implement a property that is publicly accessible and backed by a field.
@BeanProperty val/var name	public name getName() name_= (var only) setName(...) (var only)	To interoperate with JavaBeans.
private val/var name	private name name_= (var only)	To confine the field to the methods of this class, just like in Java. Use private unless you really want a public property.
private[this] val/var name	none	To confine the field to methods invoked on the same object. Not commonly used.
private[ClassName] val/var name	implementation-dependent	To grant access to an enclosing class. Not commonly used.



NOTE: If you define a field as a primary constructor parameter (see Section 5.7, “The Primary Constructor,” on page 63), and you want JavaBeans getters and setters, annotate the constructor parameter like this:

```
class Person(@BeanProperty var name: String)
```

5.6 Auxiliary Constructors

As in Java or C++, a Scala class can have as many constructors as you like. However, a Scala class has one constructor that is more important than all the others, called the *primary constructor*. In addition, a class may have any number of *auxiliary constructors*.

We discuss auxiliary constructors first because they are easier to understand. They are similar to constructors in Java or C++, with just two differences.

1. The auxiliary constructors are called `this`. (In Java or C++, constructors have the same name as the class—which is not so convenient if you rename the class.)
2. Each auxiliary constructor *must* start with a call to a previously defined auxiliary constructor or the primary constructor.

Here is a class with two auxiliary constructors:

```
class Person {
    private var name = ""
    private var age = 0

    def this(name: String) { // An auxiliary constructor
        this() // Calls primary constructor
        this.name = name
    }

    def this(name: String, age: Int) { // Another auxiliary constructor
        this(name) // Calls previous auxiliary constructor
        this.age = age
    }
}
```

We will look at the primary constructor in the next section. For now, it is sufficient to know that a class for which you don't define a primary constructor has a primary constructor with no arguments.

You can construct objects of this class in three ways:

```
val p1 = new Person // Primary constructor
val p2 = new Person("Fred") // First auxiliary constructor
val p3 = new Person("Fred", 42) // Second auxiliary constructor
```

5.7 The Primary Constructor

In Scala, every class has a primary constructor. The primary constructor is not defined with a `this` method. Instead, it is interwoven with the class definition.

1. The parameters of the primary constructor are placed *immediately after the class name*.

```
class Person(val name: String, val age: Int) {
    // Parameters of primary constructor in ...
    ...
}
```

Parameters of the primary constructor turn into fields that are initialized with the construction parameters. In our example, `name` and `age` become fields of the `Person` class. A constructor call such as `new Person("Fred", 42)` sets the `name` and `age` fields.

Half a line of Scala is the equivalent of seven lines of Java:

```
public class Person { // This is Java
    private String name; private int age; public Person(String name, int age) {
        this.name = name; this.age = age;
    }
    public String name() { return this.name; } public int age() { return this.age; }
    ...
}
```

2. The primary constructor executes *all statements in the class definition*. For example, in the following class

```
class Person(val name: String, val age: Int) {
    println("Just constructed another person")
    def description = s"$name is $age years old"
}
```

the `println` statement is a part of the primary constructor. It is executed whenever an object is constructed.

This is useful when you need to configure a field during construction. For example:

```
class MyProg {
    private val props = new Properties
    props.load(new FileReader("myprog.properties"))
    // The statement above is a part of the primary constructor
    ...
}
```



NOTE: If there are no parameters after the class name, then the class has a primary constructor with no parameters. That constructor simply executes all statements in the body of the class.



TIP: You can often eliminate auxiliary constructors by using default arguments in the primary constructor. For example:

```
class Person(val name: String = "", val age: Int = 0)
```

Primary constructor parameters can have any of the forms in Table 5–1. For example,

```
class Person(val name: String, private var age: Int)
```

declares and initializes fields

```
val name: String
private var age: Int
```

Construction parameters can also be regular method parameters, without `val` or `var`. How these parameters are processed depends on their usage inside the class.

- If a parameter without `val` or `var` is used inside at least one method, it becomes a field. For example,

```
class Person(name: String, age: Int) {
    def description = s"$name is $age years old"
}
```

declares and initializes immutable fields `name` and `age` that are object-private.

Such a field is the equivalent of a `private[this]` `val` field (see Section 5.4, “Object-Private Fields,” on page 60).

- Otherwise, the parameter is not saved as a field. It’s just a regular parameter that can be accessed in the code of the primary constructor. (Strictly speaking, this is an implementation-specific optimization.)

Table 5–2 summarizes the fields and methods that are generated for different kinds of primary constructor parameters.

Table 5–2 Fields and Methods Generated for Primary Constructor Parameters

Primary Constructor Parameter	Generated Field/Methods
<code>name: String</code>	object-private field, or no field if no method uses <code>name</code>
<code>private val/var name: String</code>	private field, private getter/setter
<code>val/var name: String</code>	private field, public getter/setter
<code>@BeanProperty val/var name: String</code>	private field, public Scala and JavaBeans getters/setters

If you find the primary constructor notation confusing, you don’t need to use it. Just provide one or more auxiliary constructors in the usual way, but remember to call `this()` if you don’t chain to another auxiliary constructor.

However, many programmers like the concise syntax. Martin Odersky suggests to think about it this way: In Scala, classes take parameters, just like methods do.



NOTE: When you think of the primary constructor’s parameters as class parameters, parameters without `val` or `var` become easier to understand. The scope of such a parameter is the entire class. Therefore, you can use the parameter in methods. If you do, it is the compiler’s job to save it in a field.



TIP: The Scala designers think that *every keystroke is precious*, so they let you combine a class with its primary constructor. When reading a Scala class, you need to disentangle the two. For example, when you see

```
class Person(val name: String) {
    var age = 0
    def description = s"$name is $age years old"
}
```

take this definition apart into a class definition:

```
class Person(val name: String) {
    var age = 0
    def description = s"$name is $age years old"
}
```

and a constructor definition:

```
class Person(val name: String) {
    var age = 0
    def description = s"$name is $age years old"
}
```



NOTE: To make the primary constructor private, place the keyword `private` like this:

```
class Person private(val id: Int) { ... }
```

A class user must then use an auxiliary constructor to construct a `Person` object.

5.8 Nested Classes

In Scala, you can nest just about anything inside anything. You can define functions inside other functions, and classes inside other classes. Here is a simple example of the latter:

```
import scala.collection.mutable.ArrayBuffer
class Network {
    class Member(val name: String) {
        val contacts = new ArrayBuffer[Member]
    }
}

private val members = new ArrayBuffer[Member]

def join(name: String) = {
    val m = new Member(name)
    members += m
}
```

```
    m
}
}
```

Consider two networks:

```
val chatter = new Network
val myFace = new Network
```

In Scala, each *instance* has its own class Member, just like each instance has its own field members. That is, `chatter.Member` and `myFace.Member` are *different classes*.



NOTE: This is different from Java, where an inner class belongs to the outer class.

The Scala approach is more regular. For example, to make a new inner object, you simply use `new` with the type name: `new chatter.Member`. In Java, you need to use a special syntax, `chatter.new Member()`.

In our network example, you can add a member within its own network, but not across networks.

```
val fred = chatter.join("Fred")
val wilma = chatter.join("Wilma")
fred.contacts += wilma // OK
val barney = myFace.join("Barney") // Has type myFace.Member
fred.contacts += barney
// No—can't add a myFace.Member to a buffer of chatter.Member elements
```

For networks of people, this behavior probably makes sense. If you don't want it, there are two solutions.

First, you can move the Member type somewhere else. A good place would be the Network companion object. (Companion objects are described in Chapter 6.)

```
object Network {
  class Member(val name: String) {
    val contacts = new ArrayBuffer[Member]
  }
}

class Network {
  private val members = new ArrayBuffer[Network.Member]
  ...
}
```

Alternatively, you can use a *type projection* `Network#Member`, which means “a Member of *any* Network.” For example,

```
class Network {
    class Member(val name: String) {
        val contacts = new ArrayBuffer[Network#Member]
    }
    ...
}
```

You would do that if you want the fine-grained “inner class per object” feature in some places of your program, but not everywhere. See Chapter 19 for more information about type projections.



NOTE: In a nested class, you can access the `this` reference of the enclosing class as `EnclosingClass.this`, like in Java. If you like, you can establish an alias for that reference with the following syntax:

```
class Network(val name: String) { outer =>
    class Member(val name: String) {
        ...
        def description = s"$name inside ${outer.name}"
    }
}
```

The `class Network { outer =>` syntax makes the variable `outer` refer to `Network.this`. You can choose any name for this variable. The name `self` is common, but perhaps confusing when used with nested classes.

This syntax is related to the “self type” syntax that you will see in Chapter 19.

Exercises

1. Improve the `Counter` class in Section 5.1, “Simple Classes and Parameterless Methods,” on page 55 so that it doesn’t turn negative at `Int.MaxValue`.
2. Write a class `BankAccount` with methods `deposit` and `withdraw`, and a read-only property `balance`.
3. Write a class `Time` with read-only properties `hours` and `minutes` and a method `before(other: Time): Boolean` that checks whether this time comes before the other. A `Time` object should be constructed as `new Time(hrs, min)`, where `hrs` is in military time format (between 0 and 23).
4. Reimplement the `Time` class from the preceding exercise so that the internal representation is the number of minutes since midnight (between 0 and $24 \times 60 - 1$). *Do not* change the public interface. That is, client code should be unaffected by your change.

5. Make a class `Student` with read-write JavaBeans properties `name` (of type `String`) and `id` (of type `Long`). What methods are generated? (Use `javap` to check.) Can you call the JavaBeans getters and setters in Scala? Should you?
6. In the `Person` class of Section 5.1, “Simple Classes and Parameterless Methods,” on page 55, provide a primary constructor that turns negative ages to 0.
7. Write a class `Person` with a primary constructor that accepts a string containing a first name, a space, and a last name, such as `new Person("Fred Smith")`. Supply read-only properties `firstName` and `lastName`. Should the primary constructor parameter be a `var`, a `val`, or a plain parameter? Why?
8. Make a class `Car` with read-only properties for manufacturer, model name, and model year, and a read-write property for the license plate. Supply four constructors. All require the manufacturer and model name. Optionally, model year and license plate can also be specified in the constructor. If not, the model year is set to -1 and the license plate to the empty string. Which constructor are you choosing as the primary constructor? Why?
9. Reimplement the class of the preceding exercise in Java, C#, or C++ (your choice). How much shorter is the Scala class?
10. Consider the class

```
class Employee(val name: String, var salary: Double) {  
    def this() { this("John Q. Public", 0.0) }  
}
```

Rewrite it to use explicit fields and a default primary constructor. Which form do you prefer? Why?

Objects

Topics in This Chapter **A1**

- 6.1 Singletons — page 71
- 6.2 Companion Objects — page 72
- 6.3 Objects Extending a Class or Trait — page 73
- 6.4 The apply Method — page 73
- 6.5 Application Objects — page 74
- 6.6 Enumerations — page 75
- Exercises — page 77

Chapter

6

In this short chapter, you will learn when to use the `object` construct in Scala. Use it when you need a class with a single instance, or when you want to find a home for miscellaneous values or functions.

The key points of this chapter are:

- Use objects for singletons and utility methods.
- A class can have a companion object with the same name.
- Objects can extend classes or traits.
- The `apply` method of an object is usually used for constructing new instances of the companion class.
- To avoid the `main` method, use an object that extends the `App` trait.
- You can implement enumerations by extending the `Enumeration` object.

6.1 Singletons

Scala has no static methods or fields. Instead, you use the `object` construct. An object defines a single instance of a class with the features that you want. For example,

```
object Accounts {  
    private var lastNumber = 0  
    def newUniqueNumber() = { lastNumber += 1; lastNumber }  
}
```

When you need a new unique account number in your application, call `Accounts.newUniqueNumber()`.

The constructor of an object is executed when the object is first used. In our example, the `Accounts` constructor is executed with the first call to `Accounts.newUniqueNumber()`. If an object is never used, its constructor is not executed.

An object can have essentially all the features of a class—it can even extend other classes or traits (see Section 6.3, “Objects Extending a Class or Trait,” on page 73). There is just one exception: You cannot provide constructor parameters.

Use an object in Scala whenever you would have used a singleton object in Java or C++:

- As a home for utility functions or constants
- When a single immutable instance can be shared efficiently
- When a single instance is required to coordinate some service (the singleton design pattern)



NOTE: Many people view the singleton design pattern with disdain. Scala gives you the tools for both good and bad design, and it is up to you to use them wisely.

6.2 Companion Objects

In Java or C++, you often have a class with both instance methods and static methods. In Scala, you can achieve this by having a class and a “companion” object of the same name. For example,

```
class Account {
    val id = Account.newUniqueNumber()
    private var balance = 0.0
    def deposit(amount: Double) { balance += amount }
    ...
}

object Account { // The companion object
    private var lastNumber = 0
    private def newUniqueNumber() = { lastNumber += 1; lastNumber }
}
```

The class and its companion object can access each other’s private features. They must be located in the *same source file*.

Note that the companion object's features are not in the scope of the class. For example, the `Account` class has to use `Account.newUniqueNumber()` and not just `newUniqueNumber()` to invoke the method of the companion object.



TIP: In the REPL, you must define the class and the object together in paste mode. Type

```
:paste
```

Then type or paste both the class and object definitions, and type `Ctrl+D`.



NOTE: A companion object contains features that accompany a class. In Chapter 7, you will see how to add features to a package using a *package object*.

6.3 Objects Extending a Class or Trait

An object can extend a class and/or one or more traits. The result is an object of a class that extends the given class and/or traits, and in addition has all of the features specified in the object definition.

One useful application is to specify default objects that can be shared. For example, consider a class for undoable actions in a program.

```
abstract class UndoableAction(val description: String) {
    def undo(): Unit
    def redo(): Unit
}
```

A useful default is the “do nothing” action. Of course, we only need one of them.

```
object DoNothingAction extends UndoableAction("Do nothing") {
    override def undo() {}
    override def redo() {}
}
```

The `DoNothingAction` object can be shared across all places that need this default.

```
val actions = Map("open" -> DoNothingAction, "save" -> DoNothingAction, ...)
// Open and save not yet implemented
```

6.4 The apply Method

It is common to have objects with an `apply` method. The `apply` method is called for expressions of the form

Object(arg1, ..., argN)

Typically, such an apply method returns an object of the companion class.

For example, the Array object defines apply methods that allow array creation with expressions such as

```
Array("Mary", "had", "a", "little", "lamb")
```

Why doesn't one just use a constructor? Not having the new keyword is handy for nested expressions, such as

```
Array(Array(1, 7), Array(2, 9))
```



CAUTION: It is easy to confuse Array(100) and new Array(100). The first expression calls apply(100), yielding an Array[Int] with a single element, the integer 100. The second expression invokes the constructor this(100). The result is an Array[Nothing] with 100 null elements.

Here is an example of defining an apply method:

```
class Account private (val id: Int, initialBalance: Double) {
    private var balance = initialBalance
    ...
}

object Account { // The companion object
    def apply(initialBalance: Double) =
        new Account(newUniqueNumber(), initialBalance)
    ...
}
```

Now you can construct an account as

```
val acct = Account(1000.0)
```

6.5 Application Objects

Each Scala program must start with an object's main method of type `Array[String]` => `Unit`:

```
object Hello {
    def main(args: Array[String]) {
        println("Hello, World!")
    }
}
```

Instead of providing a `main` method for your application, you can extend the `App` trait and place the program code into the constructor body:

```
object Hello extends App {
    println("Hello, World!")
}
```

If you need the command-line arguments, you can get them from the `args` property:

```
object Hello extends App {
    if (args.length > 0)
        println(f"Hello ${args(0)}")
    else
        println("Hello, World!")
}
```

If you invoke the application with the `scala.time` option set, then the elapsed time is displayed when the program exits.

```
$ scalac Hello.scala
$ scala -Dscala.time Hello Fred
Hello, Fred
[total 4ms]
```

All this involves a bit of magic. The `App` trait extends another trait, `DelayedInit`, that gets special handling from the compiler. All initialization code of a class with that trait is moved into a `delayedInit` method. The `main` of the `App` trait method captures the command-line arguments, calls the `delayedInit` method, and optionally prints the elapsed time.



NOTE: Older versions of Scala had an `Application` trait for the same purpose. That trait carried out the program's action in the static initializer, which is not optimized by the just-in-time compiler. Use the `App` trait instead.

6.6 Enumerations

Unlike Java or C++, Scala does not have enumerated types. However, the standard library provides an `Enumeration` helper class that you can use to produce enumerations.

Define an object that extends the `Enumeration` class and initialize each value in your enumeration with a call to the `Value` method. For example,

```
object TrafficLightColor extends Enumeration {
    val Red, Yellow, Green = Value
}
```

Here we define three fields, `Red`, `Yellow`, and `Green`, and initialize each of them with a call to `Value`. This is a shortcut for

```
val Red = Value
val Yellow = Value
val Green = Value
```

Each call to the `Value` method returns a new instance of an inner class, also called `Value`.

Alternatively, you can pass IDs, names, or both to the `Value` method:

```
val Red = Value(0, "Stop")
val Yellow = Value(10) // Name "Yellow"
val Green = Value("Go") // ID 11
```

If not specified, the ID is one more than the previously assigned one, starting with zero. The default name is the field name.

You can now refer to the enumeration values as `TrafficLightColor.Red`, `TrafficLightColor.Yellow`, and so on. If that gets too tedious, use

```
import TrafficLightColor._
```

(See Chapter 7 for information on importing members of a class or object.)

Remember that the type of the enumeration is `TrafficLightColor.Value` and *not* `TrafficLightColor`—that's the type of the object holding the values. Some people recommend that you add a type alias

```
object TrafficLightColor extends Enumeration {
    type TrafficLightColor = Value
    val Red, Yellow, Green = Value
}
```

Now the type of the enumeration is `TrafficLightColor.TrafficLightColor`, which is only an improvement if you use an `import` statement. For example,

```
import TrafficLightColor._
def doWhat(color: TrafficLightColor) = {
    if (color == Red) "stop"
    else if (color == Yellow) "hurry up"
    else "go"
}
```

The ID of an enumeration value is returned by the `id` method, and its name by the `toString` method.

The call `TrafficLightColor.values` yields a set of all values:

```
for (c <- TrafficLightColor.values) println(s"${c.id}: $c")
```

Finally, you can look up an enumeration value by its ID or name. Both of the following yield the object `TrafficLightColor.Red`:

```
TrafficLightColor(0) // Calls Enumeration.apply  
TrafficLightColor.withName("Red")
```

Exercises

1. Write an object `Conversions` with methods `inchesToCentimeters`, `gallonsToLiters`, and `milesToKilometers`.
2. The preceding problem wasn't very object-oriented. Provide a general super-class `UnitConversion` and define objects `InchesToCentimeters`, `GallonsToLiters`, and `MilesToKilometers` that extend it.
3. Define an `Origin` object that extends `java.awt.Point`. Why is this not actually a good idea? (Have a close look at the methods of the `Point` class.)
4. Define a `Point` class with a companion object so that you can construct `Point` instances as `Point(3, 4)`, without using `new`.
5. Write a Scala application, using the `App` trait, that prints its command-line arguments in reverse order, separated by spaces. For example, `scala Reverse Hello World` should print `World Hello`.
6. Write an enumeration describing the four playing card suits so that the `toString` method returns ♣, ♦, ♥, or ♠.
7. Implement a function that checks whether a card suit value from the preceding exercise is red.
8. Write an enumeration describing the eight corners of the RGB color cube. As IDs, use the color values (for example, `0xff0000` for Red).

Packages and Imports

Topics in This Chapter **A1**

- 7.1 Packages — page 80
- 7.2 Scope Rules — page 81
- 7.3 Chained Package Clauses — page 83
- 7.4 Top-of-File Notation — page 83
- 7.5 Package Objects — page 83
- 7.6 Package Visibility — page 84
- 7.7 Imports — page 85
- 7.8 Imports Can Be Anywhere — page 85
- 7.9 Renaming and Hiding Members — page 86
- 7.10 Implicit Imports — page 86
- Exercises — page 87

Chapter

7

In this chapter, you will learn how packages and import statements work in Scala. Both packages and imports are more regular than in Java; they are also a bit more flexible.

The key points of this chapter are:

- Packages nest just like inner classes.
- Package paths are *not* absolute.
- A chain `x.y.z` in a package clause leaves the intermediate packages `x` and `x.y` invisible.
- Package statements without braces at the top of the file extend to the entire file.
- A package object can hold functions and variables.
- Import statements can import packages, classes, and objects.
- Import statements can be anywhere.
- Import statements can rename and hide members.
- `java.lang`, `scala`, and `Predef` are always imported.

7.1 Packages

Packages in Scala fulfill the same purpose as packages in Java or namespaces in C++: to manage names in a large program. For example, the name `Map` can occur in the packages `scala.collection.immutable` and `scala.collection.mutable` without conflict. To access either name, you can use the fully qualified `scala.collection.immutable.Map` or `scala.collection.mutable.Map`. Alternatively, use an `import` statement to provide a shorter alias—see Section 7.7, “Imports,” on page 85.

To add items to a package, you can include them in package statements, such as:

```
package com {  
    package horstmann {  
        package impatient {  
            class Employee  
  
            ...  
        }  
    }  
}
```

Then the class name `Employee` can be accessed anywhere as `com.horstmann.impatient.Employee`.

Unlike an object or a class, a package can be defined in multiple files. The preceding code might be in a file `Employee.scala`, and a file `Manager.scala` might contain

```
package com {  
    package horstmann {  
        package impatient {  
            class Manager  
  
            ...  
        }  
    }  
}
```



NOTE: There is no enforced relationship between the directory of the source file and the package. You don't have to put `Employee.scala` and `Manager.scala` into a `com/horstmann/impatient` directory.

Conversely, you can contribute to more than one package in a single file. The file `Employee.scala` may contain

```
package com {  
    package horstmann {  
        package impatient {  
            class Employee  
            ...  
        }  
    }  
}
```

```
package net {  
    package bigjava {  
        class Counter  
        ...  
    }  
}
```

7.2 Scope Rules

In Scala, the scope rules for packages are more consistent than in Java. Scala packages nest just like all other scopes. You can access names from the enclosing scope. For example,

```
package com {  
    package horstmann {  
        object Utils {  
            def percentOf(value: Double, rate: Double) = value * rate / 100  
            ...  
        }  
  
        package impatient {  
            class Employee {  
                ...  
                def giveRaise(rate: scala.Double) {  
                    salary += Utils.percentOf(salary, rate)  
                }  
            }  
        }  
    }  
}
```

Note the `Utils.percentOf` qualifier. The `Utils` class was defined in the *parent* package. Everything in the parent package is in scope, and it is not necessary to use `com.horstmann.Utils.percentOf`. (You could, though, if you prefer—after all, `com` is also in scope.)

There is a fly in the ointment, however. Consider

```
package com {
    package horstmann {
        package impatient {
            class Manager {
                val subordinates = new collection.mutable.ArrayBuffer[Employee]
                ...
            }
        }
    }
}
```

This code takes advantage of the fact that the `scala` package is always imported. Therefore, the `collection` package is actually `scala.collection`.

And now suppose someone introduces the following package, perhaps in a different file:

```
package com {
    package horstmann {
        package collection {
            ...
        }
    }
}
```

Now the `Manager` class no longer compiles. It looks for a `mutable` member inside the `com.horstmann.collection` package and doesn't find it. The intent in the `Manager` class was the `collection` package in the top-level `scala` package, not whatever `collection` subpackage happened to be in some accessible scope.

In Java, this problem can't occur because package names are always *absolute*, starting at the root of the package hierarchy. But in Scala, package names are relative, just like inner class names. With inner classes, one doesn't usually run into problems because all the code is in one file, under control of whoever is in charge of that file. But packages are open-ended. Anyone can contribute to a package at any time.

One solution is to use absolute package names, starting with `_root_`, for example:

```
val subordinates = new _root_.scala.collection.mutable.ArrayBuffer[Employee]
```

Another approach is to use “chained” package clauses, as detailed in the next section.



NOTE: Most programmers use complete paths for package names, without the `_root_` prefix. This is safe as long as everyone avoids names `scala`, `java`, `com`, `net`, and so on, for nested packages.

7.3 Chained Package Clauses

A package clause can contain a “chain,” or path segment, for example:

```
package com.horstmann.impatient {  
    // Members of com and com.horstmann are not visible here  
    package people {  
        class Person  
  
        ...  
    }  
}
```

Such a clause limits the visible members. Now a `com.horstmann.collection` package would no longer be accessible as `collection`.

7.4 Top-of-File Notation

Instead of the nested notation that we have used up to now, you can have package clauses at the top of the file, without braces. For example:

```
package com.horstmann.impatient  
package people  
  
class Person  
  
...
```

This is equivalent to

```
package com.horstmann.impatient {  
    package people {  
        class Person  
  
        ...  
        // Until the end of the file  
    }  
}
```

This is the preferred notation if all the code in the file belongs to the same package (which is the usual case).

Note that in the example above, everything in the file belongs to the package `com.horstmann.impatient.people`, but the package `com.horstmann.impatient` has also been opened up so you can refer to its contents.

7.5 Package Objects

A package can contain classes, objects, and traits, but not the definitions of functions or variables. That’s an unfortunate limitation of the Java virtual machine.

It would make more sense to add utility functions or constants to a package than to some `Utils` object. Package objects address this limitation.

Every package can have one package object. You define it in the *parent* package, and it has the same name as the child package. For example,

```
package com.horstmann.impatient

package object people {
    val defaultName = "John Q. Public"
}

package people {
    class Person {
        var name = defaultName // A constant from the package
    }
    ...
}
```

Note that the `defaultName` value didn't need to be qualified because it was in the same package. Elsewhere, it is accessible as `com.horstmann.impatient.people.defaultName`.

Behind the scenes, the package object gets compiled into a JVM class with static methods and fields, called `package.class`, inside the package. In our example, that would be a class `com.horstmann.impatient.people.package` with a static field `defaultName`. (In the JVM, you can use `package` as a class name.)

It is a good idea to use the same naming scheme for source files. Put the package object into a file `com/horstmann/impatient/people/package.scala`. That way, anyone who wants to add functions or variables to a package can find the package object easily.

7.6 Package Visibility

In Java, a class member that isn't declared as `public`, `private`, or `protected` is visible in the package containing the class. In Scala, you can achieve the same effect with qualifiers. The following method is visible in its own package:

```
package com.horstmann.impatient.people

class Person {
    private[people] def description = s"A person with name $name"
    ...
}
```

You can extend the visibility to an enclosing package:

```
private[impatient] def description = s"A person with name $name"
```

7.7 Imports

Imports let you use short names instead of long ones. With the clause

```
import java.awt.Color
```

you can write `Color` in your code instead of `java.awt.Color`.

That is the sole purpose of imports. If you don't mind long names, you'll never need them.

You can import all members of a package as

```
import java.awt._
```

This is the same as the `*` wildcard in Java. (In Scala, `*` is a valid character for an identifier. You could define a package `com.horstmann.*.people`, but please don't.)

You can also import all members of a class or object.

```
import java.awt.Color._  
val c1 = RED // Color.RED  
val c2 = decode("#ff0000") // Color.decode
```

This is like `import static` in Java. Java programmers seem to live in fear of this variant, but in Scala it is commonly used.

Once you import a package, you can access its subpackages with shorter names. For example:

```
import java.awt._  
  
def handler(evt: event.ActionEvent) { // java.awt.event.ActionEvent  
    ...  
}
```

The `event` package is a member of `java.awt`, and the import brings it into scope.

7.8 Imports Can Be Anywhere

In Scala, an `import` statement can be anywhere, not just at the top of a file. The scope of the `import` statement extends until the end of the enclosing block. For example,

```
class Manager {  
    import scala.collection.mutable._  
    val subordinates = new ArrayBuffer[Employee]  
    ...  
}
```

This is a very useful feature, particularly with wildcard imports. It is always a bit worrisome to import lots of names from different sources. In fact, some Java programmers dislike wildcard imports so much that they never use them but let their IDE generate long lists of imported classes.

By putting the imports where they are needed, you can greatly reduce the potential for conflicts.

7.9 Renaming and Hiding Members

If you want to import more than one member from a package, use a *selector* like this:

```
import java.awt.{Color, Font}
```

The selector syntax lets you rename members:

```
import java.util.{HashMap => JavaHashMap}
import scala.collection.mutable._
```

Now `JavaHashMap` is a `java.util.HashMap` and plain `HashMap` is a `scala.collection.mutable.HashMap`.

The selector `HashMap => _` hides a member instead of renaming it. This is only useful if you import others:

```
import java.util.{HashMap => _, _}
import scala.collection.mutable._
```

Now `HashMap` unambiguously refers to `scala.collection.mutable.HashMap` since `java.util.HashMap` is hidden.

7.10 Implicit Imports

Every Scala program implicitly starts with

```
import java.lang._
import scala._
import Predef._
```

As with Java programs, `java.lang` is always imported. Next, the `scala` package is imported, but in a special way. Unlike all other imports, this one is allowed to override the preceding import. For example, `scala.StringBuilder` overrides `java.lang.StringBuilder` instead of conflicting with it.

Finally, the `Predef` object is imported. It contains commonly used types, implicit conversions, and utility methods. (The methods could equally well have been

placed into the `scala` package object, but `Predef` was introduced before Scala had package objects.)

Since the `scala` package is imported by default, you never need to write package names that start with `scala`. For example,

```
collection.mutable.HashMap
```

is just as good as

```
scala.collection.mutable.HashMap
```

Exercises

1. Write an example program to demonstrate that

```
package com.horstmann.impatient
```

is not the same as

```
package com
package horstmann
package impatient
```

2. Write a puzzler that baffles your Scala friends, using a package `com` that isn't at the top level.
3. Write a package `random` with functions `nextInt(): Int`, `nextDouble(): Double`, and `setSeed(seed: Int): Unit`. To generate random numbers, use the linear congruential generator

$$\text{next} = (\text{previous} \times a + b) \bmod 2^n,$$

where $a = 1664525$, $b = 1013904223$, $n = 32$, and the initial value of `previous` is `seed`.

4. Why do you think the Scala language designers provided the package object syntax instead of simply letting you add functions and variables to a package?
5. What is the meaning of `private[com] def giveRaise(rate: Double)`? Is it useful?
6. Write a program that copies all elements from a Java hash map into a Scala hash map. Use imports to rename both classes.
7. In the preceding exercise, move all imports into the innermost scope possible.
8. What is the effect of

```
import java._
import javax._
```

Is this a good idea?

9. Write a program that imports the `java.lang.System` class, reads the user name from the `user.name` system property, reads a password from the `StdIn` object, and prints a message to the standard error stream if the password is not "secret". Otherwise, print a greeting to the standard output stream. Do not use any other imports, and do not use any qualified names (with dots).
10. Apart from `StringBuilder`, what other members of `java.lang` does the `scala` package override?

This page intentionally left blank

Inheritance

Topics in This Chapter **A1**

- 8.1 Extending a Class — page 91
- 8.2 Overriding Methods — page 92
- 8.3 Type Checks and Casts — page 93
- 8.4 Protected Fields and Methods — page 94
- 8.5 Superclass Construction — page 94
- 8.6 Overriding Fields — page 95
- 8.7 Anonymous Subclasses — page 97
- 8.8 Abstract Classes — page 97
- 8.9 Abstract Fields — page 97
- 8.10 Construction Order and Early Definitions **L3** — page 98
- 8.11 The Scala Inheritance Hierarchy — page 100
- 8.12 Object Equality **L1** — page 102
- 8.13 Value Classes **L2** — page 103
- Exercises — page 105

Chapter 8

In this chapter, you will learn the most important ways in which inheritance in Scala differs from inheritance in Java and C++. The highlights are:

- The `extends` and `final` keywords are as in Java.
- You must use `override` when you override a method.
- Only the primary constructor can call the primary superclass constructor.
- You can override fields.

In this chapter, we only discuss the case in which a class inherits from another class. See Chapter 10 for inheriting *traits*—the Scala concept that generalizes Java interfaces.

8.1 Extending a Class

You extend a class in Scala just like you would in Java—with the `extends` keyword:

```
class Employee extends Person {  
    var salary = 0.0  
    ...  
}
```

As in Java, you specify fields and methods that are new to the subclass or that override methods in the superclass.

As in Java, you can declare a class final so that it cannot be extended. You can also declare individual methods or fields final so that they cannot be overridden. (See Section 8.6, “Overriding Fields,” on page 95 for overriding fields.) Note that this is different from Java, where a final field is immutable, similar to val in Scala.

8.2 Overriding Methods

In Scala, you *must* use the override modifier when you override a method that isn’t abstract. (See Section 8.8, “Abstract Classes,” on page 97 for abstract methods.) For example,

```
class Person {  
    ...  
    override def toString = s"${getClass.getName}[name=$name]"  
}
```

The override modifier can give useful error messages in a number of common situations, such as:

- When you misspell the name of the method that you are overriding
- When you accidentally provide a wrong parameter type in the overriding method
- When you introduce a new method in a superclass that clashes with a subclass method



NOTE: The last case is an instance of the *fragile base class problem* where a change in the superclass cannot be verified without looking at all the subclasses. Suppose programmer Alice defines a Person class, and, unbeknownst to Alice, programmer Bob defines a subclass Student with a method id yielding the student ID. Later, Alice also defines a method id that holds the person’s national ID. When Bob picks up that change, something may break in Bob’s program (but not in Alice’s test cases) since Student objects now return unexpected IDs.

In Java, one is often advised to “solve” this problem by declaring all methods as final unless they are explicitly designed to be overridden. That sounds good in theory, but programmers hate it when they can’t make even the most innocuous changes to a method (such as adding a logging call). That’s why Java eventually introduced an optional @Overrides annotation.

Invoking a superclass method in Scala works exactly like in Java, with the keyword super:

```
class Employee extends Person {
    ...
    override def toString = s"${super.toString}[salary=$salary]"
}
```

The call `super.toString` invokes the `toString` method of the superclass—that is, the `Person.toString` method.

8.3 Type Checks and Casts

To test whether an object belongs to a given class, use the `isInstanceOf` method. If the test succeeds, you can use the `asInstanceOf` method to convert a reference to a subclass reference:

```
if (p.isInstanceOf[Employee]) {
    val s = p.asInstanceOf[Employee] // s has type Employee
    ...
}
```

The `p.isInstanceOf[Employee]` test succeeds if `p` refers to an object of class `Employee` or its subclass (such as `Manager`).

If `p` is `null`, then `p.isInstanceOf[Employee]` returns `false` and `p.asInstanceOf[Employee]` returns `null`.

If `p` is not an `Employee`, then `p.asInstanceOf[Employee]` throws an exception.

If you want to test whether `p` refers to an `Employee` object, but not a subclass, use

```
if (p.getClass == getClass[Employee])
```

The `getClass` method is defined in the `scala.Predef` object that is always imported.

Table 8–1 shows the correspondence between Scala and Java type checks and casts.

Table 8–1 Type Checks and Casts in Scala and Java

Scala	Java
<code>obj.isInstanceOf[C1]</code>	<code>obj instanceof C1</code>
<code>obj.asInstanceOf[C1]</code>	<code>(C1) obj</code>
<code>getClass[C1]</code>	<code>C1.class</code>

However, pattern matching is usually a better alternative to using type checks and casts. For example,

```
p match {
    case s: Employee => ... // Process s as an Employee
    case _ => ... // p wasn't an Employee
}
```

See Chapter 14 for more information.

8.4 Protected Fields and Methods

As in Java or C++, you can declare a field or method as protected. Such a member is accessible from any subclass, but not from other locations.

Unlike in Java, a protected member is *not* visible throughout the package to which the class belongs. (If you want this visibility, you can use a package modifier—see Chapter 7.)

There is also a `protected[this]` variant that restricts access to the current object, similar to the `private[this]` variant discussed in Chapter 5.

8.5 Superclass Construction

Recall from Chapter 5 that a class has one primary constructor and any number of auxiliary constructors, and that all auxiliary constructors must start with a call to a preceding auxiliary constructor or the primary constructor.

As a consequence, an auxiliary constructor can *never* invoke a superclass constructor directly.

The auxiliary constructors of the subclass eventually call the primary constructor of the subclass. Only the primary constructor can call a superclass constructor.

Recall that the primary constructor is intertwined with the class definition. The call to the superclass constructor is similarly intertwined. Here is an example:

```
class Employee(name: String, age: Int, val salary : Double) extends
  Person(name, age)
```

This defines a subclass

```
class Employee(name: String, age: Int, val salary : Double) extends
  Person(name, age)
```

and a primary constructor that calls the superclass constructor

```
class Employee(name: String, age: Int, val salary : Double) extends
  Person(name, age)
```

Intertwining the class and the constructor makes for very concise code. You may find it helpful to think of the primary constructor parameters as parameters of

the class. Here, the `Employee` class has three parameters: `name`, `age`, and `salary`, two of which it “passes” to the superclass.

In Java, the equivalent code is quite a bit more verbose:

```
public class Employee extends Person { // Java
    private double salary;
    public Employee(String name, int age, double salary) {
        super(name, age);
        this.salary = salary;
    }
}
```



NOTE: In a Scala constructor, you can never call `super(params)`, as you would in Java, to call the superclass constructor.

A Scala class can extend a Java class. Its primary constructor must invoke one of the constructors of the Java superclass. For example,

```
class PathWriter(p: Path, cs: Charset) extends
    java.io.PrintWriter(Files.newBufferedWriter(p, cs))
```

8.6 Overriding Fields

Recall from Chapter 5 that a field in Scala consists of a private field *and* accessor/mutator methods. You can override a `val` (or a parameterless `def`) with another `val` field of the same name. The subclass has a private field and a public getter, and the getter overrides the superclass getter (or method).

For example,

```
class Person(val name: String) {
    override def toString = s"${getClass.getName}[name=$name]"
}

class SecretAgent(codename: String) extends Person(codename) {
    override val name = "secret" // Don't want to reveal name ...
    override val toString = "secret" // ... or class name
}
```

This example shows the mechanics, but it is rather artificial. A more common case is to override an abstract `def` with a `val`, like this:

```
abstract class Person { // See Section 8.8 for abstract classes
    def id: Int // Each person has an ID that is computed in some way
    ...
}

class Student(override val id: Int) extends Person
    // A student ID is simply provided in the constructor
```

Note the following restrictions (see also Table 8–2):

- A def can only override another def.
- A val can only override another val or a parameterless def.
- A var can only override an abstract var (see Section 8.8, “Abstract Classes,” on page 97).

Table 8–2 Overriding val, def, and var

	with val	with def	with var
Override val	<ul style="list-style-type: none"> • Subclass has a private field (with the same name as the superclass field—that’s OK). • Getter overrides the superclass getter. 	Error.	Error.
Override def	<ul style="list-style-type: none"> • Subclass has a private field. • Getter overrides the superclass method. 	Like in Java.	A var can override a getter/setter pair. Overriding just a getter is an error.
Override var	Error.	Error.	Only if the superclass var is abstract (see Section 8.8).



NOTE: In Chapter 5, I said that it’s OK to use a var because you can always change your mind and reimplement it as a getter/setter pair. However, the programmers extending your class do not have that choice. They cannot override a var with a getter/setter pair. In other words, if you provide a var, all subclasses are stuck with it.

8.7 Anonymous Subclasses

As in Java, you make an instance of an *anonymous* subclass if you include a block with definitions or overrides, such as

```
val alien = new Person("Fred") {
    def greeting = "Greetings, Earthling! My name is Fred."
}
```

Technically, this creates an object of a *structural type*—see Chapter 19 for details. The type is denoted as `Person{def greeting: String}`. You can use this type as a parameter type:

```
def meet(p: Person{def greeting: String}) {
    println(s"${p.name} says: ${p.greeting}")
}
```

8.8 Abstract Classes

As in Java, you can use the `abstract` keyword to denote a class that cannot be instantiated, usually because one or more of its methods are not defined. For example,

```
abstract class Person(val name: String) {
    def id: Int // No method body—this is an abstract method
}
```

Here we say that every person has an ID, but we don't know how to compute it. Each concrete subclass of `Person` needs to specify an `id` method. In Scala, unlike Java, you do not use the `abstract` keyword for an abstract method. You simply omit its body. As in Java, a class with at least one abstract method must be declared abstract.

In a subclass, you need not use the `override` keyword when you define a method that was abstract in the superclass.

```
class Employee(name: String) extends Person(name) {
    def id = name.hashCode // override keyword not required
}
```

8.9 Abstract Fields

In addition to abstract methods, a class can also have abstract fields. An abstract field is simply a field without an initial value. For example,

```
abstract class Person {
    val id: Int
        // No initializer—this is an abstract field with an abstract getter method
    var name: String
        // Another abstract field, with abstract getter and setter methods
}
```

This class defines abstract getter methods for the `id` and `name` fields, and an abstract setter for the `name` field. The generated Java class has *no fields*.

Concrete subclasses must provide concrete fields, for example:

```
class Employee(val id: Int) extends Person { // Subclass has concrete id property
    var name = "" // and concrete name property
}
```

As with methods, no `override` keyword is required in the subclass when you define a field that was abstract in the superclass.

You can always customize an abstract field by using an anonymous type:

```
val fred = new Person {
    val id = 1729
    var name = "Fred"
}
```

8.10 Construction Order and Early Definitions L3

When you override a `val` in a subclass *and* use the value in a superclass constructor, the resulting behavior is unintuitive.

Here is an example. A creature can sense a part of its environment. For simplicity, we assume the creature lives in a one-dimensional world, and the sensory data are represented as integers. A default creature can see ten units ahead.

```
class Creature {
    val range: Int = 10
    val env: Array[Int] = new Array[Int](range)
}
```

Ants, however, are near-sighted:

```
class Ant extends Creature {
    override val range = 2
}
```

Unfortunately, we now have a problem. The `range` value is used in the superclass constructor, and the superclass constructor runs *before* the subclass constructor. Specifically, here is what happens:

1. The `Ant` constructor calls the `Creature` constructor before doing its own construction.
2. The `Creature` constructor sets *its* `range` field to `10`.
3. The `Creature` constructor, in order to initialize the `env` array, calls the `range()` getter.
4. That method is overridden to yield the (as yet uninitialized) `range` field of the `Ant` class.
5. The `range` method returns `0`. (That is the initial value of all integer fields when an object is allocated.)
6. `env` is set to an array of length `0`.
7. The `Ant` constructor continues, setting its `range` field to `2`.

Even though it appears as if `range` is either `10` or `2`, `env` has been set to an array of length `0`. The moral is that you should not rely on the value of a `val` in the body of a constructor.

In Java, you have a similar issue when you call a method in a superclass constructor. The method might be overridden in a subclass, and it might not do what you want it to do. (In fact, that is the root cause of our problem—the expression `range` calls the getter method.)

There are several remedies.

- Declare the `val` as `final`. This is safe but not very flexible.
- Declare the `val` as `lazy` in the superclass (see Chapter 2). This is safe but a bit inefficient.
- Use the *early definition syntax* in the subclass—see below.

The “early definition” syntax lets you initialize `val` fields of a subclass *before* the superclass is executed. The syntax is so ugly that only a mother could love it. You place the `val` fields in a block after the `extends` keyword, like this:

```
class Ant extends { override val range = 2 } with Creature
```

Note the `with` keyword before the superclass name. This keyword is normally used with traits—see Chapter 10.

The right-hand side of an early definition can only refer to previous early definitions, not to other fields or methods of the class.



TIP: You can debug construction order problems with the `-Xcheckinit` compiler flag. This flag generates code that throws an exception (instead of yielding the default value) when an uninitialized field is accessed.



NOTE: At the root of the construction order problem lies a design decision of the Java language—namely, to allow the invocation of subclass methods in a superclass constructor. In C++, an object’s virtual function table pointer is set to the table of the superclass when the superclass constructor executes. Afterwards, the pointer is set to the subclass table. Therefore, in C++, it is not possible to modify constructor behavior through overriding. The Java designers felt that this subtlety was unnecessary, and the Java virtual machine does not adjust the virtual function table during construction.

8.11 The Scala Inheritance Hierarchy

Figure 8–1 shows the inheritance hierarchy of Scala classes. The classes that correspond to the primitive types in Java, as well as the type `Unit`, extend `AnyVal`. You can also define your own *value classes*—see Section 8.13, “Value Classes,” on page 103.

All other classes are subclasses of the `AnyRef` class. When compiling to the Java virtual machine, this is a synonym for the `java.lang.Object` class.

Both `AnyVal` and `AnyRef` extend the `Any` class, the root of the hierarchy.

The `Any` class defines methods `isInstanceOf`, `asInstanceOf`, and the methods for equality and hash codes that we will look at in Section 8.12, “Object Equality,” on page 102.

`AnyVal` does not add any methods. It is just a marker for value types.

The `AnyRef` class adds the monitor methods `wait` and `notify/notifyAll` from the `Object` class. It also provides a synchronized method with a function parameter. That method is the equivalent of a synchronized block in Java. For example,

```
account.synchronized { account.balance += amount }
```



NOTE: Just like in Java, I suggest you stay away from `wait`, `notify`, and `synchronized` unless you have a good reason to use them instead of higher-level concurrency constructs.

All Scala classes implement the marker interface `ScalaObject`, which has no methods. At the other end of the hierarchy are the `Nothing` and `Null` types.

`Null` is the type whose sole instance is the value `null`. You can assign `null` to any reference, but not to one of the value types. For example, setting an `Int` to `null` is not possible. This is better than in Java, where it would be possible to set an `Integer` wrapper to `null`.

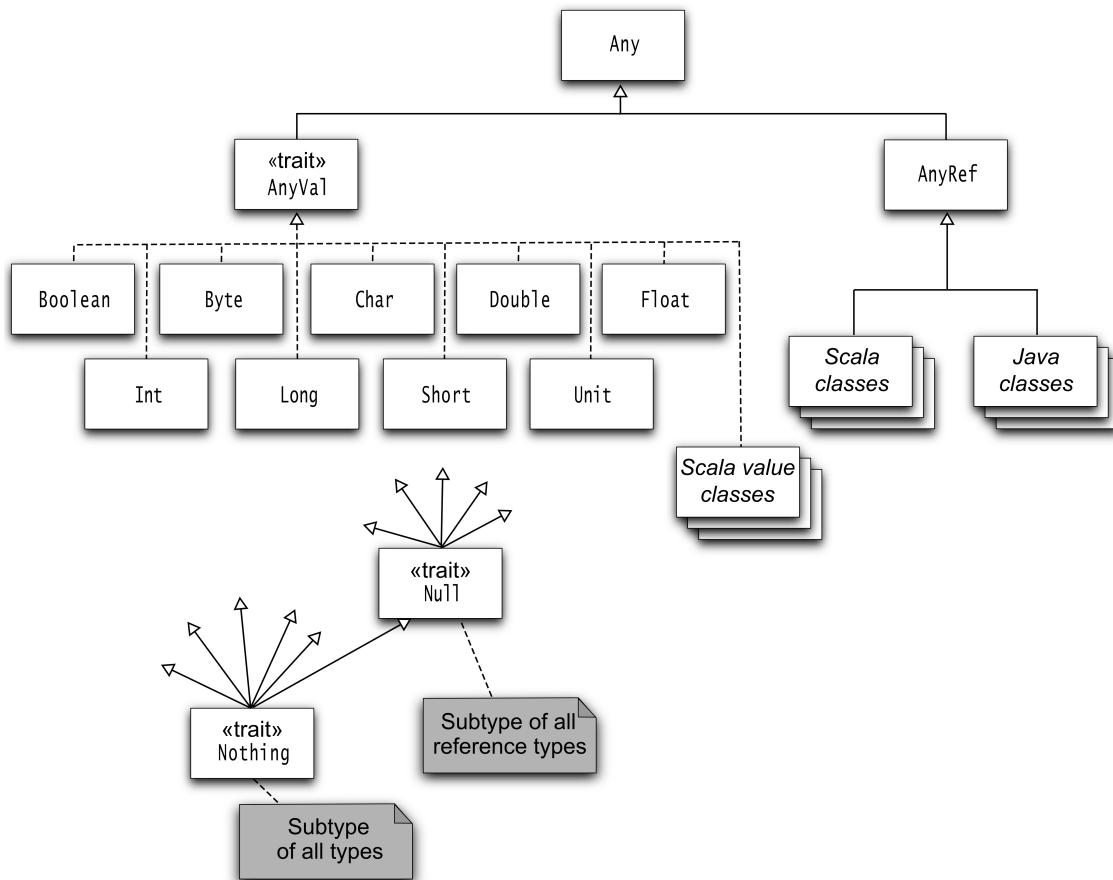


Figure 8–1 The inheritance hierarchy of Scala classes

The `Nothing` type has no instances. It is occasionally useful for generic constructs. For example, the empty list `Nil` has type `List[Nothing]`, which is a subtype of `List[T]` for any `T`.

The `???` method is declared with return type `Nothing`. It never returns but instead throws a `NotImplementedError` when invoked. You can use it for methods that you still need to implement:

```

class Person(val name: String) {
  def description = ???
}
  
```

The `Person` class compiles since `Nothing` is a subtype of every type. You can start using the class, so long as you don't call the `description` method.



CAUTION: The Nothing type is not at all the same as void in Java or C++. In Scala, void is represented by the Unit type, the type with the sole value (). Note that Unit is not a supertype of any other type. However, the compiler still allows any value to be *replaced* by a (). Consider

```
def printAny(x: Any) { println(x) }
def printUnit(x: Unit) { println(x) }
printAny("Hello") // Prints Hello
printUnit("Hello")
// Replaces "Hello" with () and calls printUnit(), which prints ()
```



CAUTION: When a method has a parameter of type Any or AnyRef, and it is called with multiple arguments, then they are placed in a tuple:

```
def show(o: Any) { println(s"${o.getClass}: $o") }
show(3) // Prints class java.lang.Integer: 3
show(3, 4, 5) // Prints class scala.Tuple3: (3,4,5)
```

If you call show() with no parameters, a Unit value is passed. However, that behavior is deprecated.

8.12 Object Equality L1

In Scala, the eq method of the AnyRef class checks whether two references refer to the same object. The equals method in AnyRef calls eq. When you implement a class, you should consider overriding the equals method to provide a natural notion of equality for your situation.

For example, if you define a class Item(val description: String, val price: Double), you might want to consider two items equal if they have the same description and price. Here is an appropriate equals method:

```
final override def equals(other: Any) = {
  other.isInstanceOf[Item] && {
    val that = other.asInstanceOf[Item]
    description == that.description && price == that.price
  }
}
```

Or better, use pattern matching:

```
final override def equals(other: Any) = other match {
  case that: Item => description == that.description && price == that.price
  case _ => false
}
```



NOTE: We defined the method as final because it is generally very difficult to correctly extend equality in a subclass. The problem is symmetry. You want `a.equals(b)` to have the same result as `b.equals(a)`, even when `b` belongs to a subclass.



CAUTION: Be sure to define the `equals` method with parameter type `Any`. The following would be wrong:

```
final def equals(other: Item) = { ... } // Don't!
```

This is a different method which does not override the `equals` method of `AnyRef`.

Also, don't supply an `==` method. You can't override the `==` method defined in `AnyRef`, but you could accidentally supply a different one with an `Item` argument:

```
final def ==(other: Item) = { ... } // Don't!
```

When you define `equals`, remember to define `hashCode` as well. The hash code should be computed only from the fields that you use in the equality check, so that equal objects have the same hash code. In the `Item` example, combine the hash codes of the fields.

```
final override def hashCode = (description, price).##
```

The `##` method is a null-safe version of the `hashCode` method that yields `0` for `null` instead of throwing an exception.



TIP: You are not compelled to override `equals` and `hashCode`. For many classes, it is appropriate to consider distinct objects unequal. For example, if you have two distinct input streams or radio buttons, you will never consider them equal.

In an application program, you don't generally call `eq` or `equals`. Simply use the `==` operator. For reference types, it calls `equals` after doing the appropriate check for `null` operands.

8.13 Value Classes L2

Some classes have a single field, such as the wrapper classes for primitive types, and the “rich” or “ops” wrappers that Scala uses to add methods to existing types. It is inefficient to allocate a new object that holds just one value. *Value classes* allow you to define classes that are “inlined,” so that the single field is used directly.

A value class has these properties:

1. The class extends `AnyVal`.
2. Its primary constructor has exactly one parameter, which is a `val`, and no body.
3. The class has no other fields or constructors.
4. The automatically provided `equals` and `hashCode` methods compare and hash the underlying value.

As an example, let us define a value class that wraps a “military time” value:

```
class MilTime(val time: Int) extends AnyVal {
    def minutes = time % 100
    def hours = time / 100
    override def toString = f"$time04d"
}
```

When you construct a new `MilTime(1230)`, the compiler doesn’t allocate a new object. Instead, it uses the underlying value, the integer 1230. You can invoke the `minutes` and `hours` methods on the value but, just as importantly, you cannot invoke `Int` methods.

```
MilTime lunch = new MilTime(1230)
println(lunch.hours) // OK
println(lunch * 2) // Error
```

To guarantee proper initialization, make the primary constructor private and provide a factory method in the companion object:

```
class MilTime private(val time: Int) extends AnyVal ...
object MilTime {
    def apply(t: Int) =
        if (0 <= t && t < 2400 && t % 100 < 60) new MilTime(t)
        else throw new IllegalArgumentException
}
```



CAUTION: In some programming languages, value types are any types that are allocated on the runtime stack, including structured types with multiple fields. In Scala, a value class can only have one field.



NOTE: If you want a value class to implement a trait (see Chapter 10), the trait must explicitly extend `Any`, and it may not have fields. Such traits are called *universal traits*.



TIP: Value types were designed to make implicit conversions efficient, but you can use them for your own overhead-free “tiny types.” For example, instead of a class `Book(val author: String, val title: String)`, you can wrap each string into a separate value class `Author` and `Title`. When the class is defined as `class Book(val author: Author, val title: Title)`, programmers who construct `Book` objects can’t accidentally switch the author and the title.

Exercises

1. Extend the following `BankAccount` class to a `CheckingAccount` class that charges \$1 for every deposit and withdrawal.

```
class BankAccount(initialBalance: Double) {
    private var balance = initialBalance
    def currentBalance = balance
    def deposit(amount: Double) = { balance += amount; balance }
    def withdraw(amount: Double) = { balance -= amount; balance }
}
```

2. Extend the `BankAccount` class of the preceding exercise into a class `SavingsAccount` that earns interest every month (when a method `earnMonthlyInterest` is called) and has three free deposits or withdrawals every month. Reset the transaction count in the `earnMonthlyInterest` method.
3. Consult your favorite Java or C++ textbook which is sure to have an example of a toy inheritance hierarchy, perhaps involving employees, pets, graphical shapes, or the like. Implement the example in Scala.
4. Define an abstract class `Item` with methods `price` and `description`. A `SimpleItem` is an item whose price and description are specified in the constructor. Take advantage of the fact that a `val` can override a `def`. A `Bundle` is an item that contains other items. Its price is the sum of the prices in the bundle. Also provide a mechanism for adding items to the bundle and a suitable `description` method.
5. Design a class `Point` whose `x` and `y` coordinate values can be provided in a constructor. Provide a subclass `LabeledPoint` whose constructor takes a label value and `x` and `y` coordinates, such as


```
new LabeledPoint("Black Thursday", 1929, 230.07)
```
6. Define an abstract class `Shape` with an abstract method `centerPoint` and subclasses `Rectangle` and `Circle`. Provide appropriate constructors for the subclasses and override the `centerPoint` method in each subclass.

7. Provide a class `Square` that extends `java.awt.Rectangle` and has three constructors: one that constructs a square with a given corner point and width, one that constructs a square with corner $(0, 0)$ and a given width, and one that constructs a square with corner $(0, 0)$ and width 0 .
8. Compile the `Person` and `SecretAgent` classes in Section 8.6, “Overriding Fields,” on page 95 and analyze the class files with `javap`. How many `name` fields are there? How many `name` getter methods are there? What do they get? (Hint: Use the `-c` and `-private` options.)
9. In the `Creature` class of Section 8.10, “Construction Order and Early Definitions,” on page 98, replace `val range` with a `def`. What happens when you also use a `def` in the `Ant` subclass? What happens when you use a `val` in the subclass? Why?
10. The file `scala/collection/immutable/Stack.scala` contains the definition

```
class Stack[A] protected (protected val elems: List[A])
```

Explain the meanings of the `protected` keywords. (Hint: Review the discussion of private constructors in Chapter 5.)

11. Define a value class `Point` that packs integer x and y coordinates into a `Long` (which you should make private).

This page intentionally left blank

Files and Regular Expressions

Topics in This Chapter **A1**

- 9.1 Reading Lines — page 109
- 9.2 Reading Characters — page 110
- 9.3 Reading Tokens and Numbers — page 111
- 9.4 Reading from URLs and Other Sources — page 111
- 9.5 Reading Binary Files — page 112
- 9.6 Writing Text Files — page 112
- 9.7 Visiting Directories — page 112
- 9.8 Serialization — page 113
- 9.9 Process Control **A2** — page 114
- 9.10 Regular Expressions — page 116
- 9.11 Regular Expression Groups — page 117
- Exercises — page 118

Chapter

9

In this chapter, you will learn how to carry out common file processing tasks, such as reading all lines or words from a file or reading a file containing numbers.

Chapter highlights:

- `Source.fromFile(...).getLines.toArray` yields all lines of a file.
- `Source.fromFile(...).mkString` yields the file contents as a string.
- To convert a string into a number, use the `toInt` or `toDouble` method.
- Use the Java `PrintWriter` to write text files.
- `"regex".r` is a Regex object.
- Use `"""..."""` if your regular expression contains backslashes or quotes.
- If a regex pattern has groups, you can extract their contents using the syntax `for (regex(var1, ..., varn) <- string)`.

9.1 Reading Lines

To read all lines from a file, call the `getLines` method on a `scala.io.Source` object:

```
import scala.io.Source
val source = Source.fromFile("myfile.txt", "UTF-8")
// The first argument can be a string or a java.io.File
// You can omit the encoding if you know that the file uses
```

```
// the default platform encoding
val lineIterator = source.getLines
```

The result is an iterator (see Chapter 13). You can use it to process the lines one at a time:

```
for (l <- lineIterator) process l
```

Or you can put the lines into an array or array buffer by applying the `toArray` or `toBuffer` method to the iterator:

```
val lines = source.getLines.toArray
```

Sometimes, you just want to read an entire file into a string. That's even simpler:

```
val contents = source.mkString
```



CAUTION: Call `close` when you are done using the `Source` object.

9.2 Reading Characters

To read individual characters from a file, you can use a `Source` object directly as an iterator since the `Source` class extends `Iterator[Char]`:

```
for (c <- source) process c
```

If you want to be able to peek at a character without consuming it (like `istream::peek` in C++ or a `PushbackInputStreamReader` in Java), call the `buffered` method on the `source` object. Then you can peek at the next input character with the `head` method without consuming it.

```
val source = Source.fromFile("myfile.txt", "UTF-8")
val iter = source.buffered
while (iter.hasNext) {
    if (iter.head is nice)
        process iter.next
    else
        ...
}
```

`source.close()`

Alternatively, if your file isn't large, you can just read it into a string and process that:

```
val contents = source.mkString
```

9.3 Reading Tokens and Numbers

Here is a quick-and-dirty way of reading all whitespace-separated tokens in a source:

```
val tokens = source.mkString.split("\\s+")
```

To convert a string into a number, use the `toInt` or `toDouble` method. For example, if you have a file containing floating-point numbers, you can read them all into an array by

```
val numbers = for (w <- tokens) yield w.toDouble
```

or

```
val numbers = tokens.map(_.toDouble)
```



TIP: Remember—you can always use the `java.util.Scanner` class to process a file that contains a mixture of text and numbers.

Finally, note that you can read numbers from `scala.io.StdIn`:

```
print("How old are you? ")
val age = Scala.io.readInt()
// Or use readDouble or readLong
```



CAUTION: These methods assume that the next input line contains a single number, without leading or trailing whitespace. Otherwise, a `NumberFormatException` occurs.

9.4 Reading from URLs and Other Sources

The `Source` object has methods to read from sources other than files:

```
val source1 = Source.fromURL("http://horstmann.com", "UTF-8")
val source2 = Source.fromString("Hello, World!")
// Reads from the given string—useful for debugging
val source3 = Source.stdin
// Reads from standard input
```



CAUTION: When you read from a URL, you need to know the character set in advance, perhaps from an HTTP header. See www.w3.org/International/O-charset for more information.

9.5 Reading Binary Files

Scala has no provision for reading binary files. You'll need to use the Java library. Here is how you can read a file into a byte array:

```
val file = new File(filename)
val in = new FileInputStream(file)
val bytes = new Array[Byte](file.length.toInt)
in.read(bytes)
in.close()
```

9.6 Writing Text Files

Scala has no built-in support for writing files. To write a text file, use a `java.io.PrintWriter`, for example:

```
val out = new PrintWriter("numbers.txt")
for (i <- 1 to 100) out.println(i)
out.close()
```

Everything works as expected, except for the `printf` method. When you pass a number to `printf`, the compiler will complain that you need to convert it to an `AnyRef`:

```
out.printf("%6d %10.2f",
           quantity.asInstanceOf[AnyRef], price.asInstanceOf[AnyRef]) // Ugh
```

Instead, use the `f` interpolator:

```
out.print(f"$quantity%6d $price%10.2f")
```

9.7 Visiting Directories

There are no “official” Scala classes for visiting all files in a directory, or for recursively traversing directories.

The simplest approach is to use the `Files.list` and `Files.walk` methods of the `java.nio.file` package. The `list` method only visits the children of a directory, and the `walk` method visits all descendants. These methods yield Java streams of `Path` objects. You can visit them as follows:

```
import java.nio.file._
String dirname = "/home/cay/scala-impatient/code"
val entries = Files.walk(Paths.get(dirname)) // or Files.list
```

```

try {
    entries.forEach(p => Process the path p)
} finally {
    entries.close()
}

```

9.8 Serialization

In Java, serialization is used to transmit objects to other virtual machines or for short-term storage. (For long-term storage, serialization can be awkward—it is tedious to deal with different object versions as classes evolve over time.)

Here is how you declare a serializable class in Java and Scala.

Java:

```

public class Person implements java.io.Serializable {
    private static final long serialVersionUID = 42L;
    ...
}

```

Scala:

```
@SerialVersionUID(42L) class Person extends Serializable
```

The `Serializable` trait is defined in the `scala` package and does not require an import.



NOTE: You can omit the `@SerialVersionUID` annotation if you are OK with the default ID.

Serialize and deserialize objects in the usual way:

```

val fred = new Person(...)
import java.io._
val out = new ObjectOutputStream(new FileOutputStream("/tmp/test.obj"))
out.writeObject(fred)
out.close()
val in = new ObjectInputStream(new FileInputStream("/tmp/test.obj"))
val savedFred = in.readObject().asInstanceOf[Person]

```

The Scala collections are serializable, so you can have them as members of your serializable classes:

```

class Person extends Serializable {
    private val friends = new ArrayBuffer[Person] // OK—ArrayBuffer is serializable
    ...
}

```

9.9 Process Control A2

Traditionally, programmers use shell scripts to carry out mundane processing tasks, such as moving files from one place to another, or combining a set of files. The shell language makes it easy to specify subsets of files and to pipe the output of one program into the input of another. However, as programming languages, most shell languages leave much to be desired.

Scala was designed to scale from humble scripting tasks to massive programs. The `scala.sys.process` package provides utilities to interact with shell programs. You can write your shell scripts in Scala, with all the power that the Scala language puts at your disposal.

Here is a simple example:

```
import scala.sys.process._  
"ls -al ..".!
```

As a result, the `ls -al ..` command is executed, showing all files in the parent directory. The result is printed to standard output.

The `scala.sys.process` package contains an implicit conversion from strings to `ProcessBuilder` objects. The `!` method *executes* the `ProcessBuilder` object.

The result of the `!` method is the exit code of the executed program: `0` if the program was successful, or a nonzero failure indicator otherwise.

If you use `!!` instead of `!`, the output is returned as a string:

```
val result = "ls -al /"!!
```



NOTE: The `!` and `!!` operators were originally intended to be used as postfix operators without the method invocation syntax:

```
"ls -al /" !!
```

However, as you will see in Chapter 11, the postfix syntax is being deprecated since it can lead to parsing errors.

You can pipe the output of one program into the input of another, using the `#|` method:

```
("ls -al /" #| "grep u")..!
```



NOTE: As you can see, the process library uses the commands of the underlying operating system. Here, I use bash commands because bash is available on Linux, Mac OS X, and Windows.

To redirect the output to a file, use the #> method:

```
("ls -al /" #> new File("filelist.txt")).!
```

To append to a file, use #>> instead:

```
("ls -al /etc" #>> new File("filelist.txt")).!
```

To redirect input from a file, use #<:

```
("grep u" #< new File("filelist.txt")).!
```

You can also redirect input from a URL:

```
("grep Scala" #< new URL("http://horstmann.com/index.html")).!
```

You can combine processes with p #&& q (execute q if p was successful) and p #|| q (execute q if p was unsuccessful). But frankly, Scala is better at control flow than the shell, so why not implement the control flow in Scala?



NOTE: The process library uses the familiar shell operators | > >> < && || |, but it prefixes them with a # so that they all have the same precedence.

If you need to run a process in a different directory, or with different environment variables, construct a `ProcessBuilder` with the `apply` method of the `Process` object. Supply the command, the starting directory, and a sequence of (*name*, *value*) pairs for environment settings:

```
val p = Process(cmd, new File(dirName), ("LANG", "en_US"))
```

Then execute it with the ! method:

```
("echo 42" #| p).!
```



NOTE: If you want to use Scala for shell scripts in a UNIX/Linux/MacOS environment, start your script files like this:

```
#!/bin/sh
exec scala "$0" "$@"
#
Scala commands
```



NOTE: You can also run Scala scripts from Java programs with the scripting integration of the `javax.script` package. To get a script engine, call

```
ScriptEngine engine =
new ScriptEngineManager().getScriptEngineByName("scala")
```

9.10 Regular Expressions

When you process input, you often want to use regular expressions to analyze it. The `scala.util.matching.Regex` class makes this simple. To construct a `Regex` object, use the `r` method of the `String` class:

```
val numPattern = "[0-9]+".r
```

If the regular expression contains backslashes or quotation marks, then it is a good idea to use the “raw” string syntax, `"""\..."""`. For example:

```
val wsnumwsPattern = """\s+[0-9]+\s+""".r
// A bit easier to read than "\s+[0-9]+\s+".r
```

The `findAllIn` method returns an `Iterator[String]` through all matches. You can use it in a `for` loop:

```
for (matchString <- numPattern.findAllIn("99 bottles, 98 bottles"))
  println(matchString)
```

Alternatively, turn the iterator into an array:

```
val matches = numPattern.findAllIn("99 bottles, 98 bottles").toArray
// Array("99", "98")
```

To find the first match in a string, use `findFirstIn`. You get an `Option[String]`. (See Chapter 14 for the `Option` class.)

```
val firstMatch = wsnumwsPattern.findFirstIn("99 bottles, 98 bottles")
// Some(" 98 ")
```



NOTE: There is no method to test whether a string matches the regex in its entirety, but you can add anchors:

```
val anchoredPattern = "^\\d+\\d+$".r
if (anchoredPattern.findFirstIn(str) != None) ...
```

Alternatively, use the `String.matches` method:

```
if (str.matches("\\d+\\d+")) ...
```

You can replace the first match, all matches, or some matches. In the latter case, supply a function `Match => Option[String]`. The `Match` class has information about the match (see the next section for details). If the function returns `Some(str)`, the match is replaced with `str`.

```
numPattern.replaceFirstIn("99 bottles, 98 bottles", "XX")
// "XX bottles, 98 bottles"
```

```
numPattern.replaceAllIn("99 bottles, 98 bottles", "XX")
// "XX bottles, XX bottles"
numPattern.replaceSomeIn("99 bottles, 98 bottles",
  m => if (m.matched.toInt % 2 == 0) Some("XX") else None)
// "99 bottles, XX bottles"
```

Here is a more useful application of the `replaceSomeIn` method. We want to replace placeholders `$0`, `$1`, and so on, in a message string with values from an argument sequence. Make a pattern for the variable with a group for the index, and then map the group to the sequence element.

```
val varPattern = """\$[0-9]+""".r
def format(message: String, vars: String*) =
  varPattern.replaceSomeIn(message, m => vars.lift(
    m.matched.tail.toInt))
format("At $1, there was $2 on $0.",
  "planet 7", "12:30 pm", "a disturbance of the force")
// At 12:30 pm, there was a disturbance of the force on planet 7.
```

The `lift` method turns a `Seq[String]` into a function. The expression `vars.lift(i)` is `Some(vars(i))` if `i` is a valid index or `None` if it is not.

9.11 Regular Expression Groups

Groups are useful to get subexpressions of regular expressions. Add parentheses around the subexpressions that you want to extract, for example:

```
val numItemPattern = "([0-9]+) ([a-z]+)".r
```

You can get the group contents from a `Match` object. The methods `findAllMatchIn` and `findFirstMatchIn` are analogs of the `findAllIn` and `findFirstIn` methods that return an `Iterator[Match]` or `Option[Match]`.

If `m` is a `Match` object, then `m.matched` is the entire match string and `m.group(i)` is the `i`th group. The start and end indices of these substrings in the original string are `m.start`, `m.end`, `m.start(i)`, and `m.end(i)`.

```
for (m <- numItemPattern.findAllMatchIn("99 bottles, 98 bottles"))
  println(m.group(1)) // Prints 99 and 98
```



CAUTION: The `Match` class has methods for retrieving groups by name.

However, this does *not* work with group names inside regular expressions, such as `"(?<num>[0-9]+) (?<item>[a-z]+)".r`. Instead, one needs to supply names to the `r` method: `"([0-9]+) ([a-z]+)".r("num", "item")`

There is another convenient way of extracting matches. Use a regular expression variable as an “extractor” (see Chapter 14), like this:

```
val numitemPattern(num, item) = "99 bottles"
// Sets num to "99", item to "bottles"
```

When you use a pattern as an extractor, it must match the string from which you extract the matches, and there must be a group for each variable.

To extract groups from multiple matches, you can use a for statement like this:

```
for (numitemPattern(num, item) <- numitemPattern.findAllIn("99 bottles, 98 bottles"))
  process num and item
```

Exercises

1. Write a Scala code snippet that reverses the lines in a file (making the last line the first one, and so on).
2. Write a Scala program that reads a file with tabs, replaces each tab with spaces so that tab stops are at n -column boundaries, and writes the result to the same file.
3. Write a Scala code snippet that reads a file and prints all words with more than 12 characters to the console. Extra credit if you can do this in a single line.
4. Write a Scala program that reads a text file containing only floating-point numbers. Print the sum, average, maximum, and minimum of the numbers in the file.
5. Write a Scala program that writes the powers of 2 and their reciprocals to a file, with the exponent ranging from 0 to 20. Line up the columns:

1	1
2	0.5
4	0.25
...	...

6. Make a regular expression searching for quoted strings "like this, maybe with \" or \\\" in a Java or C++ program. Write a Scala program that prints out all such strings in a source file.
7. Write a Scala program that reads a text file and prints all tokens in the file that are *not* floating-point numbers. Use a regular expression.
8. Write a Scala program that prints the src attributes of all img tags of a web page. Use regular expressions and groups.

9. Write a Scala program that counts how many files with .class extension are in a given directory and its subdirectories.
10. Expand the example in Section 9.8, “Serialization,” on page 113. Construct a few Person objects, make some of them friends of others, and save an Array[Person] to a file. Read the array back in and verify that the friend relations are intact.

Traits

Topics in This Chapter **L1**

- 10.1 Why No Multiple Inheritance? — page 121
- 10.2 Traits as Interfaces — page 123
- 10.3 Traits with Concrete Implementations — page 124
- 10.4 Objects with Traits — page 125
- 10.5 Layered Traits — page 125
- 10.6 Overriding Abstract Methods in Traits — page 127
- 10.7 Traits for Rich Interfaces — page 127
- 10.8 Concrete Fields in Traits — page 128
- 10.9 Abstract Fields in Traits — page 130
- 10.10 Trait Construction Order — page 130
- 10.11 Initializing Trait Fields — page 132
- 10.12 Traits Extending Classes — page 133
- 10.13 Self Types **L2** — page 134
- 10.14 What Happens under the Hood — page 135
- Exercises — page 137

Chapter 10

In this chapter, you will learn how to work with traits. A class extends one or more traits in order to take advantage of the services that the traits provide. A trait may require implementing classes to support certain features. However, unlike Java interfaces, Scala traits can supply state and behavior for these features, which makes them far more useful.

Key points of this chapter:

- A class can implement any number of traits.
- Traits can require implementing classes to have certain fields, methods, or superclasses.
- Unlike Java interfaces, a Scala trait can provide implementations of methods and fields.
- When you layer multiple traits, the order matters—the trait whose methods execute *first* goes to the back.

10.1 Why No Multiple Inheritance?

Scala, like Java, does not allow a class to inherit from multiple superclasses. At first, this seems like an unfortunate restriction. Why shouldn't a class extend multiple classes? Some programming languages, in particular C++, allow multiple inheritance—but at a surprisingly high cost.

Multiple inheritance works fine when you combine classes that have *nothing in common*. But if these classes have common methods or fields, thorny issues come up. Here is a typical example. A teaching assistant is a student and also an employee:

```
class Student {
    def id: String = ...
    ...
}

class Employee {
    def id: String = ...
    ...
}
```

Suppose we could have

```
class TeachingAssistant extends Student, Employee { // Not actual Scala code
    ...
}
```

Unfortunately, this `TeachingAssistant` class inherits *two* `id` methods. What should `myTA.id` return? The student ID? The employee ID? Both? (In C++, you need to redefine the `id` method to clarify what you want.)

Next, suppose that both `Student` and `Employee` extend a common superclass `Person`:

```
class Person {
    var name: String = _
}

class Student extends Person { ... }
class Employee extends Person { ... }
```

This leads to the *diamond inheritance* problem (see Figure 10–1). We only want one `name` field inside a `TeachingAssistant`, not two. How do the fields get merged? How does the field get constructed? In C++, you use “virtual base classes,” a complex and brittle feature, to address this issue.

Java designers were so afraid of these complexities that they took a very restrictive approach. A class can extend only one superclass; it can implement any number of *interfaces*, but interfaces can have only abstract, static, or default methods, and no fields.

Java default methods are very restricted. They can call other interface methods, but they cannot make use of object state. It is therefore common in Java to provide both an interface and an abstract base class, but that just kicks the can down the road. What if you need to extend two of those abstract base classes?

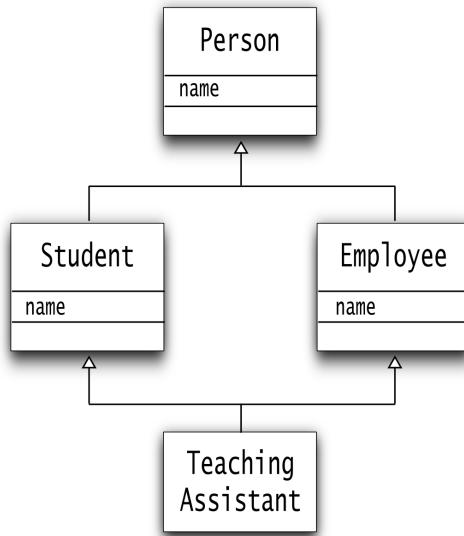


Figure 10–1 Diamond inheritance must merge common fields.

Scala has *traits* instead of interfaces. A trait can have abstract and concrete methods, as well as state, and a class can implement multiple traits. This neatly solves the problem of Java interfaces. You will see in the following sections how Scala deals with the perils of conflicting features from multiple traits.

10.2 Traits as Interfaces

Let's start with something familiar. A Scala trait can work exactly like a Java interface. For example:

```
trait Logger {
    def log(msg: String) // An abstract method
}
```

Note that you need not declare the method as abstract—an unimplemented method in a trait is automatically abstract.

A subclass can provide an implementation:

```
class ConsoleLogger extends Logger { // Use extends, not implements
    def log(msg: String) { println(msg) } // No override needed
}
```

You need not supply the `override` keyword when overriding an abstract method of a trait.



NOTE: Scala doesn't have a special keyword for implementing a trait. As you will see throughout this chapter, traits can be much more similar to classes than Java interfaces.

If you need more than one trait, add the others using the `with` keyword:

```
class ConsoleLogger extends Logger with Cloneable with Serializable
```

Here we use the `Cloneable` and `Serializable` interfaces from the Java library, just for the sake of showing the syntax. All Java interfaces can be used as Scala traits.

As in Java, a Scala class can have only one superclass but any number of traits.



NOTE: It may seem odd that you use the `extends` keyword before the first trait but `with` before all the others. But that's not the way that Scala thinks about it. In Scala, `Logger with Cloneable with Serializable` is the entity that the class extends.

10.3 Traits with Concrete Implementations

In Scala, the methods of a trait need not be abstract. For example, we can make our `ConsoleLogger` into a trait:

```
trait ConsoleLogger {
    def log(msg: String) { println(msg) }
}
```

The `ConsoleLogger` trait provides a method *with an implementation*—in this case, one that prints the logging message on the console.

Here is an example of using this trait:

```
class SavingsAccount extends Account with ConsoleLogger {
    def withdraw(amount: Double) {
        if (amount > balance) log("Insufficient funds")
        else balance -= amount
    }
    ...
}
```

Note how the `SavingsAccount` picks up a concrete implementation from the `ConsoleLogger` trait. In Java, this is also possible by using default methods in interfaces. However, as you will see shortly, traits can also have state, which would not be possible with a Java interface.

In Scala (and other programming languages that allow this), we say that the `ConsoleLogger` functionality is “mixed in” with the `SavingsAccount` class.



NOTE: Supposedly, the “mix in” term comes from the world of ice cream. In the ice cream parlor parlance, a “mix in” is an additive that is kneaded into a scoop of ice cream before dispensing it to the customer—a practice that may be delicious or disgusting depending on your point of view.

10.4 Objects with Traits

You can add a trait to an individual object when you construct it. Let’s first define this class:

```
abstract class SavingsAccount extends Account with Logger {
    def withdraw(amount: Double) {
        if (amount > balance) log("Insufficient funds")
        else ...
    }
    ...
}
```

This class is abstract since it can’t yet do any logging, which might seem pointless. But you can “mix in” a concrete logger when constructing an object.

```
trait ConsoleLogger extends Logger {
    def log(msg: String) { println(msg) }
}

val acct = new SavingsAccount with ConsoleLogger
```

When calling `log` on the `acct` object, the `log` method of the `ConsoleLogger` trait executes. Of course, another object can add in a different trait:

```
val acct2 = new SavingsAccount with FileLogger
```

10.5 Layered Traits

You can add, to a class or an object, multiple traits that invoke each other starting with the *last one*. This is useful when you need to transform a value in stages.

Here is a simple example. We may want to add a timestamp to all logging messages.

```
trait TimestampLogger extends ConsoleLogger {
    override def log(msg: String) {
        super.log(s"${java.time.Instant.now()} $msg")
    }
}
```

Also, suppose we want to truncate overly chatty log messages like this:

```
trait ShortLogger extends ConsoleLogger {
    override def log(msg: String) {
        super.log(
            if (msg.length <= 15) msg else s"${msg.substring(0, 12)}...")
    }
}
```

Note that each of the `log` methods passes a modified message to `super.log`.

With traits, `super.log` does *not* have the same meaning as it does with classes. Instead, `super.log` calls the `log` method of another trait, which depends on the order in which the traits are added.

To see how the order matters, compare the following two examples:

```
val acct1 = new SavingsAccount with TimestampLogger with ShortLogger
val acct2 = new SavingsAccount with ShortLogger with TimestampLogger
```

If we overdraw `acct1`, we get a message

```
Sun Feb 06 17:45:45 ICT 2011 Insufficient...
```

As you can see, the `ShortLogger`'s `log` method was called first, and its call to `super.log` called the `TimestampLogger`.

However, overdrawing `acct2` yields

```
Sun Feb 06 1...
```

Here, the `TimestampLogger` appeared last in the list of traits. Its `log` message was called first, and the result was subsequently shortened.

For simple mixin sequences, the “back to front” rule will give you the right intuition. See Section 10.10, “Trait Construction Order,” on page 130 for the gory details that arise when the traits form an arbitrary tree and not just a chain.



NOTE: With traits, you cannot tell from the source code which method is invoked by `super.someMethod`. The exact method depends on the ordering of the traits in the object or class that uses them. This makes super far more flexible than in plain old inheritance.



NOTE: If you want to control which trait's method is invoked, you can specify it in brackets: `super[ConsoleLogger].log(...)`. The specified type must be an immediate supertype; you can't access traits or classes that are further away in the inheritance hierarchy.

10.6 Overriding Abstract Methods in Traits

In the preceding section, the `TimestampLogger` and `ShortLogger` traits extended `ConsoleLogger`. Let's make them extend our `Logger` trait instead, where we provide *no implementation* to the `log` method.

```
trait Logger {  
    def log(msg: String) // This method is abstract  
}
```

Then, the `TimestampLogger` class no longer compiles.

```
trait TimestampLogger extends Logger {  
    override def log(msg: String) { // Overrides an abstract method  
        super.log(s"${java.time.Instant.now()} $msg") // Is super.log defined?  
    }  
}
```

The compiler flags the call to `super.log` as an error.

Under normal inheritance rules, this call could never be correct—the `Logger.log` method has no implementation. But actually, as you saw in the preceding section, there is no way of knowing which `log` method is actually being called—it depends on the order in which traits are mixed in.

Scala takes the position that `TimestampLogger.log` is still abstract—it requires a concrete `log` method to be mixed in. You therefore need to tag the method with the `abstract` keyword *and* the `override` keyword, like this:

```
abstract override def log(msg: String) {  
    super.log(s"${java.time.Instant.now()} $msg")  
}
```

10.7 Traits for Rich Interfaces

A trait can have many utility methods that depend on a few abstract ones. One example is the Scala `Iterator` trait that defines dozens of methods in terms of the abstract `next` and `hasNext` methods.

Let us enrich our rather anemic logging API. Usually, a logging API lets you specify a level for each log message to distinguish informational messages from warnings or errors. We can easily add this capability without forcing any policy for the destination of logging messages.

```
trait Logger {
    def log(msg: String)
    def info(msg: String) { log(s"INFO: $msg") }
    def warn(msg: String) { log(s"WARNING: $msg") }
    def severe(msg: String) { log(s"SEVERE: $msg") }
}
```

Note the combination of abstract and concrete methods.

A class that uses the `Logger` trait can now call any of these logging messages, for example:

```
abstract class SavingsAccount extends Account with Logger {
    def withdraw(amount: Double) {
        if (amount > balance) severe("Insufficient funds")
        else ...
    }
    ...
}
```

This use of concrete and abstract methods in a trait is very common in Scala. In Java, you can achieve the same with default methods.

10.8 Concrete Fields in Traits

A field in a trait can be concrete or abstract. If you supply an initial value, the field is concrete.

```
trait ShortLogger extends Logger {
    val maxLength = 15 // A concrete field
    abstract override def log(msg: String) {
        super.log(
            if (msg.length <= maxLength) msg
            else s"${msg.substring(0, maxLength - 3)}...")
    }
}
```

A class that mixes in this trait acquires a `maxLength` field. In general, a class gets a field for each concrete field in one of its traits. These fields are not inherited; they are simply added to the subclass. This may seem a subtle distinction, but it is important. Let us look at the process more closely, with this version of the `SavingsAccount` class:

```
class SavingsAccount extends Account with ConsoleLogger with ShortLogger {
    var interest = 0.0
    def withdraw(amount: Double) {
        if (amount > balance) log("Insufficient funds")
        else ...
    }
}
```

Note that our subclass has a field `interest`. That's a plain old field in the subclass.

Suppose `Account` has a field.

```
class Account {
    var balance = 0.0
}
```

The `SavingsAccount` class *inherits* that field in the usual way. A `SavingsAccount` object is made up of all the fields of its superclasses, together with any fields in the subclass. You can think of a `SavingsAcccount` object as “starting out” with a superclass object (see Figure 10–2).

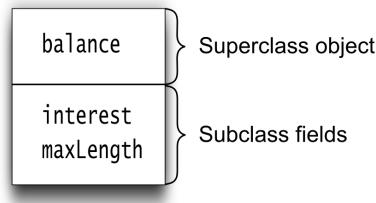


Figure 10–2 Fields from a trait are placed in the subclass.

In the JVM, a class can only extend one superclass, so the trait fields can't be inherited in the same way. Instead, the Scala compiler adds the `maxLength` field to the `SavingsAccount` class, next to the `interest` field.



CAUTION: When you extend a class and then change the superclass, the subclass doesn't have to be recompiled because the virtual machine understands inheritance. But when a trait changes, all classes that mix in that trait must be recompiled.

You can think of concrete trait fields as “assembly instructions” for the classes that use the trait. Any such fields become fields of the class.

10.9 Abstract Fields in Traits

An uninitialized field in a trait is abstract and must be overridden in a concrete subclass.

For example, the following `maxLength` field is abstract:

```
trait ShortLogger extends Logger {
    val maxLength: Int // An abstract field
    abstract override def log(msg: String) { ... }
    super.log()
        if (msg.length <= maxLength) msg
        else s"${msg.substring(0, maxLength - 3)}...")
            // The maxLength field is used in the implementation
    }
    ...
}
```

When you use this trait in a concrete class, you must supply the `maxLength` field:

```
class SavingsAccount extends Account with ConsoleLogger with ShortLogger {
    val maxLength = 20 // No override necessary
    ...
}
```

Now all logging messages are truncated after 20 characters.

This way of supplying values for trait parameters is particularly handy when you construct objects on the fly. Let's go back to our original savings account:

```
class SavingsAccount extends Account with Logger { ... }
```

Now, we can truncate the messages in an instance as follows:

```
val acct = new SavingsAccount with ConsoleLogger with ShortLogger {
    val maxLength = 20
}
```

10.10 Trait Construction Order

Just like classes, traits can have constructors, made up of field initializations and other statements in the trait's body. For example,

```
trait FileLogger extends Logger {
    val out = new PrintWriter("app.log") // Part of the trait's constructor
    out.println(s"# ${java.time.Instant.now()}") // Also part of the constructor

    def log(msg: String) { out.println(msg); out.flush() }
}
```

These statements are executed during construction of any object incorporating the trait.

Constructors execute in the following order:

1. The superclass constructor is called first.
2. Trait constructors are executed after the superclass constructor but before the class constructor.
3. Traits are constructed left-to-right.
4. Within each trait, the parents get constructed first.
5. If multiple traits share a common parent, and that parent has already been constructed, it is not constructed again.
6. After all traits are constructed, the subclass is constructed.

For example, consider this class:

```
class SavingsAccount extends Account with FileLogger with ShortLogger
```

The constructors execute in the following order:

1. Account (the superclass).
2. Logger (the parent of the first trait).
3. FileLogger (the first trait).
4. ShortLogger (the second trait). Note that its Logger parent has already been constructed.
5. SavingsAccount (the class).



NOTE: The constructor ordering is the reverse of the *linearization* of the class.

The linearization is a technical specification of all supertypes of a type. It is defined by the rule:

If C extends C_1 with C_2 with ... with C_n , then $lin(C) = C \gg lin(C_n) \gg \dots \gg lin(C_2) \gg lin(C_1)$

Here, \gg means “concatenate and remove duplicates, with the right winning out.” For example,

```
lin(SavingsAccount)
= SavingsAccount \gg lin(ShortLogger) \gg lin(FileLogger) \gg lin(Account)
= SavingsAccount \gg (ShortLogger \gg Logger) \gg (FileLogger \gg Logger) \gg lin(Account)
= SavingsAccount \gg ShortLogger \gg FileLogger \gg Logger \gg Account.
```

(For simplicity, I omitted the types ScalaObject, AnyRef, and Any that are at the end of any linearization.)

The linearization gives the order in which super is resolved in a trait. For example, calling super in a ShortLogger invokes the FileLogger method, and calling super in a FileLogger invokes the Logger method.

10.11 Initializing Trait Fields

Traits cannot have constructor parameters. Every trait has a single parameterless constructor.



NOTE: Interestingly, the absence of constructor parameters is the *only* technical difference between traits and classes. Otherwise, traits can have all the features of classes, such as concrete and abstract fields and superclasses.

This limitation can be a problem for traits that need some customization to be useful. Consider a file logger. We would like to specify the log file, but we can't use a construction parameter:

```
val acct = new SavingsAccount with FileLogger("myapp.log")
// Error: Can't have constructor parameters for traits
```

You saw one possible approach in the preceding section. The FileLogger can have an abstract field for the file name.

```
trait FileLogger extends Logger {
    val filename: String
    val out = new PrintStream(filename)
    def log(msg: String) { out.println(msg); out.flush() }
}
```

A class using this trait can override the filename field. Unfortunately, there is a pitfall. The straightforward approach does *not* work:

```
val acct = new SavingsAccount with FileLogger {
    val filename = "myapp.log" // Does not work
}
```

The problem is the construction order. The FileLogger constructor runs *before* the subclass constructor. Here, the subclass is a bit hard to see. The new statement constructs an instance of an anonymous class extending SavingsAccount (the superclass) with the FileLogger trait. The initialization of filename only happens in the anonymous subclass. Actually, it doesn't happen at all—before the subclass gets its turn, a null pointer exception is thrown in the FileLogger constructor.

One remedy is an obscure feature that we described in Chapter 8: *early definition*. Here is the correct version:

```
val acct = new { // Early definition block after new
    val filename = "myapp.log"
} with SavingsAccount with FileLogger
```

It's not pretty, but it solves our problem. The early definition happens before the regular construction sequence. When the `FileLogger` is constructed, the `filename` field is initialized.

If you need to do the same in a class, the syntax looks like this:

```
class SavingsAccount extends { // Early definition block after extends
    val filename = "savings.log"
} with Account with FileLogger {
    ... // SavingsAccount implementation
}
```

Another alternative is to use a *lazy value* in the `FileLogger` constructor, like this:

```
trait FileLogger extends Logger {
    val filename: String
    lazy val out = new PrintStream(filename)
    def log(msg: String) { out.println(msg) } // No override needed
}
```

Then the `out` field is initialized when it is first used. At that time, the `filename` field will have been set. However, lazy values are somewhat inefficient because they are checked for initialization before every use.

10.12 Traits Extending Classes

As you have seen, a trait can extend another trait, and it is common to have a hierarchy of traits. Less commonly, a trait can also extend a class. That class becomes a superclass of any class mixing in the trait.

Here is an example. The `LoggedException` trait extends the `Exception` class:

```
trait LoggedException extends Exception with ConsoleLogger {
    def log() { log(getMessage()) }
}
```

A `LoggedException` has a `log` method to log the exception's message. Note that the `log` method calls the `getMessage` method that is inherited from the `Exception` superclass.

Now let's form a class that mixes in this trait:

```
class UnhappyException extends LoggedException { // This class extends a trait
    override def getMessage() = "arggh!"
}
```

The superclass of the trait becomes the superclass of our class (see Figure 10–3).

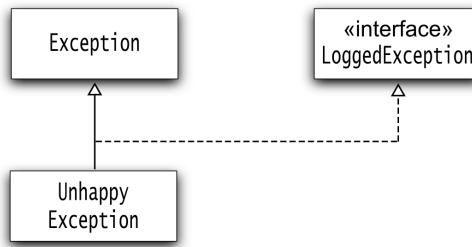


Figure 10–3 The Superclass of a trait becomes the superclass of any class mixing in the trait.

What if our class already extends another class? That's OK, as long as it's a subclass of the trait's superclass. For example,

```
class UnhappyException extends IOException with LoggedException
```

Here `UnhappyException` extends `IOException`, which already extends `Exception`. When mixing in the trait, its superclass is already present, and there is no need to add it. However, if our class extends an unrelated class, then it is not possible to mix in the trait. For example, you cannot form the following class:

```
class UnhappyFrame extends JFrame with LoggedException
// Error: Unrelated superclasses
```

It would be impossible to add both `JFrame` and `Exception` as superclasses.

10.13 Self Types L2

When a trait extends a class, there is a guarantee that the superclass is present in any class mixing in the trait. Scala has an alternate mechanism for guaranteeing this: *self types*.

When a trait starts out with

```
this: Type =>
```

then it can only be mixed into a subclass of the given type.

Let's use this feature for our `LoggedException`:

```
trait LoggedException extends ConsoleLogger {
  this: Exception =>
  def log() { log(getMessage()) }
}
```

Note that the trait does *not* extend the `Exception` class. Instead, it has a self type of `Exception`. That means it can only be mixed into subclasses of `Exception`.

In the trait's methods, we can call any methods of the self type. For example, the call to `getMessage()` in the `log` method is valid, since we know that this must be an `Exception`.

If you try to mix the trait into a class that doesn't conform to the self type, an error occurs.

```
val f = new JFrame with LoggedException
// Error: JFrame isn't a subtype of Exception, the self type of LoggedException
```

A trait with a self type is similar to a trait with a supertype. In both cases, it is ensured that a type is present in a class that mixes in the trait.

There are a few situations where the self type notation is more flexible than traits with supertypes. Self types can handle circular dependencies between traits. This can happen if you have two traits that need each other.

Self types can also handle *structural types*—types that merely specify the methods that a class must have, without naming the class. Here is the `LoggedException` using a structural type:

```
trait LoggedException extends ConsoleLogger {
  this: { def getMessage(): String } =>
  def log() { log(getMessage()) }
}
```

The trait can be mixed into any class that has a `getMessage` method.

We discuss self types and structural types in more detail in Chapter 19.

10.14 What Happens under the Hood

Scala needs to translate traits into classes and interfaces of the JVM. You are not required to know how this is done, but you may find it helpful for understanding how traits work.

A trait that has only abstract methods is simply turned into a Java interface. For example,

```
trait Logger {
  def log(msg: String)
}
```

turns into

```
public interface Logger { // Generated Java interface
  void log(String msg);
}
```

Trait methods become default methods. For example,

```
trait ConsoleLogger {
    def log(msg: String) { println(msg) }
}
```

becomes

```
public interface ConsoleLogger {
    default void log(String msg) { ... }
}
```

If the trait has fields, the Java interface has getter and setter methods.

```
trait ShortLogger extends Logger {
    val maxLength = 15 // A concrete field
    ...
}
```

is translated to

```
public interface ShortLogger extends Logger {
    int maxLength();
    void weird_prefix$maxLength$_eq(int);
    default void log(String msg) { ... } // Calls maxLength()
    default void $init$() { weird_prefix$maxLength$_eq(15); }
}
```

Of course, the interface can't have any fields, and the getter and setter methods are unimplemented. But the getter is called when the field value is needed.

The weird setter is needed to initialize the field. This happens in the `$init$` method.

When the trait is mixed into a class, the class gets a `maxLength` field, and the getter and setter are defined to get and set the field. The constructors of that class invokes the `$init$` method of the trait. For example,

```
class SavingsAccount extends Account with ConsoleLogger with ShortLogger
```

turns into

```
public class SavingsAccount extends Account
    implements ConsoleLogger, ShortLogger {
    private int maxLength;
    public int maxLength() { return maxLength; }
    public void weird_prefix$maxLength$_eq(int arg) { maxLength = arg; }
    public SavingsAccount() {
        super();
        ConsoleLogger.$init$();
        ShortLogger.$init$();
    }
    ...
}
```

If a trait extends a superclass, the trait still turns into an interface. Of course, a class mixing in the trait extends the superclass.

Exercises

1. The `java.awt.Rectangle` class has useful methods `translate` and `grow` that are unfortunately absent from classes such as `java.awt.geom.Ellipse2D`. In Scala, you can fix this problem. Define a trait `RectangleLike` with concrete methods `translate` and `grow`. Provide any abstract methods that you need for the implementation, so that you can mix in the trait like this:

```
val egg = new java.awt.geom.Ellipse2D.Double(5, 10, 20, 30) with RectangleLike
egg.translate(10, -10)
egg.grow(10, 20)
```

2. Define a class `OrderedPoint` by mixing `scala.math.Ordered[Point]` into `java.awt.Point`. Use lexicographic ordering, i.e. $(x, y) < (x', y')$ if $x < x'$ or $x = x'$ and $y < y'$.
3. Look at the `BitSet` class, and make a diagram of all its superclasses and traits. Ignore the type parameters (everything inside the [...]). Then give the linearization of the traits.
4. Provide a `CryptoLogger` trait that encrypts the log messages with the Caesar cipher. The key should be 3 by default, but it should be overridable by the user. Provide usage examples with the default key and a key of -3.
5. The JavaBeans specification has the notion of a *property change listener*, a standardized way for beans to communicate changes in their properties. The `PropertyChangeSupport` class is provided as a convenience superclass for any bean that wishes to support property change listeners. Unfortunately, a class that already has another superclass—such as `JComponent`—must reimplement the methods. Reimplement `PropertyChangeSupport` as a trait, and mix it into the `java.awt.Point` class.
6. In the Java AWT library, we have a class `Container`, a subclass of `Component` that collects multiple components. For example, a `Button` is a `Component`, but a `Panel` is a `Container`. That's the composite pattern at work. Swing has `JComponent` and `JButton`, but if you look closely, you will notice something strange. `JComponent` extends `Container`, even though it makes no sense to add other components to, say, a `JButton`. Ideally, the Swing designers would have preferred the design in Figure 10–4.

But that's not possible in Java. Explain why not. How could the design be executed in Scala with traits?

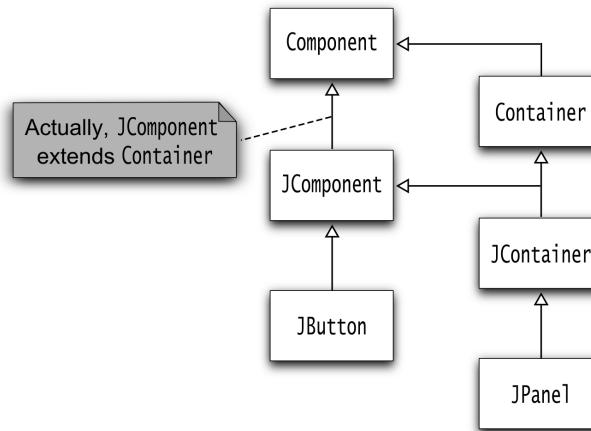


Figure 10–4 A better design for Swing containers

7. Construct an example where a class needs to be recompiled when one of the mixins changes. Start with class `SavingsAccount` extends `Account` with `ConsoleLogger`. Put each class and trait in a separate source file. Add a field to `Account`. In `Main` (also in a separate source file), construct a `SavingsAccount` and access the new field. Recompile all files *except for* `SavingsAccount` and verify that the program works. Now add a field to `ConsoleLogger` and access it in `Main`. Again, recompile all files *except for* `SavingsAccount`. What happens? Why?
8. There are dozens of Scala trait tutorials with silly examples of barking dogs or philosophizing frogs. Reading through contrived hierarchies can be tedious and not very helpful, but designing your own is very illuminating. Make your own silly trait hierarchy example that demonstrates layered traits, concrete and abstract methods, and concrete and abstract fields.
9. In the `java.io` library, you add buffering to an input stream with a `BufferedInputStream` decorator. Reimplement buffering as a trait. For simplicity, override the `read` method.
10. Using the logger traits from this chapter, add logging to the solution of the preceding problem that demonstrates buffering.
11. Implement a class `IterableInputStream` that extends `java.io.InputStream` with the trait `Iterable[Byte]`.
12. Using `javap -c -private`, analyze how the call `super.log(msg)` is translated to Java. How does the same call invoke two different methods, depending on the mixin order?

This page intentionally left blank

Operators

Topics in This Chapter [L1](#)

- 11.1 Identifiers — page 142
- 11.2 Infix Operators — page 143
- 11.3 Unary Operators — page 143
- 11.4 Assignment Operators — page 144
- 11.5 Precedence — page 144
- 11.6 Associativity — page 145
- 11.7 The `apply` and `update` Methods — page 146
- 11.8 Extractors [L2](#) — page 147
- 11.9 Extractors with One or No Arguments [L2](#) — page 149
- 11.10 The `unapplySeq` Method [L2](#) — page 149
- 11.11 Dynamic Invocation [L2](#) — page 150
- Exercises — page 153

Chapter

11

This chapter covers in detail implementing your own *operators*—methods with the same syntax as the familiar mathematical operators. Operators are often used to build *domain-specific languages*—minilanguages embedded inside Scala. *Implicit conversions* (type conversion functions that are applied automatically) are another tool facilitating the creation of domain-specific languages. This chapter also discusses the special methods `apply`, `update`, and `unapply`. We end the chapter with a discussion of *dynamic invocations*—method calls that can be intercepted at runtime, so that arbitrary actions can occur depending on the method names and arguments.

The key points of this chapter are:

- Identifiers contain either alphanumeric or operator characters.
- Unary and binary operators are method calls.
- Operator precedence depends on the first character, associativity on the last.
- The `apply` and `update` methods are called when evaluating `expr(args)`.
- Extractors extract tuples or sequences of values from an input [L2](#).
- Types extending the `Dynamic` trait can inspect the names of methods and arguments at runtime. [L2](#)

11.1 Identifiers

The names of variables, functions, classes, and so on are collectively called *identifiers*. In Scala, you have more choices for forming identifiers than in Java. Of course, you can follow the time-honored pattern: sequences of alphanumeric characters, starting with an alphabetic character or an underscore, such as `input1` or `next_token`.

As in Java, Unicode characters are allowed. For example, `quantité` or `πօօօ` are valid identifiers.

In addition, you can use *operator characters* in identifiers:

- The ASCII characters ! # % & * + - / : < = > ? @ \ ^ | ~ that are not letters, digits, underscore, the .;, punctuation marks, parentheses () [] {}, or quotation marks ' ` ".
- Unicode mathematical symbols or other symbols from the Unicode categories Sm and So.

For example, `**` and `√` are valid identifiers. With the definition

```
val √ = scala.math.sqrt _
```

you can write `√(2)` to compute a square root. This may be a good idea, provided one's programming environment makes it easy to type the symbol.



NOTE: The identifiers @ # : = _ => <- <: <% >: ⇒ ← are reserved in the specification, and you cannot redefine them.

You can also form identifiers from alphanumerical characters, followed by an underscore, and then a sequence of operator characters, such as

```
val happy_birthday_!!! = "Bonne anniversaire!!!"
```

This is probably not a good idea.

Finally, you can include just about any sequence of characters in backquotes. For example,

```
val `val` = 42
```

That example is silly, but backquotes can sometimes be an “escape hatch.” For example, in Scala, `yield` is a reserved word, but you may need to access a Java method of the same name. Backquotes to the rescue: `Thread.`yield`()`.

11.2 Infix Operators

You can write

`a identifier b`

where *identifier* denotes a method with two parameters (one implicit, one explicit). For example, the expression

`1 to 10`

is actually a method call

`1.to(10)`

This is called an *infix* expression because the operator is between the arguments. The operator can contain letters, as in `to`, or it can contain operator characters—for example,

`1 -> 10`

is a method call

`1 .->(10)`

To define an operator in your own class, simply define a method whose name is that of the desired operator. For example, here is a `Fraction` class that multiplies two fractions according to the law

$$(n_1 / d_1) \times (n_2 / d_2) = (n_1 n_2 / d_1 d_2)$$

```
class Fraction(n: Int, d: Int) {
    private val num = ...
    private val den = ...
    ...
    def *(other: Fraction) = new Fraction(num * other.num, den * other.den)
}
```

11.3 Unary Operators

Infix operators are binary operators—they have two parameters. An operator with one parameter is called a *unary operator*.

The four operators `+`, `-`, `!`, `~` are allowed as *prefix* operators, appearing before their arguments. They are converted into calls to methods with the name `unary_operator`. For example,

`-a`

means the same as `a.unary_-`.

If a unary operator follows its argument, it is a *postfix* operator. The expression

a *identifier*

is the same as the method call a.*identifier()*. For example,

42 `toString`

is the same as

`42.toString()`



CAUTION: Postfix operators can lead to parsing errors. For example, the code

```
val result = 42 toString
    println(result)
```

yields the error message “too many arguments for method `toString`”. Since parsing precedes type inference and overload resolution, the compiler does not yet know that `toString` is a unary method. Instead, the code is parsed as `val result = 42.toString(println(result))`.

For that reason, the compiler warns you if you use postfix operators. You can turn the warning off with the compiler option `-language:postfixOps`, or by adding the clause

```
import scala.language.postfixOps
```

11.4 Assignment Operators

An assignment operator has the form *operator=*, and the expression

a *operator=* b

means the same as

a = a *operator* b

For example, a `+=` b is equivalent to a = a + b.

There are a few technical details.

- `<=`, `>=`, and `!=` are not assignment operators.
- An operator starting with an = is never an assignment operator (`==`, `==>`, `=/=`, and so on).
- If a has a method called *operator=*, then that method is called directly.

11.5 Precedence

When you have two or more operators in a row without parentheses, the ones with higher *precedence* are executed first. For example, in the expression

$1 + 2 * 3$

the $*$ operator is evaluated first. Languages such as Java and C++ have a fixed number of operators, and the language standard decrees which have precedence over which. Scala can have arbitrary operators, so it uses a scheme that works for all operators, while also giving the familiar precedence order to the standard ones.

Except for assignment operators, the precedence is determined by the *first character* of the operator.

Highest precedence: An operator character other than those below

$* / %$

$+ -$

:

$< >$

$! =$

$\&$

$^$

|

A character that is not an operator character

Lowest precedence: Assignment operators

Characters in the same row yield operators with the same precedence. For example, $+$ and $-$ have the same precedence.

Postfix operators have lower precedence than infix operators:

a *infixOp* b *postfixOp*

is the same as

(a *infixOp* b) *postfixOp*

11.6 Associativity

When you have a sequence of operators of the same precedence, the *associativity* determines whether they are evaluated left-to-right or right-to-left. For example, in the expression $17 - 2 - 9$, one computes $(17 - 2) - 9$. The $-$ operator is *left-associative*.

In Scala, all operators are left-associative except for

- operators that end in a colon (:)
- assignment operators

In particular, the :: operator for constructing lists is right-associative. For example,

`1 :: 2 :: Nil`

means

`1 :: (2 :: Nil)`

This is as it should be—we first need to form the list containing 2, and that list becomes the tail of the list whose head is 1.

A right-associative binary operator is a method of its second argument. For example,

`2 :: Nil`

means

`Nil.::(2)`

11.7 The apply and update Methods

Scala lets you extend the function call syntax

`f(arg1, arg2, ...)`

to values other than functions. If f is not a function or method, then this expression is equivalent to the call

`f.apply(arg1, arg2, ...)`

unless it occurs to the left of an assignment. The expression

`f(arg1, arg2, ...) = value`

corresponds to the call

`f.update(arg1, arg2, ..., value)`

This mechanism is used in arrays and maps. For example,

```
val scores = new scala.collection.mutable.HashMap[String, Int]
scores("Bob") = 100 // Calls scores.update("Bob", 100)
val bobsScore = scores("Bob") // Calls scores.apply("Bob")
```

The apply method is also commonly used in companion objects to construct objects without calling new. For example, consider a Fraction class.

```
class Fraction(n: Int, d: Int) {
    ...
}

object Fraction {
    def apply(n: Int, d: Int) = new Fraction(n, d)
}
```

Because of the `apply` method, we can construct a fraction as `Fraction(3, 4)` instead of `new Fraction(3, 4)`. That sounds like a small thing, but if you have many `Fraction` values, it is a welcome improvement:

```
val result = Fraction(3, 4) * Fraction(2, 5)
```

11.8 Extractors L2

An extractor is an object with an `unapply` method. You can think of the `unapply` method as the opposite of the `apply` method of a companion object. An `apply` method takes construction parameters and turns them into an object. An `unapply` method takes an object and extracts values from it—usually the values from which the object was constructed.

Consider the `Fraction` class from the preceding section. The `apply` method makes a fraction from a numerator and denominator. An `unapply` method retrieves the numerator and denominator. You can use it in a variable definition

```
var Fraction(a, b) = Fraction(3, 4) * Fraction(2, 5)
// a, b are initialized with the numerator and denominator of the result
```

or a pattern match

```
case Fraction(a, b) => ... // a, b are bound to the numerator and denominator
```

(See Chapter 14 for more information about pattern matching.)

In general, a pattern match can fail. Therefore, the `unapply` method returns an `Option`. It contains a tuple with one value for each matched variable. In our case, we return an `Option[(Int, Int)]`.

```
object Fraction {
    def unapply(input: Fraction) =
        if (input.den == 0) None else Some((input.num, input.den))
}
```

Just to show the possibility, this method returns `None` when the denominator is zero, indicating no match.

A declaration

```
val Fraction(a, b) = f;
```

becomes

```
val tupleOption = Fraction.unapply(f)
if (tupleOption == None) throw new MatchError
// tupleOption is Some((t1, t2))
val a = t1
val b = t1
```



NOTE: Note that in the declaration

```
val Fraction(a, b) = f;
```

neither the `Fraction.apply` method nor the `Fraction` constructor are called. Instead, the statement means: “Initialize `a` and `b` so that if they would be passed to `Fraction.apply`, the result would be `f`.“

In the preceding example, the `apply` and `unapply` methods are inverses of one another. However, that is not a requirement. You can use extractors to extract information from an object of any type.

For example, suppose you want to extract first and last names from a string:

```
val author = "Cay Horstmann"
val Name(first, last) = author // Calls Name.unapply(author)
```

Provide an object `Name` with an `unapply` method that returns an `Option[(String, String)]`. If the match succeeds, return a pair with the first and last name. The components of the pair will be bound to the variables in the pattern. Otherwise, return `None`.

```
object Name {
    def unapply(input: String) = {
        val pos = input.indexOf(" ")
        if (pos == -1) None
        else Some((input.substring(0, pos), input.substring(pos + 1)))
    }
}
```



NOTE: In this example, there is no `Name` class. The `Name` object is an extractor for `String` objects.

Every case class automatically has `apply` and `unapply` methods. (Case classes are discussed in Chapter 14.) For example, consider

```
case class Currency(value: Double, unit: String)
```

You can construct a `Currency` instance as

```
Currency(29.95, "EUR") // Calls Currency.apply
```

You can extract values from a `Currency` object:

```
case Currency(amount, "USD") => println(s"$$amount") // Calls Currency.unapply
```

11.9 Extractors with One or No Arguments L2

In Scala, there are no tuples with one component. If the `unapply` method extracts a single value, it should just return an `Option` of the target type. For example,

```
object Number {
  def unapply(input: String): Option[Int] =
    try {
      Some(input.trim.toInt)
    } catch {
      case ex: NumberFormatException => None
    }
}
```

With this extractor, you can extract a number from a string:

```
val Number(n) = "1729"
```

An extractor can just test its input without extracting any value. In that case, the `unapply` method should return a `Boolean`. For example,

```
object IsCompound {
  def unapply(input: String) = input.contains(" ")
}
```

You can use this extractor to add a test to a pattern, for example

```
author match {
  case Name(first, IsCompound()) => ...
    // Matches if the last name is compound, such as van der Linden
  case Name(first, last) => ...
}
```

11.10 The unapplySeq Method L2

To extract an arbitrary sequence of values, the method needs to be called `unapplySeq`. It returns an `Option[Seq[A]]`, where `A` is the type of the extracted values. For example, a `Name` extractor can produce a sequence of the name's components:

```
object Name {
  def unapplySeq(input: String): Option[Seq[String]] =
    if (input.trim == "") None else Some(input.trim.split("\\s+"))
}
```

Now you can match for any number of variables:

```
author match {
    case Name(first, last) => ...
    case Name(first, middle, last) => ...
    case Name(first, "van", "der", last) => ...
    ...
}
```



CAUTION: Do not supply both an `unapply` and an `unapplySeq` methods with the same argument types.

11.11 Dynamic Invocation L2

Scala is a strongly typed language that reports type errors at compile time rather than at runtime. If you have an expression `x.f(args)`, and your program compiles, then you know for sure that `x` has a method `f` that can accept the given arguments. However, there are situations where it is desirable to define methods in a running program. This is common with object-relational mappers in dynamic languages such as Ruby or JavaScript. Objects that represent database tables have methods `findByName`, `findById`, and so on, with the method names matching the table columns. For database entities, the column names can be used to get and set fields, such as `person.lastName = "Doe"`.

In Scala, you can do this too. If a type extends the trait `scala.Dynamic`, then method calls, getters, and setters are rewritten as calls to special methods that can inspect the name of the original call and the parameters, and then take arbitrary actions.



NOTE: Dynamic types are an “exotic” feature, and the compiler wants your explicit consent when you implement such a type. You do that by adding the `import` statement

```
import scala.language.dynamics
```

Users of such types do not need to provide the `import` statement.

Here are the details of the rewriting. Consider `obj.name`, where `obj` belongs to a class that’s a subtype of `Dynamic`. Here is what the Scala compiler does with it.

1. If `name` is a known method or field of `obj`, it is processed in the usual way.
2. If `obj.name` is followed by `(arg1, arg2, ...)`,
 - a. If none of the arguments are named (of the form `name=arg`), pass the arguments on to `applyDynamic`:

```
obj.applyDynamic("name")(arg1, arg2, ...)
```

- b. If at least one of the arguments is named, pass the name/value pairs on to `applyDynamicNamed`:

```
obj.applyDynamicNamed("name")((name1, arg1), (name2, arg2), ...)
```

Here, `name1`, `name2`, and so on are strings with the argument names, or `" "` for unnamed arguments.

3. If `obj.name` is to the left of an `=`, call

```
obj.updateDynamic("name")(rightHandSide)
```

4. Otherwise call

```
obj.selectDynamic("sel")
```



NOTE: The calls to `updateDynamic`, `applyDynamic`, and `applyDynamicNamed` have two sets of parentheses, one for the selector name and one for the arguments. This construct is explained in Chapter 12.

Let's look at a few examples. Suppose `person` is an instance of a type extending `Dynamic`. A statement

```
person.lastName = "Doe"
```

is replaced with a call

```
person.updateDynamic("lastName")("Doe")
```

The `Person` class must have such a method:

```
class Person {  
  ...  
  def updateDynamic(field: String)(newValue: String) { ... }  
}
```

It is then up to you to implement the `updateDynamic` method. For example, if you are implementing an object-relational mapper, you might update the cached entity and mark it as changed, so that it can be persisted in the database.

Conversely, a statement

```
val name = person.lastName
```

turns into

```
val name = name.selectDynamic("lastName")
```

The `selectDynamic` method would simply look up the field value.

Method calls that don't involve named parameters are translated to `applyDynamic` calls. For example,

```
val does = people.findByLastName("Doe")
```

becomes

```
val does = people.applyDynamic("findByLastName")("Doe")
```

and

```
val johnDoes = people.find(lastName = "Doe", firstName = "John")
```

becomes

```
val johnDoes = people.applyDynamicNamed("find")
  ("lastName", "Doe"), ("firstName", "John"))
```

It is then up to you to implement `applyDynamic` and `applyDynamicNamed` as calls that retrieve the matching objects.

Here is a concrete example. Suppose we want to be able to dynamically look up and set elements of a `java.util.Properties` instance, using the dot notation:

```
val sysProps = new DynamicProps(System.getProperties)
sysProps.username = "Fred" // Sets the "username" property to "Fred"
val home = sysProps.java_home // Gets the "java.home" property
```

For simplicity, we replace periods in the property name with underscores. (Exercise 11 on page 154 shows how to keep the periods.)

The `DynamicProps` class extends the `Dynamic` trait and implements the `updateDynamic` and `selectDynamic` methods:

```
class DynamicProps(val props: java.util.Properties) extends Dynamic {
  def updateDynamic(name: String)(value: String) {
    props.setProperty(name.replaceAll("_", ".") , value)
  }
  def selectDynamic(name: String) =
    props.getProperty(name.replaceAll("_", "."))
}
```

As an additional enhancement, let us use the `add` method to add key/value pairs in bulk, using named arguments:

```
sysProps.add(username="Fred", password="Secret")
```

Then we need to supply the `applyDynamicNamed` method in the `DynamicProps` class. Note that the name of the method is fixed. We are only interested in arbitrary parameter names.

```
class DynamicProps(val props: java.util.Properties) extends Dynamic {
  ...
  def applyDynamicNamed(name: String)(args: (String, String)*) {
    if (name != "add") throw new IllegalArgumentException
    for ((k, v) <- args)
      props.setProperty(k.replaceAll("_", ".") , v)
  }
}
```

These examples are only meant to illustrate the mechanism—I don't think that it is a good idea to use the dot notation for map access. Like operator overloading, dynamic invocation is a feature that is best used with restraint.

Exercises

1. According to the precedence rules, how are $3 + 4 \rightarrow 5$ and $3 \rightarrow 4 + 5$ evaluated?
2. The `BigInt` class has a `pow` method, not an operator. Why didn't the Scala library designers choose `**` (as in Fortran) or `^` (as in Pascal) for a power operator?
3. Implement the `Fraction` class with operations `+` `-` `*` `/`. Normalize fractions, for example, turning $15/-6$ into $-5/2$. Divide by the greatest common divisor, like this:

```
class Fraction(n: Int, d: Int) {
  private val num: Int = if (d == 0) 1 else n * sign(d) / gcd(n, d);
  private val den: Int = if (d == 0) 0 else d * sign(d) / gcd(n, d);
  override def toString = s"$num/$den"
  def sign(a: Int) = if (a > 0) 1 else if (a < 0) -1 else 0
  def gcd(a: Int, b: Int): Int = if (b == 0) abs(a) else gcd(b, a % b)
  ...
}
```

4. Implement a class `Money` with fields for dollars and cents. Supply `+`, `-` operators as well as comparison operators `==` and `<`. For example, `Money(1, 75) + Money(0, 50) == Money(2, 25)` should be true. Should you also supply `*` and `/` operators? Why or why not?
5. Provide operators that construct an HTML table. For example,

```
Table() | "Java" | "Scala" || "Gosling" | "Odersky" || "JVM" | "JVM, .NET"
```

should produce

```
<table><tr><td>Java</td><td>Scala</td></tr><tr><td>Gosling...
```

6. Provide a class `ASCIIArt` whose objects contain figures such as

```
/\_\_\\
( ' ' )
( - - )
| | |
(_|_)
```

Supply operators for combining two `ASCIIArt` figures horizontally

```
/\_\_\\      -----
( ' ' ) / Hello \
( - - ) < Scala |
| | | \ Coder /
(_|_) -----
```

or vertically. Choose operators with appropriate precedence.

7. Implement a class `BitSequence` that stores a sequence of 64 bits packed in a `Long` value. Supply `apply` and `update` operators to get and set an individual bit.
8. Provide a class `Matrix`. Choose whether you want to implement 2×2 matrices, square matrices of any size, or $m \times n$ matrices. Supply operations `+` and `*`. The latter should also work with scalars, for example, `mat * 2`. A single element should be accessible as `mat(row, col)`.
9. Define an object `PathComponent`s with an `unapply` operation class that extracts the directory path and file name from an `java.nio.file.Path`. For example, the file `/home/cay/readme.txt` has directory path `/home/cay` and file name `readme.txt`.
10. Modify the `PathComponent`s object of the preceding exercise to instead define an `unapplySeq` operation that extracts all path segments. For example, for the file `/home/cay/readme.txt`, you should produce a sequence of three segments: `home`, `cay`, and `readme.txt`.
11. Improve the dynamic property selector in Section 11.11, “Dynamic Invocation,” on page 150 so that one doesn’t have to use underscores. For example, `sysProps.java.home` should select the property with key `"java.home"`. Use a helper class, also extending `Dynamic`, that contains partially completed paths.
12. Define a class `XMLElement` that models an XML element with a name, attributes, and child elements. Using dynamic selection and method calls, make it possible to select paths such as `rootElement.html.body.ul(id="42").li`, which should return all `li` elements inside `ul` with `id` attribute `42` inside `body` inside `html`.
13. Provide an `XMLBuilder` class for dynamically building XML elements, as `builder.ul(id="42", style="list-style: lower-alpha;")`, where the method name becomes the element name and the named arguments become the attributes. Come up with a convenient way of building nested elements.

This page intentionally left blank

Higher-Order Functions

Topics in This Chapter L1

- 12.1 Functions as Values — page 157
- 12.2 Anonymous Functions — page 159
- 12.3 Functions with Function Parameters — page 160
- 12.4 Parameter Inference — page 160
- 12.5 Useful Higher-Order Functions — page 161
- 12.6 Closures — page 162
- 12.7 SAM Conversions — page 163
- 12.8 Currying — page 164
- 12.9 Control Abstractions — page 166
- 12.10 The `return` Expression — page 167
- Exercises — page 168

Chapter 12

Scala mixes object orientation with functional features. In a functional programming language, functions are first-class citizens that can be passed around and manipulated just like any other data types. This is very useful whenever you want to pass some action detail to an algorithm. In a functional language, you just wrap that detail into a function that you pass as a parameter. In this chapter, you will see how to be productive with functions that use or return functions.

Highlights of the chapter include:

- Functions are “first-class citizens” in Scala, just like numbers.
- You can create anonymous functions, usually to give them to other functions.
- A function argument specifies behavior that should be executed later.
- Many collection methods take function parameters, applying a function to the values of the collection.
- There are syntax shortcuts that allow you to express function parameters in a way that is short and easy to read.
- You can create functions that operate on blocks of code and look much like the built-in control statements.

12.1 Functions as Values

In Scala, a function is a first-class citizen, just like a number. You can store a function in a variable:

```
import scala.math._  
val num = 3.14  
val fun = ceil _
```

This code sets `num` to 3.14 and `fun` to the `ceil` function.

The `_` behind the `ceil` function indicates that you really meant the function, and you didn't just forget to supply the arguments.

When you try this code in the REPL, the type of `num` is, not surprisingly, `Double`. The type of `fun` is reported as `(Double) => Double`—that is, a function receiving and returning a `Double`.



NOTE: Technically, the `_` turns the `ceil` *method* into a function. In Scala, you cannot manipulate methods, only functions. The type of the function is `(Double) => Double`, with an arrow. In contrast, the type of the `ceil` method is `(Double)Double`, without an arrow. There is no way for you to work with such a type, but you will find it in compiler and REPL messages.

The `_` suffix is not necessary when you use a method name in a context where a function is expected. For example, the following is legal:

```
val f: (Double) => Double = ceil // No underscore needed
```



NOTE: The `ceil` method is a method of the `scala.math` package object. If you have a method from a class, the syntax for turning it into a function is slightly different:

```
val f = (_: String).charAt(_: Int)  
// A function (String, Int) => Char
```

Alternatively, you can specify the type of the function instead of the parameter types:

```
val f: (String, Int) => Char = _.charAt(_)
```

What can you do with a function? Two things:

- Call it.
- Pass it around, by storing it in a variable or giving it to a function as a parameter.

Here is how to call the function stored in `fun`:

```
fun(num) // 4.0
```

As you can see, the normal function call syntax is used. The only difference is that `fun` is a *variable containing a function*, not a fixed function.

Here is how you can give `fun` to another function:

```
Array(3.14, 1.42, 2.0).map(fun) // Array(4.0, 2.0, 2.0)
```

The `map` method accepts a function, applies it to all values in an array, and returns an array with the function values. In this chapter, you will see many other methods that accept functions as parameters.

12.2 Anonymous Functions

In Scala, you don't have to give a name to each function, just like you don't have to give a name to each number. Here is an *anonymous function*:

```
(x: Double) => 3 * x
```

This function multiplies its argument by 3.

Of course, you can store this function in a variable:

```
val triple = (x: Double) => 3 * x
```

That's just as if you had used a `def`:

```
def triple(x: Double) = 3 * x
```

But you don't have to name the function. You can just pass it to another function:

```
Array(3.14, 1.42, 2.0).map((x: Double) => 3 * x)
// Array(9.42, 4.26, 6.0)
```

Here, we tell the `map` method: "Multiply each element by 3."



NOTE: If you prefer, you can enclose the function argument in braces instead of parentheses, for example:

```
Array(3.14, 1.42, 2.0).map{ (x: Double) => 3 * x }
```

This is more common when a method is used in infix notation (without the dot).

```
Array(3.14, 1.42, 2.0) map { (x: Double) => 3 * x }
```



NOTE: Anything defined with `def` (in the REPL or a class or object) is a method, not a function:

```
scala> def triple(x: Double) = 3 * x
triple: (x: Double)Double
```

Note the method type `(x: Double)Double`. In contrast, a function definition has a function type:

```
scala> val triple = (x: Double) => 3 * x
triple: Double => Double
```

12.3 Functions with Function Parameters

In this section, you will see how to implement a function that takes another function as a parameter. Here is an example:

```
def valueAtOneQuarter(f: (Double) => Double) = f(0.25)
```

Note that the parameter can be *any* function receiving and returning a `Double`. The `valueAtOneQuarter` function computes the value of that function at 0.25 .

For example,

```
valueAtOneQuarter(ceil _) // 1.0
valueAtOneQuarter(sqrt _) // 0.5 (because  $0.5 \times 0.5 = 0.25$ )
```

What is the type of `valueAtOneQuarter`? It is a function with one parameter, so its type is written as

$$(\text{parameterType}) \Rightarrow \text{resultType}$$

The `resultType` is clearly `Double`, and the `parameterType` is already given in the function header as `(Double) => Double`. Therefore, the type of `valueAtOneQuarter` is

$$((\text{Double}) \Rightarrow \text{Double}) \Rightarrow \text{Double}$$

Since `valueAtOneQuarter` is a function that receives a function, it is called a *higher-order function*.

A higher-order function can also *produce a function*. Here is a simple example:

```
def mulBy(factor : Double) = (x : Double) => factor * x
```

For example, `mulBy(3)` returns the function `(x : Double) => 3 * x` which you have seen in the preceding section. The power of `mulBy` is that it can deliver functions that multiply by any amount:

```
val quintuple = mulBy(5)
quintuple(20) // 100
```

The `mulBy` function has a parameter of type `Double`, and it returns a function of type `(Double) => Double`. Therefore, its type is

$$(\text{Double}) \Rightarrow ((\text{Double}) \Rightarrow \text{Double})$$

12.4 Parameter Inference

When you pass an anonymous function to another function or method, Scala helps you out by deducing types when possible. For example, you don't have to write

```
valueAtOneQuarter((x: Double) => 3 * x) // 0.75
```

Since the `valueAtOneQuarter` method knows that you will pass in a `(Double) => Double` function, you can just write

```
valueAtOneQuarter((x) => 3 * x)
```

As a special bonus, for a function that has just one parameter, you can omit the `()` around the parameter:

```
valueAtOneQuarter(x => 3 * x)
```

It gets better. If a parameter occurs only once on the right-hand side of the `=>`, you can replace it with an underscore:

```
valueAtOneQuarter(3 * _)
```

This is the ultimate in comfort, and it is also pretty easy to read: a function that multiplies something by 3.

Keep in mind that these shortcuts only work when the parameter types are known.

```
val fun = 3 * _ // Error: Can't infer types
val fun = 3 * (_: Double) // OK
val fun: (Double) => Double = 3 * _ // OK because we specified the type for fun
```

Of course, the last definition is contrived. But it shows what happens when a function is passed to a parameter (which has just such a type).



NOTE: Specifying the type of `_` is useful for turning methods into functions.

For example, `(_: String).length` is a function `String => Int`, and

`(_: String).substring(_:_Int, _: Int)` is a function `(String, Int, Int) => String`.

12.5 Useful Higher-Order Functions

A good way of becoming comfortable with higher-order functions is to practice with some common (and obviously useful) methods in the Scala collections library that take function parameters.

You have seen `map`, which applies a function to all elements of a collection and returns the result. Here is a quick way of producing a collection containing `0.1, 0.2, ..., 0.9`:

```
(1 to 9).map(0.1 * _)
```



NOTE: There is a general principle at work. If you want a sequence of values, see if you can transform it from a simpler one.

Try this to print a triangle:

```
(1 to 9).map("*" * _).foreach(println _)
```

The result is

```
*  
**  
***  
****  
*****  
*****  
*****  
*****  
*****
```

Here, we also use `foreach`, which is like `map` except that its function doesn't return a value. The `foreach` method simply applies the function to each argument.

The `filter` method yields all elements that match a particular condition. For example, here's how to get only the even numbers in a sequence:

```
(1 to 9).filter(_ % 2 == 0) // 2, 4, 6, 8
```

Of course, that's not the most efficient way of getting this result ☺.

The `reduceLeft` method takes a *binary* function—that is, a function with two parameters—and applies it to all elements of a sequence, going from left to right. For example,

```
(1 to 9).reduceLeft(_ * _)
```

is

```
1 * 2 * 3 * 4 * 5 * 6 * 7 * 8 * 9
```

or, strictly speaking,

```
((1 * 2) * 3) * ... * 9
```

Note the compact form of the multiplication function: `_ * _`. Each underscore denotes a separate parameter.

You also need a binary function for sorting. For example,

```
"Mary had a little lamb".split(" ").sortWith(_.length < _.length)
```

yields an array that is sorted by increasing length: `Array("a", "had", "Mary", "lamb", "little")`.

12.6 Closures

In Scala, you can define a function inside any scope: in a package, in a class, or even inside another function or method. In the body of a function, you can access

any variables from an enclosing scope. That may not sound so remarkable, but note that your function may be called when the variable is *no longer in scope*.

Here is an example: the `mulBy` function from Section 12.3, “Functions with Function Parameters,” on page 160.

```
def mulBy(factor : Double) = (x : Double) => factor * x
```

Consider these calls:

```
val triple = mulBy(3)
val half = mulBy(0.5)
println(s"${triple(14)} ${half(14)})" // Prints 42 7
```

Let’s look at them in slow motion.

1. The first call to `mulBy` sets the parameter variable `factor` to 3. That variable is referenced in the body of the function `(x : Double) => factor * x`, which is stored in `triple`. Then the parameter variable `factor` is popped off the runtime stack.
2. Next, `mulBy` is called again, now with `factor` set to 0.5. That variable is referenced in the body of the function `(x : Double) => factor * x`, which is stored in `half`.

Each of the returned functions has its own setting for `factor`.

Such a function is called a *closure*. A closure consists of code together with the definitions of any nonlocal variables that the code uses.

These functions are actually implemented as objects of a class, with an instance variable `factor` and an `apply` method that contains the body of the function.

It doesn’t really matter how a closure is implemented. It is the job of the Scala compiler to ensure that your functions can access nonlocal variables.



NOTE: Closures aren’t difficult or surprising if they are a natural part of the language. Many modern languages, such as JavaScript, Ruby, and Python, support closures. Java, as of version 8, has closures in the form of lambda expressions.

12.7 SAM Conversions

In Scala, you pass a function as a parameter whenever you want to tell another function what action to carry out. Prior to Java 8, this was not possible in Java without defining a class and a method for the action. For example, to implement a button callback, one had to use this Scala code (or its Java equivalent):

```

var counter = 0

val button = new JButton("Increment")
button.addActionListener(new ActionListener {
    override def actionPerformed(event: ActionEvent) {
        counter += 1
    }
})

```

In Java 8, it is possible to specify such actions with lambda expressions, which are closely related to functions in Scala. Fortunately, that means that as of Scala 2.12, one can pass Scala functions to Java code expecting a “SAM interface”—that is, any Java interface with a single **abstract method**. (Such interfaces are officially called *functional interfaces* in Java.)

Simply pass the function to the `addActionListener` method, like this:

```
button.addActionListener(event => counter += 1)
```

Note that the conversion from a Scala function to a Java SAM interface only works for *function literals*, not for variables holding functions. The following does not work:

```

val listener = (event: ActionListener) => println(counter)
button.addActionListener(listener)
// Cannot convert a nonliteral function to a Java SAM interface

```

The simplest remedy is to declare the variable holding the function as a Java SAM interface:

```

val listener: ActionListener = event => println(counter)
button.addActionListener(listener) // Ok

```

Alternatively, you can turn a function variable into a literal expression:

```

val exit = (event: ActionEvent) => if (counter > 9) System.exit(0)
button.addActionListener(exit(_))

```

12.8 Currying

Currying (named after logician Haskell Brooks Curry) is the process of turning a function that takes two arguments into a function that takes one argument. That function returns a function that consumes the second argument.

Huh? Let’s look at an example. This function takes two arguments:

```
val mul = (x: Int, y: Int) => x * y
```

This function takes one argument, yielding a function that takes one argument:

```
val mulOneAtATime = (x: Int) => ((y: Int) => x * y)
```

To multiply two numbers, you call

```
mulOneAtATime(6)(7)
```

Strictly speaking, the result of `mulOneAtATime(6)` is the function `(y: Int) => 6 * y`. That function is applied to 7, yielding 42.

When you use `def`, there is a shortcut for defining such curried methods in Scala:

```
def mulOneAtATime(x: Int)(y: Int) = x * y
```



NOTE: Recall that anything defined with `def` (in the REPL or a class or object) is a method, not a function. When defining curried methods with `def`, you can use multiple parentheses:

```
scala> def mulOneAtATime(x: Int)(y: Int) = x * y
mulOneAtATime: (x: Int)(y: Int)Int
```

Note the method type `(x: Int)(y: Int)Int`. In contrast, when you define a function, you must use multiple arrows, not multiple parentheses:

```
scala> val mulOneAtATime = (x: Int) => (y: Int) => x * y
mulOneAtATime: Int => (Int => Int)
```

As you can see, multiple parameters are just a frill, not an essential feature of a programming language. That's an amusing theoretical insight, but it has one practical use in Scala. Sometimes, you can use currying for a method parameter so that the type inferencer has more information.

Here is a typical example. The `corresponds` method can compare whether two sequences are the same under some comparison criterion. For example,

```
val a = Array("Hello", "World")
val b = Array("hello", "world")
a.corresponds(b)(_.equalsIgnoreCase(_))
```

Note that the function `_.equalsIgnoreCase(_)` is passed as a curried parameter, in a separate set of `(...)`. When you look into the Scaladoc, you will see that `corresponds` is declared as

```
def corresponds[B](that: Seq[B])(p: (A, B) => Boolean): Boolean
```

The `that` sequence and the predicate function `p` are separate curried parameters. The type inferencer can figure out what `B` is from the type of `that`, and then it can use that information when analyzing the function `p` that is passed for `p`.

In our example, `that` is a `String` sequence. Therefore, the predicate is expected to have type `(String, String) => Boolean`. With that information, the compiler can accept `_.equalsIgnoreCase(_)` as a shortcut for `(a: String, b: String) => a.equalsIgnoreCase(b)`.

12.9 Control Abstractions

In Scala, one can model a sequence of statements as a function with no parameters or return value. For example, here is a function that runs some code in a thread:

```
def runInThread(block: () => Unit) {
    new Thread {
        override def run() { block() }
    }.start()
}
```

The code is given as a function of type `() => Unit`. However, when you call this function, you need to supply an unsightly `() =>`:

```
runInThread { () => println("Hi"); Thread.sleep(10000); println("Bye") }
```

To avoid the `() =>` in the call, use the *call by name* notation: Omit the `()`, but not the `=>`, in the parameter declaration and in the call to the parameter function:

```
def runInThread(block: => Unit) {
    new Thread {
        override def run() { block }
    }.start()
}
```

Then the call becomes simply

```
runInThread { println("Hi"); Thread.sleep(10000); println("Bye") }
```

This looks pretty nice. Scala programmers can build *control abstractions*: functions that look like language keywords. For example, we can implement a function that is used *exactly* as a `while` statement. Or, we can innovate a bit and define an `until` statement that works like `while`, but with an inverted condition:

```
def until(condition: => Boolean)(block: => Unit) {
    if (!condition) {
        block
        until(condition)(block)
    }
}
```

Here is how you use `until`:

```
var x = 10
until (x == 0) {
    x -= 1
    println(x)
}
```

The technical term for such a function parameter is a *call-by-name* parameter. Unlike a regular (or call-by-value) parameter, the parameter expression is *not* evaluated when the function is called. After all, we don't want `x == 0` to evaluate to false in the call to `until`. Instead, the expression becomes the body of a function with no arguments. That function is passed as a parameter.

Look carefully at the `until` function definition. Note that it is curried: It first consumes the condition, then the block as a second parameter. Without currying, the call would look like this:

```
until(x == 0, { ... })
```

which wouldn't be as pretty.

12.10 The return Expression

In Scala, you don't use a `return` statement to return function values. The return value of a function is simply the value of the function body.

However, you can use `return` to return a value from an anonymous function to an enclosing named function. This is useful in control abstractions. For example, consider this function:

```
def indexOf(str: String, ch: Char): Int = {
    var i = 0
    until (i == str.length) {
        if (str(i) == ch) return i
        i += 1
    }
    return -1
}
```

Here, the anonymous function `{ if (str(i) == ch) return i; i += 1 }` is passed to `until`. When the `return` expression is executed, the enclosing named function `indexOf` terminates and returns the given value.

If you use `return` inside a named function, you need to specify its return type. For example, in the `indexOf` function above, the compiler was not able to infer that it returns an `Int`.

The control flow is achieved with a special exception that is thrown by the `return` expression in the anonymous function, passed out of the `until` function, and caught in the `indexOf` function.



CAUTION: If the exception is caught in a `try` block, before it is delivered to the named function, then the value will not be returned.

Exercises

1. Write a function `values(fun: (Int) => Int, low: Int, high: Int)` that yields a collection of function inputs and outputs in a given range. For example, `values(x => x * x, -5, 5)` should produce a collection of pairs $(-5, 25), (-4, 16), (-3, 9), \dots, (5, 25)$.
2. How do you get the largest element of an array with `reduceLeft`?
3. Implement the factorial function using `to` and `reduceLeft`, without a loop or recursion.
4. The previous implementation needed a special case when $n < 1$. Show how you can avoid this with `foldLeft`. (Look at the Scaladoc for `foldLeft`. It's like `reduceLeft`, except that the first value in the chain of combined values is supplied in the call.)
5. Write a function `largest(fun: (Int) => Int, inputs: Seq[Int])` that yields the largest value of a function within a given sequence of inputs. For example, `largest(x => 10 * x - x * x, 1 to 10)` should return 25. Don't use a loop or recursion.
6. Modify the previous function to return the *input* at which the output is largest. For example, `largestAt(x => 10 * x - x * x, 1 to 10)` should return 5. Don't use a loop or recursion.
7. It's easy to get a sequence of pairs, for example:

```
val pairs = (1 to 10) zip (11 to 20)
```

Now, suppose you want to do something with such a sequence—say, add up the values. But you can't do

```
pairs.map(_ + _)
```

The function `_ + _` takes two `Int` parameters, not an `(Int, Int)` pair. Write a function `adjustToPair` that receives a function of type `(Int, Int) => Int` and returns the equivalent function that operates on a pair. For example, `adjustToPair(_ * _)((6, 7))` is 42.

Then use this function in conjunction with `map` to compute the sums of the elements in `pairs`.

8. In Section 12.8, "Currying," on page 164, you saw the `corresponds` method used with two arrays of strings. Make a call to `corresponds` that checks whether the elements in an array of strings have the lengths given in an array of integers.
9. Implement `corresponds` without currying. Then try the call from the preceding exercise. What problem do you encounter?
10. Implement an `unless` control abstraction that works just like `if`, but with an inverted condition. Does the first parameter need to be a call-by-name parameter? Do you need currying?

This page intentionally left blank

Collections

Topics in This Chapter **A2**

- 13.1 The Main Collections Traits — page 172
- 13.2 Mutable and Immutable Collections — page 173
- 13.3 Sequences — page 174
- 13.4 Lists — page 175
- 13.5 Sets — page 177
- 13.6 Operators for Adding or Removing Elements — page 178
- 13.7 Common Methods — page 180
- 13.8 Mapping a Function — page 182
- 13.9 Reducing, Folding, and Scanning **A3** — page 184
- 13.10 Zipping — page 187
- 13.11 Iterators — page 188
- 13.12 Streams **A3** — page 189
- 13.13 Lazy Views **A3** — page 190
- 13.14 Interoperability with Java Collections — page 191
- 13.15 Parallel Collections — page 193
- Exercises — page 194

Chapter 13

In this chapter, you will learn about the Scala collections library from a library user's point of view. In addition to arrays and maps, which you have already encountered, you will see other useful collection types. There are many methods that can be applied to collections, and this chapter presents them in an orderly way.

The key points of this chapter are:

- All collections extend the `Iterable` trait.
- The three major categories of collections are sequences, sets, and maps.
- Scala has mutable and immutable versions of most collections.
- A Scala list is either empty, or it has a head and a tail which is again a list.
- Sets are unordered collections.
- Use a `LinkedHashSet` to retain the insertion order or a `SortedSet` to iterate in sorted order.
- `+ adds an element to an unordered collection; +: and ::+ prepend or append to a sequence; ++ concatenates two collections; - and -- remove elements.`
- The `Iterable` and `Seq` traits have dozens of useful methods for common operations. Check them out before writing tedious loops.
- Mapping, folding, and zipping are useful techniques for applying a function or operation to the elements of a collection.

13.1 The Main Collections Traits

Figure 13–1 shows the most important traits that make up the Scala collections hierarchy.

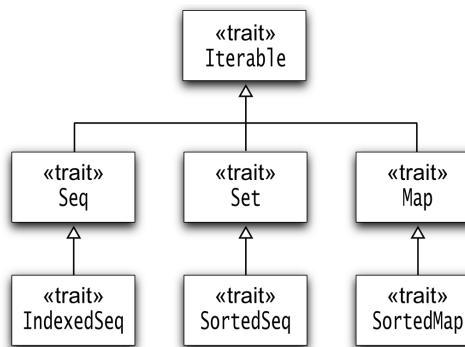


Figure 13–1 Key traits in the Scala collections hierarchy

An **Iterable** is any collection that can yield an **Iterator** with which you can access all elements in the collection:

```

val coll = ... // some Iterable
val iter = coll.iterator
while (iter.hasNext)
    do something with iter.next()
  
```

This is the most basic way of traversing a collection. However, as you will see throughout this chapter, usually there are more convenient ways.

A **Seq** is an ordered sequence of values, such as an array or list. An **IndexedSeq** allows fast random access through an integer index. For example, an **ArrayBuffer** is indexed but a linked list is not.

A **Set** is an unordered collection of values. In a **SortedSet**, elements are always visited in sorted order.

A **Map** is a set of *(key, value)* pairs. A **SortedMap** visits the entries as sorted by the keys. See Chapter 4 for more information.

This hierarchy is similar to that in Java, with a couple of welcome improvements:

1. Maps are a part of the hierarchy and not a separate hierarchy.
2. **IndexedSeq** is the supertype of arrays but not of lists, allowing you to tell the two apart.



NOTE: In Java, both `ArrayList` and `LinkedList` implement a common `List` interface, making it difficult to write efficient code when random access is preferred, for example when searching in a sorted sequence. This was a flawed design decision in the original Java collections framework. In a later version, a marker interface `RandomAccess` was added to deal with this problem.

Each Scala collection trait or class has a companion object with an `apply` method for constructing an instance of the collection. For example,

```
Iterable(0xFF, 0xFF00, 0xFF0000)
Set(Color.RED, Color.GREEN, Color.BLUE)
Map(Color.RED -> 0xFF0000, Color.GREEN -> 0xFF00, Color.BLUE -> 0xFF)
SortedSet("Hello", "World")
```

This is called the “uniform creation principle.”

There are methods `toSeq`, `toSet`, `toMap`, and so on, as well as a generic `to[C]` method, that you can use to translate between collection types.

```
val coll = Seq(1, 1, 2, 3, 5, 8, 13)
val set = coll.toSet
val buffer = coll.to[ArrayBuffer]
```



NOTE: You can use the `==` operator to compare any sequence, set, or map with another collection of the same kind. For example, `Seq(1, 2, 3) == (1 to 3)` yields true. But comparing different kinds, for example, `Seq(1, 2, 3) == Set(1, 2, 3)` always yields false. In that case, use the `sameElements` method.

13.2 Mutable and Immutable Collections

Scala supports both mutable and immutable collections. An immutable collection can never change, so you can safely share a reference to it, even in a multi-threaded program. For example, there is a `scala.collection.mutable.Map` and a `scala.collection.immutable.Map`. Both have a common supertype `scala.collection.Map` (which, of course, contains no mutation operations).



NOTE: When you have a reference to a `scala.collection.immutable.Map`, you know that *nobody* can change the map. If you have a `scala.collection.Map`, then *you* can't change it, but someone else might.

Scala gives a preference to immutable collections. The companion objects in the `scala.collection` package produce immutable collections. For example, `scala.collection.Map("Hello" -> 42)` is an immutable map.

Moreover, the `scala` package and the `Predef` object, which are always imported, have type aliases `List`, `Set`, and `Map` that refer to the immutable traits. For example, `Predef.Map` is the same as `scala.collection.immutable.Map`.



TIP: With the statement

```
import scala.collection.mutable
```

you can get an immutable map as `Map` and a mutable one as `mutable.Map`.

If you had no prior experience with immutable collections, you may wonder how you can do useful work with them. The key is that you can create new collections out of old ones. For example, if `numbers` is an immutable set, then

```
numbers + 9
```

is a new set containing the `numbers` together with 9. If 9 was already in the set, you just get a reference to the old set. This is particularly natural in recursive computations. For example, here we compute the set of all digits of an integer:

```
def digits(n: Int): Set[Int] =
  if (n < 0) digits(-n)
  else if (n < 10) Set(n)
  else digits(n / 10) + (n % 10)
```

This method starts out with a set containing a single digit. At each step, another digit is added. However, adding the digit doesn't mutate a set. Instead, in each step, a new set is constructed.

13.3 Sequences

Figure 13–2 shows the most important immutable sequences.

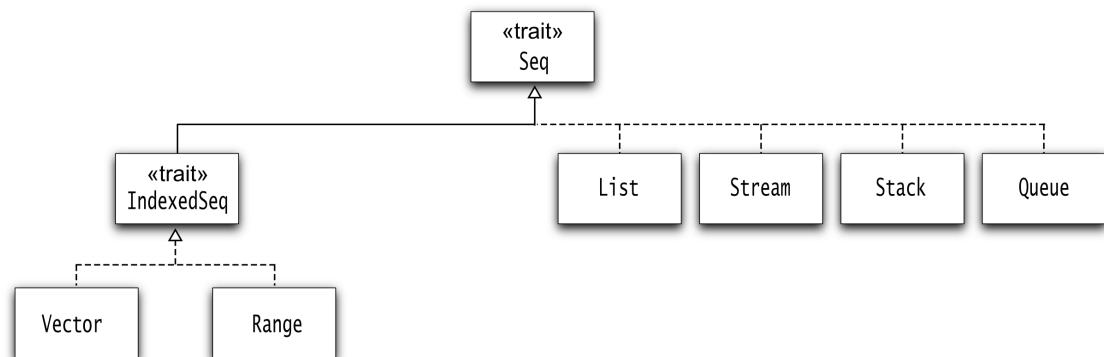


Figure 13–2 Immutable sequences

A Vector is the immutable equivalent of an ArrayBuffer: an indexed sequence with fast random access. Vectors are implemented as trees where each node has up to 32 children. For a vector with one million elements, one needs four layers of nodes. (Since $10^3 \approx 2^{10}$, $10^6 \approx 32^4$.) Accessing an element in such a list will take 4 hops, whereas in a linked list it would take an average of 500,000.

A Range represents an integer sequence, such as 0,1,2,3,4,5,6,7,8,9 or 10,20,30. Of course a Range object doesn't store all sequence values but only the start, end, and increment. You construct Range objects with the `to` and `until` methods, as described in Chapter 2.

We discuss lists in the next section, and streams in Section 13.12, “Streams,” on page 189.

See Figure 13–3 for the most useful mutable sequences.

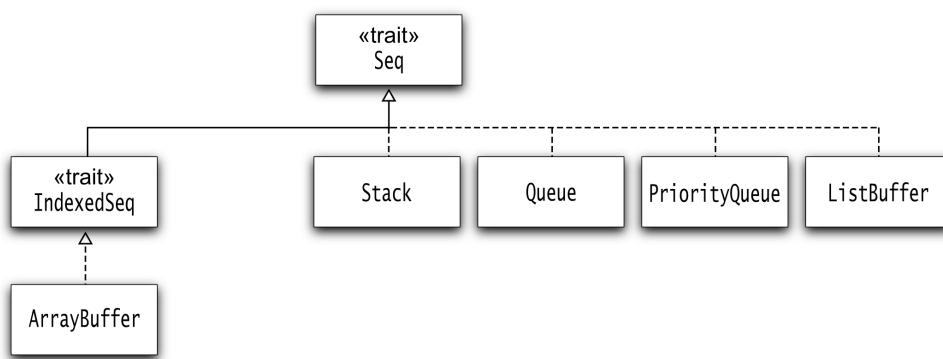


Figure 13–3 Mutable sequences

We discussed array buffers in Chapter 3. Stacks, queues, and priority queues are standard data structures that are useful for implementing certain algorithms. If you are familiar with these structures, the Scala implementations won't surprise you.

13.4 Lists

In Scala, a list is either `Nil` (that is, empty) or an object with a `head` element and a `tail` that is again a list. For example, consider the list

```
val digits = List(4, 2)
```

The value of `digits.head` is 4, and `digits.tail` is `List(2)`. Moreover, `digits.tail.head` is 2 and `digits.tail.tail` is `Nil`.

The `::` operator makes a new list from a given head and tail. For example,

```
9 :: List(4, 2)
```

is `List(9, 4, 2)`. You can also write that list as

`9 :: 4 :: 2 :: Nil`

Note that `::` is right-associative. With the `::` operator, lists are constructed from the end:

`9 :: (4 :: (2 :: Nil))`

In Java or C++, one uses an iterator to traverse a linked list. You can do this in Scala as well, but it is often more natural to use recursion. For example, the following function computes the sum of all elements in a linked list of integers:

```
def sum(lst: List[Int]): Int =  
  if (lst == Nil) 0 else lst.head + sum(lst.tail)
```

Or, if you prefer, you can use pattern matching:

```
def sum(lst: List[Int]): Int = lst match {  
  case Nil => 0  
  case h :: t => h + sum(t) // h is lst.head, t is lst.tail  
}
```

Note the `::` operator in the second pattern. It “destructures” the list into head and tail.



NOTE: Recursion works so naturally because the tail of a list is again a list.

Of course, for this particular example, you do not need to use recursion at all. The Scala library already has a `sum` method:

`List(9, 4, 2).sum // Yields 15`

If you want to mutate list elements in place, you can use a `ListBuffer`, a data structure that is backed by a linked list with a reference to the last node. This makes it efficient to add or remove elements at either end of the list.

However, adding or removing elements in the middle is not efficient. For example, suppose you want to remove every second element of a mutable list. With a Java `LinkedList`, you use an iterator and call `remove` after every second call to `next`. There is no analogous operation on a `ListBuffer`. Of course, removing multiple elements by their index positions is very inefficient in a linked list. Your best bet is to generate a new list with the result (see Exercise 3 on page 194).



NOTE: There are deprecated `LinkedList` and `DoubleLinkedList` classes and an internal `MutableList` class that you should not be using.

13.5 Sets

A set is a collection of distinct elements. Trying to add an existing element has no effect. For example,

```
Set(2, 0, 1) + 1
```

is the same as `Set(2, 0, 1)`.

Unlike lists, sets do not retain the order in which elements are inserted. By default, sets are implemented as *hash sets* in which elements are organized by the value of the `hashCode` method. (In Scala, as in Java, every object has a `hashCode` method.)

For example, if you iterate over

```
Set(1, 2, 3, 4, 5, 6)
```

the elements are visited in the order

```
5 1 6 2 3 4
```

You may wonder why sets don't retain the element order. It turns out that you can find elements much faster if you allow sets to reorder their elements. Finding an element in a hash set is *much* faster than in an array or list.

A *linked hash set* remembers the order in which elements were inserted. It keeps a linked list for this purpose. For example,

```
val weekdays = scala.collection.mutable.LinkedHashSet("Mo", "Tu", "We", "Th", "Fr")
```

If you want to iterate over elements in sorted order, use a *sorted set*:

```
val numbers = scala.collection.mutable.SortedSet(1, 2, 3, 4, 5)
```

A *bit set* is an implementation of a set of non-negative integers as a sequence of bits. The i th bit is 1 if i is present in the set. This is an efficient implementation as long as the maximum element is not too large. Scala provides both mutable and immutable `BitSet` classes.

The `contains` method checks whether a set contains a given value. The `subsetOf` method checks whether all elements of a set are contained in another set.

```
val digits = Set(1, 7, 2, 9)
digits contains 0 // false
Set(1, 2) subsetOf digits // true
```

The `union`, `intersect`, and `diff` methods carry out the usual set operations. If you prefer, you can write them as `|`, `&`, and `&~`. You can also write `union` as `++` and `difference` as `--`. For example, if we have the set

```
val primes = Set(2, 3, 5, 7)
```

then `digits union primes` is `Set(1, 2, 3, 5, 7, 9)`, `digits & primes` is `Set(2, 7)`, and `digits -- primes` is `Set(1, 9)`.

13.6 Operators for Adding or Removing Elements

When you want to add or remove an element, or a number of elements, the operators to use depend on the collection type. Table 13–1 provides a summary.

Table 13–1 Operators for Adding and Removing Elements

Operator	Description	Collection Type
<code>coll(k)</code> (i.e., <code>coll.apply(k)</code>)	The kth sequence element or the map value for key k.	Seq, Map
<code>coll ++ elem</code> <code>elem ++: coll</code>	A collection of the same type as <code>coll</code> to which <code>elem</code> has been appended or prepended.	Seq
<code>coll + elem</code> <code>coll + (e1, e2, ...)</code>	A collection of the same type as <code>coll</code> to which the given elements have been added.	Set, Map
<code>coll - elem</code> <code>coll - (e1, e2, ...)</code>	A collection of the same type as <code>coll</code> from which the given elements have been removed.	Set, Map, ArrayBuffer
<code>coll ++ coll2</code> <code>coll2 ++: coll</code>	A collection of the same type as <code>coll</code> , containing the elements of both collections.	Iterable
<code>coll -- coll2</code>	A collection of the same type as <code>coll</code> from which the elements of <code>coll2</code> have been removed. (For sequences, use <code>diff</code> .)	Set, Map, ArrayBuffer
<code>elem :: lst</code> <code>lst2 :::: lst</code>	A list with the element or given list prepended to <code>lst</code> . Same as <code>:+</code> and <code>++:</code> .	List
<code>list ::: list2</code>	Same as <code>list ++: list2</code> .	List
<code>set set2</code> <code>set & set2</code> <code>set &~ set2</code>	Set union, intersection, difference. <code> </code> is the same as <code>++</code> , and <code>&~</code> is the same as <code>--</code> .	Set
<code>coll += elem</code> <code>coll += (e1, e2, ...)</code> <code>coll ++= coll2</code> <code>coll -= elem</code> <code>coll -= (e1, e2, ...)</code> <code>coll --- coll2</code>	Modifies <code>coll</code> by adding or removing the given elements.	Mutable collections
<code>elem +=: coll</code> <code>coll2 ++=: coll</code>	Modifies <code>coll</code> by prepending the given element or collection.	ArrayBuffer

Generally, `+` is used for adding an element to an unordered collection, while `+:` and `:+` add an element to the beginning or end of an ordered collection.

```
Vector(1, 2, 3) :+ 5 // Yields Vector(1, 2, 3, 5)
1 +: Vector(1, 2, 3) // Yields Vector(1, 1, 2, 3)
```

Note that `+:`, like all operators ending in a colon, is right-associative, and that it is a method of the right operand.

These operators return new collections (of the same type as the original ones) without modifying the original. Mutable collections have a `+=` operator that mutates the left-hand side. For example,

```
val numbers = ArrayBuffer(1, 2, 3)
numbers += 5 // Adds 5 to numbers
```

With an immutable collection, you can use `+=` or `:+=` with a `var`, like this:

```
var numbers = Set(1, 2, 3)
numbers += 5 // Sets numbers to the immutable set numbers + 5
var numberVector = Vector(1, 2, 3)
numberVector :+= 5 // += does not work since vectors don't have a + operator
```

To remove an element, use the `-` operator:

```
Set(1, 2, 3) - 2 // Yields Set(1, 3)
```

You can add multiple elements with the `++` operator:

```
coll ++ coll2
```

yields a collection of the same type as `coll` that contains both `coll` and `coll2`. Similarly, the `--` operator removes multiple elements.



TIP: As you can see, Scala provides many operators for adding and removing elements. Here is a summary:

1. Append (`:+`) or prepend (`+:`) to a sequence.
2. Add (`+`) to an unordered collection.
3. Remove with the `-` operator.
4. Use `++` and `--` for bulk add and remove.
5. Mutations are `+=` `++=` `-=` `--=`.
6. For lists, many Scala programmers prefer the `::` and `:::` operators.
7. Stay away from `++:` `+=:` `++=:`.



NOTE: For lists, you can use `+` instead of `::` for consistency, with one exception: Pattern matching (case `h :: t`) does *not* work with the `+` operator.

13.7 Common Methods

Table 13–2 gives a brief overview of the most important methods of the `Iterable` trait, sorted by functionality.

Table 13–2 Important Methods of the `Iterable` Trait

Methods	Description
<code>head, last, headOption, lastOption</code>	Returns the first or last element; or, that element as an <code>Option</code> .
<code>tail, init</code>	Returns everything but the first or last element.
<code>length, isEmpty</code>	Returns the length, or true if the length is zero.
<code>map(f), flatMap(f), foreach(f), transform(f).collect(pf)</code>	Applies a function to all elements; see Section 13.8.
<code>reduceLeft(op), reduceRight(op), foldLeft(init)(op), foldRight(init)(op)</code>	Applies a binary operation to all elements in a given order; see Section 13.9.
<code>reduce(op), fold(init)(op), aggregate(init)(op, combineOp)</code>	Applies a binary operation to all elements in arbitrary order; see Section 13.15.
<code>sum, product, max, min</code>	Returns the sum or product (provided the element type can be implicitly converted to the <code>Numeric</code> trait), or the maximum or minimum (provided the element type can be converted to the <code>Ordered</code> trait).
<code>count(pred), forall(pred), exists(pred)</code>	Returns the count of elements fulfilling the predicate; true if all elements do, or at least one element does.
<code>filter(pred), filterNot(pred), partition(pred)</code>	Returns all elements fulfilling or not fulfilling the predicate; the pair of both.
<code>takeWhile(pred), dropWhile(pred), span(pred)</code>	Returns the first elements fulfilling pred; all but those elements; the pair of both.
<code>take(n), drop(n), splitAt(n)</code>	Returns the first n elements; everything but the first n elements; the pair of both.
<code>takeRight(n), dropRight(n)</code>	Returns the last n elements; everything but the last n elements.
<code>slice(from, to), view(from, to)</code>	Returns the elements in the range from until to, or a view thereto; see Section 13.13.

(Continues)

Table 13–2 Important Methods of the Iterable Trait (*Continued*)

Methods	Description
<code>zip(coll2), zipAll(coll2, fill, fill2), zipWithIndex</code>	Returns pairs of elements from this collection and another; see Section 13.10.
<code>grouped(n), sliding(n)</code>	Returns iterators of subcollections of length <code>n</code> ; <code>grouped</code> yields elements with index 0 until <code>n</code> , then with index <code>n</code> until <code>2 * n</code> , and so on; <code>sliding</code> yields elements with index 0 until <code>n</code> , then with index 1 until <code>n + 1</code> , and so on.
<code>groupBy(k)</code>	Yields a map whose keys are <code>k(x)</code> for all elements <code>x</code> . The value for each key is the collection of elements with that key.
<code>mkString(before, between, after), addString(sb, before, between, after)</code>	Makes a string of all elements, adding the given strings before the first, between each, and after the last element. The second method appends that string to a string builder.
<code>toIterable, toSeq, toIndexedSeq, toArray, toBuffer, toList, toStream, toSet, toVector, toMap, to[C]</code>	Converts the collection to a collection of the specified type.

The Seq trait adds several methods to the Iterable trait. Table 13–3 shows the most important ones.

Table 13–3 Important Methods of the Seq Trait

Methods	Description
<code>contains(elem), containsSlice(seq), startsWith(seq), endsWith(seq)</code>	Returns true if this sequence contains the given element or sequence; if it starts or ends with the given sequence.
<code>indexOf(elem), lastIndexOf(elem), indexOfSlice(seq), lastIndexOfSlice(seq)</code>	Returns the index of the first or last occurrence of the given element or element sequence.
<code>indexWhere(pred)</code>	Returns the index of the first element fulfilling <code>pred</code> .
<code>prefixLength(pred), segmentLength(pred, n)</code>	Returns the length of the longest sequence of elements fulfilling <code>pred</code> , starting with 0 or <code>n</code> .

(Continues)

Table 13–3 Important Methods of the Seq Trait (*Continued*)

Methods	Description
padTo(n, fill)	Returns a copy of this sequence, with fill appended until the length is n.
intersect(seq), diff(seq)	Returns the “multiset” intersection or difference of the sequences. For example, if a contains five 1s and b contains two, then a intersect b contains two (the smaller count), and a diff b contains three (the difference).
reverse	The reverse of this sequence.
sorted, sortWith(less), sortBy(f)	The sequence sorted using the element ordering, the binary less function, or a function f that maps each element to an ordered type.
permutations, combinations(n)	Returns an iterator over all permutations or combinations (subsequences of length n).



NOTE: Note that these methods never mutate a collection. They return a collection of the same type as the original. This is sometimes called the “uniform return type” principle.

13.8 Mapping a Function

You may want to transform all elements of a collection. The `map` method applies a function to a collection and yields a collection of the results. For example, given a list of strings

```
val names = List("Peter", "Paul", "Mary")
```

you get a list of the uppercased strings as

```
names.map(_.toUpperCase) // List("PETER", "PAUL", "MARY")
```

This is exactly the same as

```
for (n <- names) yield n.toUpperCase
```

If the function yields a collection instead of a single value, you may want to concatenate all results. In that case, use `flatMap`. For example, consider

```
def ulcase(s: String) = Vector(s.toUpperCase(), s.toLowerCase())
```

Then `names.map(ulcase)` is

```
List(Vector("PETER", "peter"), Vector("PAUL", "paul"), Vector("MARY", "mary"))
```

but `names.flatMap(uppercase)` is

```
List("PETER", "peter", "PAUL", "paul", "MARY", "mary")
```



TIP: If you use `flatMap` with a function that returns an `Option`, the resulting collection contains all values v for which the function returns `Some(v)`.



NOTE: The `map` and `flatMap` methods are important because they are used for translating for expressions. For example, the expression

```
for (i <- 1 to 10) yield i * i
```

is translated to

```
(1 to 10).map(i => i * i)
```

and

```
for (i <- 1 to 10; j <- 1 to i) yield i * j
```

becomes

```
(1 to 10).flatMap(i => (1 to i).map(j => i * j))
```

Why `flatMap`? See Exercise 9 on page 195.

The `transform` method is the in-place equivalent of `map`. It applies to mutable collections, and replaces each element with the result of a function. For example, the following code changes all buffer elements to uppercase:

```
val buffer = ArrayBuffer("Peter", "Paul", "Mary")
buffer.transform(_.toUpperCase)
```

If you just want to apply a function for its side effect and don't care about the function values, use `foreach`:

```
names.foreach(println)
```

The `collect` method works with *partial functions*—functions that may not be defined for all inputs. It yields a collection of all function values of the arguments on which it is defined. For example,

```
"-3+4".collect { case '+' => 1 ; case '-' => -1 } // Vector(-1, 1)
```

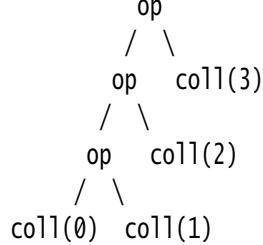
The `groupBy` method yields a map whose keys are the function values, and whose values are the collections of elements whose function value is the given key. For example,

```
val words = ...
val map = words.groupBy(_.substring(0, 1).toUpperCase)
```

builds a map that maps "A" to all words starting with A, and so on.

13.9 Reducing, Folding, and Scanning A3

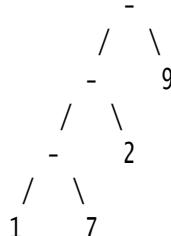
The `map` method applies a unary function to all elements of a collection. The methods that we discuss in this section combine elements with a *binary* function. The call `c.reduceLeft(op)` applies `op` to successive elements, like this:



For example,

`List(1, 7, 2, 9).reduceLeft(_ - _)`

is



or

$$(1 - 7) - 2 - 9 = 1 - 7 - 2 - 9 = -17$$

The `reduceRight` method does the same, but it starts at the end of the collection. For example,

`List(1, 7, 2, 9).reduceRight(_ - _)`

is

$$1 - (7 - (2 - 9)) = 1 - 7 + 2 - 9 = -13$$

Often, it is useful to start the computation with an initial element other than the initial element of a collection. The call `coll.foldLeft(init)(op)` computes

```

.
.
.
op
/ \
op  coll(2)
/ \
op  coll(1)
/ \
init  coll(0)

```

For example,

`List(1, 7, 2, 9).foldLeft(0)(_ - _)`

is

$$0 - 1 - 7 - 2 - 9 = -19$$



NOTE: The initial value and the operator are separate “curried” parameters so that Scala can use the type of the initial value for type inference in the operator. For example, in `List(1, 7, 2, 9).foldLeft("")(_ + _)`, the initial value is a string, so the operator must be a function `(String, Int) => String`.

You can also write the `foldLeft` operation with the `/:` operator, like this:

`(0 /: List(1, 7, 2, 9))(_ - _)`

The `/:` is supposed to remind you of the shape of the tree.



NOTE: With the `/:` operator, the initial value is the first operand. Note that, since the operator ends with a colon, it is a method of the second operand.

There is a `foldRight` or `:\ variant as well, computing`

```

.
.
.
op
/ \
coll(n-3)  op
/ \
coll(n-2)  op
/ \
coll(n-1)  init

```

These examples don't seem to be very useful. Of course, `coll.reduceLeft(_ + _)` or `coll.foldLeft(0)(_ + _)` computes the sum, but you can get that directly with `coll.sum`.

Folding is sometimes attractive as a replacement for a loop. Suppose, for example, we want to count the frequencies of the letters in a string. One way is to visit each letter and update a mutable map.

```
val freq = scala.collection.mutable.Map[Char, Int]()
for (c <- "Mississippi") freq(c) = freq.getOrElse(c, 0) + 1
// Now freq is Map('i' -> 4, 'M' -> 1, 's' -> 4, 'p' -> 2)
```

Here is another way of thinking about this process. At each step, combine the frequency map and the newly encountered letter, yielding a new frequency map. That's a fold:

```
.  
op  
/ \  
op 's'  
/ \  
op 'i'  
/ \  
empty map 'M'
```

What is `op`? The left operand is the partially filled map, and the right operand is the new letter. The result is the augmented map. It becomes the input to the next call to `op`, and at the end, the result is a map with all counts. The code is

```
(Map[Char, Int]() /: "Mississippi") {
  (m, c) => m + (c -> (m.getOrElse(c, 0) + 1))
}
```

Note that this is an immutable map. We compute a new map at each step.



NOTE: It is possible to replace any while loop with a fold. Build a data structure that combines all variables updated in the loop, and define an operation that implements one step through the loop. I am not saying that this is always a good idea, but you may find it interesting that loops and mutations can be eliminated in this way.

Finally, the `scanLeft` and `scanRight` methods combine folding and mapping. You get a collection of all intermediate results. For example,

```
(1 to 10).scanLeft(0)(_ + _)
```

yields all partial sums:

```
Vector(0, 1, 3, 6, 10, 15, 21, 28, 36, 45, 55)
```

13.10 Zipping

The methods of the preceding section apply an operation to adjacent elements in the same collection. Sometimes, you have two collections, and you want to combine corresponding elements. For example, suppose you have a list of product prices and corresponding quantities:

```
val prices = List(5.0, 20.0, 9.95)
val quantities = List(10, 2, 1)
```

The `zip` method lets you combine them into a list of pairs. For example,

```
prices zip quantities
```

is a `List[(Double, Int)]`:

```
List[(Double, Int)] = List((5.0, 10), (20.0, 2), (9.95, 1))
```

The method is called “zip” because it combines the two collections like the teeth of a zipper.

Now it is easy to apply a function to each pair.

```
(prices zip quantities) map { p => p._1 * p._2 }
```

The result is a list of prices:

```
List(50.0, 40.0, 9.95)
```

The total price of all items is then

```
((prices zip quantities) map { p => p._1 * p._2 }) sum
```

If one collection is shorter than the other, the result has as many pairs as the shorter collection. For example,

```
List(5.0, 20.0, 9.95) zip List(10, 2)
```

is

```
List((5.0, 10), (20.0, 2))
```

The `zipAll` method lets you specify defaults for the shorter list:

```
List(5.0, 20.0, 9.95).zipAll(List(10, 2), 0.0, 1)
```

is

```
List((5.0, 10), (20.0, 2), (9.95, 1))
```

The `zipWithIndex` method returns a list of pairs where the second component is the index of each element. For example,

```
"Scala".zipWithIndex  
is
```

```
Vector((‘S’, 0), (‘c’, 1), (‘a’, 2), (‘l’, 3), (‘a’, 4))
```

This can be useful if you want to compute the index of an element with a certain property. For example,

```
"Scala".zipWithIndex.max  
is (‘l’, 3). The index of the value with the largest encoding is  
"Scala".zipWithIndex.max._2
```

13.11 Iterators

You can obtain an iterator from a collection with the `iterator` method. This isn’t as common as in Java or C++ because you can usually get what you need more easily with one of the methods from the preceding sections.

However, iterators are useful for collections that are expensive to construct fully. For example, `Source.fromFile` yields an iterator because it might not be efficient to read an entire file into memory. There are a few `Iterable` methods that yield an iterator, such as `grouped` or `sliding`.

When you have an iterator, you can iterate over the elements with the `next` and `hasNext` methods.

```
while (iter.hasNext)  
  do something with iter.next()
```

If you prefer, you can use a `for` loop instead:

```
for (elem <- iter)  
  do something with elem
```

Both loops end up moving the iterator to the end of the collection, after which it is no longer usable.

Sometimes, you want to be able to look at the next element before deciding whether to consume it. In that case, use the `buffered` method to turn an `Iterator` into a `BufferedIterator`. The `head` method yields the next element without advancing the iterator.

```
val iter = scala.io.Source.fromFile(filename).buffered  
while (iter.hasNext && iter.head.isWhitespace) iter.next  
  // Now iter points to the first non-whitespace character
```

The `Iterator` class defines a number of methods that work identically to the methods on collections. In particular, all `Iterable` methods listed in Section 13.7, “Common Methods,” on page 180 are available, except for `head`, `headOption`, `last`,

`lastOption`, `tail`, `init`, `takeRight`, and `dropRight`. After calling a method such as `map`, `filter`, `count`, `sum`, or even `length`, the iterator is at the end of the collection, and you can't use it again. With other methods, such as `find` or `take`, the iterator is past the found element or the taken ones.

If you find it too tedious to work with an iterator, you can use a method such as `toArray`, `toIterable`, `toSeq`, `toSet`, or `toMap` to copy the values into a collection.

13.12 Streams A3

In the preceding sections, you saw that an iterator is a “lazy” alternative to a collection. You get the elements as you need them. If you don’t need any more elements, you don’t pay for the expense of computing the remaining ones.

However, iterators are fragile. Each call to `next` mutates the iterator. *Streams* offer an immutable alternative. A stream is an immutable list in which the tail is computed lazily—that is, only when you ask for it.

Here is a typical example:

```
def numsFrom(n: BigInt): Stream[BigInt] = n #:: numsFrom(n + 1)
```

The `#::` operator is like the `::` operator for lists, but it constructs a stream.

When you call

```
val tenOrMore = numsFrom(10)
```

you get a stream object that is displayed as

```
Stream(10, ?)
```

The tail is unevaluated. If you call

```
tenOrMore.tail.tail.tail
```

you get

```
Stream(13, ?)
```

Stream methods are executed lazily. For example,

```
val squares = numsFrom(1).map(x => x * x)
```

yields

```
Stream(1, ?)
```

You have to call `squares.tail` to force evaluation of the next entry.

If you want to get more than one answer, you can invoke `take` followed by `force`, which forces evaluation of all values. For example,

```
squares.take(5).force
```

produces `Stream(1, 4, 9, 16, 25)`.

Of course, you don't want to call

```
squares.force // No!
```

That call would attempt to evaluate all members of an infinite stream, causing an `OutOfMemoryError`.

You can construct a stream from an iterator. For example, the `Source.getLines` method returns an `Iterator[String]`. With that iterator, you can only visit the lines once. A stream caches the visited lines so you can revisit them:

```
val words = Source.fromFile("/usr/share/dict/words").getLines.toStream
words // Stream(A, ?)
words(5) // Aachen
words // Stream(A, A's, AOL, AOL's, Aachen, ?)
```



NOTE: Scala streams are quite different from streams in Java 8. However, lazy views, covered in the next section, are conceptually equivalent to Java 8 streams.

13.13 Lazy Views A3

In the preceding section, you saw that stream methods are computed lazily, delivering results only when they are needed. You can get a similar effect with other collections by applying the `view` method. This method yields a collection on which methods are applied lazily. For example,

```
val palindromicSquares = (1 to 1000000).view
  .map(x => x * x)
  .filter(x => x.toString == x.toString.reverse)
```

yields a collection that is unevaluated. (Unlike a stream, not even the first element is evaluated.) When you call

```
palindromicSquares.take(10).mkString(",")
```

then enough squares are generated until ten palindromes have been found, and then the computation stops. Unlike streams, views do not cache any values. If you call `palindromicSquares.take(10).mkString(",")` again, the computation starts over.

As with streams, use the `force` method to force evaluation of a lazy view. You get back a collection of the same type as the original.



CAUTION: The apply method forces evaluation of the entire view. Instead of calling `lazyView(i)`, call `lazyView.take(i).last`.

When you obtain a view into a slice of a mutable collection, any mutations affect the original collection. For example,

```
ArrayBuffer buffer = ...
buffer.view(10, 20).transform(x => 0)
```

clears the given slice and leaves the other elements unchanged.

13.14 Interoperability with Java Collections

At times you may need to use a Java collection, and you will likely miss the rich set of methods that you get with Scala collections. Conversely, you may want to build up a Scala collection and then pass it to Java code. The `JavaConversions` object provides a set of conversions between Scala and Java collections.

Give the target value an explicit type to trigger the conversion. For example,

```
import scala.collection.JavaConversions...
val props: scala.collection.mutable.Map[String, String] = System.getProperties()
```

If you are worried about unwanted implicit conversions, just import the ones you need, for example:

```
import scala.collection.JavaConversions.propertiesAsScalaMap
```

Table 13–4 shows the conversions from Scala to Java collections.

And Table 13–5 shows the opposite conversions from Java to Scala collections.

Note that the conversions yield wrappers that let you use the target interface to access the original type. For example, if you use

```
val props: scala.collection.mutable.Map[String, String] = System.getProperties()
```

then `props` is a wrapper whose methods call the methods of the underlying Java object. If you call

```
props("com.horstmann.scala") = "impatient"
```

then the wrapper calls `put("com.horstmann.scala", "impatient")` on the underlying `Properties` object.

Table 13–4 Conversions from Scala Collections to Java Collections

Implicit Function	From Type in <code>scala.collection</code>	To Type in <code>java.util</code>
<code>asJavaCollection</code>	<code>Iterable</code>	<code>Collection</code>
<code>asJavaIterable</code>	<code>Iterable</code>	<code>Iterable</code>
<code>asJavaIterator</code>	<code>Iterator</code>	<code>Iterator</code>
<code>asJavaEnumeration</code>	<code>Iterator</code>	<code>Enumeration</code>
<code>seqAsJavaList</code>	<code>Seq</code>	<code>List</code>
<code>mutableSeqAsJavaList</code>	<code>mutable.Seq</code>	<code>List</code>
<code>bufferAsJavaList</code>	<code>mutable.Buffer</code>	<code>List</code>
<code>setAsJavaSet</code>	<code>Set</code>	<code>Set</code>
<code>mutableSetAsJavaSet</code>	<code>mutable.Set</code>	<code>Set</code>
<code>mapAsJavaMap</code>	<code>Map</code>	<code>Map</code>
<code>mutableMapAsJavaMap</code>	<code>mutable.Map</code>	<code>Map</code>
<code>asJavaDictionary</code>	<code>Map</code>	<code>Dictionary</code>
<code>asJavaConcurrentMap</code>	<code>mutable.ConcurrentMap</code>	<code>concurrent.ConcurrentMap</code>

Table 13–5 Conversions from Java Collections to Scala Collections

Implicit Function	From Type in <code>java.util</code>	To Type in <code>scala.collection</code>
<code>collectionAsScalaIterable</code>	<code>Collection</code>	<code>Iterable</code>
<code>iterableAsScalaIterable</code>	<code>Iterable</code>	<code>Iterable</code>
<code>asScalaIterator</code>	<code>Iterator</code>	<code>Iterator</code>
<code>enumerationAsScalaIterator</code>	<code>Enumeration</code>	<code>Iterator</code>
<code>asScalaBuffer</code>	<code>List</code>	<code>mutable.Buffer</code>
<code>asScalaSet</code>	<code>Set</code>	<code>mutable.Set</code>
<code>mapAsScalaMap</code>	<code>Map</code>	<code>mutable.Map</code>
<code>dictionaryAsScalaMap</code>	<code>Dictionary</code>	<code>mutable.Map</code>
<code>propertiesAsScalaMap</code>	<code>Properties</code>	<code>mutable.Map</code>
<code>asScalaConcurrentMap</code>	<code>concurrent.ConcurrentMap</code>	<code>mutable.ConcurrentMap</code>

13.15 Parallel Collections

It is hard to write correct concurrent programs, yet concurrency is often required nowadays to keep all processors of a computer busy. Scala offers a particularly attractive solution for manipulating large collections. Such tasks often parallelize naturally. For example, to compute the sum of all elements, multiple threads can concurrently compute the sums of different sections; in the end, these partial results are summed up. Of course it is troublesome to schedule these concurrent activities—but with Scala, you don't have to. If `coll` is a large collection, then

```
coll.par.sum
```

computes the sum concurrently. The `par` method produces a *parallel implementation* of the collection. That implementation parallelizes the collection methods whenever possible. For example,

```
coll.par.count(_ % 2 == 0)
```

counts the even numbers in `coll` by evaluating the predicate on subcollections in parallel and combining the results.

You can parallelize a `for` loop by applying `.par` to the collection over which you iterate, like this:

```
for (i <- (0 until 100000).par) print(s" $i")
```

Try it out—the numbers are printed in the order they are produced by the threads working on the task.

In a `for/yield` loop, the results are assembled in order. Try this:

```
(for (i <- (0 until 100000).par) yield i) == (0 until 100000)
```



CAUTION: If parallel computations mutate shared variables, the result is unpredictable. For example, do not update a shared counter:

```
var count = 0
for (c <- coll.par) { if (c % 2 == 0) count += 1 } // Error!
```

The parallel collections returned by the `par` method belong to types that extend the `ParSeq`, `ParSet`, or `ParMap` traits. These are *not* subtypes of `Seq`, `Set`, or `Map`, and you cannot pass a parallel collection to a method that expects a sequential collection.

You can convert a parallel collection back to a sequential one with the `seq` method.

```
val result = coll.par.filter(p).seq
```

Not all methods can be parallelized. For example, `reduceLeft` and `reduceRight` require that each operator is applied in sequence. There is an alternate method, `reduce`, that operates on parts of the collection and combines the results. For this

to work, the operator must be *associative*—it must fulfill $(a \text{ op } b) \text{ op } c = a \text{ op } (b \text{ op } c)$. For example, addition is associative but subtraction is not: $(a - b) - c \neq a - (b - c)$.

Similarly, there is a `fold` method that operates on parts of the collection. Unfortunately, it is not as flexible as `foldLeft` or `foldRight`—both arguments of the operator must be elements. That is, you can do `coll.par.fold(0)(_ + _)`, but you cannot do a more complex fold such as the one at the end of Section 13.9, “Reducing, Folding, and Scanning,” on page 184.

To solve this problem, there is an even more general aggregate that applies an operator to parts of the collection, and then uses another operator to combine the results. For example, `str.par.aggregate(Set[Char]())(_ + _, _ ++ _)` is the equivalent of `str.foldLeft(Set[Char]())(_ + _)`, forming a set of all distinct characters in `str`.



NOTE: By default, parallel collections use a global fork-join pool, which is well suited for processor-intensive calculations. If you carry out parallel computation steps that block, you should choose a different “execution context”—see Chapter 17.

Exercises

1. Write a function that, given a string, produces a map of the indexes of all characters. For example, `indexes("Mississippi")` should return a map associating 'M' with the set {0}, 'i' with the set {1, 4, 7, 10}, and so on. Use a mutable map of characters to mutable sets. How can you ensure that the set is sorted?
2. Repeat the preceding exercise, using an immutable map of characters to lists.
3. Write a function that removes every second element from a `ListBuffer`. Try it two ways. Call `remove(i)` for all even `i` starting at the end of the list. Copy every second element to a new list. Compare the performance.
4. Write a function that receives a collection of strings and a map from strings to integers. Return a collection of integers that are values of the map corresponding to one of the strings in the collection. For example, given `Array("Tom", "Fred", "Harry")` and `Map("Tom" -> 3, "Dick" -> 4, "Harry" -> 5)`, return `Array(3, 5)`. Hint: Use `flatMap` to combine the `Option` values returned by `get`.
5. Implement a function that works just like `mkString`, using `reduceLeft`.
6. Given a list of integers `lst`, what is `(lst :\ List[Int]())(_ :: _)`? (`List[Int]()` /: `lst`)`(_ :: _)`? How can you modify one of them to reverse the list?
7. In Section 13.10, “Zipping,” on page 187, the expression `(prices zip quantities) map { p => p._1 * p._2 }` is a bit inelegant. We can’t do `(prices zip quantities) map { _ * _ }` because `_ * _` is a function with two arguments, and we need a function with one argument that is a tuple. The `tupled` method of the `Function`

object changes a function with two arguments to one that takes a tuple. Apply `tupled` to the multiplication function so you can map it over the list of pairs.

8. Write a function that turns an array of Double values into a two-dimensional array. Pass the number of columns as a parameter. For example, with `Array(1, 2, 3, 4, 5, 6)` and three columns, return `Array(Array(1, 2, 3), Array(4, 5, 6))`. Use the `grouped` method.
9. The Scala compiler transforms a `for/yield` expression

```
for (i <- 1 to 10; j <- 1 to i) yield i * j
```

to invocations of `flatMap` and `map`, like this:

```
(1 to 10).flatMap(i => (1 to i).map(j => i * j))
```

Explain the use of `flatMap`. Hint: What is `(1 to i).map(j => i * j)` when `i` is 1, 2, 3?

What happens when there are three generators in the `for/yield` expression?

10. The method `java.util.TimeZone.getAvailableIDs` yields time zones such as Africa/Cairo and Asia/Chungking. Which continent has the most time zones? Hint: `groupBy`.
11. Harry Hacker reads a file into a string and wants to use a parallel collection to update the letter frequencies concurrently on portions of the string. He uses the following code:

```
val frequencies = new scala.collection.mutable.HashMap[Char, Int]
for (c <- str.par) frequencies(c) = frequencies.getOrElse(c, 0) + 1
```

Why is this a terrible idea? How can he really parallelize the computation? (Hint: Use `aggregate`.)

Pattern Matching and Case Classes

Topics in This Chapter **A2**

- 14.1 A Better Switch — page 198
- 14.2 Guards — page 199
- 14.3 Variables in Patterns — page 199
- 14.4 Type Patterns — page 200
- 14.5 Matching Arrays, Lists, and Tuples — page 201
- 14.6 Extractors — page 202
- 14.7 Patterns in Variable Declarations — page 203
- 14.8 Patterns in for Expressions — page 204
- 14.9 Case Classes — page 205
- 14.10 The copy Method and Named Parameters — page 205
- 14.11 Infix Notation in case Clauses — page 206
- 14.12 Matching Nested Structures — page 207
- 14.13 Are Case Classes Evil? — page 208
- 14.14 Sealed Classes — page 209
- 14.15 Simulating Enumerations — page 209
- 14.16 The Option Type — page 210
- 14.17 Partial Functions **L2** — page 211
- Exercises — page 212

Chapter 14

Pattern matching is a powerful mechanism that has a number of applications: switch statements, type inquiry, and “destructuring” (getting at the parts of complex expressions). Case classes are optimized to work with pattern matching.

The key points of this chapter are:

- The `match` expression is a better `switch`, without fall-through.
- If no pattern matches, a `MatchError` is thrown. Use the `case _` pattern to avoid that.
- A pattern can include an arbitrary condition, called a guard.
- You can match on the type of an expression; prefer this over `isInstanceOf`/`asInstanceOf`.
- You can match patterns of arrays, tuples, and case classes, and bind parts of the pattern to variables.
- In a `for` expression, nonmatches are silently skipped.
- A case class is a class for which the compiler automatically produces the methods that are needed for pattern matching.
- The common superclass in a case class hierarchy should be sealed.
- Use the `Option` type for values that may or may not be present—it is safer than using `null`.

14.1 A Better Switch

Here is the equivalent of the C-style switch statement in Scala:

```
var sign = ...
val ch: Char = ...

ch match {
  case '+' => sign = 1
  case '-' => sign = -1
  case _ => sign = 0
}
```

The equivalent of default is the catch-all case _ pattern. It is a good idea to have such a catch-all pattern. If no pattern matches, a MatchError is thrown.

Unlike the switch statement, Scala pattern matching does not suffer from the “fall-through” problem. (In C and its derivatives, you must use explicit break statements to exit a switch at the end of each branch, or you will fall through to the next branch. This is annoying and error-prone.)



NOTE: In his entertaining book *Deep C Secrets*, Peter van der Linden reports a study of a large body of C code in which the fall-through behavior was unwanted in 97% of the cases.

Similar to if, match is an expression, not a statement. The preceding code can be simplified to

```
sign = ch match {
  case '+' => 1
  case '-' => -1
  case _ => 0
}
```

Use | to separate multiple alternatives:

```
prefix match {
  case "0" | "0x" | "0X" => ...
  ...
}
```

You can use the match statement with any types. For example:

```
color match {
  case Color.RED => ...
  case Color.BLACK => ...
  ...
}
```

14.2 Guards

Suppose we want to extend our example to match all digits. In a C-style switch statement, you would simply add multiple case labels, for example case '0': case '1': ... case '9':. (Except that, of course, you can't use ... but must write out all ten cases explicitly.) In Scala, you add a *guard clause* to a pattern, like this:

```
ch match {
  case _ if Character.isDigit(ch) => digit = Character.digit(ch, 10)
  case '+' => sign = 1
  case '-' => sign = -1
  case _ => sign = 0
}
```

The guard clause can be any Boolean condition.

Patterns are always matched top-to-bottom. If the pattern with the guard clause doesn't match, the case '+' pattern is attempted next.

14.3 Variables in Patterns

If the case keyword is followed by a variable name, then the match expression is assigned to that variable. For example:

```
str(i) match {
  case '+' => sign = 1
  case '-' => sign = -1
  case ch => digit = Character.digit(ch, 10)
}
```

You can think of case _ as a special case of this feature, where the variable name is _.

You can use the variable name in a guard:

```
str(i) match {
  case ch if Character.isDigit(ch) => digit = Character.digit(ch, 10)
  ...
}
```



CAUTION: Unfortunately, variable patterns can conflict with constants, for example:

```
import scala.math._  
0.5 * c / r match {  
    case Pi => ... // If 0.5 * c / r equals Pi ...  
    case x => ... // Otherwise set x to 0.5 * c / r ...  
}
```

How does Scala know that Pi is a constant, not a variable? The rule is that a variable must start with a *lowercase* letter.

If you have an expression that starts with a lowercase letter, enclose it in backquotes:

```
import java.io.File._ // Imports java.io.File.pathSeparator  
str match {  
    case `pathSeparator` => ... // If str == pathSeparator ...  
    case pathSeparator => ...  
        // Caution—declares a new variable pathSeparator  
}
```

14.4 Type Patterns

You can match on the type of an expression, for example:

```
obj match {  
    case x: Int => x  
    case s: String => Integer.parseInt(s)  
    case _: BigInt => Int.MaxValue  
    case _ => 0  
}
```

In Scala, this form is preferred to using the `isInstanceOf` operator.

Note the variable names in the patterns. In the first pattern, the match is bound to `x` as an `Int`, and in the second pattern, it is bound to `s` as a `String`. No `asInstanceOf` casts are needed!



CAUTION: When you match against a type, you must supply a variable name. Otherwise, you match the *object*:

```
obj match {  
    case _: BigInt => Int.MaxValue // Matches any object of type BigInt  
    case BigInt => -1 // Matches the BigInt object of type Class  
}
```



CAUTION: Matches occur at runtime, and generic types are erased in the Java virtual machine. For that reason, you cannot make a type match for a specific Map type.

```
case m: Map[String, Int] => ... // Don't
```

You can match a generic map:

```
case m: Map[_, _) => ... // OK
```

However, arrays are not erased. You can match an Array[Int].

14.5 Matching Arrays, Lists, and Tuples

To match an array against its contents, use Array expressions in the patterns, like this:

```
arr match {
  case Array(0) => "0"
  case Array(x, y) => s"$x $y"
  case Array(0, _) => "0 ..."
  case _ => "something else"
}
```

The first pattern matches the array containing 0. The second pattern matches any array with two elements, and it binds the variables x and y to the elements. The third pattern matches any array starting with zero.

If you want to bind a variable argument match `_*` to a variable, use the `@` notation like this:

```
case Array(x, rest @ _) => rest.min
```

You can match lists in the same way, with List expressions. Alternatively, you can use the `::` operator:

```
lst match {
  case 0 :: Nil => "0"
  case x :: y :: Nil => s"$x $y"
  case 0 :: tail => "0 ..."
  case _ => "something else"
}
```

With tuples, use the tuple notation in the pattern:

```
pair match {
  case (0, _) => "0 ..."
  case (y, 0) => s"$y 0"
  case _ => "neither is 0"
}
```

Again, note how the variables are bound to parts of the list or tuple. Since these bindings give you easy access to parts of a complex structure, this operation is called *destructuring*.



CAUTION: The same warning applies as in Section 14.3, “Variables in Patterns,” on page 199. The variable names that you use in the pattern must start with a *lowercase* letter. In a match against `Array(X, Y)`, `X` and `Y` are deemed constants, not variables.



NOTE: If a pattern has alternatives, you cannot use variables other than an underscore. For example,

```
pair match {
    case (_, 0) | (0, _) => ... // OK, matches if one component is zero
    case (x, 0) | (0, x) => ... // Error—cannot bind with alternatives
}
```

14.6 Extractors

In the preceding section, you have seen how patterns can match arrays, lists, and tuples. These capabilities are provided by *extractors*—objects with an `unapply` or `unapplySeq` method that extract values from an object. The implementation of these methods is covered in Chapter 11. The `unapply` method is provided to extract a fixed number of objects, while `unapplySeq` extracts a sequence whose length can vary.

For example, consider the expression

```
arr match {
    case Array(x, 0) => x
    case Array(x, rest @ _) => rest.min
    ...
}
```

The `Array` companion object is an extractor—it defines an `unapplySeq` method. That method is called *with the expression that is being matched*, not with what appear to be the parameters in the pattern. The call `Array.unapplySeq(arr)`, when successful, results in a sequence of values, namely the two values in the array. In the first case, the match succeeds if the array has length 2 and the second element is zero. In that case, the initial array element is assigned to `x`.

Regular expressions provide another good use of extractors. When a regular expression has groups, you can match each group with an extractor pattern. For example:

```
val pattern = "([0-9]+) ([a-z]+)".r
"99 bottles" match {
  case pattern(num, item) => ...
    // Sets num to "99", item to "bottles"
}
```

The call `pattern.unapplySeq("99 bottles")` yields a sequence of strings that match the groups. These are assigned to the variables `num` and `item`.

Note that here the extractor isn't a companion object but a regular expression object.

14.7 Patterns in Variable Declarations

In the preceding sections, you have seen how patterns can contain variables. You can use these patterns inside variable declarations. For example,

```
val (x, y) = (1, 2)
```

simultaneously defines `x` as 1 and `y` as 2. That is useful for functions that return a pair, for example:

```
val (q, r) = BigInt(10) % 3
```

The `%` method returns a pair containing the quotient and the remainder, which are captured in the variables `q` and `r`.

The same syntax works for any patterns with variable names. For example,

```
val Array(first, second, rest @_*) = arr
```

assigns the `first` and `second` element of the array `arr` to the variables `first` and `second` and `rest` to a `Seq` of the remaining elements.



CAUTION: The same warning applies as in Section 14.3, “Variables in Patterns,” on page 199. The variable names that you use in the pattern must start with a *lowercase* letter. In a declaration `val Array(E, x) = arr`, `E` is deemed a constant and `x` a variable that becomes `arr(1)` if `arr` has length 2 and `arr(0) == E`.



NOTE: The expression

```
val p(x1, ..., xn) = e
```

is, by definition, exactly the same as

```
val $result = e match { case p(x1, ..., xn) => (x1, ..., xn) }
val x1 = $result._1
...
val xn = $result._n
```

where x_1, \dots, x_n are the free variables in the pattern p .

This definition holds even when there are no free variables in the pattern. For example,

```
val 2 = x
```

is perfectly legal Scala code, provided x has been defined elsewhere. When you apply the definition, you get

```
val $result = x match { case 2 => () }
```

followed by no assignments. In other words, it is equivalent to

```
if (!(2 == x)) throw new MatchError
```

14.8 Patterns in for Expressions

You can use patterns with variables in for expressions. For each traversed value, the variables are bound. This makes it possible to traverse a map:

```
import scala.collection.JavaConversions.propertiesAsScalaMap
// Converts Java Properties to a Scala map—just to get an interesting example
for ((k, v) <- System.getProperties())
  println(s"$k -> $v")
```

For each (*key, value*) pair in the map, k is bound to the key and v to the value.

In a for expression, match failures are silently ignored. For example, the following loop prints all keys with empty value, skipping over all others:

```
for ((k, "") <- System.getProperties())
  println(k)
```

You can also use a guard. Note that the if goes after the $<-$ symbol:

```
for ((k, v) <- System.getProperties() if v == "")
  println(k)
```

14.9 Case Classes

Case classes are a special kind of classes that are optimized for use in pattern matching. In this example, we have two case classes that extend a regular (noncase) class:

```
abstract class Amount
case class Dollar(value: Double) extends Amount
case class Currency(value: Double, unit: String) extends Amount
```

You can also have case objects for singletons:

```
case object Nothing extends Amount
```

When we have an object of type `Amount`, we can use pattern matching to match the amount type and bind the property values to variables:

```
amt match {
  case Dollar(v) => s"$$${v}"
  case Currency(_, u) => s"Oh noes, I got $u"
  case Nothing => ""
}
```



NOTE: Use `()` with case class instances, no parentheses with case objects.

When you declare a case class, several things happen automatically.

- Each of the constructor parameters becomes a `val` unless it is explicitly declared as a `var` (which is not recommended).
- An `apply` method is provided for the companion object that lets you construct objects without `new`, such as `Dollar(29.95)` or `Currency(29.95, "EUR")`.
- An `unapply` method is provided that makes pattern matching work—see Chapter 11 for the details. (You don't really need to know these details to use case classes for pattern matching.)
- Methods `toString`, `equals`, `hashCode`, and `copy` are generated unless they are explicitly provided.

Otherwise, case classes are just like any other classes. You can add methods and fields to them, extend them, and so on.

14.10 The copy Method and Named Parameters

The `copy` method of a case class makes a new object with the same values as an existing one. For example,

```
val amt = Currency(29.95, "EUR")
val price = amt.copy()
```

By itself, that isn't very useful—after all, a `Currency` object is immutable, and one can just share the object reference. However, you can use named parameters to modify some of the properties:

```
val price = amt.copy(value = 19.95) // Currency(19.95, "EUR")
```

or

```
val price = amt.copy(unit = "CHF") // Currency(29.95, "CHF")
```

14.11 Infix Notation in case Clauses

When an `unapply` method yields a pair, you can use infix notation in the case clause. In particular, you can use infix notation with a case class that has two parameters. For example:

```
amt match { case a Currency u => ... } // Same as case Currency(a, u)
```

Of course, that is a silly example. The feature is meant for matching sequences. For example, every `List` object is either `Nil` or an object of the case class `::`, defined as

```
case class ::[E](head: E, tail: List[E]) extends List[E]
```

Therefore, you can write

```
lst match { case h :: t => ... }
// Same as case ::(h, t), which calls ::.unapply(lst)
```

In Chapter 20, you will encounter the `~` case class for combining pairs of parse results. It is also intended for use as an infix expression in case clauses:

```
result match { case p ~ q => ... } // Same as case ~(p, q)
```

These infix expressions are easier to read when you have more than one. For example,

```
result match { case p ~ q ~ r => ... }
```

is nicer than `~(~(p, q), r)`.

If the operator ends in a colon, it associates right-to-left. For example,

```
case first :: second :: rest
```

means

```
case ::(first, ::(second, rest))
```



NOTE: Infix notation works with any unapply method that returns a pair. Here is an example:

```
case object +: {
    def unapply[T](input: List[T]) =
        if (input.isEmpty) None else Some((input.head, input.tail))
}
```

Now you can destructure lists using +:.

```
1 +: 7 +: 2 +: 9 +: Nil match {
    case first +: second +: rest => first + second + rest.length
}
```

14.12 Matching Nested Structures

Case classes are often used for nested structures. Consider, for example, items that a store sells. Sometimes, we bundle items together for a discount.

```
abstract class Item
case class Article(description: String, price: Double) extends Item
case class Bundle(description: String, discount: Double, items: Item*) extends Item
```

Not having to use new makes it easy to specify nested objects:

```
Bundle("Father's day special", 20.0,
      Article("Scala for the Impatient", 39.95),
      Bundle("Anchor Distillery Sampler", 10.0,
             Article("Old Potrero Straight Rye Whiskey", 79.95),
             Article("Junípero Gin", 32.95)))
```

Patterns can match specific nestings, for example

```
case Bundle(_, _, Article(descr, _), _) => ...
```

binds descr to the description of the first article in a bundle.

You can bind a nested value to a variable with the @ notation:

```
case Bundle(_, _, art @ Article(_, _), rest @ _) => ...
```

Now art is the first article in a bundle and rest is the sequence of the other items.

Note that the @_ is required in this example. The pattern

```
case Bundle(_, _, art @ Article(_, _), rest) => ...
```

would match a bundle with an article and exactly one additional item, bound to rest.

As an application, here is a function that computes the price of an item:

```
def price(it: Item): Double = it match {
    case Article(_, p) => p
    case Bundle(_, disc, its @ _*) => its.map(price _).sum - disc
}
```

14.13 Are Case Classes Evil?

The example in the preceding section can enrage OO purists. Shouldn't `price` be a method of the superclass? Shouldn't each subclass override it? Isn't polymorphism better than making a switch on each type?

In many situations, this is true. If someone comes up with another kind of `Item`, you'll need to revisit all those `match` clauses. In such a situation, case classes are not the right solution.

Case classes work well for structures whose makeup doesn't change. For example, the Scala `List` is implemented with case classes. Simplifying things a bit, a list is essentially

```
abstract class List
case object Nil extends List
case class ::(head: Any, tail: List) extends List
```

A list is either empty, or it has a head and a tail (which may be empty or not). Nobody is ever going to add a third case. (You'll see in the next section how to stop anyone from trying.)

When they are appropriate, case classes are quite convenient, for the following reasons:

- Pattern matching often leads to more concise code than inheritance.
- It is easier to read compound objects that are constructed without `new`.
- You get `toString`, `equals`, `hashCode`, and `copy` for free.

Those automatically generated methods do what you think they do—print, compare, hash, and copy each field. See Section 14.10, “The `copy` Method and Named Parameters,” on page 205 for more information about the `copy` method.

For certain kinds of classes, case classes give you exactly the right semantics. Some people call them *value classes*. For example, consider the `Currency` class:

```
case class Currency(value: Double, unit: String)
```

A `Currency(10, "EUR")` is the same as any other `Currency(10, "EUR")`, and that's how `equals` and `hashCode` are implemented. Typically, such classes are immutable.

Case classes with variable fields are somewhat suspect, at least with respect to the hash code. With mutable classes, one should always derive the hash code from fields that are never mutated, such as an ID.



CAUTION: The `toString`, `equals`, `hashCode`, and `copy` methods are *not* generated for case classes that extend other case classes. You get a compiler warning if one case class inherits from another. A future version of Scala may outlaw such inheritance altogether. If you need multiple levels of inheritance to factor out common behavior of case classes, make only the leaves of the inheritance tree into case classes.

14.14 Sealed Classes

When you use pattern matching with case classes, you would like the compiler to check that you exhausted all alternatives. To achieve this, declare the common superclass as *sealed*:

```
sealed abstract class Amount
case class Dollar(value: Double) extends Amount
case class Currency(value: Double, unit: String) extends Amount
```

All subclasses of a sealed class must be defined in the same file as the class itself. For example, if someone wants to add another class for euros,

```
case class Euro(value: Double) extends Amount
```

they must do so in the file in which `Amount` is declared.

When a class is sealed, all of its subclasses are known at compile time, enabling the compiler to check pattern clauses for completeness. It is a good idea for all case classes to extend a sealed class or trait.

14.15 Simulating Enumerations

Case classes let you simulate enumerated types in Scala.

```
sealed abstract class TrafficLightColor
case object Red extends TrafficLightColor
case object Yellow extends TrafficLightColor
case object Green extends TrafficLightColor
```

```
color match {
  case Red => "stop"
  case Yellow => "hurry up"
  case Green => "go"
}
```

Note that the superclass was declared as sealed, enabling the compiler to check that the match clause is complete.

If you find this a bit heavyweight, you may prefer the `Enumeration` helper class that was described in Chapter 6.

14.16 The Option Type

The `Option` type in the standard library uses case classes to express values that might or might not be present. The case subclass `Some` wraps a value, for example `Some("Fred")`. The case object `None` indicates that there is no value.

This is less ambiguous than using an empty string and safer than using `null` for a missing value.

`Option` is a generic type. For example, `Some("Fred")` is an `Option[String]`.

The `get` method of the `Map` class returns an `Option`. If there is no value for a given key, `get` returns `None`. Otherwise, it wraps the value inside `Some`.

You can use pattern matching to analyze such a value.

```
val alicesScore = scores.get("Alice")
alicesScore match {
  case Some(score) => println(score)
  case None => println("No score")
}
```

But frankly, that is tedious. Alternatively, you can use the `isEmpty` and `get`:

```
if (alicesScore.isEmpty) println("No score")
else println(alicesScore.get)
```

That's tedious too. It is better to use the `getOrElse` method:

```
println(alicesScore.getOrElse("No score"))
```

If `alicesScore` is `None`, then `getOrElse` returns "No score".

A more powerful way of working with options is to consider them as collections that have zero or one element. You can visit the element with a `for` loop:

```
for (score <- alicesScore) println(score)
```

If `alicesScore` is `None`, nothing happens. If it is a `Some`, then the loop executes once, with `score` bound to the contents of the option.

You can also use methods such as `map`, `filter`, or `foreach`. For example,

```
val biggerScore = alicesScore.map(_ + 1) // Some(score + 1) or None
val acceptableScore = alicesScore.filter(_ > 5) // Some(score) if score > 5 or None
alicesScore.foreach(println _) // Prints the score if it exists
```



TIP: When you create an `Option` from a value that may be `null`, you can simply use `Option(value)`. The result is `None` if `value` is `null` and `Some(value)` otherwise.

14.17 Partial Functions L2

A set of case clauses enclosed in braces is a *partial function*—a function which may not be defined for all inputs. It is an instance of a class `PartialFunction[A, B]`. (`A` is the parameter type, `B` the return type.) That class has two methods: `apply`, which computes the function value from the matching pattern, and `isDefinedAt`, which returns true if the input matches at least one of the patterns.

For example,

```
val f: PartialFunction[Char, Int] = { case '+' => 1 ; case '-' => -1 }
f('-') // Calls f.apply(' -'), returns -1
f.isDefinedAt('0') // false
f('0') // Throws MatchError
```

Some methods accept a `PartialFunction` as a parameter. For example, the `collect` method of the `GenTraversable` trait applies a partial function to all elements where it is defined, and returns a sequence of the results.

```
"-3+4".collect { case '+' => 1 ; case '-' => -1 } // Vector(-1, 1)
```

The partial function expression must be in a context where the compiler can infer the return type. This is the case when you assign it to a typed variable or pass it as an argument.



NOTE: An exhaustive set of case clauses defines a `Function1`, not just a `PartialFunction`, that you can pass whenever a function is expected.

```
"-3+4".map { case '+' => 1 ; case '-' => -1; case _ => 0 }
// Vector(-1, 0, 1, 0)
```

A `Seq[A]` is a `PartialFunction[Int, A]`, and a `Map[K, V]` is a `PartialFunction[K, V]`. For example, you can pass a map to `collect`:

```
val names = Array("Alice", "Bob", "Carmen")
val scores = Map("Alice" -> 10, "Carmen" -> 7)
names.collect(scores) // Yields Array(10, 7)
```

The `lift` method turns a `PartialFunction[T, R]` into a regular function with return type `Option[R]`.

```
val f: PartialFunction[Char, Int] = { case '+' => 1 ; case '-' => -1 }
val g = f.lift // A function with type Char => Option[Int]
```

Now `g('-')` is `Some(-1)` and `g('*')` is `None`.

In Chapter 9, you saw that the `Regex.replaceSomeIn` method requires a function `String => Option[String]` for the replacement. If you have a map (or some other `PartialFunction`), you can use `lift` to produce such a function:

```
val varPattern = """\{([0-9]+)\}""".r
val message = "At {1}, there was {2} on {0}"
val vars = Map("{0}" -> "planet 7", "{1}" -> "12:30 pm",
               "{2}" -> "a disturbance of the force.")
val result = varPattern.replaceSomeIn(message, m => vars.lift(m.matched))
```

Conversely, you can turn a function returning Option[R] into a partial function by calling Function.unlift.



NOTE: The catch clause of the try statement is a partial function. You can even use a variable holding a function:

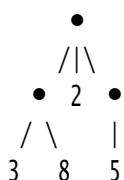
```
def tryCatch[T](b: => T, catcher: PartialFunction[Throwable, T]) =
  try { b } catch catcher
```

Then you can supply a custom catch clause like this:

```
val result = tryCatch(str.toInt,
                      { case _: NumberFormatException => -1 })
```

Exercises

1. Your Java Development Kit distribution has the source code for much of the JDK in the `src.zip` file. Unzip and search for case labels (regular expression `case [^:]++:`). Then look for comments starting with `//` and containing `[Ff]alls?` `thr` to catch comments such as `// Falls through` or `// just fall thru`. Assuming the JDK programmers follow the Java code convention, which requires such a comment, what percentage of cases falls through?
2. Using pattern matching, write a function `swap` that receives a pair of integers and returns the pair with the components swapped.
3. Using pattern matching, write a function `swap` that swaps the first two elements of an array provided its length is at least two.
4. Add a case class `Multiple` that is a subclass of the `Item` class. For example, `Multiple(10, Article("Blackwell Toaster", 29.95))` describes ten toasters. Of course, you should be able to handle any items, such as bundles or multiples, in the second argument. Extend the `price` function to handle this new case.
5. One can use lists to model trees that store values only in the leaves. For example, the list `((3 8) 2 (5))` describes the tree



However, some of the list elements are numbers and others are lists. In Scala, you cannot have heterogeneous lists, so you have to use a `List[Any]`. Write a `leafSum` function to compute the sum of all elements in the leaves, using pattern matching to differentiate between numbers and lists.

6. A better way of modeling such trees is with case classes. Let's start with binary trees.

```
sealed abstract class BinaryTree
case class Leaf(value: Int) extends BinaryTree
case class Node(left: BinaryTree, right: BinaryTree) extends BinaryTree
```

Write a function to compute the sum of all elements in the leaves.

7. Extend the tree in the preceding exercise so that each node can have an arbitrary number of children, and reimplement the `leafSum` function. The tree in Exercise 5 should be expressible as

```
Node(Node(Leaf(3), Leaf(8)), Leaf(2), Node(Leaf(5)))
```

8. Extend the tree in the preceding exercise so that each nonleaf node stores an operator in addition to the child nodes. Then write a function `eval` that computes the value. For example, the tree

```

+
/|\ \
* 2 - 
/ \   |
3 8 5

```

has value $(3 \times 8) + 2 + (-5) = 21$.

Pay attention to the unary minus.

9. Write a function that computes the sum of the non-`None` values in a `List[Option[Int]]`. Don't use a `match` statement.
10. Write a function that composes two functions of type `Double => Option[Double]`, yielding another function of the same type. The composition should yield `None` if either function does. For example,

```
def f(x: Double) = if (x != 1) Some(1 / (x - 1)) else None
def g(x: Double) = if (x >= 0) Some(sqrt(x)) else None
val h = compose(g, f) // h(x) should be g(f(x))
```

Then `h(2)` is `Some(1)`, and `h(1)` and `h(0)` are `None`.

Annotations

Topics in This Chapter **A2**

- 15.1 What Are Annotations? — page 216
- 15.2 What Can Be Annotated? — page 216
- 15.3 Annotation Arguments — page 217
- 15.4 Annotation Implementations — page 218
- 15.5 Annotations for Java Features — page 219
 - Java Modifiers. Marker Interfaces. Checked Exceptions.
 - Variable Arguments. JavaBeans.
- 15.6 Annotations for Optimizations — page 222
 - Tail Recursion. Jump Table Generation and Inlining.
 - Eliding Methods. Specialization for Primitive Types.
- 15.7 Annotations for Errors and Warnings — page 226
- Exercises — page 227

Chapter

15

Annotations let you add information to program items. This information can be processed by the compiler or by external tools. In this chapter, you will learn how to interoperate with Java annotations and how to use the annotations that are specific to Scala.

The key points of this chapter are:

- You can annotate classes, methods, fields, local variables, parameters, expressions, type parameters, and types.
- With expressions and types, the annotation follows the annotated item.
- Annotations have the form `@Annotation`, `@Annotation(value)`, or `@Annotation(name1 = value1, ...)`.
- `@volatile`, `@transient`, `@strictfp`, and `@native` generate the equivalent Java modifiers.
- Use `@throws` to generate Java-compatible throws specifications.
- The `@tailrec` annotation lets you verify that a recursive function uses tail call optimization.
- The `assert` function takes advantage of the `@elidable` annotation. You can optionally remove assertions from your Scala programs.
- Use the `@deprecated` annotation to mark deprecated features.

15.1 What Are Annotations?

Annotations are tags that you insert into your source code so that some tools can process them. These tools can operate at the source level, or they can process the class files into which the compiler has placed your annotations.

Annotations are widely used in Java, for example by testing tools such as JUnit 4 and enterprise technologies such as Java EE.

The syntax is just like in Java. For example:

```
@Test(timeout = 100) def testSomeFeature() { ... }

@Entity class Credentials {
    @Id @BeanProperty var username : String = _
    @BeanProperty var password : String = _
}
```

You can use Java annotations with Scala classes. The annotations in the preceding examples are from JUnit and JPA, two Java frameworks that have no particular knowledge of Scala.

You can also use Scala annotations. These annotations are specific to Scala and are usually processed by the Scala compiler or a compiler plugin. (Implementing a compiler plugin is a nontrivial undertaking that is not covered in this book.)

Java annotations do not affect how the compiler translates source code into bytecode; they merely add data to the bytecode that can be harvested by external tools. In Scala, annotations can affect the compilation process. For example, the `@BeanProperty` annotation that you saw in Chapter 5 causes the generation of getter and setter methods.

15.2 What Can Be Annotated?

In Scala, you can annotate classes, methods, fields, local variables, and parameters, just like in Java.

```
@Entity class Credentials
@Test def testSomeFeature() {}
@BeanProperty var username = _
def doSomething(@NotNull message: String) {}
```

You can apply multiple annotations. The order doesn't matter.

```
@BeanProperty @Id var username = _
```

When annotating the primary constructor, place the annotation before the constructor, and add a pair of parentheses if the annotation has no arguments.

```
class Credentials @Inject() (var username: String, var password: String)
```

You can also annotate expressions. Add a colon followed by the annotation, for example:

```
(myMap.get(key): @unchecked) match { ... }
// The expression myMap.get(key) is annotated
```

You can annotate type parameters:

```
class MyContainer[@specialized T]
```

Annotations on an actual type are placed *after* the type, like this:

```
def country: String @Localized
```

Here, the `String` type is annotated. The method returns a localized string.

15.3 Annotation Arguments

Java annotations can have named arguments, such as

```
@Test(timeout = 100, expected = classOf[IOException])
```

However, if the argument name is `value`, it can be omitted. For example:

```
@Named("creds") var credentials: Credentials = _
// The value argument is "creds"
```

If the annotation has no arguments, the parentheses can be omitted:

```
@Entity class Credentials
```

Most annotation arguments have defaults. For example, the `timeout` argument of the JUnit `@Test` annotation has a default value of `0`, indicating no timeout. The `expected` argument has as default a dummy class to signify that no exception is expected. If you use

```
@Test def testSomeFeature() { ... }
```

this annotation is equivalent to

```
@Test(timeout = 0, expected = classOf[org.junit.Test.None])
def testSomeFeature() { ... }
```

Arguments of Java annotations are restricted to the following types:

- Numeric literals
- Strings
- Class literals
- Java enumerations
- Other annotations
- Arrays of the above (but not arrays of arrays)

Arguments of Scala annotations can be of arbitrary types, but only a couple of the Scala annotations take advantage of this added flexibility. For instance, the `@deprecatedName` annotation has an argument of type `Symbol`.

15.4 Annotation Implementations

I don't expect that many readers of this book will feel the urge to implement their own Scala annotations. The main point of this section is to be able to decipher the implementation of the existing annotation classes.

An annotation must extend the `Annotation` trait. For example, the `unchecked` annotation is defined as follows:

```
class unchecked extends annotation.Annotation
```

A type annotation must extend the `TypeAnnotation` trait:

```
class Localized extends StaticAnnotation with TypeConstraint
```



CAUTION: If you want to implement a new Java annotation, you need to write the annotation class in Java. You can, of course, use that annotation for your Scala classes.

Generally, an annotation describes the expression, variable, field, method, class, or type to which it is applied. For example, the annotation

```
def check(@NotNull password: String)
```

applies to the parameter variable `password`.

However, field definitions in Scala can give rise to multiple features in Java, all of which can potentially be annotated. For example, consider

```
class Credentials(@NotNull @BeanProperty var username: String)
```

Here, there are six items that can be annotation targets:

- The constructor parameter
- The private instance field
- The accessor method `username`
- The mutator method `username_=`
- The bean accessor `getUsername`
- The bean mutator `setUsername`

By default, constructor parameter annotations are only applied to the parameter itself, and field annotations are only applied to the field. The meta-annotations

`@param`, `@field`, `@getter`, `@setter`, `@beanGetter`, and `@beanSetter` cause an annotation to be attached elsewhere. For example, the `@deprecated` annotation is defined as:

```
@getter @setter @beanGetter @beanSetter
class deprecated(message: String = "", since: String = "")
  extends annotation.StaticAnnotation
```

You can also apply these annotations in an ad-hoc fashion:

```
@Entity class Credentials {
  @Id @beanGetter) @BeanProperty var id = 0
  ...
}
```

In this situation, the `@Id` annotation is applied to the Java `getId` method, which is a JPA requirement for property access.

15.5 Annotations for Java Features

The Scala library provides annotations for interoperating with Java. They are presented in the following sections.

15.5.1 Java Modifiers

Scala uses annotations instead of modifier keywords for some of the less commonly used Java features.

The `@volatile` annotation marks a field as volatile:

```
@volatile var done = false // Becomes a volatile field in the JVM
```

A volatile field can be updated in multiple threads.

The `@transient` annotation marks a field as transient:

```
@transient var recentLookups = new HashMap[String, String]
// Becomes a transient field in the JVM
```

A transient field is not serialized. This makes sense for cache data that need not be saved, or data that can easily be recomputed.

The `@strictfp` annotation is the analog of the Java `strictfp` modifier:

```
@strictfp def calculate(x: Double) = ...
```

This method does its floating-point calculations with IEEE double values, not using the 80 bit extended precision (which Intel processors use by default). The result is slower and less precise but more portable.

The `@native` annotation marks methods that are implemented in C or C++ code. It is the analog of the native modifier in Java.

```
@native def win32RegKeys(root: Int, path: String): Array[String]
```

15.5.2 Marker Interfaces

Scala uses annotations `@cloneable` and `@remote` instead of the `Cloneable` and `java.rmi.Remote` marker interfaces for cloneable and remote objects.

```
@cloneable class Employee
```

With serializable classes, you can use the `@SerialVersionUID` annotation to specify the serial version:

```
@SerialVersionUID(6157032470129070425L)
class Employee extends Person with Serializable
```



NOTE: For more information about Java concepts such as volatile fields, cloning, or serialization, see C. Horstmann, *Core Java®, Tenth Edition* (Prentice Hall, 2016).

15.5.3 Checked Exceptions

Unlike Scala, the Java compiler tracks checked exceptions. If you call a Scala method from Java code, its signature should include the checked exceptions that can be thrown. Use the `@throws` annotation to generate the correct signature. For example,

```
class Book {
    @throws(classOf[IOException]) def read(filename: String) { ... }
    ...
}
```

The Java signature is

```
void read(String filename) throws IOException
```

Without the `@throws` annotation, the Java code would not be able to catch the exception.

```
try { // This is Java
    book.read("war-and-peace.txt");
} catch (IOException ex) {
    ...
}
```

The Java compiler needs to know that the `read` method can throw an `IOException`, or it will refuse to catch it.

15.5.4 Variable Arguments

The `@varargs` annotation lets you call a Scala variable-argument method from Java. By default, if you supply a method such as

```
def process(args: String*)
```

the Scala compiler translates the variable argument into a sequence

```
def process(args: Seq[String])
```

That is cumbersome to use in Java. If you add `@varargs`,

```
@varargs def process(args: String*)
```

then a Java method

```
void process(String... args) // Java bridge method
```

is generated that wraps the `args` array into a `Seq` and calls the Scala method.

15.5.5 JavaBeans

You have seen the `@BeanProperty` annotation in Chapter 5. When you annotate a field with `@scala.reflect.BeanProperty`, the compiler generates JavaBeans-style getter and setter methods. For example,

```
class Person {  
    @BeanProperty var name : String = _  
}
```

generates methods

```
getName() : String  
setName(newValue : String) : Unit
```

in addition to the Scala getter and setter.

The `@BooleanBeanProperty` annotation generates a getter with an `is` prefix for a Boolean method.



NOTE: The annotations `@BeanDescription`, `@BeanDisplayName`, `@BeanInfo`, `@BeanInfoSkip` let you control some of the more obscure features of the JavaBeans specifications. Very few programmers need to worry about these. If you are among them, you'll figure out what to do from the Scaladoc descriptions.

15.6 Annotations for Optimizations

Several annotations in the Scala library let you control compiler optimizations. They are discussed in the following sections.

15.6.1 Tail Recursion

A recursive call can sometimes be turned into a loop, which conserves stack space. This is important in functional programming where it is common to write recursive methods for traversing collections.

Consider this method that computes the sum of a sequence of integers using recursion:

```
object Util {
    def sum(xs: Seq[Int]): BigInt =
        if (xs.isEmpty) 0 else xs.head + sum(xs.tail)
    ...
}
```

This method cannot be optimized because the last step of the computation is addition, not the recursive call. But a slight transformation can be optimized:

```
def sum2(xs: Seq[Int], partial: BigInt): BigInt =
    if (xs.isEmpty) partial else sum2(xs.tail, xs.head + partial)
```

The partial sum is passed as a parameter; call this method as `sum2(xs, 0)`. Since the *last* step of the computation is a recursive call to the same method, it can be transformed into a loop to the top of the method. The Scala compiler automatically applies the “tail recursion” optimization to the second method. If you try

```
sum(1 to 1000000)
```

you will get a stack overflow error (at least with the default stack size of the JVM), but

```
sum2(1 to 1000000, 0)
```

returns the sum 500000500000.

Even though the Scala compiler will try to use tail recursion optimization, it is sometimes blocked from doing so for nonobvious reasons. If you rely on the compiler to remove the recursion, you should annotate your method with `@tailrec`. Then, if the compiler cannot apply the optimization, it will report an error.

For example, suppose the `sum2` method is in a class instead of an object:

```
class Util {
  @tailrec def sum2(xs: Seq[Int], partial: BigInt): BigInt =
    if (xs.isEmpty) partial else sum2(xs.tail, xs.head + partial)
  ...
}
```

Now the program fails with an error message "could not optimize @tailrec annotated method sum2: it is neither private nor final so can be overridden". In this situation, you can move the method into an object, or you can declare it as private or final.



NOTE: A more general mechanism for recursion elimination is “trampolining”

A trampoline implementation runs a loop that keeps calling functions. Each function returns the next function to be called. Tail recursion is a special case where each function returns itself. The more general mechanism allows for mutual calls—see the example that follows.

Scala has a utility object called `TailCalls` that makes it easy to implement a trampoline. The mutually recursive functions have return type `TailRec[A]` and return either `done(result)` or `tailcall(fun)` where `fun` is the next function to be called. This needs to be a parameterless function that also returns a `TailRec[A]`. Here is a simple example:

```
import scala.util.control.TailCalls._
def evenLength(xs: Seq[Int]): TailRec[Boolean] =
  if (xs.isEmpty) done(true) else tailcall(oddLength(xs.tail))
def oddLength(xs: Seq[Int]): TailRec[Boolean] =
  if (xs.isEmpty) done(false) else tailcall(evenLength(xs.tail))
```

To obtain the final result from the `TailRec` object, use the `result` method:

```
evenLength(1 to 1000000).result
```

15.6.2 Jump Table Generation and Inlining

In C++ or Java, a `switch` statement can often be compiled into a jump table, which is more efficient than a sequence of `if/else` expressions. Scala attempts to generate jump tables for `match` clauses as well. The `@switch` annotation lets you check whether a Scala `match` clause is indeed compiled into one. Apply the annotation to the expression preceding a `match` clause:

```
(n: @switch) match {
  case 0 => "Zero"
  case 1 => "One"
  case _ => "?"
}
```

A common optimization is method inlining—replacing a method call with the method body. You can tag methods with `@inline` to suggest inlining, or `@noinline` to suggest not to inline. Generally, inlining is done in the JVM, whose “just in time” compiler does a good job without any annotations. The `@inline` and `@noinline` annotations let you direct the Scala compiler, in case you perceive the need to do so.

15.6.3 Eliding Methods

The `@elidable` annotation flags methods that can be removed in production code. For example,

```
@elidable(500) def dump(props: Map[String, String]) { ... }
```

If you compile with

```
scalac -Xelide-below 800 myprog.scala
```

then the method code will not be generated. The `elidable` object defines the following numerical constants:

- MAXIMUM or OFF = `Int.MaxValue`
- ASSERTION = 2000
- SEVERE = 1000
- WARNING = 900
- INFO = 800
- CONFIG = 700
- FINE = 500
- FINER = 400
- FINEST = 300
- MINIMUM or ALL = `Int.MinValue`

You can use one of these constants in the annotation:

```
import scala.annotation.elidable._

@elidable(FINE) def dump(props: Map[String, String]) { ... }
```

You can also use these names in the command line:

```
scalac -Xelide-below INFO myprog.scala
```

If you don’t specify the `-Xelide-below` flag, annotated methods with values below 1000 are elided, leaving SEVERE methods and assertions, but removing warnings.



NOTE: The levels ALL and OFF are potentially confusing. The annotation @elide(ALL) means that the method is always elided, and @elide(OFF) means that it is never elided. But -Xelide-below OFF means to elide everything, and -Xelide-below ALL means to elide nothing. That's why MAXIMUM and MINIMUM have been added.

The Predef object defines an elidable assert method. For example,

```
def makeMap(keys: Seq[String], values: Seq[String]) = {
  assert(keys.length == values.length, "lengths don't match")
  ...
}
```

If the method is called with mismatched arguments, the assert method throws an AssertionError with message assertion failed: lengths don't match.

To disable assertions, compile with -Xelide-below 2001 or -Xelide-below MAXIMUM. Note that by default assertions are *not* disabled. This is a welcome improvement over Java assertions.



CAUTION: Calls to elided methods are replaced with Unit objects. If you use the return value of an elided method, a ClassCastException is thrown. It is best to use the @elidable annotation only with methods that don't return a value.

15.6.4 Specialization for Primitive Types

It is inefficient to wrap and unwrap primitive type values—but in generic code, this often happens. Consider, for example,

```
def allDifferent[T](x: T, y: T, z: T) = x != y && x != z && y != z
```

If you call allDifferent(3, 4, 5), each integer is wrapped into a java.lang.Integer before the method is called. Of course, one can manually supply an overloaded version

```
def allDifferent(x: Int, y: Int, z: Int) = ...
```

as well as seven more methods for the other primitive types.

You can generate these methods automatically by annotating the type parameter with @specialized:

```
def allDifferent[@specialized T](x: T, y: T, z: T) = ...
```

You can restrict specialization to a subset of types:

```
def allDifferent[@specialized(Long, Double) T](x: T, y: T, z: T) = ...
```

In the annotation constructor, you can provide any subset of Unit, Boolean, Byte, Short, Char, Int, Long, Float, Double.

15.7 Annotations for Errors and Warnings

If you mark a feature with the @deprecated annotation, the compiler generates a warning whenever the feature is used. The annotation has two optional arguments, message and since.

```
@deprecated(message = "Use factorial(n: BigInt) instead")
def factorial(n: Int): Int = ...
```

The @deprecatedName is applied to a parameter, and it specifies a former name for the parameter.

```
def draw(@deprecatedName('sz) size: Int, style: Int = NORMAL)
```

You can still call `draw(sz = 12)` but you will get a deprecation warning.



NOTE: The constructor argument is a *symbol*—a name preceded by a single quote. Symbols with the same name are guaranteed to be unique. Therefore, comparing symbols is a bit more efficient than comparing strings. More importantly, there is a semantic distinction: A symbol denotes a name of some item in a program.

The @deprecatedInheritance and @deprecatedOverriding annotations generate warnings that inheriting from a class or overriding a method is now deprecated.

The @implicitNotFound and @implicitAmbiguous annotations generates meaningful error messages when an implicit value is not available or ambiguous. See Chapter 21 for details about implicits.

The @unchecked annotation suppresses a warning that a match is not exhaustive. For example, suppose we know that a given list is never empty:

```
(lst: @unchecked) match {
    case head :: tail => ...
}
```

The compiler won't complain that there is no `Nil` option. Of course, if `lst` is `Nil`, an exception is thrown at runtime.

The @uncheckedVariance annotation suppresses a variance error message. For example, it would make sense for `java.util.Comparator` to be contravariant. If `Student` is a subtype of `Person`, then a `Comparator[Person]` can be used when a `Comparator[Student]` is required. However, Java generics have no variance. We can fix this with the @uncheckedVariance annotation:

```
trait Comparator[-T] extends  
    java.lang.Comparator[T @uncheckedVariance]
```

Exercises

1. Write four JUnit test cases that use the `@Test` annotation with and without each of its arguments. Run the tests with JUnit.
2. Make an example class that shows every possible position of an annotation. Use `@deprecated` as your sample annotation.
3. Which annotations from the Scala library use one of the meta-annotations `@param`, `@field`, `@getter`, `@setter`, `@beanGetter`, or `@beanSetter`?
4. Write a Scala method `sum` with variable integer arguments that returns the sum of its arguments. Call it from Java.
5. Write a Scala method that returns a string containing all lines of a file. Call it from Java.
6. Write a Scala object with a volatile Boolean field. Have one thread sleep for some time, then set the field to true, print a message, and exit. Another thread will keep checking whether the field is true. If so, it prints a message and exits. If not, it sleeps for a short time and tries again. What happens if the variable is not volatile?
7. Give an example to show that the tail recursion optimization is not valid when a method can be overridden.
8. Add the `allDifferent` method to an object, compile and look at the bytecode. What methods did the `@specialized` annotation generate?
9. The `Range.foreach` method is annotated as `@specialized(Unit)`. Why? Look at the bytecode by running

```
javap -classpath /path/to/scala/lib/scala-library.jar  
      scala.collection.immutable.Range
```

and consider the `@specialized` annotations on `Function1`. Click on the `Function1.scala` link in Scaladoc to see them.

10. Add `assert(n >= 0)` to a `factorial` method. Compile with assertions enabled and verify that `factorial(-1)` throws an exception. Compile without assertions. What happens? Use `javap` to check what happened to the assertion call.

XML Processing

Topics in This Chapter A2

- 16.1 XML Literals — page 230
- 16.2 XML Nodes — page 230
- 16.3 Element Attributes — page 232
- 16.4 Embedded Expressions — page 233
- 16.5 Expressions in Attributes — page 234
- 16.6 Uncommon Node Types — page 235
- 16.7 XPath-like Expressions — page 235
- 16.8 Pattern Matching — page 237
- 16.9 Modifying Elements and Attributes — page 238
- 16.10 Transforming XML — page 239
- 16.11 Loading and Saving — page 239
- 16.12 Namespaces — page 242
- Exercises — page 243

Chapter 16

Scala has built-in support for XML literals that makes it easy to generate XML fragments in your programs. The Scala library includes support for common XML processing tasks. In this chapter, you will learn how to put these features to use for reading, analyzing, creating, and writing XML.



NOTE: The XML support in Scala is brilliant because it makes it possible to slice and dice XML data easily and conveniently, in the REPL and in Scala programs. It is also flawed because of some unfortunate design decisions and lack of maintenance. It is difficult to remedy these issues because XML is so tightly integrated into Scala. A future version of Scala will likely abandon this tight integration and instead rely on string interpolation and third-party libraries. But for now, we can enjoy a uniquely powerful way of processing XML.



NOTE: The API documentation for Scala XML is at www.scala-lang.org/api/current/scala-xml.

The key points of this chapter are:

- XML literals <like>this</like> are of type `NodeSeq`.
- You can embed Scala code inside XML literals.
- The `child` property of a `Node` yields the child nodes.

- The `attributes` property of a `Node` yields a `MetaData` object containing the node attributes.
- The `\` and `\\"` operators carry out XPath-like matches.
- You can match node patterns with XML literals in case clauses.
- Use the `RuleTransformer` with `RewriteRule` instances to transform descendants of a node.
- The `XML` object interfaces with Java XML methods for loading and saving.
- The `ConstructingParser` is an alternate parser that preserves comments and `CDATA` sections.

16.1 XML Literals

In Scala, you can define XML *literals*, simply by using the XML code:

```
val doc = <html><head><title>Fred's Memoirs</title></head><body>...</body></html>
```

In this case, `doc` becomes a value of type `scala.xml.Elem`, representing an XML element.

An XML literal can also be a sequence of nodes. For example,

```
val items = <li>Fred</li><li>Wilma</li>
```

yields a `scala.xml.NodeSeq`. We will discuss the `Elem` and `NodeSeq` classes in the next section.



CAUTION: Sometimes, the compiler suspects XML literals when none are intended. For example,

```
val (x, y) = (1, 2)
x < y // OK
x <y // Error—unclosed XML literal
```

In this case, the remedy is to add a space after the `<`.

16.2 XML Nodes

The `Node` class is the ancestor of all XML node types. Its two most important subclasses are `Text` and `Elem`. Figure 16–1 shows the complete hierarchy.

The `Elem` class describes an XML element, such as

```
val elem = <a href="http://scala-lang.org">The <em>Scala</em> language</a>
```

The `label` property yields the tag name (here, "a"), and `child` is the child node sequence (two `Text` nodes and an `Elem` node in this example).



CAUTION: Unfortunately, unlike a DOM node, a Scala Node retains no information about its parent.

Node sequences are of type `NodeSeq`, a subtype of `Seq[Node]` that adds support for XPath-like operators (see Section 16.7, “XPath-like Expressions,” on page 235). You can use any of the `Seq` operations described in Chapter 13 with XML node sequences. To traverse a sequence, simply use a `for` loop, for example:

```
for (n <- elem.child) process n
```



NOTE: The `Node` class extends `NodeSeq`. A single node is a sequence of length 1.

This is supposed to make it easier to deal with functions that can return a single node or a sequence. (It actually creates as many problems as it solves, so I don’t recommend using this trick in your own designs.)

There are also node classes for XML comments (`<!-- ... -->`), entity references (`&...;`), and processing instructions (`<? ... ?>`). Figure 16–1 shows all node types.

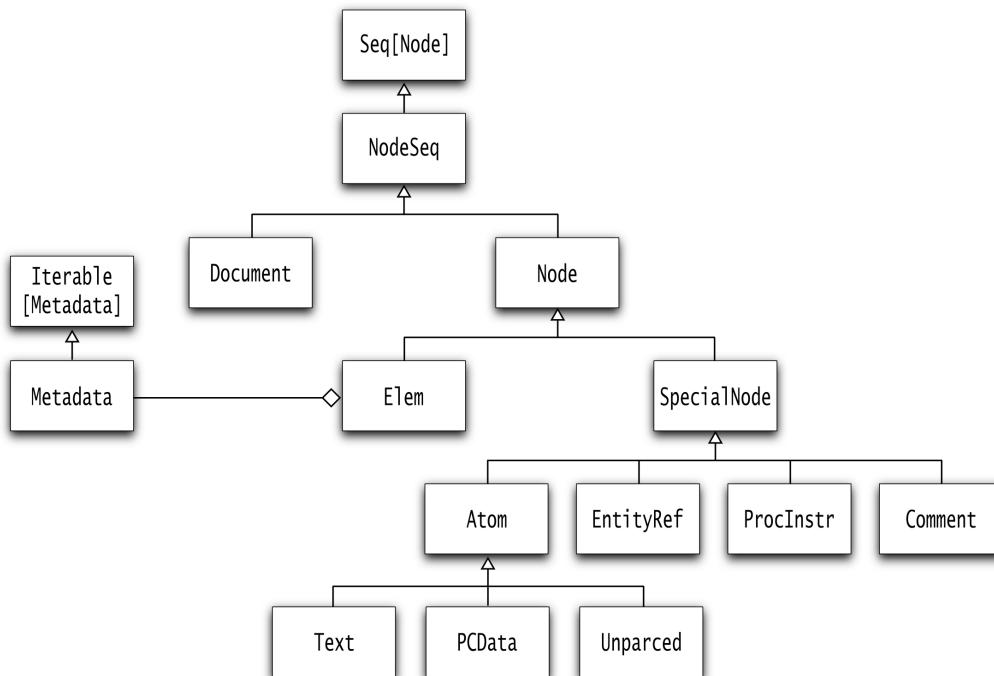


Figure 16–1 XML node types

If you build node sequences programmatically, you can use a `NodeBuffer`, a subclass of `ArrayBuffer[Node]`.

```
val items = new NodeBuffer
items += <li>Fred</li>
items += <li>Wilma</li>
val nodes: NodeSeq = items
```



CAUTION: A NodeBuffer is a Seq[Node]. It can be implicitly converted to a NodeSeq. Once this conversion has occurred, you need to be careful not to mutate the node buffer any longer since XML node sequences are supposed to be immutable.

16.3 Element Attributes

To process the attribute keys and values of an element, use the attributes property. It yields an object of type MetaData which is almost, but not quite, a Map from attribute keys to values. You can use the () operator to access the value for a given key:

```
val elem = <a href="http://scala-lang.org">The Scala language</a>
val url = elem.attributes("href")
```

Unfortunately, this yields a node sequence, not a string, because the attribute might contain entity references. For example, consider

```
val image = 
val alt = image.attributes("alt")
```

Here, the value for the key "alt" is a node sequence consisting of a text node for "San Jos", an EntityRef for é, and another text node for " State University Logo".

Why not resolve the entity reference? There is no way to know what é means. In XHTML it means é (the code for é), but in another document type it can be defined as something else.



TIP: If you find it inconvenient to deal with entity references in XML literals, you can use character references instead: .

If you are certain you don't have unresolved entities in your attributes, you can simply call the text method to turn the node sequence into a string:

```
val url = elem.attributes("href").text
```

If an attribute is not present, the () operator returns null. If you dislike working with null, use the get method, which returns an Option[Seq[Node]].

Unfortunately, the `MetaData` class has no `getOrElse` method, but you can apply `getOrElse` to the `Option` that `get` returns:

```
val url = elem.attributes.get("href").getOrElse(Text(""))
```

To iterate over all attributes, use

```
for (attr <- elem.attributes)
  process attr.key and attr.value.text
```

Alternatively, call the `asAttrMap` method:

```
val image = 
val map = image.attributes.asAttrMap // Map("alt" -> "TODO", "src" -> "hamster.jpg")
```

16.4 Embedded Expressions

You can include blocks of Scala code inside XML literals to dynamically compute items. For example:

```
<ul><li>{items(0)}</li><li>{items(1)}</li></ul>
```

Each block is evaluated, and its result is spliced into the XML tree.

If the block yields a node sequence, the nodes are simply added to the XML. Everything else is turned into an `Atom[T]`, a container for a type `T`. In this way, you can store arbitrary values in an XML tree. Retrieve the value from an `Atom` node with the `data` property.

In many cases, one doesn't care about retrieving the items from the atoms. When the XML document gets saved, each atom is turned into a string, by calling `toString` on the `data` property.



CAUTION: Embedded strings do not get turned into `Text` nodes but into `Atom[String]` nodes. That is not quite the same—`Text` is a subclass of `Atom[String]`. It doesn't matter when saving a document. But if you later do pattern matching on `Text` nodes, the match will fail. In that case, you should insert `Text` nodes instead of strings:

```
<li>{Text("Another item")}</li>
```

Not only can you nest Scala inside XML, but the nested Scala code can again contain XML literals. For example, if you have a list of items, you will want to place each item inside an `li` element:

```
<ul>{for (i <- items) yield <li>{i}</li>}</ul>
```

We have a Scala block `{...}` inside the `ul` element. That block yields a sequence of XML expressions.

```
for (i <- items) yield an XML literal
```

That XML literal `...` contains another Scala block!

```
<li>{i}</li>
```

It's Scala inside XML inside Scala inside XML. The mind reels if you think about it. But if you don't, it's a very natural construction: make a `ul` that contains an `li` for every element of `items`.



NOTE: To place an opening or closing brace into an XML literal, use two braces:

```
<h1>The Natural Numbers {{1, 2, 3, ...}}</h1>
```

This produces

```
<h1>The Natural Numbers {1, 2, 3, ...}</h1>
```

16.5 Expressions in Attributes

You can compute attribute values with Scala expressions, for example:

```
<img src={makeURL(fileName)} />
```

Here, the `makeURL` function returns a string that becomes the attribute value.



CAUTION: Braces inside quoted strings are *not* evaluated. For example,

```

```

sets the `src` attribute to the string "`{makeURL(fileName)}`", which is probably not what you want.

The embedded block can also yield a node sequence. This is potentially useful if you want to include entity references or atoms in an attribute:

```
<a id={new Atom(1)} ... />
```

If the embedded block returns `null` or `None`, the attribute is not set. For example:

```
<img alt={if (description == "TODO") null else description} ... />
```

If `description` is the string "TODO" or `null`, the element will have no `alt` attribute.

You can get the same effect with an `Option[Seq[Node]]`. For example,

```
<img alt={if (description == "TODO" || description == null) None  
else Some(Text(description))} ... />
```



CAUTION: It is a syntax error if the block yields something other than a String, a Seq[Node], or an Option[Seq[Node]]. This is inconsistent with blocks inside elements, where the result would be wrapped in an Atom. If you want an atom in an attribute value, you must construct it yourself.

16.6 Uncommon Node Types

Sometimes, you need to include non-XML text into an XML document. A typical example is JavaScript code in an XHTML page. You can use CDATA markup in XML literals:

```
val js = <script><![CDATA[if (temp < 0) alert("Cold!")]]></script>
```

However, the parser does not retain the fact that the text was marked up with CDATA. What you get is a node with a Text child. If you want the CDATA in the output, include a PCDATA node, like this:

```
val code = """if (temp < 0) alert("Cold!")"""
val js = <script>{PCData(code)}</script>
```

You can include arbitrary text in an Unparsed node. It is saved as is. You can generate such nodes as literals or programmatically:

```
val n1 = <xm1:unparsed>&</xm1:unparsed>
val n2 = Unparsed("&")
```

I don't recommend this technique since you can easily end up with malformed XML.

Finally, you can group a node sequence into a single "group" node.

```
val g1 = <xm1:group><li>Item 1</li><li>Item 2</li></xm1:group>
val g2 = Group(Seq(<li>Item 1</li>, <li>Item 2</li>))
```

Group nodes are "ungrouped" when you iterate over them. Contrast:

```
val items = <li>Item 1</li><li>Item 2</li>
for (n <- <xm1:group>{items}</xm1:group>) yield n
  // Yields two li elements
for (n <- <ol>{items}</ol>) yield n
  // Yields one ol element
```

16.7 XPath-like Expressions

The NodeSeq class provides methods that resemble the / and // operators in XPath (XML Path Language, www.w3.org/TR/xpath). Since // denotes comments and is therefore not a valid operator, Scala uses \ and \\ instead.

The \ operator locates direct descendants of a node or node sequence. For example,

```
val list = <dl><dt>Java</dt><dd>Gosling</dd><dt>Scala</dt><dd>Odersky</dd></dl>
val languages = list \ "dt"
```

sets languages to a node sequence containing <dt>Java</dt> and <dt>Scala</dt>.

A wildcard matches any element. For example,

```
doc \ "body" \ "_" \ "li"
```

finds all li elements, whether they are contained in a ul, an ol, or any other element inside the body.

The \\ operator locates descendants at any depth. For example,

```
doc \\ "img"
```

locates all img elements anywhere inside the doc.

A string starting with @ locates attributes. For example,

```
img \ "@alt"
```

returns the value of the alt attribute of the given node, and

```
doc \\ "@alt"
```

locates all alt attributes of any elements inside doc.



NOTE: There is no wildcard notation for attributes; img \ "@_" does *not* return all attributes.



CAUTION: Unlike XPath, you cannot use a single \ to extract attributes from multiple nodes. For example, doc \\ "img" \ "@src" will not work if the document contains more than one img element. Use doc \\ "img" \\ "@src" instead.

The result of \ or \\ is a node sequence. It might be a single node, but unless you know that for sure, you should traverse the sequence. For example,

```
for (n <- doc \\ "img") process n
```

If you simply call text on a result of \ or \\, all texts of the result sequence will be concatenated. For example,

```
() \\ "@src).text
```

returns a string "hamster.jpgfrog.jpg".

16.8 Pattern Matching

You can use XML literals in pattern matching expressions. For example,

```
node match {
    case <img/> => ...
    ...
}
```

The first match succeeds if `node` is an `img` element with *any* attributes and *no* child elements.

To deal with child elements is a little tricky. You can match a single child with

```
case <li>{_}</li> => ...
```

However, if `li` has more than one child, for example `An important item`, then the match fails. To match any number of items, use

```
case <li>{_*}</li> => ...
```

Note the braces—they might remind you of the embedded code notation for XML literals. However, inside XML patterns, braces indicate code patterns, not code to be evaluated.

Instead of the wildcard indicators, you can use variable names. The match is bound to the variable.

```
case <li>{child}</li> => child.text
```

To match a text node, use a case class match like this:

```
case <li>{Text(item)}</li> => item
```

To bind a variable to a node sequence, use the following syntax:

```
case <li>{children @ _*}</li> => for (c <- children) yield c
```



CAUTION: In such a match, `children` is a `Seq[Node]` and *not* a `NodeSeq`.

You can only use one node in the case clause. For example, the following is not legal:

```
case <p>{_*}</p><br/> => ... // Not legal
```

XML patterns can't have attributes.

```
case <img alt="TODO"/> => ... // Not legal
```

To match an attribute, use a guard:

```
case n @ <img/> if (n.attributes("alt").text == "TODO") => ...
```

16.9 Modifying Elements and Attributes

In Scala, XML nodes and node sequences are immutable. If you want to edit a node, you have to create a copy of it, making any needed changes and copying what hasn't changed.

To copy an `Elem` node, use the `copy` method. It has five named parameters: the familiar `label`, `attributes`, and `child`, as well as `prefix` and `scope` which are used for namespaces (see Section 16.12, “Namespaces,” on page 242). Any parameters that you don’t specify are copied from the original element. For example,

```
val list = <ul><li>Fred</li><li>Wilma</li></ul>
val list2 = list.copy(label = "ol")
```

makes a copy of `list`, changing the label from `ul` to `ol`. The children are shared, but that’s OK since node sequences are immutable.

To add a child, make a call to `copy` like this:

```
list.copy(child = list.child ++ <li>Another item</li>)
```

To add or change an attribute, use the `%` operator:

```
val image = 
val image2 = image % Attribute(null, "alt", "An image of a hamster", Null)
```

The first argument is the namespace. The last one is a list of additional metadata. Just like `Node` extends `NodeSeq`, the `Attribute` trait extends `MetaData`. To add more than one attribute, you can chain them like this:

```
val image3 = image % Attribute(null, "alt", "An image of a frog",
    Attribute(null, "src", "frog.jpg", Null))
```



CAUTION: Here, `scala.xml.Null` is an empty attribute list. It is *not* the `scala.Null` type.

Adding an attribute with the same key replaces the existing one. The `image3` element has a single attribute with key “`src`”; its value is “`frog.jpg`”.

16.10 Transforming XML

Sometimes, you need to rewrite all descendants that match a particular condition. The XML library provides a `RuleTransformer` class that applies one or more `RewriteRule` instances to a node and its descendants.

For example, suppose you want to change all `ul` nodes in a document to `ol`. Define a `RewriteRule` that overrides the `transform` method:

```
val rule1 = new RewriteRule {
  override def transform(n: Node) = n match {
    case e @ <ul>{_*}</ul> => e.asInstanceOf[Elem].copy(label = "ol")
    case _ => n
  }
}
```

Then you can transform a tree with the command

```
val transformed = new RuleTransformer(rule1).transform(root)
```

You can supply any number of rules in the constructor of the `RuleTransformer`:

```
val transformer = new RuleTransformer(rule1, rule2, rule3);
```

The `transform` method traverses the descendants of a node, applies all rules, and returns the transformed tree.

16.11 Loading and Saving

To load an XML document from a file, call the `loadFile` method of the `XML` object:

```
import scala.xml.XML
val root = XML.loadFile("myfile.xml")
```

You can also load from a `java.io.InputStream`, a `java.io.Reader`, or a URL:

```
val root2 = XML.load(new FileInputStream("myfile.xml"))
val root3 = XML.load(new InputStreamReader(
  new FileInputStream("myfile.xml"), "UTF-8"))
val root4 = XML.load(new URL("http://horstmann.com/index.html"))
```

The document is loaded using the standard SAX parser from the Java library. Unfortunately, the Document Type Definition (DTD) is not made available.



CAUTION: This parser suffers from a problem that is inherited from the Java library. It does not read DTDs from a local catalog. In particular, fetching an XHTML file can take a very long time, or fail altogether, as the parser retrieves the DTDs from the www.w3.org site.

To use a local catalog, you need the `CatalogResolver` class that is available in the `com.sun.org.apache.xml.internal.resolver.tools` package of the JDK or, if you are squeamish about using a class outside the official API, in the Apache Commons Resolver project (<http://xml.apache.org/commons/components/resolver/resolver-article.html>).

Unfortunately, the `XML` object has no API for installing an entity resolver. Here is how you can do it through the back door:

```
val res = new CatalogResolver
val doc = new factory.XMLLoader[Elem] {
    override def adapter = new parsing.NoBindingFactoryAdapter() {
        override def resolveEntity(publicId: String, systemId: String) = {
            res.resolveEntity(publicId, systemId)
        }
    }
}.load(new URL("http://horstmann.com/index.html"))
```

There is another parser that preserves comments, CDATA sections, and, optionally, whitespace:

```
import scala.xml.parsing.ConstructingParser
import java.io.File
val parser = ConstructingParser.fromFile(new File("myfile.xml"), preserveWS = true)
val doc = parser.document
val root = doc.documentElement
```

Note that the `ConstructingParser` returns a node of type `Document`. Call its `docElem` method to get the document root.

If your document has a DTD and you need it (for example, when saving the document), it is available as `doc.dtd`.



CAUTION: By default, the `ConstructingParser` does not resolve entities but converts them into useless comments, such as

```
<!-- unknown entity nbsp; -->
```

If you happen to be reading an XHTML file, you can use the `XhtmlParser` subclass:

```
val parser = new XhtmlParser(scala.io.Source.fromFile("myfile.html"))
val doc = parser.initialize.document
```

Otherwise, you need to add entities to the parser's entity map. For example,

```
parser.ent ::= List(
  "nbsp" -> ParsedEntityDecl("nbsp", IntDef("\u00A0")),
  "acute" -> ParsedEntityDecl("acute", IntDef("\u00E9")))
```

To save an XML document to a file, use the `save` method:

```
XML.save("myfile.xml", root)
```

This method takes three optional parameters:

- `enc` specifies the character encoding (default is "ISO-8859-1").
- `xmlDecl` specifies whether to emit an XML declaration (`<?xml...?>`) at the beginning of the output (default is `false`).
- `doctype` is an object of the case class `scala.xml.dtd.DocType` (default is `null`).

For example, to write an XHTML file, you might use

```
XML.save("myfile.xhtml", root,
  enc = "UTF-8",
  xmlDecl = true,
  doctype = DocType("html",
    PublicID("-//W3C//DTD XHTML 1.0 Strict//EN",
      "http://www.w3.org/TR/xhtml1/DTD/xhtml1-strict.dtd"),
    Nil))
```

The last parameter in the `DocType` constructor lets you specify internal DTD declarations—an obscure XML feature that I won't discuss here.

You can also save to a `java.io.Writer`, but then you must specify all parameters.

```
XML.write(writer, root, "UTF-8", false, null)
```



NOTE: When saving an XML file, elements without content are not written with self-closing tags. For example:

```
</img>
```

If you prefer

```

```

use

```
val str = xml.Utility.toXML(node, minimizeTags = true)
```



TIP: If you want your XML code to line up prettily, use the PrettyPrinter class:

```
val printer = new PrettyPrinter(width = 100, step = 4)
val str = printer.formatNodes(nodeSeq)
```

16.12 Namespaces

In XML, namespaces are used to avoid name clashes, similar to packages in Java or Scala. However, an XML namespace is a URI (and usually a URL), such as

<http://www.w3.org/1999/xhtml>

The `xmllns` attribute declares a namespace, for example:

```
<html xmllns="http://www.w3.org/1999/xhtml">
  <head>...</head>
  <body>...</body>
</html>
```

The `html` element and its descendants (`head`, `body`, and so on) are placed in this namespace.

A descendant can introduce its own namespace, for example:

```
<svg xmllns="http://www.w3.org/2000/svg" width="100" height="100">
  <rect x="25" y="25" width="50" height="50" fill="#ff0000"/>
</svg>
```

In Scala, each element has a `scope` property of type `NamespaceBinding`. The `uri` property of that class yields the namespace URI.

If you want to mix elements from multiple namespaces, it is tedious to work with namespace URLs. An alternative is a *namespace prefix*. For example, the tag

```
<html xmllns="http://www.w3.org/1999/xhtml"
      xmllns:svg="http://www.w3.org/2000/svg">
```

introduces the prefix `svg` for the namespace `http://www.w3.org/2000/svg`. All elements prefixed with `svg:` belong to that namespace. For example,

```
<svg:svg width="100" height="100">
  <svg:rect x="25" y="25" width="50" height="50" fill="#ff0000"/>
</svg:svg>
```

As mentioned in Section 16.9, “Modifying Elements and Attributes,” on page 238, each `Elem` object has `prefix` and `scope` values. The parser automatically computes these values. To find out the namespace of an element, look at the `scope.uri` value. If there are multiple namespaces, they are chained together through `scope.parent` links. This method retrieves them all:

```
def namespaces(node: Node) = {
  def namespaces(scope: NamespaceBinding): List[(String, String)] =
    if (scope == null) List()
    else namespaces(scope.parent) :+ ((scope.prefix, scope.uri))
  namespaces(node.scope)
}
```

To get the namespace of an attribute, use the `prefixedKey` method.

When you produce XML elements programmatically, you need to set prefixes and scopes. For example,

```
val scope = new NamespaceBinding("svg", "http://www.w3.org/2000/svg", TopScope)
val attrs = Attribute(null, "width", "100",
  Attribute(null, "height", "100", Null))
val elem = Elem(null, "body", Null, TopScope,
  Elem("svg", "svg", attrs, scope))
```

Exercises

1. What is `<fred/>(0)? <fred/>(0)(0)?` Why?
2. What is the result of

```
<ul>
  <li>Opening bracket: [</li>
  <li>Closing bracket: ]</li>
  <li>Opening brace: {</li>
  <li>Closing brace: }</li>
</ul>
```

How do you fix it?

3. Contrast

```
<li>Fred</li> match { case <li>{Text(t)}</li> => t }
```

and

```
<li>{"Fred"}</li> match { case <li>{Text(t)}</li> => t }
```

Why do they act differently?

4. Read an XHTML file and print all `img` elements that don't have an `alt` attribute.
5. Print the names of all images in an XHTML file. That is, print all `src` attribute values inside `img` elements.
6. Read an XHTML file and print a table of all hyperlinks in the file, together with their URLs. That is, print the child text and the `href` attribute of each `a` element.
7. Write a function that has a parameter of type `Map[String, String]` and returns a `dl` element with a `dt` for each key and `dd` for each value. For example,

```
Map("A" -> "1", "B" -> "2")
```

should yield `<dl><dt>A</dt><dd>1</dd><dt>B</dt><dd>2</dd></dl>`.

8. Write a function that takes a `dl` element and turns it into a `Map[String, String]`. This function should be the inverse of the function in the preceding exercise, provided all `dt` children are distinct.
9. Transform an XHTML document by adding an `alt="TODO"` attribute to all `img` elements without an `alt` attribute, preserving everything else.
10. Write a function that reads an XHTML document, carries out the transformation of the preceding exercise, and saves the result. Be sure to preserve the DTD and any CDATA sections.

This page intentionally left blank

Futures

Topics in This Chapter A2

- 17.1 Running Tasks in the Future — page 248
- 17.2 Waiting for Results — page 250
- 17.3 The Try Class — page 251
- 17.4 Callbacks — page 251
- 17.5 Composing Future Tasks — page 252
- 17.6 Other Future Transformations — page 255
- 17.7 Methods in the Future Object — page 256
- 17.8 Promises — page 258
- 17.9 Execution Contexts — page 260
- Exercises — page 260

Chapter

17

Writing concurrent applications that work correctly and with high performance is very challenging. The traditional approach, in which concurrent tasks have side effects that mutate shared data, is tedious and error-prone. Scala encourages you to think of a computation in a functional way. A computation yields a value, sometime in the future. As long as the computations don't have side effects, you can let them run concurrently and combine the results when they become available. In this chapter, you will see how to use the Future and Promise traits to organize such computations.

The key points of this chapter are:

- A block of code wrapped in a Future { ... } executes concurrently.
- A future succeeds with a result or fails with an exception.
- You can wait for a future to complete, but you don't usually want to.
- You can use callbacks to get notified when a future completes, but that gets tedious when chaining callbacks.
- Use methods such as `map`/`flatMap`, or the equivalent for expressions, to compose futures.
- A promise has a future whose value can be set (once), which gives added flexibility for implementing tasks that produce results.
- Pick an execution context that is suitable for the concurrent workload of your computation.

17.1 Running Tasks in the Future

The `scala.concurrent.Future` object can execute a block of code “in the future.”

```
import java.time._  
import scala.concurrent._  
import ExecutionContext.Implicits.global  
  
Future {  
    Thread.sleep(10000)  
    println(s"This is the future at ${LocalTime.now}")  
}  
println(s"This is the present at ${LocalTime.now}")
```

When running this code, a line similar to the following is printed:

This is the present at 13:01:19.400

About ten seconds later, a second line appears:

This is the future at 13:01:29.140

When you create a `Future`, its code is run on some thread. One could of course create a new thread for each task, but thread creation is not free. It is better to keep some pre-created threads around and use them to execute tasks as needed. A data structure that assigns tasks to threads is usually called a *thread pool*. In Java, the `Executor` interface describes such a data structure. Scala uses the `ExecutionContext` trait instead.

Each `Future` must be constructed with a reference to an `ExecutionContext`. The simplest way is to import

```
import ExecutionContext.Implicits.global
```

Then the tasks execute on a global thread pool. This is fine for demos, but in a real program, you should make another choice if your tasks block. See Section 17.9, “Execution Contexts,” on page 260 for more information.

When you construct multiple futures, they can execute concurrently. For example, try running

```
Future { for (i <- 1 to 100) { print("A"); Thread.sleep(10) } }  
Future { for (i <- 1 to 100) { print("B"); Thread.sleep(10) } }
```

You will get an output that looks somewhat like

ABABABABABABABABABABABABA...AABABBBABABABABABABBBBBBBBBBBBBB

A future can—and normally will—have a result:

```
val f = Future {
    Thread.sleep(10000)
    42
}
```

When you evaluate `f` in the REPL immediately after the definition, you will get this output:

```
res12: scala.concurrent.Future[Int] = Future(<not completed>)
```

Wait ten seconds and try again:

```
res13: scala.concurrent.Future[Int] = Future(Success(42))
```

Alternatively, something bad may happen in the future:

```
val f2 = Future {
    if (LocalTime.now.getHour > 12)
        throw new Exception("too late")
    42
}
```

If it is after noon, the task terminates with an exception. In the REPL, you will see

```
res14: scala.concurrent.Future[Int] = Future(Failure(java.lang.Exception: too late))
```

Now you know what a Future is. It is an object that will give you a result (or failure) at some point in the future. In the next section, you will see one way of harvesting the result of a Future.



NOTE: In the Play web framework, you are encouraged to return Future objects in the “action” methods that react to web requests. Then you don’t have to worry how the results are harvested—that’s the job of the framework.



NOTE: The `java.util.concurrent` package has a Future interface that is much more limited than the Scala Future trait. A Scala future is equivalent to the `CompletionStage` interface in Java 8.



NOTE: The Scala language imposes no restrictions on what you can do in concurrent tasks. However, you should stay away from computations with side effects. It is best if you don’t increment shared counters—even atomic ones. Don’t populate shared maps—even threadsafe ones. Instead, have each future compute a value. Then you can combine the computed values after all contributing futures have completed. That way, each value is only owned by one task at a time, and it is easy to reason about the correctness of the computation.

17.2 Waiting for Results

When you have a Future, you can use the `isCompleted` method to check whether it is completed. But of course you don't want to wait for completion in a loop.

You can make a blocking call that waits for the result.

```
import scala.concurrent.duration._
val f = Future { Thread.sleep(10000); 42 }
val result = Await.result(f, 10.seconds)
```

The call to `Await.result` blocks for ten seconds and then yields the result of the future.

The second argument of the `Await.result` method is a `Duration` object. Importing `scala.concurrent.duration._` enables conversion methods from integers to `Duration` objects, called `seconds`, `millis`, and so on.

If the task is not ready by the allotted time, the `Await.ready` method throws a `TimeoutException`.

If the task throws an exception, it is rethrown in the call to `Await.result`. To avoid that, you can call `Await.ready` and then get the result.

```
val f = Future { ... }
Await.ready(f, 10.seconds)
val Some(t) = f.value
```

The `value` method returns an `Option[Try[T]]`, which is `None` when the future is not completed and `Some(t)` when it is. Here, `t` is an object of the `Try` class, which holds either the result or the exception that caused the task to fail. In our situation, the `value` method is only executed if the future is completed, so we can use an extractor to get the `Try` object. You will see how to look inside it in the next section.



NOTE: In practice, you won't use the `Await.result` or `Await.ready` methods much. You run tasks concurrently when they are time-consuming and your program can do something more useful than waiting for the result. Section 17.4, "Callbacks," on page 251 shows you how you can harvest the results without blocking.



CAUTION: In this section, we used the `result` and `ready` methods of the `Await` object. The `Future` class also has `result` and `ready` methods, but you should not call them. If the execution context uses a small number of threads (which is the case for the default fork-join pool), you don't want them all to block. Unlike the `Future` methods, the `Await` methods notify the execution context so that it can adjust the pooled threads.



NOTE: Not all exceptions that occur during execution of the future are stored in the result. Virtual machine errors and the `InterruptedException` are allowed to propagate in the usual way.

17.3 The Try Class

A `Try[T]` instance is either a `Success(v)`, where `v` is a value of type `T` or a `Failure(ex)`, where `ex` is a `Throwable`. One way of processing it is with a `match` statement.

```
t match {
  case Success(v) => println(s"The answer is $v")
  case Failure(ex) => println(ex.getMessage)
}
```

Alternatively, you can use the `isSuccess` or `isFailure` methods to find out whether the `Try` object represents success or failure. In the case of success, you can obtain the value with the `get` method:

```
if (t.isSuccess) println(s"The answer is ${t.get}")
```

To get the exception in case of failure, first apply the `failed` method which turns the failed `Try[T]` object into a `Try[Throwable]` wrapping the exception. Then call `get` to get the exception object.

```
if (t.isFailure) println(t.failed.get.getMessage)
```

You can also turn a `Try` object into an `Option` with the `toOption` method if you want to pass it on to a method that expects an option. This turns `Success` into `Some` and `Failure` into `None`.

To construct a `Try` object, call `Try(block)` with some block of code. For example,

```
val result = Try(str.toInt)
```

is either a `Success` object with the parsed integer, or a `Failure` wrapping a `NumberFormatException`.

There are several methods for composing and transforming `Try` objects. However, analogous methods exist for futures, where they are more commonly used. You will see how to work with multiple futures in Section 17.5, “Composing Future Tasks,” on page 252. At the end of that section, you will see how those techniques apply to `Try` objects.

17.4 Callbacks

As already mentioned, one does not usually use a blocking wait to get the result of a future. For better performance, the future should report its result to a callback function.

This is easy to arrange with the `onComplete` method.

```
f.onComplete(t => ...)
```

When the future has completed, either successfully or with a failure, it calls the given function with a `Try` object.

You can then react to the success or failure, for example by passing a `match` function to the `onComplete` method.

```
val f = Future { Thread.sleep(10000)
  if (random() < 0.5) throw new Exception
  42
}
f.onComplete {
  case Success(v) => println(s"The answer is $v")
  case Failure(ex) => println(ex.getMessage)
}
```

By using a callback, we avoid blocking. Unfortunately, we now have another problem. In all likelihood, the long computation in one `Future` task will be followed by another computation, and another. It is possible to nest callbacks within callbacks, but it is profoundly unpleasant. (This technique is sometimes called “callback hell”.)

A better approach is to think of futures as entities that can be composed, similar to functions. You compose two functions by calling the first one, then passing its result to the second one. In the next section, you will see how to do the same with futures.



NOTE: There are callback methods `onSuccess` and `onFailure` that are only called on success or failure of a future. However, these are deprecated because they are even bigger contributors to callback hell.

17.5 Composing Future Tasks

Suppose we need to get some information from two web services and then combine the two. Each task is long-running and should be executed in a `Future`. It is possible to link them together with callbacks:

```
val future1 = Future { getData1() }
val future2 = Future { getData2() }
```

```

future1 onComplete {
  case Success(n1) =>
    future2 onComplete {
      case Success(n2) => {
        val n = n1 + n2
        println(s"Result: $n")
      }
      case Failure(ex) => ...
    }
  case Failure(ex) => ...
}

```

Even though the callbacks are ordered sequentially, the tasks run concurrently. Each task starts after the `Future.apply` method executes or soon afterwards. We don't know which of `f1` and `f2` completes first, and it doesn't matter. We can't process the result until both tasks complete. Once `f1` completes, its completion handler registers a completion handler on `f2`. If `f2` has already completed, the second handler is called right away. Otherwise, it is called when `f2` finally completes.

Even though this chaining of the futures works, it looks very messy, and it will look worse with each additional level of processing.

Instead of nesting callbacks, we will use an approach that you already know from working with Scala collections. Think of a `Future` as a collection with (hopefully, eventually) one element. You know how to transform the values of a collection—with `map`:

```

val future1 = Future { getData1() }
val combined = future1.map(n1 => n1 + getData2())

```

Here `future1` is a `Future[Int]`—a collection of (hopefully, eventually) one value. We map a function `Int => Int` and get another `Future[Int]`—a collection of (hopefully, eventually) one integer.

But wait—that's not quite the same as in the callback code. The call to `getData2` is running *after* `getData1`, not concurrently. Let's fix that with a second `map`:

```

val future1 = Future { getData1() }
val future2 = Future { getData2() }
val combined = future1.map(n1 => future2.map(n2 => n1 + n2))

```

When `future1` and `future2` have delivered their results, the sum is computed.

Unfortunately, now `combined` is a `Future[Future[Int]]`, which isn't so good. That's what `flatMap` is for:

```
val combined = f1.flatMap(n1 => f2.map(n2 => n1 + n2))
```

This looks much nicer when you use a `for` expression instead of chaining `flatMap` and `map`:

```
val combined = for (n1 <- future1; n2 <- future2) yield n1 + n2
```

This is exactly the same code since `for` expressions are translated to chains of `map` and `flatMap`.

You can also apply guards in the `for` expression:

```
val combined =
  for (int n1 <- future1; n2 <- future2 if n1 != n2) yield n1 + n2
```

If the guard fails, the computation fails with a `NoSuchElementException`.

What if something goes wrong? The `map` and `flatMap` implementations take care of all that. As soon as one of the tasks fails, the entire pipeline fails, and the exception is captured. In contrast, when you manually combine callbacks, you have to deal with failure at every step.



NOTE: If you find the `for/yield` construct unnatural, check out the Scala Async library at <http://github.com/scala/async>. It uses Scala Macros to let you express the flow more naturally as

```
val combined = async { await(future1) + await(future2) }
```

So far, you have seen how to run two tasks concurrently. Sometimes, you need one task to run after another. A `Future` starts execution immediately when it is created. To delay the creation, use functions.

```
def future1 = Future { getData() }
def future2 = Future { getMoreData() } // def, not val
val combined = for (n1 <- future1; n2 <- future2) yield n1 + n2
```

Now `future2` is only evaluated when `future1` has completed.

It doesn't matter whether you use `val` or `def` for `future1`. If you use `def`, its creation is slightly delayed to the start of the `for` expression.

This is particularly useful if the second step depends on the output of the first:

```
def future1 = Future { getData() }
def future2(arg: Int) = Future { getMoreData(arg) }
val combined = for (n1 <- future1; n2 <- future2(n1)) yield n1 + n2
```



NOTE: Like the Future trait, the Try class from Section 17.3, “The Try Class,” on page 251 has `map` and `flatMap` methods. A `Try[T]` is a collection of, hopefully, one element. It is just like a `Future[T]`, except you don’t have to wait. You can apply `map` with a function that changes that one element, or `flatMap` if you have Try-valued function and want to flatten the result. And you can use `for` expressions. For example, here is how to compute the sum of two function calls that might fail:

```
def readInt(prompt: String) = Try(StdIn.readLine(s"$prompt: ").toInt)
val t = for (n1 <- readInt("n1"); n2 <- readInt("n2")) yield n1 + n2
```

In this way, you can compose Try-valued computations and you don’t need to deal with the boring part of error handling.

17.6 Other Future Transformations

The `map` and `flatMap` methods that you saw in the preceding section are the most fundamental transformation of Future objects.

Table 17–1 shows several ways of applying functions to the contents of a future that differ in subtle details.

The `foreach` method works exactly like it does for collections, applying a method for its side effect. The method is applied to the single value in the future. It is convenient for harvesting the answer when it materializes.

```
val combined = for (n1 <- future1; n2 <- future2) yield n1 + n2
combined.foreach(n => println(s"Result: $n"))
```

The `recover` method accepts a partial function that can turn an exception into a successful result. Consider this call:

```
val f = Future { persist(data) } recover { case e: SQLException => 0 }
```

If a `SQLException` occurs, the future succeeds with result 0.

The `fallbackTo` method provides a different recovery mechanism. When you call `f.fallbackTo(f2)`, then `f2` is executed if `f` fails, and its value becomes the value of the future. However, `f2` cannot inspect the reason for the failure.

The `failed` method turns a failed `Future[T]` into a successful `Future[Throwable]`, just like the `Try.failed` method. You can retrieve the failure in a `for` expression like this:

```
val f = Future { ... }
for (ex <- f.failed) println(ex)
```

Finally, you can zip two futures together. The call `f1.zip(f2)` yields a future whose result is a pair `(v, w)` if `v` was the result of `f1` and `w` the result of `f2`, or an exception if either `f1` or `f2` failed. (If both fail, the exception of `f1` is reported.)

Table 17–1 Transformations on a Future[T] with success value v or exception ex

Method	Result	Description
collect(pf: PartialFunction[T, S])	Future[S]	Like map, but with a partial function. The result fails with a NoSuchElementException if pf(v) is not defined.
foreach(f: T => U)	Unit	Calls f(v) like map, but only for its side effect.
andThen(pf: PartialFunction[Try[T], U])	Future[T]	Calls pf(v) for its side effect and returns a future with v.
filter(p: T => Boolean)	Future[T]	Calls p(v) and returns a future with v or a NoSuchElementException.
recover(pf: PartialFunction[Throwable, U]) recoverWith(pf: PartialFunction[Throwable, Future[U]])	Future[U] (where U is a supertype of T)	A future with value v or pf(ex), flattened in the asynchronous case.
fallbackTo(f2: Future[U])	Future[U] (where U is a supertype of T)	A future with value v, or if this future failed, with the value of f2, or if that also failed, with exception ex.
failed	Future[Throwable]	A future with value ex.
transform(s: T => S, f: Throwable => Throwable) transform(f: Try[T] => Try[S]) transformWith(f: Try[T] => Future[Try[S]])	Future[S]	Transforms both the success and failure.

The zipWith method is similar, but it takes a method to combine the two results instead of returning a pair. For example, here is another way of obtaining the sum of two computations:

```
val future1 = Future { getData1() }
val future2 = Future { getData2() }
val combined = future1.zipWith(future2)(_ + _)
```

17.7 Methods in the Future Object

The Future companion object contains useful methods for working on collections of futures.

Suppose that, as you are computing a result, you organize the work so that you can concurrently work on different parts. For example, each part might be a range of the inputs. Make a future for each part:

```
val futures = parts.map(p => Future { compute result in p })
```

Now you have a collection of futures. Often, you want to combine the results. By using the `Future.sequence` method, you can get a collection of all results for further processing:

```
val result = Future.sequence(futures);
```

Note that the call doesn't block—it gives you a future to a collection. For example, assume `futures` is a `Set[Future[T]]`. Then the result is a `Future[Set[T]]`. When the results for all elements of `futures` are available, the result future will complete with a set of the results.

If any of the futures fail, then the resulting future fails as well with the exception of the leftmost failed future. If multiple futures fail, you don't get to see the remaining failures.

The `traverse` method combines the `map` and `sequence` steps. Instead of

```
val futures = parts.map(p => Future { compute result in p })
val result = Future.sequence(futures);
```

you can call

```
val result = Future.traverse(parts)(p => Future { compute result in p })
```

The function in the second curried argument is applied to each element of `parts`. You get a future to a collection of all results.

There are also `reduceLeft` and `foldLeft` operations that are analogous to the reductions and folds described in Section 13.9, "Reducing, Folding, and Scanning," on page 184. You supply an operation that combines the results of all futures as they become available. For example, here is how you can compute the sum of the results:

```
val result = Future.reduceLeft(futures)(_ + _)
// Yields a future to the sum of the results of all futures
```

So far, we have collected the results from all futures. Suppose you are willing to accept a result from any of the parts. Then call

```
Future[T] result = Future.firstCompletedOf(futures)
```

You get a future that, when it completes, has the result or failure of the first completed element of futures.

The `find` method is similar, but you also supply a predicate.

```
val result = Future.find(futures)(predicate)
// Yields a Future[Option[T]]
```

You get a future that, when it completes successfully, yields `Some(r)`, where `r` is the result of one of the given futures that fulfills the predicate. Failed futures are ignored. If all futures complete but none yields a result that matches the predicate, then `find` returns `None`.



CAUTION: A potential problem with `firstCompletedOf` and `find` is that the other computations keep on going even when the result has been determined. Scala futures do not have a mechanism for cancellation. If you want to stop unnecessary work, you have to provide your own mechanism.

Finally, the `Future` object provides convenience methods for generating simple futures:

- `Future.successful(r)` is an already completed future with result `r`.
- `Future.failed(e)` is an already completed future with exception `e`.
- `Future.fromTry(t)` is an already completed future with the result or exception given in the `Try` object `t`.
- `Future.unit` is an already completed future with `Unit` result.
- `Future.never` is a future that never completes.

17.8 Promises

A `Future` object is read-only. The value of the future is set implicitly when the task has finished or failed. A `Promise` is similar, but the value can be set explicitly.

Consider this method that yields a `Future`:

```
def computeAnswer(arg: String) = Future {
    val n = workHard(arg)
    n
}
```

With a `Promise`, it looks like this:

```
def computeAnswer(arg: String) = {
    val p = Promise[Int]()
    Future {
        val n = workHard(arg)
        p.success(n)
        workOnSomethingElse()
    }
    p.future
}
```

Calling `future` on a promise yields the associated `Future` object. Note that the method returns the `Future` right away, immediately after starting the task that will eventually yield the result. That task is run in another `Future`, defined by the expression `Future { ... }`, that is unrelated to the promise's `future`.

Calling `success` on a promise sets the result. Alternatively, you can call `failure` with an exception to make the promise fail. As soon as one of these methods is called, the associated `future` is completed, and neither method can be called again. (An `IllegalStateException` is thrown otherwise.)

From the point of view of the consumer (that is, caller of the `computeAnswer` method), there is no difference between the two approaches. Either way, the consumer has a `Future` and eventually gets the result.

The producer, however, has more flexibility when using a `Promise`. As suggested in the code sample, the producer can do other work besides fulfilling this promise. For example, the producer might work on fulfilling multiple promises.

```
val p1 = Promise[Int]()
val p2 = Promise[Int]()
Future {
    val n1 = getData1()
    p1.success(n1)
    val n2 = getData2()
    p2.success(n2)
}
```

It is also possible to have multiple tasks that work concurrently to fulfill a single promise. When one of the tasks has a result, it calls `trySuccess` on the promise. Unlike the `success` method, that method accepts the result and returns `true` if the promise has not yet completed; otherwise it returns `false` and ignores the result.

```
val p = Promise[Int]()
Future {
    var n = workHard(arg)
    p.trySuccess(n)
}
Future {
    var n = workSmart(arg)
    p.trySuccess(n)
}
```

The promise is completed by the first task that manages to produce the result. With this approach, the tasks might want to periodically call `p.isCompleted` to check whether they should continue.



NOTE: Scala promises are equivalent to the `CompletableFuture` class in Java 8.

17.9 Execution Contexts

By default, Scala futures are executed on the global fork-join pool. That works well for computationally intensive tasks. However, the fork-join pool only manages a small number of threads (by default, equal to the number of cores of all processors). This is a problem when tasks have to wait, for example when communicating with a remote resource. A program could exhaust all available threads, waiting for results.

You can notify the execution context that you are about to block, by placing the blocking code inside `blocking { ... }`:

```
val f = Future {
    val url = ...
    blocking {
        val contents = Source.fromURL(url).mkString
        ...
    }
}
```

The execution context may then increase the number of threads. The fork-join pool does exactly that, but it isn't designed for perform well for many blocking threads. If you do input/output or connect to databases, you are better off using a different thread pool. The `Executors` class from the Java concurrency library gives you several choices. A cached thread pool works well for I/O intensive workloads. You can pass it explicitly to the `Future.apply` method, or you can set it as the implicit execution context:

```
val pool = Executors.newCachedThreadPool()
implicit val ec = ExecutionContext.fromExecutor(pool)
```

Now this pool is used by all futures when `ec` is in scope.

Exercises

1. Consider the expression

```
for (n1 <- Future { Thread.sleep(1000) ; 2 })
    n2 <- Future { Thread.sleep(1000); 40 }
    println(n1 + n2)
```

How is the expression translated to `map` and `flatMap` calls? Are the two futures executed concurrently or one after the other? In which thread does the call to `println` occur?

2. Write a function `doInOrder` that, given two functions $f: T \Rightarrow Future[U]$ and $g: U \Rightarrow Future[V]$, produces a function $T \Rightarrow Future[V]$ that, for a given t , eventually yields $g(f(t))$.
3. Repeat the preceding exercise for any sequence of functions of type $T \Rightarrow Future[T]$.
4. Write a function `doTogether` that, given two functions $f: T \Rightarrow Future[U]$ and $g: U \Rightarrow Future[V]$, produces a function $T \Rightarrow Future[(U, V)]$, running the two computations in parallel and, for a given t , eventually yielding $(f(t), g(t))$.
5. Write a function that receives a sequence of futures and returns a future that eventually yields a sequence of all results.
6. Write a method

```
Future[T] repeat(action: => T, until: T => Boolean)
```

that asynchronously repeats the action until it produces a value that is accepted by the `until` predicate, which should also run asynchronously. Test with a function that reads a password from the console, and a function that simulates a validity check by sleeping for a second and then checking that the password is "secret". Hint: Use recursion.

7. Write a program that counts the prime numbers between 1 and n , as reported by `BigInt.isProbablePrime`. Divide the interval into p parts, where p is the number of available processors. Count the primes in each part in concurrent futures and combine the results.
8. Write a program that asks the user for a URL, reads the web page at that URL, and displays all the hyperlinks. Use a separate `Future` for each of these three steps.
9. Write a program that asks the user for a URL, reads the web page at that URL, finds all the hyperlinks, visits each of them concurrently, and locates the Server HTTP header for each of them. Finally, print a table of which servers were found how often. The futures that visit each page should return the header.
10. Change the preceding exercise where the futures that visit each header update a shared Java `ConcurrentHashMap` or Scala `TrieMap`. This isn't as easy as it sounds. A threadsafe data structure is safe in the sense that you cannot corrupt its implementation, but you have to make sure that sequences of reads and updates are atomic.

11. Using futures, run four tasks that each sleep for ten seconds and then print the current time. If you have a reasonably modern computer, it is very likely that it reports four available processors to the JVM, and the futures should all complete at around the same time. Now repeat with forty tasks. What happens? Why? Replace the execution context with a cached thread pool. What happens now? (Be careful to define the futures *after* replacing the implicit execution context.)
12. Write a method that, given a URL, locates all hyperlinks, makes a promise for each of them, starts a task in which it will eventually fulfill all promises, and returns a sequence of futures for the promises. Why would it not be a good idea to return a sequence of promises?
13. Use a promise for implementing cancellation. Given a range of big integers, split the range into subranges that you concurrently search for palindromic primes. When such a prime is found, set it as the value of the future. All tasks should periodically check whether the promise is completed, in which case they should terminate.

This page intentionally left blank

Type Parameters

Topics in This Chapter [L2](#)

- 18.1 Generic Classes — page 266
- 18.2 Generic Functions — page 266
- 18.3 Bounds for Type Variables — page 266
- 18.4 View Bounds — page 268
- 18.5 Context Bounds — page 268
- 18.6 The ClassTag Context Bound — page 269
- 18.7 Multiple Bounds — page 269
- 18.8 Type Constraints [L3](#) — page 269
- 18.9 Variance — page 271
- 18.10 Co- and Contravariant Positions — page 272
- 18.11 Objects Can't Be Generic — page 274
- 18.12 Wildcards — page 275
- Exercises — page 275

Chapter 18

In Scala, you can use type parameters to implement classes and functions that work with multiple types. For example, an `Array[T]` stores elements of an arbitrary type `T`. The basic idea is very simple, but the details can get tricky. Sometimes, you need to place restrictions on the type. For example, to sort elements, `T` must provide an ordering. Furthermore, if the parameter type varies, what should happen with the parameterized type? For example, can you pass an `Array[String]` to a function that expects an `Array[Any]`? In Scala, you specify how your types should vary depending on their parameters.

The key points of this chapter are:

- Classes, traits, methods, and functions can have type parameters.
- Place the type parameters after the name, enclosed in square brackets.
- Type bounds have the form `T <: UpperBound, T >: LowerBound, T : ContextBound`.
- You can restrict a method with a type constraint such as `(implicit ev: T <: UpperBound)`.
- Use `+T` (covariance) to indicate that a generic type's subtype relationship is in the same direction as the parameter `T`, or `-T` (contravariance) to indicate the reverse direction.
- Covariance is appropriate for parameters that denote outputs, such as elements in an immutable collection.

- Contravariance is appropriate for parameters that denote inputs, such as function arguments.

18.1 Generic Classes

As in Java or C++, classes and traits can have type parameters. In Scala, you use square brackets for type parameters, for example:

```
class Pair[T, S](val first: T, val second: S)
```

This defines a class with two type parameters `T` and `S`. You use the type parameters in the class definition to define the types of variables, method parameters, and return values.

A class with one or more type parameters is *generic*. If you substitute actual types for the type parameters, you get an ordinary class, such as `Pair[Int, String]`.

Pleasantly, Scala attempts to infer the actual types from the construction parameters:

```
val p = new Pair(42, "String") // It's a Pair[Int, String]
```

You can also specify the types yourself:

```
val p2 = new Pair[Any, Any](42, "String")
```

18.2 Generic Functions

Functions and methods can also have type parameters. Here is a simple example:

```
def getMiddle[T](a: Array[T]) = a(a.length / 2)
```

As with generic classes, you place the type parameter after the name.

Scala infers the actual types from the arguments in the call.

```
getMiddle(Array("Mary", "had", "a", "little", "lamb")) // Calls getMiddle[String]
```

If you need to, you can specify the type:

```
val f = getMiddle[String] _ // The function, saved in f
```

18.3 Bounds for Type Variables

Sometimes, you need to place restrictions on type variables. Consider a `Pair` type where both components have the same type, like this:

```
class Pair[T](val first: T, val second: T)
```

Now we want to add a method that produces the smaller value:

```
class Pair[T](val first: T, val second: T) {
    def smaller = if (first.compareTo(second) < 0) first else second // Error
}
```

That's wrong—we don't know if `first` has a `compareTo` method. To solve this, we can add an *upper bound* `T <: Comparable[T]`.

```
class Pair[T <: Comparable[T]](val first: T, val second: T) {
    def smaller = if (first.compareTo(second) < 0) first else second
}
```

This means that `T` must be a subtype of `Comparable[T]`.

Now we can instantiate `Pair[java.lang.String]` but not `Pair[java.net.URL]`, since `String` is a subtype of `Comparable[String]` but `URL` does not implement `Comparable[URL]`. For example:

```
val p = new Pair("Fred", "Brooks")
println(p.smaller) // Prints Brooks
```



CAUTION: This example is a bit simplistic. If you try a new `Pair(4, 2)`, you will be told that for `T = Int`, the bound `T <: Comparable[T]` is not fulfilled. See Section 18.4, “View Bounds,” on page 268 for a remedy.

You can also specify a lower bound for a type. For example, suppose we want to define a method that replaces the first component of a pair with another value. Our pairs are immutable, so we need to return a new pair. Here is a first attempt:

```
class Pair[T](val first: T, val second: T) {
    def replaceFirst(newFirst: T) = new Pair[T](newFirst, second)
}
```

But we can do better than that. Suppose we have a `Pair[Student]`. It should be possible to replace the first component with a `Person`. Of course, then the result must be a `Pair[Person]`. In general, the replacement type must be a supertype of the pair's component type.

```
def replaceFirst[R >: T](newFirst: R) = new Pair[R](newFirst, second)
```

Here, I included the type parameter in the returned pair for greater clarity. You can also write

```
def replaceFirst[R >: T](newFirst: R) = new Pair(newFirst, second)
```

Then the return type is correctly inferred as `new Pair[R]`.



CAUTION: If you omit the lower bound,

```
def replaceFirst[R](newFirst: R) = new Pair(newFirst, second)
```

the method will compile, but it will return a `Pair[Any]`.

18.4 View Bounds

In the preceding section, we had an example of an upper bound:

```
class Pair[T <: Comparable[T]]
```

Unfortunately, if you try constructing a new `Pair(4, 2)`, the compiler complains that `Int` is not a subtype of `Comparable[Int]`. Unlike the `java.lang.Integer` wrapper type, the Scala `Int` type does not implement `Comparable`. However, `RichInt` does implement `Comparable[Int]`, and there is an implicit conversion from `Int` to `RichInt`. (See Chapter 21 for more information on implicit conversions.)

A solution is to use a “view bound” like this:

```
class Pair[T <% Comparable[T]]
```

The `<%` relation means that `T` can be converted to a `Comparable[T]` through an implicit conversion.

However, view bounds are on their way out in Scala. If you compile with the `-future` flag, you’ll get a warning when you use them. You can replace a view bound with a “type constraint” (see Section 18.8, “Type Constraints,” on page 269), like this:

```
class Pair[T](val first: T, val second: T)(implicit ev: T => Comparable[T]) {
    def smaller = if (first.compareTo(second) < 0) first else second
    ...
}
```

18.5 Context Bounds

A view bound `T <% V` requires the existence of an implicit conversion from `T` to `V`. A *context bound* has the form `T : M`, where `M` is another generic type. It requires that there is an “implicit value” of type `M[T]`. We discuss implicit values in detail in Chapter 21.

For example,

```
class Pair[T : Ordering]
```

requires that there is an implicit value of type `Ordering[T]`. That implicit value can then be used in the methods of the class. When you declare a method that uses the implicit value, you have to add an “implicit parameter.” Here is an example:

```
class Pair[T : Ordering](val first: T, val second: T) {
  def smaller(implicit ord: Ordering[T]) =
    if (ord.compare(first, second) < 0) first else second
}
```

As you will see in Chapter 21, implicit values are more flexible than implicit conversions.

18.6 The ClassTag Context Bound

To instantiate a generic `Array[T]`, one needs a `ClassTag[T]` object. This is required for primitive type arrays to work correctly. For example, if `T` is `Int`, you want an `int[]` array in the virtual machine. If you write a generic function that constructs a generic array, you need to help it out and pass that class tag object. Use a context bound, like this:

```
import scala.reflect._
def makePair[T : ClassTag](first: T, second: T) = {
  val r = new Array[T](2); r(0) = first; r(1) = second; r
}
```

If you call `makePair(4, 9)`, the compiler locates the implicit `ClassTag[Int]` and actually calls `makePair(4, 9)(classTag)`. Then the `new` operator is translated to a call `classTag newArray`, which in the case of a `ClassTag[Int]` constructs a primitive array `int[2]`.

Why all this complexity? In the virtual machine, generic types are erased. There is only a single `makePair` method that needs to work for *all* types `T`.

18.7 Multiple Bounds

A type variable can have both an upper and a lower bound. The syntax is this:

```
T >: Lower <: Upper
```

You can't have multiple upper or lower bounds. However, you can still require that a type implements multiple traits, like this:

```
T <: Comparable[T] with Serializable with Cloneable
```

You can have more than one context bound:

```
T : Ordering : ClassTag
```

18.8 Type Constraints L3

Type constraints give you another way of restricting types. There are three relationships that you can use:

$T ::= U$
 $T <: U$
 $T \Rightarrow U$

These constraints test whether T equals U , is a subtype of U , or is convertible to U . To use such a constraint, you add an “implicit evidence parameter” like this:

```
class Pair[T](val first: T, val second: T)(implicit ev: T <: Comparable[T])
```



NOTE: See Chapter 21 for an explanation of the curious syntax, and for an analysis of the inner workings of the type constraints.

In the example above, there is no advantage to using a type constraint over a type bound class `Pair[T <: Comparable[T]]`. However, type constraints are useful in some specialized circumstances. In this section, you will see two uses of type constraints.

Type constraints let you supply a method in a generic class that can be used only under certain conditions. Here is an example:

```
class Pair[T](val first: T, val second: T) {
    def smaller(implicit ev: T <: Ordered[T]) =
        if (first < second) first else second
}
```

You can form a `Pair[URL]`, even though `URL` is not ordered. You will get an error only if you invoke the `smaller` method.

Another example is the `orNull` method in the `Option` class:

```
val friends = Map("Fred" -> "Barney", ...)
val friendOpt = friends.get("Wilma") // An Option[String]
val friendOrNull = friendOpt.orNull // A String or null
```

The `orNull` method can be useful when working with Java code where it is common to encode missing values as `null`. But it can't be applied to value types such as `Int` that don't have `null` as a valid value. Because `orNull` is implemented using a constraint `Null <: A`, you can still instantiate `Option[Int]`, as long as you stay away from `orNull` for those instances.

Another use of type constraints is for improving type inference. Consider

```
def firstLast[A, C <: Iterable[A]](it: C) = (it.head, it.last)
```

When you call

```
firstLast(List(1, 2, 3))
```

you get a message that the inferred type arguments `[Nothing, List[Int]]` don't conform to `[A, C <: Iterable[A]]`. Why `Nothing`? The type inferencer cannot figure

out what A is from looking at List(1, 2, 3), because it matches A and C in a single step. To help it along, first match C and then A:

```
def firstLast[A, C](it: C)(implicit ev: C <:< Iterable[A]) =  
(it.head, it.last)
```



NOTE: You saw a similar trick in Chapter 12. The `corresponds` method checks whether two sequences have corresponding entries:

```
def corresponds[B](that: Seq[B])(match: (A, B) => Boolean): Boolean
```

The `match` predicate is a curried parameter so that the type inferencer can first determine the type of B and then use that information to analyze `match`. In the call

```
Array("Hello", "Fred").corresponds(Array(5, 4))(_.length == _)
```

the compiler can infer that B is Int. Then it can make sense of `_.length == _`.

18.9 Variance

Suppose we have a function that does something with a `Pair[Person]`:

```
def makeFriends(p: Pair[Person])
```

If `Student` is a subclass of `Person`, can I call `makeFriend` with a `Pair[Student]`? By default, this is an error. Even though `Student` is a subtype of `Person`, there is *no* relationship between `Pair[Student]` and `Pair[Person]`.

If you want such a relationship, you have to indicate it when you define the `Pair` class:

```
class Pair[+T](val first: T, val second: T)
```

The `+` means that the type is *covariant* in T—that is, it varies in the same direction. Since `Student` is a subtype of `Person`, a `Pair[Student]` is now a subtype of `Pair[Person]`.

It is also possible to have variance in the other direction. Consider a generic type `Friend[T]`, which denotes someone who is willing to befriend anyone of type T.

```
trait Friend[-T] {  
    def befriend(someone: T)  
}
```

Now suppose you have a function

```
def makeFriendWith(s: Student, f: Friend[Student]) { f.befriend(s) }
```

Can you call it with a `Friend[Person]`? That is, if you have

```
class Person extends Friend[Person] { ... }
class Student extends Person
val susan = new Student
val fred = new Person
```

will the call `makeFriendWith(susan, fred)` succeed? It seems like it should. If Fred is willing to befriend any person, he'll surely like to be friends with Susan.

Note that the type varies in the opposite direction of the subtype relationship. `Student` is a subtype of `Person`, but `Friend[Student]` is a supertype of `Friend[Person]`. In that case, you declare the type parameter to be *contravariant*:

```
trait Friend[-T] {
    def befriend(someone: T)
}
```

You can have both variance types in a single generic type. For example, single-argument functions have the type `Function1[-A, +R]`. To see why these are the appropriate variances, consider a function

```
def friends(students: Array[Student], find: Function1[Student, Person]) =
    // You can write the second parameter as find: Student => Person
    for (s <- students) yield find(s)
```

Suppose you have a function

```
def findStudent(p: Person) : Student
```

Can you call `friends` with that function? Of course you can. It's willing to take any person, so surely it will take a `Student`. It yields `Student` results, which can be put into an `Array[Person]`.

18.10 Co- and Contravariant Positions

In the preceding section, you saw that functions are contravariant in their arguments and covariant in their results. Generally, it makes sense to use contravariance for the values an object consumes, and covariance for the values it produces. (Aide-mémoire: **contravariance consumes**.)

If an object does both, then the type should be left *invariant*. This is generally the case for mutable data structures. For example, in Scala, arrays are invariant. You can't convert an `Array[Student]` to an `Array[Person]` or the other way around. This would not be safe. Consider the following:

```
val students = new Array[Student](length)
val people: Array[Person] = students // Not legal, but suppose it was ...
people(0) = new Person("Fred") // Oh no! Now students(0) isn't a Student
```

Conversely,

```
val people = new Array[Person](length)
val students: Array[Student] = people // Not legal, but suppose it was . .
people(0) = new Person("Fred") // Oh no! Now students(0) isn't a Student
```



NOTE: In Java, it is possible to convert a `Student[]` array to a `Person[]` array, but if you try to add a nonstudent into such an array, an `ArrayStoreException` is thrown. In Scala, the compiler rejects programs that could cause type errors.

Suppose we tried to declare a covariant *mutable* pair. This wouldn't work. It would be like an array with two elements, and one could produce the same kind of error that you just saw.

Indeed, if you try

```
class Pair[+T](var first: T, var second: T) // Error
```

you get an error complaining that the covariant type `T` occurs in a *contravariant position* in the setter

```
first_=(value: T)
```

Parameters are contravariant positions, and return types are covariant.

However, inside a function parameter, the variance flips—its parameters are covariant. For example, look at the `foldLeft` method of `Iterable[+A]`:

```
foldLeft[B](z: B)(op: (B, A) => B): B
      -      +  +      -      +
```

Note that `A` is now in a covariant position.

These position rules are simple and safe, but they sometimes get in the way of doing something that would be risk-free. Consider the `replaceFirst` method from Section 18.3, “Bounds for Type Variables,” on page 266 in an immutable pair:

```
class Pair[+T](val first: T, val second: T) {
  def replaceFirst(newFirst: T) = new Pair[T](newFirst, second) // Error
}
```

The compiler rejects this, because the parameter type `T` is in a contravariant position. Yet this method cannot damage the pair—it returns a new pair.

The remedy is to come up with a second type parameter for the method, like this:

```
def replaceFirst[R >: T](newFirst: R) = new Pair[R](newFirst, second)
```

Now the method is a generic method with another type parameter `R`. But `R` is *invariant*, so it doesn't matter that it appears in a contravariant position.

18.11 Objects Can't Be Generic

It is not possible to add type parameters to objects. Consider, for example, immutable lists. A list with element type `T` is either empty, or it is a node with a head of type `T` and a tail of type `List[T]`:

```
abstract class List[+T] {
    def isEmpty: Boolean
    def head: T
    def tail: List[T]
}

class Node[T](val head: T, val tail: List[T]) extends List[T] {
    def isEmpty = false
}

class Empty[T] extends List[T] {
    def isEmpty = true
    def head = throw new UnsupportedOperationException
    def tail = throw new UnsupportedOperationException
}
```



NOTE: Here I use `Node` and `Empty` to make the discussion easier to follow for Java programmers. If you are experienced with Scala lists, just substitute `::` and `Nil` in your mind.

It seems silly to define `Empty` as a class. It has no state. But you can't simply turn it into an object:

```
object Empty[T] extends List[T] // Error
```

You can't add a parameterized type to an object. In this case, a remedy is to inherit `List[Nothing]`:

```
object Empty extends List[Nothing]
```

Recall from Chapter 8 that the `Nothing` type is a subtype of all types. Thus, when we make a one-element list

```
val lst = new Node(42, Empty)
```

type checking is successful. Due to covariance, a `List[Nothing]` is convertible into a `List[Int]`, and the `Node[Int]` constructor can be invoked.

18.12 Wildcards

In Java, all generic types are invariant. However, you can vary the types where you use them, using wildcards. For example, a method

```
void makeFriends(List<? extends Person> people) // This is Java
```

can be called with a `List<Student>`.

You can use wildcards in Scala too. They look like this:

```
def process(people: java.util.List[_ <: Person]) // This is Scala
```

In Scala, you don't need the wildcard for a covariant `Pair` class. But suppose `Pair` is invariant:

```
class Pair[T](var first: T, var second: T)
```

Then you can define

```
def makeFriends(p: Pair[_ <: Person]) // OK to call with a Pair[Student]
```

You can also use wildcards for contravariance:

```
import java.util.Comparator
def min[T](p: Pair[T])(comp: Comparator[_ >: T])
```

Wildcards are “syntactic sugar” for existential types, which we will discuss in detail in Chapter 19.



CAUTION: In certain complex situations, Scala wildcards are still a work in progress. For example, the following declaration does not work in Scala 2.12:

```
def min[T <: Comparable[_ >: T]](p: Pair[T]) = ...
```

A workaround is the following:

```
type SuperComparable[T] = Comparable[_ >: T]
def min[T <: SuperComparable[T]](p: Pair[T]) = ...
```

Exercises

1. Define an immutable class `Pair[T, S]` with a method `swap` that returns a new pair with the components swapped.
2. Define a mutable class `Pair[T]` with a method `swap` that swaps the components of the pair.

3. Given a class `Pair[T, S]`, write a generic method `swap` that takes a pair as its argument and returns a new pair with the components swapped.
4. Why don't we need a lower bound for the `replaceFirst` method in Section 18.3, "Bounds for Type Variables," on page 266 if we want to replace the first component of a `Pair[Person]` with a `Student`?
5. Why does `RichInt` implement `Comparable[Int]` and not `Comparable[RichInt]`?
6. Write a generic method `middle` that returns the middle element from any `Iterable[T]`. For example, `middle("World")` is 'r'.
7. Look through the methods of the `Iterable[+A]` trait. Which methods use the type parameter `A`? Why is it in a covariant position in these methods?
8. In Section 18.10, "Co- and Contravariant Positions," on page 272, the `replaceFirst` method has a type bound. Why can't you define an equivalent method on a mutable `Pair[T]`?

```
def replaceFirst[R >: T](newFirst: R) { first = newFirst } // Error
```

9. It may seem strange to restrict method parameters in an immutable class `Pair[+T]`. However, suppose you could define

```
def replaceFirst(newFirst: T)
```

in a `Pair[+T]`. The problem is that this method can be overridden in an unsound way. Construct an example of the problem. Define a subclass `NastyDoublePair` of `Pair[Double]` that overrides `replaceFirst` so that it makes a pair with the square root of `newFirst`. Then construct the call `replaceFirst("Hello")` on a `Pair[Any]` that is actually a `NastyDoublePair`.

10. Given a mutable `Pair[S, T]` class, use a type constraint to define a `swap` method that can be called if the type parameters are the same.

This page intentionally left blank

Advanced Types

Topics in This Chapter **L2**

- 19.1 Singleton Types — page 280
- 19.2 Type Projections — page 281
- 19.3 Paths — page 282
- 19.4 Type Aliases — page 283
- 19.5 Structural Types — page 283
- 19.6 Compound Types — page 284
- 19.7 Infix Types — page 285
- 19.8 Existential Types — page 286
- 19.9 The Scala Type System — page 287
- 19.10 Self Types — page 288
- 19.11 Dependency Injection — page 289
- 19.12 Abstract Types **L3** — page 291
- 19.13 Family Polymorphism **L3** — page 293
- 19.14 Higher-Kinded Types **L3** — page 296
- Exercises — page 299

Chapter 19

In this chapter, you will see all the types that Scala has to offer, including some of the more technical ones. We will end with a discussion of self types and dependency injection.

The key points of this chapter are:

- Singleton types are useful for method chaining and methods with object parameters.
- A type projection includes inner class instances for all objects of an outer class.
- A type alias gives a short name for a type.
- Structural types are equivalent to “duck typing.”
- Existential types provide the formalism for wildcard parameters of generic types.
- Use a self type declaration to indicate that a trait requires another type.
- The “cake pattern” uses self types to implement a dependency injection mechanism.
- An abstract type must be made concrete in a subclass.
- A higher-kinded type has a type parameter that is itself a parameterized type.

19.1 Singleton Types

Given any reference `v`, you can form the type `v.type`, which has two values: `v` and `null`. This sounds like a curious type, but it has a couple of useful applications.

First, consider a method that returns this so you can chain method calls:

```
class Document {
    def setTitle(title: String) = { ...; this }
    def setAuthor(author: String) = { ...; this }
    ...
}
```

You can then call

```
article.setTitle("Whatever Floats Your Boat").setAuthor("Cay Horstmann")
```

However, if you have a subclass, there is a problem:

```
class Book extends Document {
    def addChapter(chapter: String) = { ...; this }
    ...
}

val book = new Book()
book.setTitle("Scala for the Impatient").addChapter(chapter1) // Error
```

Since the `setTitle` method returns `this`, Scala infers the return type as `Document`. But `Document` doesn't have an `addChapter` method.

The remedy is to declare the return type of `setTitle` as `this.type`:

```
def setTitle(title: String): this.type = { ...; this }
```

Now the return type of `book.setTitle("...")` is `book.type`, and since `book` has an `addChapter` method, the chaining works.

You can also use a singleton type if you want to define a method that takes an object instance as parameter. You may wonder why you would ever want to do that—after all, if there is just one instance, the method could simply use it instead of making the caller pass it.

However, some people like to construct “fluent interfaces” that read like English, for example:

```
book.set Title to "Scala for the Impatient"
```

This is parsed as

```
book.set(Title).to("Scala for the Impatient")
```

For this to work, `set` is a method whose argument is the singleton `Title`:

```

object Title

class Document {
    private var useNextArgAs: Any = null
    def set(obj: Title.type): this.type = { useNextArgAs = obj; this }
    def to(arg: String) = if (useNextArgAs == Title) title = arg; else ...
    ...
}

```

Note the `Title.type` parameter. You can't use

```
def set(obj: Title) ... // Error
```

since `Title` denotes the singleton *object*, not a type.

19.2 Type Projections

In Chapter 5, you saw that a nested class belongs to the *object* in which it is nested. Here is the example:

```

import scala.collection.mutable.ArrayBuffer
class Network {
    class Member(val name: String) {
        val contacts = new ArrayBuffer[Member]
    }

    private val members = new ArrayBuffer[Member]

    def join(name: String) = {
        val m = new Member(name)
        members += m
        m
    }
}

```

Each `Network` instance has its own `Member` class. For example, here are two networks:

```
val chatter = new Network
val myFace = new Network
```

Now `chatter.Member` and `myFace.Member` are *different classes*.

You can't add a member from one network to another:

```
val fred = chatter.join("Fred") // Has type chatter.Member
val barney = myFace.join("Barney") // Has type myFace.Member
fred.contacts += barney // Error
```

If you don't want this restriction, you should simply move the `Member` type outside the `Network` class. A good place would be the `Network` companion object.

If what you want is fine-grained classes, with an occasional loose interpretation, use a *type projection* `Network#Member`, which means “a Member of *any* Network.”

```
class Network {
    class Member(val name: String) {
        val contacts = new ArrayBuffer[Network#Member]
    }
    ...
}
```

You would do that if you want the fine-grained “inner class per object” feature in some places of your program, but not everywhere.



CAUTION: A type projection such as `Network#Member` is not considered a “path,” and you cannot import it. We discuss paths in the next section.

19.3 Paths

Consider a type such as

```
com.horstmann.impatient.chatter.Member
```

or, if we nest `Member` inside the companion object,

```
com.horstmann.impatient.Network.Member
```

Such an expression is called a *path*.

Each component of the path before the final type must be “stable,” that is, it must specify a single, definite scope. Each such component is one of the following:

- A package
- An object
- A `val`
- `this`, `super`, `super[S]`, `C.this`, `C.super`, or `C.super[S]`

A path component can’t be a class because, as you have seen, a nested class isn’t a single type, but it gives rise to a separate type for each instance.

Moreover, a path element can’t be a var. For example,

```
var chatter = new Network
...
val fred = new chatter.Member // Error—chatter is not stable
```

Since you might assign a different value to `chatter`, the compiler can’t assign a definite meaning to the type `chatter.Member`.



NOTE: Internally, the compiler translates all nested type expressions `a.b.c.T` to type projections `a.b.c.type#T`. For example, `chatter.Member` becomes `chatter.type#Member`—any `Member` inside the singleton `chatter.type`. That is not something you generally need to worry about. However, sometimes you will see an error message with a type of the form `a.b.c.type#T`. Just translate it back to `a.b.c.T`.

19.4 Type Aliases

You can create a simple *alias* for a complicated type with the `type` keyword, like this:

```
class Book {  
    import scala.collection.mutable._  
    type Index = HashMap[String, (Int, Int)]  
    ...  
}
```

Then you can refer to `Book.Index` instead of the cumbersome type `scala.collection.mutable.HashMap[String, (Int, Int)]`.

A type alias must be nested inside a class or object. It cannot appear at the top level of a Scala file. However, in the REPL, you can declare a type at the top level, since everything in the REPL is implicitly contained in a top-level object.



NOTE: The `type` keyword is also used for *abstract types* that are made concrete in a subclass, for example:

```
abstract class Reader {  
    type Contents  
    def read(fileName: String): Contents  
}
```

We will discuss abstract types in Section 19.12, “Abstract Types,” on page 291.

/

19.5 Structural Types

A “structural type” is a specification of abstract methods, fields, and types that a conforming type should possess. For example, this method has a structural type parameter:

```
def appendLines(target: { def append(str: String): Any },
               lines: Iterable[String]) {
    for (l <- lines) { target.append(l); target.append("\n") }
}
```

You can call the `appendLines` method with an instance of *any* class that has an `append` method. This is more flexible than defining a `Appendable` trait, because you might not always be able to add that trait to the classes you are using.

Under the hood, Scala uses reflection to make the calls to `target.append(...)`. Structural typing gives you a safe and convenient way of making such calls.

However, a reflective call is *much* more expensive than a regular method call. For that reason, you should only use structural typing when you model common behavior from classes that cannot share a trait.



NOTE: Structural types are similar to “duck typing” in dynamically typed programming languages such as JavaScript or Ruby. In those languages, variables have no type. When you write `obj.quack()`, the runtime figures out whether the particular object to which `obj` refers at this point has a `quack` method. In other words, you don’t have to declare `obj` as a Duck as long as it walks and quacks like one.

19.6 Compound Types

A compound type has the form

T_1 with T_2 with $T_3 \dots$

where T_1 , T_2 , T_3 , and so on are types. In order to belong to the compound type, a value must belong to all of the individual types. Therefore, such a type is also called an intersection type.

You can use a compound type to manipulate values that must provide multiple traits. For example,

```
val image = new ArrayBuffer[java.awt.Shape with java.io.Serializable]
```

You can draw the `image` object as for `(s <- image)` `graphics.draw(s)`. You can serialize the `image` object because you know that all elements are serializable.

Of course, you can only add elements that are both shapes and serializable objects:

```
val rect = new Rectangle(5, 10, 20, 30)
image += rect // OK—Rectangle is Serializable
image += new Area(rect) // Error—Area is a Shape but not Serializable
```



NOTE: When you have a declaration

```
trait ImageShape extends Shape with Serializable
```

this means that ImageShape extends the intersection type Shape with Serializable.

You can add a structural type declaration to a simple or compound type. For example,

```
Shape with Serializable { def contains(p: Point): Boolean }
```

An instance of this type must be a subtype of Shape and Serializable, and it must have a contains method with a Point parameter.

Technically, the structural type

```
{ def append(str: String): Any }
```

is an abbreviation for

```
AnyRef { def append(str: String): Any }
```

and the compound type

```
Shape with Serializable
```

is a shortcut for

```
Shape with Serializable {}
```

19.7 Infix Types

An infix type is a type with two type parameters, written in “infix” syntax, with the type name between the type parameters. For example, you can write

```
String Map Int
```

instead of

```
Map[String, Int]
```

The infix notation is common in mathematics. For example, $A \times B = \{ (a, b) \mid a \in A, b \in B \}$ is the set of pairs with components of types A and B. In Scala, this type is written as (A, B). If you prefer the mathematical notation, you can define

```
type ×[A, B] = (A, B)
```

Then you can write String × Int instead of (String, Int).

All infix type operators have the same precedence. As with regular operators, they are left-associative unless their names end in :. For example,

```
String × Int × Int
```

means `((String, Int), Int)`. This type is similar to, but not the same, as `(String, Int, Int)`, which could not be written in infix form in Scala.



NOTE: An infix type name can be any sequence of operator characters, except for a single `*`. This rule avoids confusion with variable argument declarations `T*`.

19.8 Existential Types

Existential types were added to Scala for compatibility with Java wildcards. An existential type is a type expression followed by `forSome { ... }`, where the braces contain type and `val` declarations. For example,

```
Array[T] forSome { type T <: JComponent }
```

This is the same as the wildcard type

```
Array[_ <: JComponent]
```

that you saw in Chapter 18.

Scala wildcards are syntactic sugar for existential types. For example,

```
Array[_]
```

is the same as

```
Array[T] forSome { type T }
```

and

```
Map[_, _]
```

is the same as

```
Map[T, U] forSome { type T; type U }
```

The `forSome` notation allows for more complex relationships than wildcards can express, for example:

```
Map[T, U] forSome { type T; type U <: T }
```

You can use `val` declarations in the `forSome` block because a `val` can have its own nested types (see Section 19.2, “Type Projections,” on page 281). Here is an example:

```
n.Member forSome { val n: Network }
```

By itself, that’s not so interesting—you could just use a type projection `Network#Member`. But consider

```
def process[M <: n.Member forSome { val n: Network }](m1: M, m2: M) = (m1, m2)
```

This method will accept members from the same network, but reject members from different ones:

```
val chatter = new Network
val myFace = new Network
val fred = chatter.join("Fred")
val wilma = chatter.join("Wilma")
val barney = myFace.join("Barney")
process(fred, wilma) // OK
process(fred, barney) // Error
```



NOTE: To use existential types without warnings, you must import `scala.language.existentials` or use the compiler option `-language:existentials`.

19.9 The Scala Type System

The Scala language reference gives an exhaustive list of all Scala types, which is reproduced in Table 19–1, with brief explanations for each type.

Table 19–1 Scala Types

Type	Syntax	Notes
Class or trait	<code>class C ...</code> , <code>trait C ...</code>	See Chapter 5, Chapter 10
Tuple type	(T_1, \dots, T_n)	Section 4.7
Function type	$(T_1, \dots, T_n) \Rightarrow T$	
Annotated type	$T @A$	See Chapter 15
Parameterized type	$A[T_1, \dots, T_n]$	See Chapter 18
Singleton type	$value.type$	See Section 19.1
Type projection	$O#I$	See Section 19.2
Compound type	$T_1 \text{ with } T_2 \text{ with } \dots \text{ with } T_n \{ declarations \}$	See Section 19.6
Infix type	$T_1 \wedge T_2$	See Section 19.7
Existential type	$T \text{ forSome } \{ type\ and\ val\ declarations \}$	See Section 19.8



NOTE: This table shows the types that you, the programmer, can declare. There are a few types that the Scala compiler uses internally. For example, a *method type* is denoted by $(T_1, \dots, T_n)T$ without a \Rightarrow . You will occasionally see such types. For example, when you enter

```
def square(x: Int) = x * x
```

in the Scala REPL, it responds with

```
square (x: Int)Int
```

This is different from

```
val triple = (x: Int) => 3 * x
```

which yields

```
triple: Int => Int
```

You can turn a method into a function by appending a `_`. The type of

```
square _
```

is `Int => Int`.

19.10 Self Types

In Chapter 10, you saw how a trait can require that it is mixed into a class that extends another type. You define the trait with a *self type* declaration:

```
this: Type =>
```

Such a trait can only be mixed into a subclass of the given type. In the following example, the `LoggedException` trait can only be mixed into a class that extends `Exception`:

```
trait Logged {
    def log(msg: String)
}

trait LoggedException extends Logged {
    this: Exception =>
    def log() { log(getMessage()) }
    // OK to call getMessage because this is an Exception
}
```

If you try to mix the trait into a class that doesn't conform to the self type, an error occurs:

```
val f = new JFrame with LoggedException
// Error: JFrame isn't a subtype of Exception, the self type of LoggedException
```

If you want to require multiple types, use a compound type:

this: T with U with ... =>



NOTE: You can combine the self type syntax with the “alias for enclosing this” syntax that I briefly introduced in Chapter 5. If you give a name other than this to the variable, then it can be used in subtypes by that name. For example,

```
trait Group {
    outer: Network =>
    class Member {
        ...
    }
}
```

The Group trait requires that it is added to a subtype of Network, and inside Member, you can refer to Group.this as outer.

This syntax seems to have grown organically over time; unfortunately, it introduces a great deal of confusion for a small amount of added functionality.



CAUTION: Self types do not automatically inherit. If you define

```
trait ManagedException extends LoggedException { ... }
```

you get an error that ManagedException doesn’t supply Exception. In this situation, you need to repeat the self type:

```
trait ManagedException extends LoggedException {
    this: Exception =>
    ...
}
```

19.11 Dependency Injection

When building a large system out of components, with different implementations for each component, one needs to assemble the component choices. For example, there may be a mock database and a real database, or console logging and file logging. A particular implementation may want the real database and console logging for running an experiment, or the mock database and file logging for running an automated test script.

Usually, there is some dependency among the components. For example, the data access component may require logging.

Java has several tools that allow programmers to express dependencies through frameworks such as Spring or module systems such as OSGi. Each component describes on which other component interfaces it depends. References to actual component implementations are “injected” when the application is assembled.

In Scala, you can achieve a simple form of dependency injection with traits and self types.

For logging, suppose we have a trait

```
trait Logger { def log(msg: String) }
```

with implementations `ConsoleLogger` and `FileLogger`.

The user authentication trait has a logging dependency to log authentication failures:

```
trait Auth {
    this: Logger =>
    def login(id: String, password: String): Boolean
}
```

The application depends on both:

```
trait App {
    this: Logger with Auth =>
    ...
}
```

Now we can assemble an application as

```
object MyApp extends App with FileLogger("test.log") with MockAuth("users.txt")
```

It's a bit awkward to use trait composition in this way. After all, an application isn't an authenticator and a file logger. It has these components, and it is more natural to use instance variables for the components than to glue them all into one huge type. A better design is given by the *cake pattern*. In this pattern, you supply a component trait for each service that contains

- Any dependent components, expressed as self types
- A trait describing the service interface
- An abstract `val` that will be instantiated with an instance of the service
- Optionally, implementations of the service interface

```

trait LoggerComponent {
  trait Logger { ... }
  val logger: Logger
  class FileLogger(file: String) extends Logger { ... }
  ...
}

trait AuthComponent {
  this: LoggerComponent => // Gives access to logger

  trait Auth { ... }
  val auth: Auth
  class MockAuth(file: String) extends Auth { ... }
  ...
}

```

Note the use of the self type to indicate that the authentication component requires the logger component.

Now the component configuration can happen in one central place:

```

object AppComponents extends LoggerComponent with AuthComponent {
  val logger = new FileLogger("test.log")
  val auth = new MockAuth("users.txt")
}

```

Either approach is better than component wiring in an XML file because the compiler can verify that the module dependencies are satisfied.

19.12 Abstract Types L3

A class or trait can define an *abstract type* that is made concrete in a subclass. For example:

```

trait Reader {
  type Contents
  def read(fileName: String): Contents
}

```

Here, the type *Contents* is abstract. A concrete subclass needs to specify the type:

```

class StringReader extends Reader {
    type Contents = String
    def read(fileName: String) = Source.fromFile(fileName, "UTF-8").mkString
}

class ImageReader extends Reader {
    type Contents = BufferedImage
    def read(fileName: String) = ImageIO.read(new File(fileName))
}

```

The same effect could be achieved with a type parameter:

```

trait Reader[C] {
    def read(fileName: String): C
}

class StringReader extends Reader[String] {
    def read(fileName: String) = Source.fromFile(fileName, "UTF-8").mkString
}

class ImageReader extends Reader[BufferedImage] {
    def read(fileName: String) = ImageIO.read(new File(fileName))
}

```

Which is better? In Scala, the rule of thumb is:

- Use type parameters when the types are supplied as the class is instantiated. For example, when you construct a `HashMap[String, Int]`, you want control over the types.
- Use abstract types when the types are expected to be supplied in a subclass. That is the case in our `Reader` example.

Nothing bad will happen if you specify type parameters as you form a subclass. But abstract types can work better when there are many type dependencies—you avoid long lists of type parameters. For example,

```

trait Reader {
    type In
    type Contents
    def read(in: In): Contents
}

class ImageReader extends Reader {
    type In = File
    type Contents = BufferedImage
    def read(file: In) = ImageIO.read(file)
}

```

With type parameters, `ImageReader` would extend `Reader[File, BufferedImage]`. That's still OK, but you can see that this technique doesn't scale so well in more complex situations.

Also, abstract types can express subtle interdependencies between types. The next section has an example.

Abstract types can have type bounds, just like type parameters. For example:

```
trait Listener {
    class EventObject { ... }
    type Event <: EventObject
    ...
}
```

A subclass must provide a compatible type, for example:

```
trait ActionListener extends Listener {
    class ActionEvent extends EventObject { ... }
    type Event = ActionEvent // OK, it's a subtype
    ...
}
```

Note that this example could not have been done with type parameters since the bound is an inner class.

19.13 Family Polymorphism L3

It is a challenge to model families of types that vary together, share common code, and preserve type safety. For example, consider event handling in client-side Java. There are different types of events (such as `ActionEvent`, `ChangeEvent`, and so on). Each type has a separate listener interface (`ActionListener`, `ChangeListener`, and so on). This is an example of "family polymorphism."

Let's design a generic mechanism for listener management. We'll start using generic types, and then switch to abstract types.

In Java, each listener interface has a different name for the method called when an event occurs: `actionPerformed`, `stateChanged`, `itemStateChanged`, and so on. We'll unify those:

```
trait Listener[E] {
    def occurred(e: E): Unit
}
```

An event source needs a collection of listeners, and a method to fire them all:

```
trait Source[E, L <: Listener[E]] {
    private val listeners = new ArrayBuffer[L]
    def add(l: L) { listeners += l }
    def remove(l: L) { listeners -= l }
    def fire(e: E) {
        for (l <- listeners) l.occurred(e)
    }
}
```

Now consider a button firing action events. We define a listener type

```
trait ActionListener extends Listener[ActionEvent]
```

The Button class can mix in the Source trait:

```
class Button extends Source[ActionEvent, ActionListener] {
    def click() {
        fire(new ActionEvent(this, ActionEvent.ACTION_PERFORMED, "click"))
    }
}
```

Mission accomplished: The Button class didn't need to replicate the code for listener management, and the listeners are typesafe. You can't add a ChangeListener to a button.

The ActionEvent class sets the event source to this, but the type of the event source is Object. We can make this typesafe with a self type:

```
trait Event[S] {
    var source: S = _
}

trait Listener[S, E <: Event[S]] {
    def occurred(e: E): Unit
}

trait Source[S, E <: Event[S], L <: Listener[S, E]] {
    this: S =>
    private val listeners = new ArrayBuffer[L]
    def add(l: L) { listeners += l }
    def remove(l: L) { listeners -= l }
    def fire(e: E) {
        e.source = this // Self type needed here
        for (l <- listeners) l.occurred(e)
    }
}
```

Note the self type `this: S =>` that is required for setting the source to this. Otherwise, this would only be some `Source`, not necessarily the one required by `Event[S]`.

Here is how you define a button:

```
class ButtonEvent extends Event[Button]

trait ButtonListener extends Listener[Button, ButtonEvent]

class Button extends Source[Button, ButtonEvent, ButtonListener] {
  def click() { fire(new ButtonEvent) }
}
```

You can see the proliferation of the type parameters. With abstract types, this looks a little nicer.

```
trait ListenerSupport {
  type S <: Source
  type E <: Event
  type L <: Listener

  trait Event {
    var source: S = _
  }

  trait Listener {
    def occurred(e: E): Unit
  }

  trait Source {
    this: S =>
    private val listeners = new ArrayBuffer[L]
    def add(l: L) { listeners += l }
    def remove(l: L) { listeners -= l }
    def fire(e: E) {
      e.source = this
      for (l <- listeners) l.occurred(e)
    }
  }
}
```

But there is a price to pay. You can't have top-level type declarations. That's why everything is wrapped in the `ListenerSupport` trait.

Now when you want to define a button with a button event and a button listener, enclose the definitions in a module that extends this trait:

```

object ButtonModule extends ListenerSupport {
    type S = Button
    type E = ButtonEvent
    type L = ButtonListener

    class ButtonEvent extends Event
    trait ButtonListener extends Listener
    class Button extends Source {
        def click() { fire(new ButtonEvent) }
    }
}

```

To use the button, import the module:

```

object Main {
    import ButtonModule._

    def main(args: Array[String]) {
        val b = new Button
        b.add(new ButtonListener {
            override def occurred(e: ButtonEvent) { println(e) }
        })
        b.click()
    }
}

```



NOTE: In this example, I used single-letter names for the abstract types, to show the analogy with the version that uses type parameters. It is common in Scala to use more descriptive type names, which leads to more self-documenting code:

```

object ButtonModule extends ListenerSupport {
    type SourceType = Button
    type EventType = ButtonEvent
    type ListenerType = ButtonListener
    ...
}

```

19.14 Higher-Kinded Types L3

The generic type `List` depends on a type `T` and produces a type. For example, given the type `Int`, you get the type `List[Int]`. For that reason, a generic type such as `List` is sometimes called a *type constructor*. In Scala, you can go up another level and define types that depend on types that depend on types.

To see why this can be useful, consider the following simplified `Iterable` trait:

```
trait Iterable[E] {
  def iterator(): Iterator[E]
  def map[F](f: (E) => F): Iterable[F]
}
```

Now consider a class implementing this trait:

```
class Buffer[E] extends Iterable[E] {
  def iterator(): Iterator[E] = ...
  def map[F](f: (E) => F): Buffer[F] = ...
}
```

For a buffer, we expect that `map` returns a `Buffer`, not a mere `Iterable`. That means we cannot implement `map` in the `Iterable` trait itself, which we would like to do. A remedy is to parameterize the `Iterable` with a type constructor, like this:

```
trait Iterable[E, C[_]] {
  def iterator(): Iterator[E]
  def build[F](): C[F]
  def map[F](f: (E) => F): C[F]
}
```

Now an `Iterable` depends on a type constructor for the result, denoted as `C[_]`. This makes `Iterable` a higher-kinded type.

The type returned by `map` may or may not be the same as the type of the `Iterable` on which `map` was invoked. For example, if you invoke `map` on a `Range`, the result is not generally a range, so `map` must construct a different type such as a `Buffer[F]`. Such a `Range` type is declared as

```
class Range extends Iterable[Int, Buffer]
```

Note that the second parameter is the type constructor `Buffer`.

To implement `map` in `Iterable`, we need a bit more support. An `Iterable` needs to be able to produce a container that holds values of any type `F`. Let's define a `Container` trait—it is something to which you can add values:

```
trait Container[E] {
  def +=(e: E): Unit
}
```

The `build` method is required to yield such an object:

```
trait Iterable[E, C[X] <: Container[X]] {
  def build[F](): C[F]
  ...
}
```

The type constructor C has now been constrained to be a Container, so we know that we can add items to the object that `build` returns. We can no longer use a wildcard for the parameter of C since we need to indicate that $C[X]$ is a container for the same X .



NOTE: The Container trait is a simpler version of the builder mechanism that is used in the Scala collections library.

The `map` method can be implemented in the `Iterable` trait:

```
def map[F](f : (E) => F) : C[F] = {
    val res = build[F]()
    val iter = iterator()
    while (iter.hasNext) res += f(iter.next())
    res
}
```

`Iterable` classes no longer need to supply their own `map`. Here is the `Range` class:

```
class Range(val low: Int, val high: Int) extends Iterable[Int, Buffer] {
    def iterator() = new Iterator[Int] {
        private var i = low
        def hasNext = i <= high
        def next() = { i += 1; i - 1 }
    }

    def build[F]() = new Buffer[F]
}
```

Note that a `Range` is an `Iterable`: You can iterate over its contents. But it is not a `Container`: You can't add values to it.

A `Buffer`, on the other hand, is both:

```
class Buffer[E : ClassTag] extends Iterable[E, Buffer] with Container[E] {
    private var capacity = 10
    private var length = 0
    private var elems = new Array[E](capacity) // See note

    def iterator() = new Iterator[E] {
        private var i = 0
        def hasNext = i < length
        def next() = { i += 1; elems(i - 1) }
    }

    def build[F : ClassTag]() = new Buffer[F]
```

```

def +=(e: E) {
  if (length == capacity) {
    capacity = 2 * capacity
    val nelems = new Array[E](capacity) // See note
    for (i <- 0 until length) nelems(i) = elems(i)
    elems = nelems
  }
  elems(length) = e
  length += 1
}
}

```



NOTE: There is one additional complexity in this example, and it has nothing to do with higher-kinded types. In order to construct a generic `Array[E]`, the type `E` must fulfill the `ClassTag` context bound that was discussed in Chapter 18.



NOTE: To use higher-kinded types without a warning, add the statement `import scala.language.higherKinds` or use the compiler option `-language:higherKinds`.

This example showed a typical use of higher-kinded types. An `Iterator` depends on `Container`, but `Container` isn't a type—it is a mechanism for making types.

The `Iterable` trait of the Scala collections library doesn't have an explicit parameter for making collections. Instead, Scala uses an *implicit parameter* to conjure up an object for building the target collection. See Chapter 21 for more information.

Exercises

1. Implement a `Bug` class modeling a bug that moves along a horizontal line. The `move` method moves in the current direction, the `turn` method makes the bug turn around, and the `show` method prints the current position. Make these methods chainable. For example,

```
bugsy.move(4).show().move(6).show().turn().move(5).show()
```

should display 4 10 5.

2. Provide a fluent interface for the `Bug` class of the preceding exercise, so that one can write

```
bugsy move 4 and show and then move 6 and show turn around move 5 and show
```

3. Complete the fluent interface in Section 19.1, “Singleton Types,” on page 280 so that one can call

```
book.set Title to "Scala for the Impatient" set Author to "Cay Horstmann"
```

4. Implement the `equals` method for the `Member` class that is nested inside the `Network` class in Section 19.2, “Type Projections,” on page 281. For two members to be equal, they need to be in the same network.

5. Consider the type alias

```
type NetworkMember = n.Member forSome { val n: Network }
```

and the function

```
def process(m1: NetworkMember, m2: NetworkMember) = (m1, m2)
```

How does this differ from the `process` function in Section 19.8, “Existential Types,” on page 286?

6. The `Either` type in the Scala library can be used for algorithms that return either a result or some failure information. Write a function that takes two parameters: a sorted array of integers and an integer value. Return either the index of the value in the array or the index of the element that is closest to the value. Use an infix type as the return type.

7. Implement a method that receives an object of any class that has a method

```
def close(): Unit
```

together with a function that processes that object. Call the function and invoke the `close` method upon completion, or when any exception occurs.

8. Write a function `printValues` with three parameters `f`, `from`, `to` that prints all values of `f` with inputs from the given range. Here, `f` should be any object with an `apply` method that consumes and yields an `Int`. For example,

```
printValues((x: Int) => x * x, 3, 6) // Prints 9 16 25 36
printValues(Array(1, 1, 2, 3, 5, 8, 13, 21, 34, 55), 3, 6) // Prints 3 5 8 13
```

9. Consider this class that models a physical dimension:

```
abstract class Dim[T](val value: Double, val name: String) {
  protected def create(v: Double): T
  def +(other: Dim[T]) = create(value + other.value)
  override def toString() = s"$value $name"
}
```

Here is a concrete subclass:

```
class Seconds(v: Double) extends Dim[Seconds](v, "s") {
  override def create(v: Double) = new Seconds(v)
}
```

But now a knucklehead could define

```
class Meters(v: Double) extends Dim[Seconds](v, "m") {  
    override def create(v: Double) = new Seconds(v)  
}
```

allowing meters and seconds to be added. Use a self type to prevent that.

10. Self types can usually be replaced with traits that extend classes, but there can be situations where using self types changes the initialization and override orders. Construct such an example.

Parsing

Topics in This Chapter **A3**

- 20.1 Grammars — page 304
- 20.2 Combining Parser Operations — page 305
- 20.3 Transforming Parser Results — page 307
- 20.4 Discarding Tokens — page 308
- 20.5 Generating Parse Trees — page 309
- 20.6 Avoiding Left Recursion — page 310
- 20.7 More Combinators — page 311
- 20.8 Avoiding Backtracking — page 314
- 20.9 Packrat Parsers — page 314
- 20.10 What Exactly Are Parsers? — page 315
- 20.11 Regex Parsers — page 316
- 20.12 Token-Based Parsers — page 317
- 20.13 Error Handling — page 319
- Exercises — page 320

Chapter 20

In this chapter, you will see how to use the “parser combinators” library to analyze data with fixed structure. Examples of such data are programs in a programming language or data in formats such as HTTP or JSON. Not everyone needs to write parsers for these languages, so you may not find this chapter useful for your work. If you are familiar with the basic concepts of grammars and parsers, glance through the chapter anyway because the Scala parser library is a good example of a sophisticated *domain-specific language* embedded in the Scala language.



NOTE: The API documentation for Scala parser combinators is at www.scala-lang.org/api/current/scala-parser-combinators.

The key points of this chapter are:

- Alternatives, concatenation, options, and repetitions in a grammar turn into `|`, `~`, `opt`, and `rep` in Scala combinator parsers.
- With `RegexParsers`, literal strings and regular expressions match tokens.
- Use `^` to process parse results.
- Use pattern matching in a function supplied to `^` to take apart `~` results.
- Use `~>` and `<~` to discard tokens that are no longer needed after matching.
- The `repsep` combinator handles the common case of repeated items with a separator.

- A token-based parser is useful for parsing languages with reserved words and operators. Be prepared to define your own lexer.
- Parsers are functions that consume a reader and yield a parse result: success, failure, or error.
- For a practical parser, you need to implement robust error reporting.
- Thanks to operator symbols, implicit conversions, and pattern matching, the parser combinator library makes parser writing easy for anyone who understands context-free grammars. Even if you don't feel the urge to write your own parsers, you may find this an interesting case study for an effective domain-specific language.

20.1 Grammars

To understand the Scala parsing library, you need to know a few concepts from the theory of formal languages. A *grammar* is a set of rules for producing all strings that follow a particular format. For example, we can say that an arithmetic expression is given by the following rules:

- Each whole number is an arithmetic expression.
- `+` `-` `*` are operators.
- If *left* and *right* are arithmetic expressions and *op* is an operator, then *left op right* is an arithmetic expression.
- If *expr* is an arithmetic expression, then `(expr)` is an arithmetic expression.

According to these rules, `3+4` and `(3+4)*5` are arithmetic expressions, but `3+` or `3*x` are not.

A grammar is usually written in a notation called Backus-Naur Form (BNF). Here is the BNF for our expression language:

```
op ::= "+" | "-" | "*"
expr ::= number | expr op expr | "(" expr ")"
```

Here, `number` is undefined. We could define it as

```
digit ::= "0" | "1" | "2" | "3" | "4" | "5" | "6" | "7" | "8" | "9"
number ::= digit | digit number
```

But in practice, it is more efficient to collect numbers before parsing starts, in a separate step called *lexical analysis*. A *lexer* discards whitespace and comments, and forms *tokens*—identifiers, numbers, or symbols. In our expression language, tokens are `number` and the symbols `+` `-` `*` `(` `)`.

Note that `op` and `expr` are not tokens. They are structural elements that were invented by the author of the grammar, in order to produce correct token sequences.

Such symbols are called *nonterminal symbols*. One of the nonterminal symbols is at the root of the hierarchy; in our case, that is `expr`. It is called the *start symbol*. To produce correctly formatted strings, you start with the start symbol and apply the grammar rules until all nonterminals have been replaced and only tokens remain. For example, the derivation

```
expr -> expr op expr -> number op expr ->
-> number "+" expr -> number "+" number
```

shows that `3+4` is a valid expression.

The most often used “extended Backus-Naur form,” or EBNF, allows specifying optional elements and repetition. I will use the familiar regex operators `? *` for 0 or 1, 0 or more, 1 or more, correspondingly. For example, a comma-separated list of numbers can be described with the grammar

```
numberList ::= number ( "," numberList )?
```

or with

```
numberList ::= number ( "," number )*
```

As another example of EBNF, let’s make an improvement to the grammar for arithmetic expressions to support operator precedence. Here is the revised grammar:

```
expr ::= term ( ( "+" | "-" ) expr )?
term ::= factor ( "*" factor )*
factor ::= number | "(" expr ")"
```

20.2 Combining Parser Operations

To use the Scala parsing library, provide a class that extends the `Parsers` trait and defines parsing operations that are combined from primitive operations, such as

- Matching a token
- Choosing between two operations (`|`)
- Performing two operations in sequence (`~`)
- Repeating an operation (`rep`)
- Optionally performing an operation (`opt`)

The following parser recognizes arithmetic expressions. It extends `RegexParsers`, a subtrait of `Parsers` that can match tokens against regular expressions. Here, we specify `number` with the regular expression `"[0-9]+".r`:

```
class ExprParser extends RegexParsers {
    val number = "[0-9]+".r

    def expr: Parser[Any] = term ~ opt(("+" | "-") ~ expr)
    def term: Parser[Any] = factor ~ rep("*" ~ factor)
    def factor: Parser[Any] = number | "(" ~ expr ~ ")"
}
```

Note that the parser is a straightforward translation from the EBNF of the preceding section.

Simply use the `~` operator to join the parts, and use `opt` and `rep` instead of `?` and `*`.

In our example, each function has return type `Parser[Any]`. This type isn't very useful, and we will improve it in the next section.

To run the parser, invoke the inherited `parse` method, for example:

```
val parser = new ExprParser
val result = parser.parseAll(parser.expr, "3-4*5")
if (result.successful) println(result.get)
```

The `parseAll` method receives the method to be invoked—that is, the method associated with the grammar's start symbol—and the string to be parsed.



NOTE: There is also a `parse` method that parses a prefix of a string, stopping when it can't find another match. That method isn't very useful; for example, `parser.parse(parser.expr, "3-4/5")` parses 3-4, then quietly stops at the / which it cannot handle.

The output of the program snippet is

```
((3~List())~Some((~((4~List((*~5))~None))))
```

To interpret this output, you need to know the following:

- Literal strings and regular expressions return `String` values.
- `p ~ q` returns an instance of the `~` case class, which is very similar to a pair.
- `opt(p)` returns an `Option`, either `Some(...)` or `None`.
- `rep(p)` returns a `List`.

The call to `expr` returns the result from `term` (shown in bold) joined with something optional—the `Some(...)` part which I won't analyze.

Since `term` is defined as

```
def term = factor ~ rep(("*" | "/" ) ~ factor)
```

it returns the result from factor joined with a List. That's an empty list because there is no * in the subexpression to the left of the -.

Of course, this result is quite tedious. In the next section, you will see how to transform it to something more useful.

20.3 Transforming Parser Results

Instead of having a parser build up a complex structure of ~, options, and lists, you should transform intermediate outputs to a useful form. Consider, for example, the arithmetic expression parser. If it is our goal to evaluate the expression, then each of the functions expr, term, and factor should return the value of the parsed subexpression. Let's start with

```
def factor: Parser[Any] = number | "(" ~ expr ~ ")"
```

We want it to return an Int:

```
def factor: Parser[Int] = ...
```

When a whole number is received, we want its integer value:

```
def factor: Parser[Int] = number ^^ { _.toInt } | ...
```

Here, the ^^ operator applies the function { _.toInt } to the result of number.



NOTE: There is no particular significance to the ^^ symbol. It conveniently has lower precedence than ~ but higher precedence than |.

Assuming that expr has been changed to return a Parser[Int], we can evaluate "(" ~ expr ~ ")" simply by returning expr, which yields an Int. Here is one way of doing that (you'll see a simpler one in the next section):

```
def factor: Parser[Int] = ... | "(" ~ expr ~ ")" ^^ {
  case _ ~ e ~ _ => e
}
```

In this case, the argument of the ^^ operator is the partial function { case _ ~ e ~ _ => e }.



NOTE: The ~ combinator returns an instance of the ~ case class instead of a pair to make matching easier. If ~ returned a pair, then you would have to write case ((_, e), _) instead of case _ ~ e ~ _.

A similar pattern match yields the sum or difference. Note that opt yields an Option: either None or Some(...).

```
def expr: Parser[Int] = term ~ opt(("+" | "-") ~ expr) ^^ {
    case t ~ None => t
    case t ~ Some("+ " ~ e) => t + e
    case t ~ Some("- " ~ e) => t - e
}
```

Finally, to multiply the factors, note that `rep("*" ~ factor)` yields a `List` of items of the form `"*" ~ f`, where `f` is an `Int`. We extract the second component of each `~` pair and compute their product:

```
def term: Parser[Int] = factor ~ rep("*" ~ factor) ^^ {
    case f ~ r => f * r.map(_._2).product
}
```

In this example, we simply computed the value of the expression. When building a compiler or interpreter, the usual goal is to build a *parse tree*—a tree structure that describes the parsed result; see Section 20.5, “Generating Parse Trees,” on page 309.



CAUTION: If you turn off the warning against postfix operators, you can write `p?` instead of `opt(p)` and `p*` instead of `rep(p)`:

```
def expr: Parser[Any] = term ~ (( "+" | "-" ) ~ expr)?
def term: Parser[Any] = factor ~ ("*" ~ factor)*
```

It seems a good idea to use these familiar operators, but they conflict with the `^^` operator. You’ll have to add another set of parentheses, such as

```
def term: Parser[Any] = factor ~ (( "*" ~ factor )*) ^^ { ... }
```

For that reason, I prefer `opt` and `rep`.

20.4 Discarding Tokens

As you saw in the preceding section, it can be tedious to deal with tokens when analyzing a match. The tokens are required for parsing, but they can often be discarded after they have been matched. The `~>` and `<~` operators are used to match and discard a token. For example, the result of `"*" ~> factor` is just the result of `factor`, not a value of the form `"*" ~ f`. With that notation, we can simplify the `term` function to

```
def term = factor ~ rep("*" ~> factor) ^^ {
    case f ~ r => f * r.product
}
```

Similarly, we can discard the parentheses around an expression, like this:

```
def factor = number ^& { _.toInt } | "(" ~> expr <~ ")"
```

A transformation is no longer required in the expression "`" ~> expr <~ "`", since the value is now simply `e`, which already yields an `Int`.

Note that the “arrow tip” of the `~>` or `<~` operator points to the part that is retained.



CAUTION: You need to be careful when using multiple `~`, `~>`, and `<~` in the same expression. `<~` has a lower precedence than `~` and `~>`. Consider, for example:

```
"if" ~> "(" ~> expr <~ ")" ~ expr
```

Unfortunately, this doesn’t just discard `")"`, but the subexpression `")" ~ expr`. The remedy is to use parentheses: `"if" ~> "(" ~> (expr <~ ")"")" ~ expr`.

20.5 Generating Parse Trees

The parsers of the preceding examples simply computed numeric results. When you build an interpreter or compiler, you will want to build up a parse tree instead. This is usually done with case classes. For example, the following classes can represent an arithmetic expression:

```
class Expr
case class Number(value: Int) extends Expr
case class Operator(op: String, left: Expr, right: Expr) extends Expr
```

The parser’s job is to turn an input such as `3+4*5` into a value

```
Operator("+", Number(3), Operator("*", Number(4), Number(5)))
```

In an interpreter, such an expression can be evaluated. In a compiler, it can be used for generating code.

To generate a parse tree, use `^&` with functions that yield tree nodes. For example,

```
class ExprParser extends RegexParsers {
  ...
  def term: Parser[Expr] = (factor ~ opt("*" ~> term)) ^& {
    case a ~ None => a
    case a ~ Some(b) => Operator("*", a, b)
  }
  def factor: Parser[Expr] = wholeNumber ^& (n => Number(n.toInt)) |
    "(" ~> expr <~ ")"
}
```

20.6 Avoiding Left Recursion

If a parser function calls itself without first consuming some input, the recursion will never stop. Consider this function that is supposed to consume any sequence of ones:

```
def ones: Parser[Any] = ones ~ "1" | "1"
```

Such a function is called *left-recursive*. To avoid the recursion, you can reformulate the grammar. Here are two alternatives:

```
def ones: Parser[Any] = "1" ~ ones | "1"
```

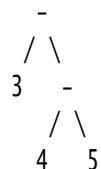
or

```
def ones: Parser[Any] = rep1("1")
```

This problem occurs commonly in practice. For example, consider our arithmetic expression parser:

```
def expr: Parser[Any] = term ~ opt(("+" | "-") ~ expr)
```

The rule for `expr` has an unfortunate effect with subtraction. The expressions are grouped in the wrong order. When the input is 3-4-5, the expression is parsed as



That is, 3 is accepted as `term`, and -4-5 as `"-" ~ expr`. This yields a wrong answer of 4 instead of -6.

Could we turn the grammar around?

```
def expr: Parser[Any] = expr ~ opt(("+" | "-") ~ term)
```

Then we would get the correct parse tree. But that doesn't work—this `expr` function is left-recursive.

The original version eliminates the left recursion, but at a cost—it is harder to compute the parse result. You need to collect the intermediate results and then combine them in the correct order.

Collecting intermediate results is easier if you can use a repetition, which yields a List of the collected values. For example, an `expr` is a sequence of `term` values, joined by + or -:

```
def expr: Parser[Any] = term ~ rep(("+" | "-") ~ term)
```

To evaluate the expression, replace each $s \sim t$ in the repetition with t or $-t$, depending on whether s is "+" or "-". Then compute the sum of the list.

```
def expr: Parser[Int] = term ~ rep(
  ("+" | "-") ~ term ^^ {
    case "+" ~ t => t
    case "-" ~ t => -t
  }) ^^ { case t ~ r => t + r.sum }
```

If rewriting the grammar is too cumbersome, see Section 20.9, "Packrat Parsers," on page 314 for another remedy.

20.7 More Combinators

The `rep` method matches zero or more repetitions. Table 20–1 shows several variations of this combinator. The most commonly used among them is `repsep`. For example, a list of comma-separated numbers can be defined as

```
def numberList = number ~ rep(", " ~> number)
```

or more concisely as

```
def numberList = repsep(number, ",")
```

Table 20–2 shows additional combinators that are occasionally useful. The `into` combinator can come in handy to store information from an earlier combinator in a variable so that it can be used later. For example, in the grammar rule

```
def term: Parser[Any] = factor ~ rep("*" ~> factor)
```

you can store the first factor in a variable, like this:

```
def term: Parser[Int] = factor into { first =>
  rep("*" ~> factor) ^^ { first * _.product }
}
```

The `log` combinator can help debug a grammar. Replace a parser `p` with `log(p)(str)`, and you get a printout whenever `p` is called. For example,

```
def factor: Parser[Int] = log(number)("number") ^^ { _.toInt } | ...
```

yields outputs such as

```
trying number at scala.util.parsing.input.CharSequenceReader@76f7c5
number --> [1.2] parsed: 3
```

Table 20–1 Combinators for Repetitions

Combinator	Description	Notes
<code>rep(p)</code>	0 or more matches of <code>p</code> .	
<code>rep1(p)</code>	1 or more matches of <code>p</code> .	<code>rep1("[" ~> expr <~ "]")</code> yields a list of expressions that were included inside brackets—for example, to specify bounds of a multidimensional array.
<code>rep1(p, q), p and q are Parser[P]</code>	1 match of <code>p</code> followed by 0 or more matches of <code>q</code> .	
<code>repN(n, p)</code>	<code>n</code> matches of <code>p</code> .	<code>repN(4, number)</code> matches a sequence of four numbers—for example, to specify a rectangle.
<code>repsep(p, s) rep1sep(p, s) p is a Parser[P]</code>	0 or more/1 or more matches of <code>p</code> , separated by matches of <code>s</code> . The result is a <code>List[P]</code> ; the <code>s</code> are discarded.	<code>repsep(expr, ",")</code> yields a list of expressions that were separated by commas. Useful for parsing the arguments to a function call.
<code>chainl1(p, s)</code>	Like <code>rep1sep</code> , but <code>s</code> must, upon matching each separator, produce a binary function that is used to combine neighboring values. If <code>p</code> produces values v_0, v_1, v_2, \dots and <code>s</code> produces functions f_1, f_2, \dots , then the result is $(v_0 f_1 v_1) f_2 v_2 \dots$	<code>chainl1(number ^& { _.toInt }, "*" ^& { _ * _})</code> computes the product of a sequence of integers separated by <code>*</code> .

Table 20–2 Additional Combinators

Combinator	Description	Notes
<code>p ^&& v</code>	Like <code>^&</code> , but returns a constant result.	Useful for parsing literals: "true" ^& true.
<code>p into f or p >> f</code>	<code>f</code> is a function whose argument is the result of <code>p</code> . Useful to bind the result of <code>p</code> to a variable.	<code>(number ^& { _.toInt }) >> { n => repN(n, number) }</code> parses a sequence of numbers, where the first number indicates how many numbers follow.

(Continues)

Table 20–2 Additional Combinators (*Continued*)

Combinator	Description	Notes
<code>p ^? f</code>	Like <code>^&</code> , but takes a partial function <code>f</code> . Fails if <code>f</code> can't be applied to the result of <code>p</code> . In the second version, <code>error</code> is a function from the result type of <code>p</code> , yielding an error message string.	<code>ident ^? (symbols, "undefined symbol " + _)</code> looks up <code>ident</code> in the map <code>symbols</code> and reports an error if it is not contained in the map. Note that a map can be converted to a partial function.
<code>log(p)(str)</code>	Executes <code>p</code> and prints a logging message.	<code>log(number)("number") ^& { _.toInt }</code> prints a message whenever a number is parsed.
<code>guard(p)</code>	Calls <code>p</code> , succeeds or fails, then restores the input as if <code>p</code> had not been called.	Useful for looking ahead. For example, to distinguish between a variable and a function call, you can use a <code>guard(ident ~ "(")</code> .
<code>not(p)</code>	Calls <code>p</code> and succeeds if <code>p</code> fails, or fails if <code>p</code> succeeds.	
<code>p ~! q</code>	Like <code>~</code> , but if the second match fails, the failure turns into an error, inhibiting backtracking in an enclosing <code> </code> .	See Section 20.8.
<code>accept(descr, f)</code>	Accepts an item that is accepted by the partial function <code>f</code> , returning the function result. The string <code>descr</code> is used to describe the expected item in the failure message.	<code>accept("string literal", { case t: lexical.StringLit => t.chars })</code>
<code>success(v)</code>	Always succeeds with value <code>v</code> .	Can be used to add a value <code>v</code> to a result.
<code>failure(msg)</code> <code>err(msg)</code>	Fails with the given error message.	See Section 20.13 on how to improve error messages.
<code>phrase(p)</code>	Succeeds if <code>p</code> succeeds and no input is left over.	Useful for defining a <code>parseAll</code> method; see, for example, Section 20.12.
<code>positioned(p)</code>	Adds position to the result of <code>p</code> (must extend <code>Positional</code>).	Useful for reporting errors after parsing has finished.

20.8 Avoiding Backtracking

Whenever an alternative $p \mid q$ is parsed and p fails, the parser tries q on the same input. This is called *backtracking*. Backtracking also happens when there is a failure in an `opt` or `rep`. Clearly, backtracking can be inefficient. For example, consider an arithmetic expression parser with the rules

```
def expr: Parser[Any] = term ~ ("+" | "-") ~ expr | term
def term: Parser[Any] = factor ~ "*" ~ term | factor
def factor: Parser[Any] = "(" ~ expr ~ ")" | number
```

If the expression $(3+4)*5$ is parsed, then `term` matches the entire input. Then the match for `+` or `-` fails, and the compiler backtracks to the second alternative, parsing `term` again.

It is often possible to rearrange the grammar rules to avoid backtracking. For example:

```
def expr: Parser[Any] = term ~ opt(("+" | "-") ~ expr)
def term: Parser[Any] = factor ~ rep("*" ~ factor)
```

You can then use the `~!` operator instead of `~` to express that there is no need to backtrack.

```
def expr: Parser[Any] = term ~ opt(("+" | "-") ~! expr)
def term: Parser[Any] = factor ~ rep("*" ~! factor)
def factor: Parser[Any] = "(" ~! expr ~! ")" | number
```

When $p \sim! q$ is evaluated and q fails, no other alternatives are tried in an enclosing `l`, `opt`, or `rep`. For example, if `factor` finds a `("` and then `expr` doesn't match, the parser won't even try matching `number`.

20.9 Packrat Parsers

A packrat parser uses an efficient parsing algorithm that caches previous parse results. This has two advantages:

- Parse time is guaranteed to be proportional to the length of the input.
- The parser can accept left-recursive grammars.

In order to use packrat parsing in Scala, follow these steps:

1. Mix the `PackratParsers` trait into your parser.
2. Use `val` or `lazy val`, not `def`, for each parser function. This is important because the parser caches these values, and it relies on them being identical. (A `def` would return a different value each time it is called.)
3. Have each parser function return `PackratParser[T]` instead of `Parser[T]`.

4. Use a PackratReader and supply a `parseAll` method (which is annoyingly missing from the `PackratParsers` trait).

For example,

```
class OnesPackratParser extends RegexParsers with PackratParsers {
  lazy val ones: PackratParser[Any] = ones ~ "1" | "1"

  def parseAll[T](p: Parser[T], input: String) =
    phrase(p)(new PackratReader(new CharSequenceReader(input)))
}
```

20.10 What Exactly Are Parsers?

Technically, a `Parser[T]` is a function with one argument, of type `Reader[Elem]`, and a return value of type `ParseResult[T]`. In this section, we will have a closer look at these types.

The type `Elem` is an abstract type of the `Parsers` trait. (See Section 19.12, “Abstract Types,” on page 291 for more information about abstract types.) The `RegexParsers` trait defines `Elem` as `Char`, and the `StdTokenParsers` trait defines `Elem` as `Token`. (We will have a look at token-based parsing in Section 20.12, “Token-Based Parsers,” on page 317.)

A `Reader[Elem]` reads a sequence of `Elem` values (that is, characters or tokens) from some source and tracks their positions for error reporting.

When a `Parser[T]` is invoked on a reader, it returns an object of one of three subclasses of `ParseResult[T]`: `Success[T]`, `Failure`, or `Error`.

An `Error` terminates the parser and anything that called it. It can arise in one of these circumstances:

- A parser `p ~! q` fails to match `q`.
- A `commit(p)` fails.
- The `err(msg)` combinator is encountered.

A `Failure` simply arises from a failure to match; it normally triggers alternatives in an enclosing `|`.

A `Success[T]` has, most importantly, a `result` of type `T`. It also has a `Reader[Elem]` called `next`, containing the input beyond the match that is yet to be consumed.

Consider this part of our arithmetic expression parser:

```
val number = "[0-9]+".r
def expr = number | "(" ~ expr ~ ")"
```

Our parser extends `RegexParsers`, which has an implicit conversion from a `Regex` to a `Parser[String]`. The regular expression `number` is converted into such a parser—a function that consumes a `Reader[Char]`.

If the initial characters in the reader match the regular expression, the function returns a `Success[String]`. The `result` property of the returned object is the matched input, and the `next` property is the reader with the match removed.

If the initial characters in the reader don't match the regular expression, the parser function returns a `Failure` object.

The `|` method combines two parsers. That is, if `p` and `q` are functions, then `p | q` is again a function. The combined function consumes a reader, say `r`. It calls `p(r)`. If that call returns `Success` or `Error`, then that's the return value of `p | q`. Otherwise, the return value is the result of `q(r)`.

20.11 Regex Parsers

The `RegexParsers` trait, which we have used in all examples up to this point, provides two implicit conversions for defining parsers:

- `literal` makes a `Parser[String]` from a literal string (such as `"+"`).
- `regex` makes a `Parser[String]` from a regular expression (such as `"[0-9]".r`).

By default, regex parsers skip whitespace. If your notion of whitespace is different from the default of `"""\s+""".r` (for example, if you want to skip comments), override `whiteSpace` with your definition. If you don't want whitespace skipped, use

```
override val whiteSpace = """.r
```

The `JavaTokenParsers` trait extends `RegexParsers` and specifies five tokens, shown in Table 20–3. None of them correspond exactly to their Java forms, which makes that trait of limited utility.

Table 20–3 Predefined Tokens in `JavaTokenParsers`

Token	Regular Expression
<code>ident</code>	<code>[a-zA-Z_]\w*</code>
<code>wholeNumber</code>	<code>-?\d+</code>
<code>decimalNumber</code>	<code>\d+(\.\d*)? \d*\.\d+</code>
<code>stringLiteral</code>	<code">([^\p{Cntrl}"] \\[\\"/bf\nrt] \\u[a-fA-F0-9]{4})*</code">
<code>floatingPointNumber</code>	<code>-?(\d+(.\d*)? \d*\.\d+)([eE][+-]?\d+)?[fFdD]?</code>

20.12 Token-Based Parsers

Token-based parsers use a Reader[Token] instead of a Reader[Char]. The Token type is defined in the trait `scala.util.parsing.combinator.token.Tokens`. The `StdTokens` subtrait defines four types of tokens that one commonly finds when parsing a programming language:

- Identifier
- Keyword
- NumericLit
- StringLit

The `StandardTokenParsers` class provides a parser that produces these tokens. Identifiers consist of letters, digits, or `_` but don't start with a digit.



CAUTION: The rules for letters and digits are subtly different from those in Java or Scala. Digits in any script are supported, but letters in the “supplementary” range (above U+FFFF) are excluded.

Numeric literals are sequences of digits. String literals are enclosed in `"..."` or `'...'`, with no escapes. Comments, enclosed in `/* ... */` or from `//` to the end of the line, are considered whitespace.

When you extend this parser, add any reserved words and special tokens to the `lexical.reserved` and `lexical.delimiters` sets:

```
class MyLanguageParser extends StandardTokenParser {
    lexical.reserved += ("auto", "break", "case", "char", "const", ...)
    lexical.delimiters += ("=", "<", "<=", ">", ">=", "==", "!=")
    ...
}
```

When a reserved word is encountered, it becomes a `Keyword`, not an `Identifier`.

The parser sorts the delimiters according to the “maximum munch” rule. For example, when the input contains `<=`, you will get that as a single token, not as a sequence of tokens `<` and `=`.

The `ident` function parses an identifier; `numericLit` and `stringLit` parse literals.

For example, here is our arithmetic expression grammar, using `StandardTokenParsers`:

```
class ExprParser extends StandardTokenParsers {
    lexical.delimiters += ("+", "-", "*", "(", ")")

    def expr: Parser[Any] = term ~ rep(("+" | "-") ~ term)
    def term: Parser[Any] = factor ~ rep("*" ~> factor)
```

```
def factor: Parser[Any] = numericLit | "(" ~> expr <~ ")"

def parseAll[T](p: Parser[T], in: String): ParseResult[T] =
  phrase(p)(new lexical.Scanner(in))
}
```

Note that you need to supply a `parseAll` method, which is annoyingly missing from the `StandardTokenParsers` class. In that method, you use a `lexical.Scanner`, which is the `Reader[Token]` supplied by the `StdLexical` trait.



TIP: If you need to process a language with different tokens, it is easy to adapt the token parser. Extend `StdLexical` and override the `token` method to recognize the token types you need. Consult the source code of `StdLexical` for guidance—it is quite short. Then extend `StdTokenParsers` and override `lexical`:

```
class MyParser extends StdTokenParsers {
  val lexical = new MyLexical
  ...
}
```



TIP: The `token` method in `StdLexical` is a bit tedious. It's nicer to define tokens with regular expressions. Add this definition when you extend `StdLexical`:

```
def regex(r: Regex): Parser[String] = new Parser[String] {
  def apply(in: Input) = r.findPrefixMatchOf(
    in.source.subSequence(in.offset, in.source.length)) match {
    case Some(matched) =>
      Success(in.source.subSequence(in.offset,
        in.offset + matched.end).toString, in.drop(matched.end))
    case None =>
      Failure("string matching regex '$r' expected but
        ${in.first} found", in)
  }
}
```

Then you can use regular expressions in your `token` method, like this:

```
override def token: Parser[Token] = {
  regex("[a-zA-Z][a-zA-Z0-9]*".r) ^^ { processIdent(_) } |
  regex("0|[1-9][0-9]*".r) ^^ { NumericLit(_) } |
  ...
}
```

20.13 Error Handling

When a parser can't accept an input, you want to get an accurate message indicating where the failure occurred.

The parser generates an error message that describes the position at which the parser was unable to continue. If there were several failure points, the one that was visited last is reported.

You may want to keep error reporting in mind when you order alternatives. For example, if you have a rule

```
def value: Parser[Any] = numericLit | "true" | "false"
```

and the parser doesn't match any of them, then it's not so useful to know that the input failed to match "false". You can add a failure clause with an explicit error message:

```
def value: Parser[Any] = numericLit | "true" | "false" |
  failure("Not a valid value")
```

The failure combinator only reports an error when it is visited. It does not change the error messages that another combinator reports. For example, a RegexParser has an error message such as

```
string matching regex '\d+' expected but 'x' found
```

Then use the withFailureMessage method, like this:

```
def value = opt(sign) ~ digits withFailureMessage "Not a valid number"
```

When the parser fails, the parseAll method returns a Failure result. Its msg property is an error message that you can display to the user. The next property is the Reader that points to the unconsumed input at the point of failure. You will want to display the line number and column, which are available as next.pos.line and next.pos.column.

Finally, next.first is the lexical element at which the failure occurred. If you use the RegexParsers trait, that element is a Char, which is not very useful for error reporting. But with a token parser, next.first is a token, which is worth reporting.



TIP: If you want to report errors that you detect after a successful parse (such as type errors in a programming language), use the positioned combinator to add a position to a parse result. The result type must extend the Positional trait. For example,

```
def vardec1 = "var" ~ positioned(ident ^^ { Ident(_) }) ~ "=" ~ value
```

Exercises

1. Add / and % operations to the arithmetic expression evaluator.
2. Add a \wedge operator to the arithmetic expression evaluator. As in mathematics, \wedge should have a higher precedence than multiplication, and it should be right-associative. That is, $4\wedge 2\wedge 3$ should be $4\wedge(2\wedge 3)$, or 65536.
3. Write a parser that parses a list of integers (such as (1, 23, -79)) into a List[Int].
4. Write a parser that can parse date and time expressions in ISO 8601. Your parser should return a java.time.LocalDateTime object.
5. Write a parser that parses a subset of XML. Handle tags of the form `<ident> ... </ident>` or `<ident/>`. Tags can be nested. Handle attributes inside tags. Attribute values can be delimited by single or double quotes. You don't need to deal with character data (that is, text inside tags or CDATA sections). Your parser should return a Scala XML Elem value. The challenge is to reject mismatched tags. Hint: into, accept.
6. Assume that the parser in Section 20.5, “Generating Parse Trees,” on page 309 is completed with

```
class ExprParser extends RegexParsers {
    def expr: Parser[Expr] = (term ~ opt(("+" | "-") ~ expr)) ^^ {
        case a ~ None => a
        case a ~ Some(op ~ b) => Operator(op, a, b)
    }
    ...
}
```

Unfortunately, this parser computes an incorrect expression tree—operators with the same precedence are evaluated right-to-left. Modify the parser so that the expression tree is correct. For example, 3-4-5 should yield an `Operator("-", Operator("-", 3, 4), 5)`.

7. Suppose in Section 20.6, “Avoiding Left Recursion,” on page 310, we first parse an `expr` into a list of `~` with operations and values:

```
def expr: Parser[Int] = term ~ rep(("+" | "-") ~ term) ^^ {...}
```

To evaluate the result, we need to compute $((t_0 \pm t_1) \pm t_2) \pm \dots$. Implement this computation as a fold (see Chapter 13).

8. Add variables and assignment to the calculator program. Variables are created when they are first used. Uninitialized variables are zero. To print a value, assign it to the special variable `out`.

9. Extend the preceding exercise into a parser for a programming language that has variable assignments, Boolean expressions, and if/else and while statements.
10. Add function definitions to the programming language of the preceding exercise.

Implicits

Topics in This Chapter L3

- 21.1 Implicit Conversions — page 324
- 21.2 Using Implicits for Enriching Existing Classes — page 324
- 21.3 Importing Implicits — page 325
- 21.4 Rules for Implicit Conversions — page 326
- 21.5 Implicit Parameters — page 328
- 21.6 Implicit Conversions with Implicit Parameters — page 329
- 21.7 Context Bounds — page 329
- 21.8 Type Classes — page 331
- 21.9 Evidence — page 333
- 21.10 The `@implicitNotFound` Annotation — page 334
- 21.11 `CanBuildFrom` Demystified — page 334
- Exercises — page 336

Chapter

21

Implicit conversions and implicit parameters are Scala's power tools that do useful work behind the scenes. In this chapter, you will learn how implicit conversions can be used to enrich existing classes, and how implicit objects are summoned automatically to carry out conversions or other tasks. With implicits, you can provide elegant libraries that hide tedious details from library users.

The key points of this chapter are:

- Implicit conversions are used to convert between types.
- You must import implicit conversions so that they are in scope.
- An implicit parameter list requests objects of a given type. They can be obtained from implicit objects that are in scope, or from the companion object of the desired type.
- If an implicit parameter is a single-argument function, it is also used as an implicit conversion.
- A context bound of a type parameter requires the existence of an implicit object of the given type.
- If it is possible to locate an implicit object, this can serve as evidence that a type conversion is valid.

21.1 Implicit Conversions

An *implicit conversion function* is a function with a single parameter that is declared with the `implicit` keyword. As the name suggests, such a function is automatically applied to convert values from one type to another.

Consider the `Fraction` class from Section 11.2, “Infix Operators,” on page 143 with a method `*` for multiplying a fraction with another. We want to convert integers n to fractions $n / 1$.

```
implicit def int2Fraction(n: Int) = Fraction(n, 1)
```

Now we can evaluate

```
val result = 3 * Fraction(4, 5) // Calls int2Fraction(3)
```

The implicit conversion function turns the integer 3 into a `Fraction` object. That object is then multiplied by `Fraction(4, 5)`.

You can give any name to the conversion function. Since you don’t call it explicitly, you may be tempted to use something short such as `i2f`. But, as you will see in Section 21.3, “Importing Implicits,” on page 325, sometimes it is useful to import a conversion function. I suggest that you stick with the *source2target* convention.

Scala is not the first language that allows the programmer to provide automatic conversions. However, Scala gives programmers a great deal of control over when to apply these conversions. In the following sections, we will discuss exactly when the conversions happen, and how you can fine-tune the process.



NOTE: Even though Scala gives you tools to fine-tune implicit conversions, the language designers realize that implicit conversions are potentially problematic. To avoid a warning when using implicit functions, add the statement `import scala.language.implicitConversions` or the compiler option `-language:implicitConversions`.



NOTE: In C++, you specify implicit conversions as one-argument constructors or member functions with the name operator `Type()`. However, in C++, you cannot selectively allow or disallow these functions, and it is common to run into unwanted conversions.

21.2 Using Implicits for Enriching Existing Classes

Did you ever wish that a class had a method its creator failed to provide? For example, wouldn’t it be nice if the `java.io.File` class had a `read` method for reading a file:

```
val contents = new File("README").read
```

As a Java programmer, your only recourse is to petition Oracle Corporation to add that method. Good luck!

In Scala, you can define an enriched class that provides what you want:

```
class RichFile(val from: File) {  
    def read = Source.fromFile(from.getPath).mkString  
}
```

Then, provide an implicit conversion to that type:

```
implicit def file2RichFile(from: File) = new RichFile(from)
```

Now it is possible to call `read` on a `File` object. It is implicitly converted to a `RichFile`.

Instead of providing a conversion function, you can declare `RichFile` as an implicit class:

```
implicit class RichFile(val from: File) { ... }
```

An implicit class must have a primary constructor with exactly one argument. That constructor becomes the implicit conversion function.

It is a good idea to declare the enriched class as a value class:

```
implicit class RichFile(val from: File) extends AnyVal { ... }
```

In that case, no `RichFile` objects are created. A call `file.read` is directly compiled into a static method call `RichFile$.read$extension(file)`.



CAUTION: An implicit class cannot be a top-level class. You can place it inside the class that uses the type conversion, or inside another object or class that you import, as explained in the next section.

21.3 Importing Implicits

Scala will consider the following implicit conversion functions:

1. Implicit functions or classes in the companion object of the source or target type
2. Implicit functions or classes that are in scope

For example, consider the `int2Fraction` function. We can place it into the `Fraction` companion object, and it will be available for converting fractions.

Alternatively, let's suppose we put it inside a `FractionConversions` object, which we define in the `com.horstmann.impatient` package. If you want to use the conversion, import the `FractionConversions` object, like this:

```
import com.horstmann.impatient.FractionConversions._
```

For an implicit conversion to be in scope, it must be imported *without a prefix*. For example, if you import `com.horstmann.impatient.FractionConversions` or `com.horstmann.impatient.FractionConversions.int2Fraction`, then the `int2Fraction` method is in scope as `FractionConversions.int2Fraction` or `impatient.FractionConversions.int2Fraction`, to anyone who wants to call it explicitly. But if the function is not available as `int2Fraction`, without a prefix, the compiler won't use it implicitly.



TIP: In the REPL, type `:implicits` to see all implicits that have been imported from a source other than Predef, or `:implicits -v` to see all implicits.

You can localize the import to minimize unintended conversions. For example,

```
object Main extends App {
    import com.horstmann.impatient.FractionConversions.-
    val result = 3 * Fraction(4, 5) // Uses imported conversion
    println(result)
}
```

You can even select the specific conversions that you want. Suppose you have a second conversion

```
object FractionConversions {
    ...
    implicit def fraction2Double(f: Fraction) = f.num * 1.0 / f.den
}
```

If you prefer this conversion over `int2Fraction`, you can import it:

```
import com.horstmann.impatient.FractionConversions.fraction2Double
val result = 3 * Fraction(4, 5) // result is 2.4
```

You can also exclude a specific conversion if it causes you trouble:

```
import com.horstmann.impatient.FractionConversions.{fraction2Double => _, _}
// Imports everything but fraction2Double
```



TIP: If you want to find out why the compiler *doesn't* use an implicit conversion that you think it should use, try adding it explicitly, for example by calling `fraction2Double(3) * Fraction(4, 5)`. You may get an error message that shows the problem.

21.4 Rules for Implicit Conversions

In this section, you will see when implicit conversions are attempted. To illustrate the rules, we again use the `Fraction` class and assume that the implicit conversions `int2Fraction` and `fraction2Double` are both available.

Implicit conversions are considered in three distinct situations:

- If the type of an expression differs from the expected type:

```
3 * Fraction(4, 5) // Calls fraction2Double
```

The `Int` class doesn't have a method `*(Fraction)`, but it has a method `*(Double)`.

- If an object accesses a nonexistent member:

```
3.den // Calls int2Fraction
```

The `Int` class doesn't have a `den` member but the `Fraction` class does.

- If an object invokes a method whose parameters don't match the given arguments:

```
Fraction(4, 5) * 3
// Calls int2Fraction
```

The `*` method of `Fraction` doesn't accept an `Int` but it accepts a `Fraction`.

On the other hand, there are three situations when an implicit conversion is *not* attempted:

- No implicit conversion is used if the code compiles without it. For example, if `a * b` compiles, the compiler won't try `a * convert(b)` or `convert(a) * b`.
- The compiler will never attempt multiple conversions, such as `convert1(convert2(a)) * b`.
- Ambiguous conversions are an error. For example, if both `convert1(a) * b` and `convert2(a) * b` are valid, the compiler will report an error.



CAUTION: It is *not* an ambiguity that

```
3 * Fraction(4, 5)
```

could be either

```
3 * fraction2Double(Fraction(4, 5))
```

or

```
int2Fraction(3) * Fraction(4, 5)
```

The first conversion wins over the second, since it does not require modification of the object to which the `*` method is applied.



TIP: If you want to find out which implicit conversion the compiler uses, compile your program as

```
scalac -Xprint:typer MyProg.scala
```

You will see the source after implicit conversions have been added.

21.5 Implicit Parameters

A function or method can have a parameter list that is marked `implicit`. In that case, the compiler will look for default values to supply with the function call. Here is a simple example:

```
case class Delimiters(left: String, right: String)

def quote(what: String)(implicit delims: Delimiters) =
  delims.left + what + delims.right
```

You can call the `quote` method with an explicit `Delimiters` object, like this:

```
quote("Bonjour le monde")(Delimiters("«", "»")) // Returns «Bonjour le monde»
```

Note that there are two argument lists. This function is “curried”—see Chapter 12.

You can also omit the implicit parameter list:

```
quote("Bonjour le monde")
```

In that case, the compiler will look for an implicit value of type `Delimiters`. This must be a value that is declared as `implicit`. The compiler looks for such an object in two places:

- Among all `val` and `def` of the desired type that are in scope without a prefix.
- In the companion object of a type that is *associated* with the desired type. Associated types include the desired type itself and, if it is a parameterized type, its type parameters.

In our example, it is useful to make an object, such as

```
object FrenchPunctuation {
  implicit val quoteDelimiters = Delimiters("«", "»")
  ...
}
```

Then one imports all values from the object:

```
import FrenchPunctuation._
```

or just the specific value:

```
import FrenchPunctuation.quoteDelimiters
```

Now the French delimiters are supplied implicitly to the `quote` function.



NOTE: There can only be one implicit value for a given data type. Thus, it is not a good idea to use implicit parameters of common types. For example,

```
def quote(what: String)(implicit left: String, right: String) // No!
```

would not work—one could not supply two different strings.

21.6 Implicit Conversions with Implicit Parameters

An implicit function parameter is also usable as an implicit conversion. To understand the significance, consider first this simple generic function:

```
def smaller[T](a: T, b: T) = if (a < b) a else b // Not quite
```

That doesn't actually work. The compiler won't accept the function because it doesn't know that `a` and `b` belong to a type with a `<` operator.

We can supply a conversion function for that purpose:

```
def smaller[T](a: T, b: T)(implicit order: T => Ordered[T])
  = if (order(a) < b) a else b
```

Since the `Ordered[T]` trait has a `<` operator that consumes a `T`, this version is correct.

As it happens, this is such a common situation that the `Predef` object defines implicit values of type `T => Ordered[T]` for a large number of types, including all types that already implement `Ordered[T]` or `Comparable[T]`. Therefore, you can call

```
smaller(40, 2)
```

and

```
smaller("Hello", "World")
```

If you want to call

```
smaller(Fraction(1, 7), Fraction(2, 9))
```

then you need to define a function `Fraction => Ordered[Fraction]` and either supply it in the call or make it available as an `implicit val`. I leave this as an exercise because it moves us too far from the point that I want to make in this section.

Here, finally, is the point. Look again at

```
def smaller[T](a: T, b: T)(implicit order: T => Ordered[T])
```

Note that `order` is a function that is tagged `implicit` and is in scope. Therefore, *it is an implicit conversion*, in addition to being an implicit parameter. So, we can omit the call to `order` in the body of the function:

```
def smaller[T](a: T, b: T)(implicit order: T => Ordered[T])
  = if (a < b) a else b // Calls order(a) < b if a doesn't have a < operator
```

21.7 Context Bounds

A type parameter can have a *context bound* of the form `T : M`, where `M` is another generic type. It requires that there is an implicit value of type `M[T]` in scope.

For example,

```
class Pair[T : Ordering]
```

requires that there is an implicit value of type `Ordering[T]`. That implicit value can then be used in the methods of the class. Consider this example:

```
class Pair[T : Ordering](val first: T, val second: T) {
    def smaller(implicit ord: Ordering[T]) =
        if (ord.compare(first, second) < 0) first else second
}
```

If we form a new `Pair(40, 2)`, then the compiler infers that we want a `Pair[Int]`. Since there is an implicit value of type `Ordering[Int]` in the `Ordering` companion object, `Int` fulfills the context bound. That ordering becomes a field of the class, and it is passed to the methods that need it.

If you prefer, you can retrieve the ordering with the `implicitly` method in the `Predef` class:

```
class Pair[T : Ordering](val first: T, val second: T) {
    def smaller =
        if (implicitly[Ordering[T]].compare(first, second) < 0) first else second
}
```

The `implicitly` function is defined as follows in `Predef.scala`:

```
def implicitly[T](implicit e: T) = e
// For summoning implicit values from the nether world
```



NOTE: The comment is apt—the implicit objects live in the “nether world” and are invisibly added to methods.

Alternatively, you can take advantage of the fact that the `Ordered` trait defines an implicit conversion from `Ordering` to `Ordered`. If you import that conversion, you can use relational operators:

```
class Pair[T : Ordering](val first: T, val second: T) {
    def smaller = {
        import Ordered._;
        if (first < second) first else second
    }
}
```

These are just minor variations; the important point is that you can instantiate `Pair[T]` whenever there is an implicit value of type `Ordering[T]`. For example, if you want a `Pair[Point]`, arrange for an implicit `Ordering[Point]` value:

```
implicit object PointOrdering extends Ordering[Point] {
    def compare(a: Point, b: Point) = ...
}
```

21.8 Type Classes

Have another look at the `Ordering` trait in the preceding section. We had an algorithm that required the parameters to have an ordering. Normally, in object-oriented programming, we would require the parameter types to extend a trait. But that's not what happened here. In order to make a class usable with the algorithm, it is not necessary to modify the class at all. Instead, one provides an implicit conversion. This is much more flexible than the object-oriented approach.

A trait such as `Ordering` is called a *type class*. A type class defines some behavior, and a type can join the class by providing that behavior. (The term comes from Haskell, and “class” is not used in the same way as in object-oriented programming. Think of “class” as in a “class action”—types banding together for a common purpose.)

To see how a type joins a type class, let us look at a simple example. We want to compute averages, $(x_1 + \dots + x_n) / n$. To do that, we need to be able to add two values and divide a value by an integer. There is a type class `Numeric` in the Scala library, which requires that values can be added, multiplied, and compared. But it doesn't have any requirement that values can be divided by an integer. Therefore, let's define our own:

```
trait NumberLike[T] {
    def plus(x: T, y: T): T
    def divideBy(x: T, n: Int): T
}
```

Next, to make sure that the type class is useful out of the gate, let's add some common types as members. That's easy to do by providing implicit objects in the companion object:

```
object NumberLike {
    implicit object NumberLikeDouble extends NumberLike[Double] {
        def plus(x: Double, y: Double) = x + y
        def divideBy(x: Double, n: Int) = x / n
    }

    implicit object NumberLikeBigDecimal extends NumberLike[BigDecimal] {
        def plus(x: BigDecimal, y: BigDecimal) = x + y
        def divideBy(x: BigDecimal, n: Int) = x / n
    }
}
```

Now we are ready to put the type class to use. In the average method, we need an instance of the type class so that we can call `plus` and `divideBy`. (Note that these are methods of the type class, not the member types.)

Here, we'll just compute the average of two values. The general case is left as an exercise. There are two ways in which we can supply the type class instance: as an implicit parameter, or with a context bound. Here is the first approach:

```
def average[T](x: T, y: T)(implicit ev: NumberLike[T]) =
    ev.divideBy(ev.plus(x, y), 2)
```

The parameter name `ev` is shorthand for "evidence"—see the next section.

When using a context bound, we retrieve the implicit object from the "nether world."

```
def average[T : NumberLike](x: T, y: T) = {
    val ev = implicitly[NumberLike[T]]
    ev.divideBy(ev.plus(x, y), 2)
}
```

That's all there is to it. Finally, let's see what a type needs to do to join the `NumberLike` type class. It must provide an implicit object, just like the `NumberLikeDouble` and `NumberLikeBigDecimal` objects that we provided out of the gate. Here is how `Point` can join:

```
class Point(val x: Double, val y: Double) {
    ...
}

object Point {
    def apply(x: Double, y: Double) = new Point(x, y)
    implicit object NumberLikePoint extends NumberLike[Point] {
        def plus(p: Point, q: Point) = Point(p.x + q.x, p.y + q.y)
        def divideBy(p: Point, n: Int) = Point(p.x * 1.0 / n, p.y * 1.0 / n)
    }
}
```

Here we added the implicit object to the companion object of `Point`. If you can't modify the `Point` class, you can put the implicit object elsewhere and import it as needed.

The standard Scala library provides useful type classes, such as `Equiv`, `Ordering`, `Numeric`, `Fractional`, `Hashing`, `IsTraversableOnce`, `IsTraversableLike`. As you have seen, it is easy to provide your own.

The important point about type classes is that they provide an "ad-hoc" way of providing polymorphism that is less rigid than inheritance.

21.9 Evidence

In Chapter 18, you saw the type constraints

```
T =:= U
T <:< U
T => U
```

The constraints test whether `T` equals `U`, is a subtype of `U`, or is convertible to `U`. To use such a type constraint, you supply an implicit parameter, such as

```
def firstLast[A, C](it: C)(implicit ev: C <:< Iterable[A]) =
  (it.head, it.last)
```

The `=:=` and `<:<` are classes with implicit values, defined in the `Predef` object. For example, `<:<` is essentially

```
abstract class <:<[-From, +To] extends Function1[From, To]

object <:< {
  implicit def conforms[A] = new (A <:< A) { def apply(x: A) = x }
}
```

Suppose the compiler processes a constraint `implicit ev: String <:< AnyRef`. It looks in the companion object for an implicit object of type `String <:< AnyRef`. Note that `<:<` is contravariant in `From` and covariant in `To`. Therefore the object

```
<:<.conforms[String]
```

is usable as a `String <:< AnyRef` instance. (The `<:<.conforms[AnyRef]` object is also usable, but it is less specific and therefore not considered.)

We call `ev` an “evidence object”—its existence is evidence of the fact that, in this case, `String` is a subtype of `AnyRef`.

Here, the evidence object is the identity function. To see why the identity function is required, have a closer look at

```
def firstLast[A, C](it: C)(implicit ev: C <:< Iterable[A]) =
  (it.head, it.last)
```

The compiler doesn’t actually know that `C` is an `Iterable[A]`—recall that `<:<` is not a feature of the language, but just a class. So, the calls `it.head` and `it.last` are not valid. But `ev` is a function with one parameter, and therefore an implicit conversion from `C` to `Iterable[A]`. The compiler applies it, computing `ev(it).head` and `ev(it).last`.



TIP: To test whether a generic implicit object exists, you can call the `implicitly` function in the REPL. For example, type `implicitly[String <:< AnyRef]` in the REPL, and you get a result (which happens to be a function). But `implicitly[AnyRef <:< String]` fails with an error message.

21.10 The @implicitNotFound Annotation

The `@implicitNotFound` annotation raises an error message when the compiler cannot construct an implicit parameter of the annotated type. The intent is to give a useful error message to the programmer. For example, the `<:<` class is annotated as

```
@implicitNotFound(msg = "Cannot prove that ${From} <:< ${To}.")
abstract class <:<[-From, +To] extends Function1[From, To]
```

For example, if you call

```
firstLast[String, List[Int]](List(1, 2, 3))
```

then the error message is

```
Cannot prove that List[Int] <:< Iterable[String]
```

That is more likely to give the programmer a hint than the default

```
Could not find implicit value for parameter ev: <:<[List[Int],Iterable[String]]
```

Note that `${From}` and `${To}` in the error message are replaced with the type parameters `From` and `To` of the annotated class.

21.11 CanBuildFrom Demystified

In Chapter 1, I wrote that you should ignore the implicit `CanBuildFrom` parameter. Now you are finally ready to understand how it works.

Consider the `map` method. Simplifying slightly, `map` is a method of `Iterable[A, Repr]` with the following implementation:

```
def map[B, That](f : (A) => B)(implicit bf: CanBuildFrom[Repr, B, That]): That = {
    val builder = bf()
    val iter = iterator()
    while (iter.hasNext) builder += f(iter.next())
    builder.result
}
```

Here, `Repr` is the “representation type.” `That` parameter will enable us to select appropriate builder factories for unusual collections such as `Range` or `String`.



NOTE: In the Scala library, `map` is actually defined in the `TraversableLike[A, Repr]` trait. That way, the more commonly used `Iterable` trait doesn't need to carry with it the `Repr` type parameter.

The `CanBuildFrom[From, E, To]` trait provides evidence that it is possible to create a collection of type `To`, holding values of type `E`, that is compatible with type `From`. Before discussing how these evidence objects are generated, let's see what they do.

The `CanBuildFrom` trait has an `apply` method that yields an object of type `Builder[-E, +To]`. A `Builder` has methods `+=` for adding elements into an internal buffer, and `result` for producing the desired collection.

```
trait Builder[-E, +To] {
  def +=(e: E): Unit
  def result(): To
}

trait CanBuildFrom[-From, -E, +To] {
  def apply(): Builder[E, To]
}
```

Therefore, the `map` method simply constructs a builder for the target type, fills the builder with the values of the function `f`, and yields the resulting collection.

Each collection provides an implicit `CanBuildFrom` object in its companion object. Consider a simplified version of the standard `ArrayBuffer` class:

```
class Buffer[E : ClassTag] extends Iterable[E, Buffer[E]]
  with Builder[E, Buffer[E]] {
  private var elems = new Array[E](10)

  ...
  def iterator() = ...
  private var i = 0
  def hasNext = i < length
  def next() = { i += 1; elems(i - 1) }
}
def +=(e: E) { ... }
def result() = this
}

object Buffer {
  implicit def canBuildFrom[E : ClassTag] =
    new CanBuildFrom[Buffer[_], E, Buffer[E]] {
      def apply() = new Buffer[E]
    }
}
```

Consider a call `buffer.map(f)`, where `f` is a function of type `A => B`. The implicit `bf` parameter is obtained by calling the `canBuildFrom[B]` method in the `Buffer` companion object. Its `apply` method returns the builder, in this case a `Buffer[B]`.

As it happens, the `Buffer` class already has a `+=` method, and its `result` method is defined to return itself. Therefore, a `Buffer` is its own builder.

However, a builder for the `Range` class doesn't return a `Range`, and clearly it can't. For example, `(1 to 10).map(x => x * x)` has a value that isn't a `Range`. In the actual

Scala library, Range extends IndexedSeq[Int], and the IndexedSeq companion object defines a builder that constructs a Vector.

Here is a simplified Range class that provides a Buffer as builder:

```
class Range(val low: Int, val high: Int) extends Iterable[Int, Range] {
    def iterator() = ...
}

object Range {
    implicit def canBuildFrom[E : ClassTag] = new CanBuildFrom[Range, E, Buffer[E]] {
        def apply() = new Buffer[E]
    }
}
```

Now consider a call `Range(1, 10).map(f)`. That method needs an `implicit bf: CanBuildFrom[Repr, B, That]`. Since Repr is Range, the associated types are `CanBuildFrom`, `Range`, `B`, and the unknown `That`. The Range object yields a match by calling its `canBuildFrom[B]` method, which returns a `CanBuildFrom[Range, B, Buffer[B]]`. That object is `bf`, and its `apply` method yields a `Buffer[B]` for building the result.

As you just saw, the `implicit CanBuildFrom[Repr, B, That]` parameter locates a factory object that can produce a builder for the target collection. The builder factory is defined as `implicit` in the companion object of `Repr`.

Exercises

- How does `->` work? That is, how can "Hello" \rightarrow 42 and 42 \rightarrow "Hello" be pairs ("Hello", 42) and (42, "Hello")? Hint: `Predef.ArrowAssoc`.
- Define an operator `+%` that adds a given percentage to a value. For example, $120 +\% 10$ should be 132. Use an `implicit` class.
- Define a `!` operator that computes the factorial of an integer. For example, $5.!$ is 120. Use an `implicit` class.
- Some people are fond of “fluent APIs” that read vaguely like English sentences. Create such an API for reading integers, floating-point numbers, and strings from the console. For example: Read in `aString` askingFor “Your name” and `aInt` askingFor “Your age” and `aDouble` askingFor “Your weight”.
- Provide the machinery that is needed to compute

```
smaller(Fraction(1, 7), Fraction(2, 9))
```

with the `Fraction` class of Chapter 11. Supply an `implicit` class `RichFraction` that extends `Ordered[Fraction]`.

- Compare objects of the class `java.awt.Point` by lexicographic comparison.

7. Continue the previous exercise, comparing two points according to their distance to the origin. How can you switch between the two orderings?
8. Use the `implicitly` command in the REPL to summon the implicit objects described in Section 21.5, “Implicit Parameters,” on page 328 and Section 21.6, “Implicit Conversions with Implicit Parameters,” on page 329. What objects do you get?
9. Explain why `Ordering` is a type class and why `Ordered` is not.
10. Generalize the `average` method in Section 21.8, “Type Classes,” on page 331 to a `Seq[T]`.
11. Make `String` a member of the `NumberLike` type class in Section 21.8, “Type Classes,” on page 331. The `divBy` method should retain every `n`th letter, so that `average("Hello", "World")` becomes `"Hlool"`.
12. Look up the `=:=` object in `Predef.scala`. Explain how it works.
13. The result of `"abc".map(_.toUpperCase)` is a `String`, but the result of `"abc".map(_.toInt)` is a `Vector`. Find out why.

Index

Symbols and Numbers

- (minus sign)
 - in identifiers, 142
 - operator:
 - arithmetic, 6
 - for collections, 178–179
 - for maps, 49
 - for type parameters, 271
 - left-associative, 145
 - precedence of, 145
 - unary (negation), 14, 143
- operator, 177–179
 - not used for arithmetic decrements, 7
- = operator
 - arithmetic, 7
 - for collections, 178–179
 - for maps, 49
- = operator, 178–179
- _ (underscore)
 - as wildcard:
 - for XML elements, 236
 - in case clauses, 30, 198–199, 237
 - in imports, 8, 76, 85–86
 - in tuples, 52
 - for function calls, 158, 288
 - for function parameters, 161
 - in identifiers, 142, 317
- _* syntax
 - for arrays, 201
 - for nested structures, 207
 - in function parameters, 27
 - in pattern matching, 237
- _=, in setter methods, 57
- _1, _2, _3 methods (tuples), 51
- ;(semicolon)
 - after statements, 5, 18–20
 - inside loops, 24–25

- : (colon)
 - followed by annotations, 217
 - in case clauses, 200–201
 - in identifiers, 142
 - in implicits, 329–330
 - in operator names, 285
 - and precedence, 145
 - right-associative, 146, 185
 - in type parameters, 268–269
- :: operator
 - for lists, 175–176, 178–179
 - in case clauses, 201, 206
 - right-associative, 146, 176
- :::, ::+ operators, 178–179
- :\\, \:/ operators, 185
- :+= operator, 179
- ! (exclamation mark)
 - in identifiers, 142
 - operator:
 - in shell scripts, 114–115
 - precedence of, 145
 - unary, 143
- !! operator, in shell scripts, 114
- != operator, 144
- ? (question mark)
 - in identifiers, 142
 - in parsers, 308
- ? : operator, 18
- ??? method, 101
- / (slash)
 - in identifiers, 142
 - in XPath, 235
 - operator, 6
 - precedence of, 145
- //
for comments, 317
in XPath, 235
- /* ... */ comments, parsing, 317

% method (`BigInt`), 7, 203
`...` (backquotes), for identifiers, 142, 200
^ (caret)
 in identifiers, 142
 in Pascal, 153
 operator, 6
 precedence of, 145
^? operator, 313
^^ operator, 307–309, 312
^^^ operator, 312
~ (tilde)
 in identifiers, 142
operator:
 in case clauses, 206
 in parsers, 305–311, 313–314
 unary, 143
~! operator, 313–314
~> operator, 308–309, 312
' (single quote)
 in symbols, 226
 parsing, 317
" (double quote), parsing, 317
"""", in regular expressions, 116
() (parentheses)
 as value of `Unit`, 18–21
 discarding, in parsers, 309
 for annotations, 216
 for apply method, 9–10
 for array elements, 36
 for curried parameters, 165
 for functions, 158–161, 166
 for maps, 48
 for primary constructors, 63
 for tuples, 51, 287
 in case clauses, 201, 205
 in method declarations, 8, 56, 59
 in regular expressions, 117
 to access XML attributes, 232
[] (square brackets)
 for arrays, 36
 for methods in traits, 126
 for type parameters, 14, 266, 287
{} (braces)
 for block expressions, 20–21
 for existential types, 286
 for function arguments, 159
 for structural types, 283–285
 in formatted strings, 21
 in imports, 86
 in package clauses, 83
 in pattern matching, 211–212, 237
 in REPL, 19
 in XML literals, 234
 Kernighan & Ritchie style for, 20
@ (at)
 for XML attributes, 236
 in case clauses, 207
 in identifiers, 142
\$ (dollar sign), in formatted strings, 21
* (asterisk)
 as wildcard in Java, 8, 85
 in identifiers, 142
 in parsers, 308
 operator, 6, 327
 no infix notation for, 286
 precedence of, 145
**
 in Fortran, 153
 in identifiers, 142
\ (backslash)
 in identifiers, 142
 operator, 235–236
\\ operator, 235–236
& (ampersand)
 in identifiers, 142
operator:
 arithmetic, 6
 for sets, 177–179
 precedence of, 145
&...; (XML), 231
&~ operator, 177–178
&#...; (XML), 232
(number sign)
 for type projections, 67, 281–283, 286
 in identifiers, 142
#: operator, 189
&&, #<, #>, #>>, #|| operators (shell scripts), 115
#| operator (shell scripts), 114
% (percent sign)
 for XML attributes, 238
 in identifiers, 142
operator, 6
 precedence of, 145
+ (plus sign)
 in identifiers, 142
operator:
 arithmetic, 6
 for collections, 178–179
 for maps, 49
 for type parameters, 271
 precedence of, 145
 unary, 143

- +: operator
 - for collections, 178–179
 - in case clauses, 207
 - right-associative, 146, 179
- ++ operator, 177–179
 - not used for arithmetic increments, 7
- ++:, ++=:, +=: operators, 178–179
- += operator
 - for array buffers, 36
 - for collections, 178–179
- += operator
 - arithmetic, 7
 - assignment, 144
 - for array buffers, 36, 42
 - for collections, 178–179, 335
 - for maps, 49
 - of Buffer, 335
- < (left angle bracket)
 - in identifiers, 142
 - in XML literals, 230
- operator:
 - and implicits, 329
 - precedence of, 145
- <- operator, 23–24, 204
- <: operator, 267–271, 286, 293
- <:< operator, 269, 333
- <!---->, <?...?> comments (XML), 231
- <?xml...?> (XML), 241
- <~ operator, 308–309, 312
- % operator, 268
- << operator, 6
- <= operator, 144
- = (equal sign)
 - in identifiers, 142
 - operator, 49, 144
 - precedence of, 145
- =:= operator, 269, 333
- == operator, 144
 - for collections, 173
 - for reference types, 103
- ==, /= operators, 144
- > operator, 269
 - for functions, 166, 287–288
 - for self types, 134–135, 288–289, 295
 - in case clauses, 198
- > (right angle bracket)
 - in identifiers, 142
 - operator, 145
- >- operator, 48
 - precedence of, 145
- >: operator, 267, 269
- >= operator, 144
- >> operator
 - arithmetic, 6
 - in parsers, 312
- | (vertical bar)
 - in case clauses, 198
 - in identifiers, 142
- operator:
 - arithmetic, 6
 - for sets, 177–178
 - in parsers, 304–319
 - precedence of, 145
- √ (square root), 142
- 80 bit extended precision, 219

A

- abstract keyword, 97, 123, 127
- abstract types, 291–293, 315
 - bounds for, 293
 - made concrete in subclass, 283, 291
 - naming, 296
 - vs. type parameters, 292
- accept combinator, 313
- ActionEvent class (Java), 293–294
- actionPerformed method (listeners), 293
- addString method
 - of Iterable, 181
 - of Iterator, 189
- aggregate method
 - of Iterable, 180
 - of Iterator, 189
 - of parallel collections, 194
- aliases
 - for this, 68, 289
 - for types, 174, 283
- andThen method (Future), 256
- Annotation trait, 218
- annotations, 215–226, 287
 - arguments of, 217–218
 - for compiler optimizations, 222–226
 - for errors and warnings, 226
 - implementing, 218–219
 - in Java, 216–221
 - meta-annotations for, 219
 - order of, 216
- anonymous functions, 25, 159–160, 167
- Any class, 100, 103
 - asInstanceOf, isInstanceOf methods, 93, 100
- AnyRef class, 100, 112
 - == method, 103
 - eq, equals methods, 102
 - notify, notifyAll, synchronized, wait methods, 100

- AnyVal class, 100, 104
- Apache Commons Resolver project, 240
- App trait, 75
- append, appendAll methods, 41
- Application trait, 75
- application objects, 74
- apply method, 9–10, 73–74, 146–148
 - of Array, 10, 74
 - of BigInt, 10
 - of Buffer, CanBuildFrom, 335
 - of case classes, 205
 - of collections, 173
 - of Future, 253, 260
 - of PartialFunction, 211
 - of Process, 115
 - of StringOps, 9
- applyDynamic method (Dynamic), 150–152
- applyDynamicNamed method (Dynamic), 151–152
- args property, 75
- array buffers, 36–37
 - adding/removing elements of, 36–37
 - appending collections to, 36
 - converting to arrays, 37
 - displaying contents of, 40
 - empty, 36
 - largest/smallest elements in, 40
 - sorting, 40
 - transforming, 38–40
 - traversing, 37–38
- Array class, 35–36, 269
 - apply method, 74
 - corresponds method, 165, 271
 - mkString, quickSort methods, 40
 - ofDim method, 42
 - toBuffer method, 37
- Array companion object, 10, 202
- ArrayBuffer class, 36–37, 172, 335
 - max, min, mkString methods, 40
 - mutable, 175
 - serializing, 113
 - sorted, sortWith methods, 40
 - subclasses of, 42
 - toArray method, 37
- ArrayList class (Java), 36, 43, 173
- arrays, 35–44
 - converting to array buffers, 37
 - displaying contents of, 40
 - fixed-length, 35–36
 - function call syntax for, 146
 - generic, 269
 - interoperating with Java, 43–44
 - invariant, 272
- largest/smallest elements in, 40
- multidimensional, 42–43, 74
- pattern matching for, 201
- ragged, 43
- sorting, 40
- transforming, 38–40
- traversing, 37–38
- vs. lists, 172
- Arrays class (Java), 43
- ArrayStoreException, 273
- asAttrMap method, 233
- ASCII characters, 142
- asInstanceOf method (Any), 93, 100
- asJavaXxx, asScalaXxx functions (JavaConversions), 192
- assert method (Predef), 225
- AssertionError, 225
- assignments, 20–21, 144
 - no chaining of, 21
 - precedence of, 145
 - right-associative, 146, 179
 - value of, 20–21
- Async library, 254
- Atom class, 233–235
- Attribute trait, 238
- attributes (XML), 232–233
 - atoms in, 234
 - entity references in, 232, 234
 - expressions in, 234–235
 - in pattern matching, 238
 - iterating over, 233
 - modifying, 238
 - namespace of, 243
 - no wildcard notation for, 236
- attributes property, 232
- automatic conversions. *See* implicit conversions
- Await object, 250

B

- backtracking, 313–314
- balanced trees, 50
- bash shell, 114
- bean properties, 61–62
 - @BeanDescription, @BeanDisplayName, @BeanInfo, @BeanInfoSkip annotations, 221
 - @beanGetter, @beanSetter annotations, 219
 - @BeanProperty annotation, 61–62, 216, 221
 - generated methods for, 65
- BigDecimal class, 6–7
- BigInt class, 6–9
- % method, 7, 203

- BigInt class (*continued*)**
 - pow method, 153
 - unary_- method, 14
 - BigInt companion object**
 - apply method, 10
 - probablePrime method, 9
 - binary files**, reading, 112
 - binary functions**, 162, 184
 - binary operators**, 143–146
 - BitSet class**, 177
 - blocking keyword**, 260
 - blocking calls**, 250
 - blocks**, 20–21
 - BNF (Backus-Naur Form)**, 304
 - Boolean type**, 5
 - reading, 22
 - @BooleanBeanProperty annotation**, 221
 - Breaks object**, break method, 23
 - Buffer class**, 298, 335
 - bufferAsJavaList function (`JavaConversions`), 192
 - buffered method (`Source`), 110, 188
 - BufferedInputStream class (Java)**, 138
 - BufferedIterator trait**, 188
 - Builder trait**, 335
 - Byte type**, 5
 - arrays of, 112
 - reading, 22
- C**
- C++ programming language**
 - ? operator in, 18
 - arrays in, 36
 - assignments in, 21
 - construction order in, 100
 - expressions in, 17
 - functions in, 25
 - implicit conversions in, 324
 - linked lists in, 176
 - loops in, 22, 38
 - methods in, 72, 94, 219
 - multiple inheritance in, 121–122
 - namespaces in, 80
 - operators in, 145
 - protected fields in, 94
 - reading files in, 110
 - singleton objects in, 72
 - statements in, 17–20
 - switch in, 198–199, 223
 - virtual base classes in, 122
 - void in, 18, 20, 102
 - cached thread pool, 260
 - cake pattern**, 290
 - callbacks**, 251–252
 - call-by-name parameters**, 167
 - case clause**, 198–212
 - _ in, 199, 237
 - : in, 200–201
 - :: in, 201, 206
 - '...' in, 200
 - ~ in, 206
 - () in, 201
 - {_*} in, 237
 - @, +: in, 207
 - =>, | in, 198
 - catch-all pattern for, 198–199
 - enclosed in braces, 211–212, 252
 - followed by variable, 199–200
 - infix notation in, 206–207
 - variables in, 199–200
 - XML literals in, 237
 - case classes**, 205–211
 - applicability of, 208–209
 - declaring, 205
 - default methods of, 148, 205, 208–209
 - extending other case classes, 209
 - in parsers, 306, 309
 - modifying properties in, 205
 - sealed, 209
 - with variable fields, 208
 - case objects**, 205
 - casts**, 93–94
 - CatalogResolver class (Java)**, 240
 - catch statement**, 30–31
 - as a partial function, 212
 - CDATA markup**, 235, 240
 - ceil method (`scala.math`)**, 158
 - chaining**
 - assignments, 21
 - auxiliary constructors, 65
 - method calls, 42
 - packages, 82–83
 - chain11 method (Parsers)**, 312
 - ChangeEvent class (Java)**, 293
 - Char type**, 5, 22, 315
 - character references**, 232
 - character sets**, 111
 - characters**
 - common, in two strings, 6
 - in identifiers, 142, 317
 - reading, 22, 110
 - sequences of, 14
 - uppercase, 14

circular dependencies, 29, 135
class keyword, 55, 287
ClassCastException, 225
 classes, 10, 55–68, 287
 abstract, 97, 125
 abstract types in, 291
 and primitive types, 5
 annotated, 216
 combined with primary constructor, 66
 companion objects to, 9, 67, 72–73, 146, 281, 325
 concrete, 130
 definitions of, 56
 executing statements in, 64
 using traits in, 125
 enriching, 324–325
 equality checks in, 102
 extending, 73, 91–92
 Java classes, 95
 multiple traits, 124
 only one superclass, 121, 129
 final, 92
 generic, 266, 270
 granting access to, 61–62
 immutable, 7, 208
 implementing, 265
 importing members of, 76, 85
 inheritance hierarchy of, 100–102
 inlined, 103–105
 interoperating with Java, 57
 linearization of, 131
 mutable, 208
 names of, 142
 nested, 66–68, 281
 properties of, 57, 59
 recompiling, 129
 serializable, 113, 220
 type aliases in, 283
 visibility of, 56
 vs. traits, 132
classOf method (`scala.Predef`), 93
ClassTag trait, 269, 299
Cloneable interface (Java), 124, 220
@cloneable annotation, 220
close method (`Source`), 110
 closures, 162–163
collect method
 of `Future`, 256
 of `GenTraversable`, 211
 of `Iterable`, 180
 of `Iterator`, 189
 with partial functions, 183
collectionAsScalaIterable function (`JavaConversions`), 192
collections, 171–194
 adding/removing elements of, 178–179
 appending to array buffers, 36
 applying functions to all elements of, 161, 180–183
 combining, 187–188
 companion objects of, 173, 335
 constructing in a loop, 24, 38
 converting to specific type, 181
 filtering, 180
 folding, 180, 184–186
 hierarchy of, 41, 172–173
 immutable, 173–174
 instances of, 173
 interoperating with Java, 191–192
 iterators for, 188–189
 methods for, 180–182
 mutable, 173–174, 179, 191
 ordered/unordered, 172, 179
 parallel implementations of, 193–194
 reducing, 180, 184
 scanning, 180, 186
 serializing, 113
 sorted, 172
 traits for, 172–173
 traversing, 23, 38, 172, 222–223
 unevaluated, 190
 unordered, 179
com.**s**un.org.apache.xml.internal.resolver.tools package, 240
combinations method (`Seq`), 182
combinators, 311–313
command-line arguments, 75
comma-separated lists, 311
comments
 in lexical analysis, 304
 in XML, 231
 parsing, 240, 316–317
companion objects, 9, 67, 72–73, 281, 328
 apply method in, 146
 implicits in, 325
 of collections, 173, 335
Comparable interface (Java), 42, 267–268, 329
Comparator interface (Java), 226
compareTo method, 267
compiler
 -future flag, 268
 implicits in, 327, 334
 internal types in, 288
 -language:existentials flag, 287

- compiler (*continued*)
 - language:higherKinds flag, 299
 - language:implicitConversions flag, 324
 - language:postfixOps flag, 144
 - nested type expressions in, 283
 - optimizations in, 222–226
 - Scala annotations in, 216
 - Xcheckinit flag, 99
 - Xprint flag, 327
- compiler plugin, 216
- compile-time errors, 22
- CompletableFuture class (Java), 260
- CompletionStage interface (Java), 249
- Component class (Java), 137
- compound types, 284–287
- comprehensions, 24
- concurrency, 193
- console
 - input from, 22, 111
 - printing to, 21
- constants. *See* values
- ConstructingParser class, 240–241
- constructors
 - auxiliary, 62–63, 94
 - chaining, 65
 - eliminating, 64
 - executing, 72
 - order of, 98–100, 130–132
 - parameterless, 64, 132
 - parameters of, 62–66
 - annotated, 218–219
 - implicit, 269
 - primary, 62–66, 94
 - annotated, 216
 - default parameters in, 64
 - private, 66
 - superclass, 94–95
 - vals in, 99
- Container trait, 298
- Container class (Java), 137
- contains method, 48
 - of BitSet, 177
 - of Iterator, 189
 - of Seq, 181
- containsSlice method, 42
 - of Iterator, 189
 - of Seq, 181
 - of StringOps, 14
- context bounds, 268–269, 329–330
- control abstractions, 166–167
- copy method, 238
 - of case classes, 205, 208–209
 - copyToArray method, 42
 - of Iterable, 181
 - of Iterator, 189
 - copyToBuffer method
 - of Iterable, 181
 - of Iterator, 189
 - corresponds method (Array), 165, 271
 - count method, 41
 - of Iterable, 180
 - of Iterator, 189
 - of StringOps, 14
 - Curry, Haskell Brooks, 164
 - currying, 164–165
- D**
- debugging
 - implicit conversions, 326
 - reading from strings for, 111
 - reporting types for, 41
- decrements, 7
- def keyword, 25, 159, 165
 - abstract, 95
 - in parsers, 314
 - overriding, 95–96
 - parameterless, 95
 - return value of, 314
- default statement, 198
- definitions, 24–25
- DelayedInit trait, 75
- Delimiters type, 328
- dependency injections, 289–291
- @deprecated annotation, 219, 226
- @deprecatedInheritance, @deprecatedOverriding annotations, 226
- @deprecatedName annotation, 218, 226
- destructuring
 - of lists, 207
 - of tuples, 202
- diamond inheritance problem, 122–123
- dictionaryAsScalaMap function (JavaConversions), 192
- diff method
 - of Iterator, 189
 - of Seq, 182
 - of sets, 177
- directories
 - and packages, 80
 - naming, 15
 - traversing, 112–113
- division, quotient and remainder of, 7, 203
- do loop, 22
- docElem method (ConstructingParser), 240

DocType class, 241
 domain-specific languages, 141, 303
 Double type, 5
 reading, 22
 DoubleLinkedList class (deprecated), 176
 drop, dropWhile methods
 of Iterable, 180
 of Iterator, 189
 dropRight method (Iterable), 180
 DTDs (Document Type Definitions), 239–241
 duck typing, 284
 Duration object, 250
 Dynamic trait, 150
 dynamic invocation, 150–153
 dynamically typed languages, 284

E

early definitions, 99, 132–133
 EBNF (Extended Backus-Naur Form), 305–306
 Eclipse-based Scala IDE, 2
 Eiffel programming language, 58
 Either type, 300
 Elem type, 230, 315
 prefix, scope values, 238, 243
 elements (XML), 230
 attributes of. *See* attributes (XML)
 child, 237–238
 empty, 242
 matching, 236
 modifying, 238
 endsWith method
 of Iterator, 189
 of Seq, 181
 entity references, 231
 in attributes, 232, 234
 resolving, 241
 EntityRef class, 232
 Enumeration class, 75–77
 enumerationAsScalaIterator function (JavaConversions),
 192
 enumerations, 75–77
 simulating, 209
 values of, 76–77
 eq method (AnyRef), 102
 equals method
 of AnyRef, 102–103
 of case classes, 205, 208–209
 of value classes, 104
 overriding, 103
 parameter type of, 103
 Equiv type, 332

err combinator, 313
 Error class, 315
 error messages
 explicit, 319
 for implicit conversions, 326, 334
 generated with annotations, 226
 suppressed, 226
 type projections in, 283
 with override modifier, 92
 escape hatch, 142
 evidence objects, 333
 Exception trait, 288
 exceptions, 29–31
 catching, 30
 checked, in Java, 29, 220
 execution contexts, 260
 ExecutionContext trait, 248
 Executor interface (Java), 248
 Executors, 260
 existential types, 287
 exists method
 of Iterable, 180
 of Iterator, 189
 expressions
 annotated, 217
 conditional, 18–19
 traversing values of, 23
 type of, 18
 vs. statements, 17
 extends keyword, 91, 99, 123–124
 extractors, 118, 147–149, 202–203

F

f prefix, in formatted strings, 21–22, 112
 failed method
 of Future object, 258
 of Future trait, 255–256
 failure method (Promise), 259
 Failure class, 251, 315
 failure combinator, 313
 fallbackTo method (Future), 255
 fall-through problem, 198
 family polymorphism, 293–296
 @field annotation, 219
 fields
 abstract, 97–98, 130, 132
 accessing uninitialized, 99
 annotated, 216
 comparing, 208
 concrete, 98, 128–129
 copying, 208

- fields (*continued*)
 - for primary constructor parameters, 62, 65
 - getter/setter methods for, 57, 61–62, 65
 - hash codes of, 103, 208
 - immutable, 65
 - object-private, 60–61, 65
 - overriding, 91, 95–96, 130, 132
 - printing, 208
 - private, private final, 59
 - protected, 94
 - public, 56
 - static, 71
 - transient, volatile, 219
- `FileInputStream` class (Java), 112
- files
 - and packages, 80
 - appending, 115
 - binary, 112
 - naming, 15
 - processing, 109–115
 - reading, 109–111
 - redirecting input/output from/to, 115
 - saving, 241–242
 - writing, 112
- `Files` class (Java), 112–113
- filter method, 39, 162, 210
 - of Future, 256
 - of Iterable, 180
 - of Iterator, 189
- final keyword, 59, 92
- finally statement, 30–31
- find, `firstCompletedOf` methods (Future object), 257–258
- findAllIn, `findFirstIn` methods, 116
- flatMap method, 182–183
 - of Future, 254
 - of Iterable, 180
 - of Iterator, 189
 - of Try, 255
- Float type, 5
 - reading, 22
- floating-point calculations, 219
- fluent interfaces, 280–281
- fold method
 - of Iterable, 180
 - of Iterator, 189
 - of parallel collections, 194
- foldLeft method, 168, 184–185
 - of Future object, 257
 - of Iterable, 180, 273
 - of Iterator, 189
 - of parallel collections, 194
- foldRight method, 185
 - of Iterable, 180
 - of Iterator, 189
 - of parallel collections, 194
- for loop, 22–25
 - constructing collections in, 24, 38
 - enhanced (Java), 38
 - for arrays, 37–40
 - for futures, 254
 - for maps, 50
 - for regex groups, 118
 - guards in, 204, 254
 - pattern matching in, 204
 - range-based (C++), 38
 - regular expressions in, 116
 - with Option type, 210
- forall method
 - of Iterable, 180
 - of Iterator, 189
- force method, 190
- foreach method, 162, 183, 210
 - of Future, 255–256
 - of Iterable, 180
 - of Iterator, 189
- fork-join pool, 260
- formatted strings, 21–22, 112
- Fortran programming language, 153
- Fraction class, 146–147, 324–327
- Fraction companion object, 325
- Fractional type, 332
- fragile base class problem, 92
- French delimiters, 328
- fromString, `fromURL` methods (Source), 111
- fromTry method (Future), 258
- functional programming languages, 157
- functions, 25–26, 157–167, 287
 - anonymous, 25, 159–160, 167
 - as method parameters, 14, 159
 - binary, 162, 184
 - calling, 8–9, 158
 - converting to Java interfaces, 164
 - curried, 164–167, 328
 - exiting immediately, 25
 - for all elements of a collection, 161
 - from methods, 288
 - generic, 266, 269
 - higher-order, 160–162
 - identity, 333
 - implementing, 265
 - left-recursive, 310
 - mapping, 182–183
 - names of, 14, 142, 324

nested, 23
 parameterless, 166–167
 parameters of, 25
 call-by-name, 167
 default, 26
 named, 26
 other functions as, 160
 type deduction in, 160
 variable, 26–27
 partial, 183, 211–212, 313
 passing to another function, 159–160, 163
 recursive, 25–27
 return type of, 5, 25, 28
 return value of, 166–167
 scope of, 163
 single-argument, 161, 272
 storing in variables, 157–159
 syntax of, 146–147
 variance of, 272
 vs. variables, in parsers, 313
future method (`Promise`), 259
Future companion object, 256–258
 `apply` method, 253, 260
 `find, firstCompletedOf` methods, 257–258
 `foldLeft, reduceLeft, traverse` methods, 257
 `fromTry, never, successful, unit` methods, 258
Future interface (Java), 249
Future trait, 248–260
 `andThen, collect` methods, 256
 `failed, fallbackTo` method, 255–256
 `filter` method, 256
 `flatMap` method, 254
 `foreach` method, 255–256
 `isCompleted` method, 250
 `map` method, 253–254
 `onComplete, onSuccess, onFailure` methods, 252
 `ready, result` methods, 250
 `recover, recoverWith` methods, 255–256
 `transform, transformWith` methods, 256
 `value` method, 250
 `zip, zipWith` methods, 255
futures
 blocking waits for, 251
 chaining, 253
 delaying creation of, 254
 execution context of, 260
 tasks of:
 composing, 252–255
 failing, 250
 running, 248–249
 transformations of, 255–256
 vs. promises, 259

G

generators, 24–25
 generic arrays, 269
 generic classes, 266
 conditionally used methods in, 270
 generic functions, 266, 269
 generic methods, 266
 generic types
 erased, in JVM, 269
 variance of, 271–275
`GenTraversable` trait, 211
`GenXxx` sequences, 14
`get` method (`Try`), 251
`get, getOrElse` methods (`Map`), 48, 210, 232
`getLines` method (`Source`), 109, 190
`@getter` annotation, 219
`getXxx` methods, 57, 61, 221
 grammars, 304–305
 left-recursive, 314
 Group type, 235
`groupBy` method, 183
 grouped method
 of `Iterable`, 181, 188
 of `Iterator`, 189
 guard combinator, 313
 guards, 24–25, 38, 199
 for pattern matching, 238
 in `for` statements, 204, 254
 variables in, 199

H

hash codes, 100, 103
 hash sets, 177
 hash tables, 47, 50
`hashCode` method, 177
 of `AnyRef`, 103
 of case classes, 205, 208–209
 of value classes, 104
 overriding, 103
 Hashing type, 332
 Haskell programming language, 26, 331
`hasNext` method
 of `Iterable`, 188
 of `Iterator`, 127
`head` method, 110
 of `Iterable`, 180, 188
 of lists, 175–176
`headOption` method (`Iterable`), 180
 Hindley-Milner algorithm, 26
 HTTP (Hypertext Transfer Protocol), 111, 303

I

- id method, 76
- @Id annotation, 219
- ident method, 317
- Identifier type, 317
- identifiers, 142, 317
- identity functions, 333
- IEEE double values, 219
- if expression, in loop generators, 24, 38
- if/else expression, 18–19, 30
- IllegalStateException, 259
- immutability, 7
- implements keyword, 123
- implicit conversions, 42–43, 141, 323–336
 - ambiguous or multiple, 327
 - debugging, 326
 - for parsers, 316
 - for strings to ProcessBuilder objects, 114
 - for type parameters, 268
 - importing, 325–326, 330
 - naming, 324
 - rules for, 326–327
 - unwanted, 191, 324, 326
 - uses of, 324–325
- implicit evidence parameter, 270
- implicit keyword, 324, 328–329
- implicit parameters, 9, 14, 268, 299, 328–336
 - not available, 226, 334
 - of common types, 328
- implicit values, 268–269
- @implicitAmbiguous annotation, 226
- implicitly method (Predef), 330, 333
- @implicitNotFound annotation, 226, 334
- :implicits in REPL, 326
- import statement, 76, 80, 85–87
 - implicit, 86–87
 - location of, 85
 - overriding, 86
 - selectors for, 86
 - wildcards in, 8, 85–86
- increments, 7
- IndexedSeq companion object, 336
- IndexedSeq trait, 172, 174–175, 336
- indexXxx methods
 - of Iterator, 189
 - of Seq, 181
- infix notation
 - for operators, 143–145
 - precedence of, 285
 - for types, 285–287
 - in case clauses, 206–207
- in math, 285
- with anonymous functions, 159
- inheritance, 91–105
 - diamond, 122–123
 - hierarchy of, 100–102
 - multiple, 121–123
- init method (Iterable), 180
- @inline annotation, 224
- inlining, 224
- InputStream class (Java), 239
- Int type, 5, 268, 270
 - immutability of, 7
 - no null value in, 100
 - reading, 22
- int2Fraction method, 325
- Integer class (Java), 225
- interfaces
 - fluent, 280–281
 - rich, 127
- interpolated strings. *See* formatted strings
- InterruptedException, 251
- intersect method
 - of Iterator, 189
 - of Seq, 182
 - of sets, 177
 - of StringOps, 6
- intersection types. *See* compound types
- intertwining classes and constructors, 94–95
- into combinator, 311–312
- isCompleted method
 - of Future, 250
 - of Promise, 259
- isDefinedAt method (PartialFunction), 211
- isEmpty method
 - of Iterable, 180
 - of Iterator, 189
- isInstanceOf method (Any), 93, 100, 200
- isSuccess, isFailure methods (Try), 251
- IsTraversableXxx type classes, 332
- istream::peek function (C++), 110
- itemStateChanged method (listeners), 293
- Iterable trait, 41, 172, 296–299
 - methods of, 180–182, 188, 273
- iterableAsScalaIterable function (JavaConversions), 192
- iterator method (collections), 188
- Iterator trait, 127, 172
 - methods of, 189
- iterators, 110, 188–189
 - constructing streams from, 190
 - fragility of, 189
 - turning into arrays, 116

J

Java programming language

- ? operator in, 18
- annotations in, 216–221
- arrays in, 36, 43, 173, 273
- assertions in, 225
- assignments in, 21
- checked exceptions in, 29, 220
- classes in, 91–92
 - hierarchy of, 67
 - serializable, 113
- closures in, 163
- collections in, 172
- completable futures in, 260
- construction order in, 100
- dependencies in, 290
- event handling in, 293
- expressions in, 17
- fields in:
 - protected, 94
 - public, 56
- functional (SAM) interfaces in, 164
- futures in, 249
- generic types in, 275
- hashCode method for each object in, 177
- identifiers in, 142
- interfaces in, 121–124, 135, 137
- interoperating with Scala:
 - arrays, 43–44
 - classes, 57, 95, 216
 - collections, 191–192
 - maps, 50–51, 204
 - methods, 221
 - traits, 135, 137
- lambda expressions in, 164
- linked lists in, 173, 176
- loops in, 22, 38
- maps in, 173
- methods in, 72, 92, 94
 - abstract, 97
 - overriding, 99
 - static, 9, 25
 - with variable arguments, 27
- missing values in, 270
- modifiers in, 219
- no multiple inheritance in, 121
- no operator overloading in, 8
- no variance in, 226
- null value in, 100
- operators in, 145
- packages in, 80, 82, 84

primitive types in, 36, 100

- reading files in, 110–112
- singleton objects in, 72
- statements in, 17–20
- superclass constructors in, 95
- switch in, 223
- synchronized in, 100
- toString method in, 41
- traversing directories in, 112–113
- type checks and casts in, 93
- void in, 18, 20, 102
- wildcards in, 85, 275, 286
- Java AWT library, 137
- Java EE (Java Platform, Enterprise Edition), 216
- java.awt.Component, java.awt.Container classes, 137
- java.io.BufferedReader class, 138
- java.io.FileInputStream class, 112
- java.io.InputStream class, 239
- java.io.PrintWriter class, 112
- java.io.Reader class, 239
- java.io.Serializable interface, 124
- java.io.Writer class, 241
- java.lang package, 86–87
- java.lang.Cloneable interface, 124, 220
- java.lang.Comparable interface, 42
- java.lang.Integer class, 225
- java.lang.Object class, 100
- java.lang.ProcessBuilder class, 44
- java.lang.String class, 6, 36
- java.lang.Throwable class, 29
- java.math.BigDecimal classes, 6
- java.nio.file.Files class, 112–113
- java.rmi.Remote interface, 220
- java.util package, 192
- java.util.ArrayList class, 36, 43, 173
- java.util.Arrays class, 43
- java.util.Comparator class, 226
- java.util.concurrent.CompletableFuture class, 260
- java.util.concurrent.CompletionStage interface, 249
- java.util.concurrent.Executor interface, 248
- java.util.concurrent.Executors class, 260
- java.util.concurrent.Future interface, 249
- java.util.LinkedList class, 173
- java.util.List interface, 43, 173
- java.util.Properties class, 51, 204
- java.util.RandomAccess interface, 173
- java.util.Scanner class, 53, 111
- JavaBeans, 61–62, 221
- property change listeners in, 137

- JavaConversions class, 43, 51, 191–192
 Javadoc, 10
 javap command, 58
 JavaScript, 235
 - closures in, 163
 - duck typing in, 284
 - translating from Scala, 10
 JavaTokenParsers trait, 316
 JButton, JComponent classes (Swing), 137
 JDK (Java Development Kit), source code for, 212
 JSON (JavaScript Object Notation), 303
 jump tables, 223
 JUnit framework, 216–217
 JVM (Java Virtual Machine)
 - arrays in, 36
 - classes in, 129
 - generic types in, 269
 - Inlining in, 224
 - stack size of, 222
 - transient/volatile fields in, 219
- ## K
- Kernighan & Ritchie brace style, 20
 keySet method (Map), 50
 Keyword type, 317
- ## L
- last, lastOption methods (Iterable), 180
 lastIndexof, lastIndexofslice methods
 - of Iterator, 189
 - of Seq, 181
 lazy keyword, 28–29, 133
 length method
 - of Iterable, 180
 - of Iterator, 189
 lexical analysis, 304
 li element (XML), 233–234
 lines
 - iterating over, 110
 - reading, 109
 linked hash sets, 177
 LinkedHashMap class, 50
 LinkedList class (deprecated), 176
 LinkedList class (Java), 173
 List class, 296, 306
 - immutable, 174
 - implemented with case classes, 208
 List interface (Java), 43, 173
 list method (java.nio.file.Files), 112
- ListBuffer class, 175
 lists, 175–176
 - adding/removing elements of, 178–179
 - constructing, 146, 175
 - destructuring, 176, 207
 - empty, 101
 - heterogeneous, 213
 - immutable, 189, 274
 - linked, 172
 - pattern matching for, 201–202, 206
 - traversing, 176
 - vs. arrays, 172
 literals. *See XML literals*
 loadFile method (XML), 239
 log combinator, 311, 313
 log messages
 - adding timestamp to, 125
 - printing, 313
 - truncating, 125, 130
 - types of, 127
 LoggedException trait, 135, 288–289
 Logger trait, 128
 Long type, 5
 - reading, 22
 loops, 22–25
 - breaking out of, 23
 - for collections, 23
 - variables within, 24
 - vs. folding, 186
- ## M
- main method, 74
 makeURL function, 234
 ManagedException trait, 289
 map method, 39, 161, 210
 - of Future, 253–254
 - of Iterable, 180–184, 297–298
 - of Iterator, 189
 - of Try, 255
 Map trait, 48, 172
 - get, getOrElse methods, 48, 210, 232
 - immutable, 174
 - keySet, values methods, 50
 - mapAsXxxMap functions (JavaConversions), 51, 192
 maps, 47–52
 - blank, 48
 - constructing, 48
 - from collection of pairs, 52
 - function call syntax for, 146
 - interoperating with Java, 50–51

- iterating over, 50
- keys of:
 - checking, 48
 - removing, 49
 - visiting in insertion order, 50
- mutable/imutable, 48–49, 173
- reversing, 50
- sorted, 50
- traversing, 204
- updating values of, 49
- values of, 48
- `match` expression, 198–210, 224, 271
 - annotated, 223
 - processing `Try` instances with, 251
- `Match` class, 117
- `MatchError`, 198
- `Math` class, 8
- mathematical functions, 8, 14
- `max` method, 42
 - of `ArrayBuffer`, 40
 - of `Iterable`, 180
 - of `Iterator`, 189
- maximum munch rule, 317
- `MessageFormat.format` method (Java), 27
- `MetaData` type, 232–233, 238
- method types (in compiler), 288
- methods
 - abstract, 95–98, 123, 127, 135
 - abundance of, 10, 14
 - accessor, 56
 - annotated, 216
 - as arithmetic operators, 6
 - calling, 3, 5, 8–9, 56–57, 59, 126
 - chained, 280
 - clashes of, in sub- and superclasses, 92
 - co-/contravariant, 333
 - curried, 165, 271
 - declaring, 56
 - executed lazily, 189
 - final, 92, 103, 223
 - for primary constructor parameters, 65
 - generic, 266
 - getter/setter, 56–60, 98, 216, 221
 - Inlining, 224
 - modifiers for, 84
 - mutator, 56
 - names of, 7
 - misspelled, 92
 - overriding, 91–93, 95–96, 127
 - parallelized, 193
 - parameterless, 8, 56, 95
 - parameters of, 266, 273, 280
 - implicit, 14, 268
 - using functions for, 14, 159
 - wrong type, 92
 - private, 59, 223
 - protected, 94
 - public, 57
 - return type of, 273, 280
 - return value of, 266
 - static, 71
 - turning into functions, 158, 288
 - used under certain conditions, 270
 - variable-argument, 27, 221
 - with two parameters, 7
- Meyer, Bertrand, 58
- `min` method
 - of `ArrayBuffer`, 40
 - of `Iterable`, 180
 - of `Iterator`, 189
 - of `scala.math`, 8
- `mkString` method
 - of `Array`, `ArrayBuffer`, 40
 - of `Iterable`, 181
 - of `Iterator`, 189
 - of `Source`, 110–111
- ML programming language, 26
- `mulBy` function, 160, 163
- multiple inheritance, 121–123
- `MutableList` class (deprecated), 176
- `mutableXxxAsJavaXxx` functions (`JavaConversions`), 192

N

- `NamespaceBinding` class, 242
- namespaces, 242–243
- `@native` annotation, 219
- negation operator, 14
- `never` method (`Future`), 258
- `new` keyword, 66–67
 - and nested expressions, 74
 - constructing objects without, 146
 - omitting, 205, 207–208
- newline character
 - in long statements, 20
 - in printed values, 21
 - inside loops, 25
- `next` method
 - of `Iterable`, 188
 - of `Iterator`, 127
- `Nil` list, 101, 175–176, 226

node sequences, 230
 binding variables to, 237
 building programmatically, 231
 descendants of, 236
 grouping, 235
 immutability of, 232, 238
 in embedded blocks, 234
 traversing, 231
 turning into strings, 232
 Node type, 230–232
 NodeBuffer class, 231–232
 NodeSeq type, 230–232
 XPath-like expressions in, 235–236
 @noinline annotation, 224
 None object, 210, 306–307
 nonterminal symbols, 305
 NoSuchElementException, 254
 not combinator, 313
 Nothing type, 30, 100, 271, 274
 notify, notifyAll methods (*AnyRef*), 100
 NotImplementedException, 101
 @NotNull annotation, 218
 Null type, 100, 238
 null value, 100, 270
 NumberFormatException, 111, 251
 numbers
 average value of, 332
 classes for, 14
 converting:
 between numeric types, 6, 10
 to arrays, 111
 greatest common divisor of, 153
 in identifiers, 317
 invoking methods on, 5
 parsing, 312, 317
 random, 9
 ranges of, 14
 reading, 22, 111
 sums of, 40
 Numeric type, 331
 numericLit method, NumericLit type, 317

O

Object class, 100
 Object class (Java), 100
 object keyword, 71–77, 281
 objects, 71–77
 adding traits to, 125
 cloneable, 220
 companion, 9, 67, 72–73, 146, 173, 281, 325,
 328, 335

compound, 208
 constructing, 10, 56, 72, 125
 default methods for, 177
 equality of, 102–103
 extending class or trait, 73
 extracting values from, 202
 importing members of, 76, 85
 nested, 207
 nested classes in, 66–68, 281
 no type parameters for, 274
 of a given class, testing for, 93–94
 pattern matching for, 200
 remote, 220
 scope of, 282
 serializing, 113, 284
 type aliases in, 283
 variance of, 272
 ofDim method (*Array*), 42
 onComplete, onSuccess, onFailure methods (*Future*),
 252
 operators, 141–149
 arithmetic, 6–8
 assignment, 144–146
 associativity of, 145, 194
 binary, 143–146
 for adding/removing elements, 178–179
 infix, 143, 145
 overloading, 8
 parsing, 317
 postfix, 143–145
 precedence of, 144–145, 285, 307
 unary, 143–144
 opt method (*Parsers*), 305–306
 optimization, 222–226
 Option class, 48, 116, 147, 149, 180, 210, 233, 235,
 306–307
 orNull method, 270
 using with for, 210
 Ordered trait, 40, 42, 329–330
 Ordering type, 42, 329–332
 OSGi (Open Services Gateway initiative
 framework), 290
 OutOfMemoryError, 190
 override keyword, 92–93, 95–96, 123, 127
 omitted, 97–98
 @Overrides annotation, 92

P

package objects, 8, 83–84
 packages, 80–87
 adding items to, 80

chained, 82–83
 defined in multiple files, 80
 importing, 8, 85–87
 always, 86–87
 selected members of, 86
 modifiers for, 84
 naming, 82, 87
 nested, 81–82
 scope of, 282
 top-of-file notation for, 83
 packrat parsers, 314–315
 PackratParsers trait, 314
 PackratReader class, 315
 padTo method
 of Iterator, 189
 of Seq, 182
 Pair class, 275
 par method (parallel collections), 193
 parallel implementations, 193
 @param annotation, 219
 parameters
 annotated, 216
 curried, 271
 deprecated, 226
 implicit, 9, 14, 299
 named, 205
 ParMap trait, 193
 parse method, 306
 parse trees, 308–309
 parseAll method, 306, 313, 315, 318–319
 ParSeq trait, 193
 parsers, 303–319
 backtracking in, 313–314
 combining operations in, 305–307
 entity map of, 241
 error handling in, 319
 implicit conversions for, 316
 numbers in, 312
 output of, 307–308
 packrat, 314–315
 regex, 316, 319
 skipping comments in, 316
 strings in, 312
 whitespace in, 316
 Parsers trait, 305, 315–319
 ParSet trait, 193
 partial functions, 183, 211–212
 PartialFunction class, 211–212
 partition method
 of Iterable, 180
 of Iterator, 189
 of StringOps, 52
 Pascal programming language, 153
 paths, 282–283
 pattern matching, 197–211
 and +: operator, 179
 by type, 200–201
 classes for. *See* case classes
 extractors in, 147
 failed, 147
 for arrays, 201
 for lists, 176, 201–202, 206
 for maps, 50
 for objects, 200
 for tuples, 52, 201–202
 guards in, 199, 238
 in for expressions, 204
 in variable declarations, 203–204
 in XML, 237–238
 jump tables for, 223
 not exhaustive, 226
 of nested structures, 207
 regular expressions in, 203
 variables in, 199–200
 vs. type checks and casts, 93–94
 with extractors, 202
 with Option type, 210
 with partial functions, 211–212
 PCData type, 235
 permutations method
 of Iterator, 189
 of Seq, 182
 phrase combinator, 313
 piping, 114
 Play web framework, 249
 Point class, 332
 polymorphism, 208
 Positional trait, 313, 319
 positioned combinator, 313, 319
 postfix operators, 143–145
 in parsers, 308
 pow method (`scala.math`), 8, 153
 Predef object
 always imported, 86–87, 93, 174
 assert method, 225
 classOf method, 93
 implicitly method, 330, 333
 implicits in, 329–333
 prefixedKey method, 243
 prefixLength method
 of Iterator, 189
 of Seq, 181
 PrettyPrinter class, 242
 primitive types, 5, 36, 225

- print function, 21, 111
- printf function, 21, 112
- println function, 21
- PrintStream.printf method (Java), 27
- PrintWriter class (Java), 112
- PriorityQueue class, 175
- private keyword, 57–68, 84
- probablePrime method (BigInt), 9
- procedures, 28
- process control, 114–115
- Process object, 115
- ProcessBuilder class (Java), 44
 - constructing, 115
 - implicit conversions to, 114
- processing instructions, 231
- product method
 - of Iterable, 180
 - of Iterator, 189
- programs
 - concurrent, 193
 - displaying elapsed time for, 75
 - implicit imports in, 86–87
 - piping, 114
 - readability of, 8
 - self-documenting, 296
- Promise trait, 258–260
- promises, 258–260
- properties, 57
 - in Java. *See* bean properties
 - read-only, 59
 - write-only (not possible), 60
- Properties class (Java), 51, 204
- propertiesAsScalaMap function (JavaConversions), 51, 192
- propertiesAsScalaMap method (JavaConversions), 204
- property change listeners, 137
- protected keyword, 84, 94
- public keyword, 56, 84
- PushbackInputStreamReader class (Java), 110
- Python programming languages, closures in, 163

- Q**
- Queue class, 174–175
- quickSort method (Array), 40

- R**
- r method (String), 116–117
- Random object, 9
- RandomAccess interface (Java), 173
- Range class, 5, 14, 297–298, 336
 - immutable, 174–175
 - traversing, 23
- raw prefix, in formatted strings, 22
- raw string syntax, 116
- readBoolean, readByte, readChar, readFloat, readLine, readShort methods (scala.io.StdIn), 22
- readDouble, readInt, readLong methods (scala.io.StdIn), 22, 111
- Reader class (Java), 239, 315
- ready, result methods (Await), 250
- recover, recoverWith methods (Future), 255–256
- recursions, 174
 - for lists, 176
 - left, 310–311
 - tail, 223
- reduce method
 - of Iterable, 180
 - of Iterator, 189
 - of parallel collections, 193
- reduceLeft method, 162, 184
 - of Future object, 257
 - of Iterable, 180
 - of Iterator, 189
 - of parallel collections, 193
- reduceRight method, 184
 - of Iterable, 180
 - of Iterator, 189
 - of parallel collections, 193
- reference types
 - == operator for, 103
 - assigning null to, 100
- reflective calls, 284
- Regex class, 116–117
 - replaceXxxIn methods, 116, 211
- RegexParsers trait, 305, 315–316, 319
- regular expressions, 116–117
 - for extractors, 203
 - grouping, 117–118
 - in parsers, 316
 - matching tokens against, 305
 - raw string syntax in, 116
 - return value of, 306
- Remote interface (Java), 220
- @remote annotation, 220
- rep method (Parsers), 305–306, 311–312
- REPL (read-eval-print loop), 3–4
 - implicits in, 326, 333
 - nearsightedness of, 19
 - paste mode in, 19, 73
 - types in, 158, 283
- replaceXxxIn methods (Regex), 116, 211

`repXxx` methods (`Parsers`), 312

`result` method, 223

`return` keyword, 25, 167

`reverse` method

- of `Iterator`, 189

- of `Seq`, 182

`RewriteRule` class, 239

`rich` interfaces, 127

`RichChar` class, 6

`RichDouble` class, 6, 14

`RichFile` class, 325

`RichInt` class, 6, 14, 22, 268

`_root_` prefix, in package names, 82

`Ruby` programming language

- closures in, 163

- duck typing in, 284

`RuleTransformer` class, 239

S

`s` prefix, in formatted strings, 22

`SAM` (single abstract method) conversions, 163

`save` method (`XML`), 241

`SAX` parser, 239

`scala` package, always imported, 82, 86–87, 113, 174

`Scala` programming language

- embedded languages in, 141, 303

- interoperating with shell programs, 114

- interpreter of, 1–4

- older versions of, 75

- strongly typed, 150

- translating to `JavaScript`, 10

`scala/bin` directory, 1

`scala.` prefix, in package names, 8, 87

`scala.collection` package, 173, 192

`Scala.js` project, 10

`scala.language.existentials`, importing, 287

`scala.language.higherKinds`, importing, 299

`scala.language.implicitConversions`, importing, 324

`scala.math` package, 8, 14

- `ceil` method, 158

`scala.sys.process` package, 114

`scala.util` package, 9

`scala-ARM` library, 31

`scalac` program. *See compiler*

`Scaladoc`, 6, 10–15, 41–42, 221

`ScalaObject` interface, 100

`scanLeft`, `scanRight` methods, 186

`Scanner` class (`Java`), 53, 111

`sealed` keyword, 209

`segmentLength` method

- of `Iterator`, 189

- of `Seq`, 181

`selectDynamic` method (`Dynamic`), 151–152

`selectors`, for imports, 86

`self types`, 68, 134–135, 288–295

- dependency injections in, 290–291

- no automatic inheritance for, 289

- structural types in, 135

- typesafe, 294

- vs. traits with supertypes, 135

`Seq` trait, 27, 41, 172–175, 235

- methods of, 181–182

`seq` method (parallel collections), 193

`Seq[Char]` class, 14

`Seq[Node]` class, 231–232

`seqAsJavaList` function (`JavaConversions`), 192

`sequences`, 14

- adding/removing elements of, 179

- as function parameters, 27

- comparing, 165, 271

- extracting values from, 149–150

- filtering, 162

- integer, 175

- mutable/imutable, 174–175

- of characters, 14

- reversing, 182

- sorting, 162, 182

- with fast random access, 175

`Serializable` interface (`Java`), 124

`Serializable` trait, 113

`serialization`, 113

`@SerialVersionUID` annotation, 113, 220

`Set` trait, 172

- immutable, 174

`setAsJavaSet` function (`JavaConversions`), 192

`sets`, 177

- difference, intersection, union operations on, 177–178

- elements of, 177–179

`@setter` annotation, 219

`setXxx` methods, 57, 61, 221

`shadowing` rule, 24

`shell scripts`, 114–115

`Short` type, 5

- reading, 22

`singleton` objects, 8–9, 71–72, 281

- case objects for, 205

`singleton` types, 280–281, 283, 287

`slice`, `span`, `splitAt` methods

- of `Iterable`, 180

- of `Iterator`, 189

- sliding method
 - of Iterable, 181, 188
 - of Iterator, 189
- SmallTalk programming language, 60
- Some class, 210, 306–307
- sortBy method
 - of Iterator, 189
 - of Seq, 182
- sorted method
 - of ArrayBuffer, 40
 - of Iterator, 189
 - of Seq, 182
 - of StringOps, 8, 14
- sorted sets, 177
- SortedMap trait, 50, 172
- SortedSet trait, 172
- sortWith method, 162
 - of ArrayBuffer, 40
 - of Iterator, 189
 - of Seq, 182
- Source object, 109–111
 - buffered, close, toArray, toBuffer methods, 110
 - fromString, fromURL methods, 111
 - getLines method, 109, 190
 - mkString method, 110–111
- @specialized annotation, 225–226
- Spring framework, 290
- sqrt method (`scala.math`), 8
- src.zip file (JDK), 212
- Stack class, 174–175
- stack overflow, 222
- standard input, 111
- StandardTokenParsers class, 317
- start symbol, 305–306
- startsWith method
 - of Iterator, 189
 - of Seq, 181
- stateChanged method (listeners), 293
- statements
 - line breaks in, 20
 - terminating, 19–20
 - vs. expressions, 17
- stdin method, 111
- StdIn object, 22, 111
- StdLexical trait, 318
- StdTokenParsers trait, 315, 318
- StdTokens trait, 317
- Stream class, 174
- streams, 189–190
- @strictfp annotation, 219
- String class, 116–117
- String class (Java), 6, 36
- string interpolation, 21
- stringLiteral method, `StringLit` type, 317
- StringOps class, 6, 14, 52
 - apply method, 9
 - containsSlice method, 14
 - sorted method, 8, 14
- strings, 6
 - characters in:
 - common, 6
 - uppercase, 14
 - classes for, 14
 - converting:
 - from any objects, 6
 - to numbers, 10, 111
 - to ProcessBuilder objects, 114
 - formatted, 21–22, 112
 - parsing, 312, 317
 - sorting, 8
 - testing for sequences in, 14
 - traversing, 23
 - vs. symbols, 226
- structural types, 97, 135, 283–284
 - adding to compound types, 285
- subclasses
 - anonymous, 97
 - concrete, 98
 - early definitions in, 99
 - equality in, 103
 - implementing abstract methods in, 123
 - mixing traits into, 288–289
- subsetOf method (`BitSet`), 177
- Success class, 251, 315
- success combinator, 313
- success method (`Promise`), 259
- successful method (`Future`), 258
- sum method
 - of ArrayBuffer, 40
 - of Iterable, 180
 - of Iterator, 189
- super keyword, 92–93, 126
- super keyword (Java), 95
- superclasses, 133–134
 - abstract fields in, 98
 - constructing, 94–95
 - extending, 137
- methods of:
 - abstract, 97
 - clashing with subclass methods, 92
 - overriding, 96
- no multiple inheritance of, 121, 124, 129
- scope of, 282
- sealed, 209

supertypes, 18, 42, 102

Swing toolkit, 137–138

switch statement, 19, 198

@switch annotation, 223–224

symbols, 226

synchronized method (*AnyRef*), 100

syntactic sugar, 275, 286

T

tab completion, 3

tail method

- of Iterable, 180

- of lists, 175–176

TailCalls, TailRec objects, 223

@tailrec annotation, 222

take, takeWhile methods

- of Iterable, 180

- of Iterator, 189

takeRight method (Iterable), 180

@Test annotation, 217

Text class, 233

- pattern matching for, 237

text method, 232

this keyword, 42, 62, 65, 134–135, 280, 288–289, 295

- aliases for, 68, 289

- scope of, 282

- with private, 59

- with protected, 94

thread pool, 248

threads

- assigning tasks to, 248

- blocking, 250–251

throw expression, 30

Throwable class (Java), 29

@throws annotation, 220

to method (RichInt), 6, 22, 175

toArray method

- of ArrayBuffer, 37

- of Iterable, 181

- of Iterator, 189

- of Source, 110

toBuffer method

- of Array, 37

- of Source, 110

toChar method, 6

toIndexedSeq, toIterable methods

- of Iterable, 181

- of Iterator, 189

toInt,toDouble methods, 6, 111

token method (StdLexical), 318

Token type, 315

tokens, 304

- discarding, 308–309

- matching against regular expressions, 305

Tokens trait, 317

toList, toStream methods

- of Iterable, 181

- of Iterator, 189

toMap method, 52, 173

- of Iterable, 181

- of Iterator, 189

toSeq, toSet methods, 173

- of Iterable, 181

- of Iterator, 189

toString method, 6, 41, 76, 233

- of case classes, 205, 208–209

trait keyword, 123, 287

traits, 123–137, 287

- abstract types in, 291

- adding to objects, 125

- construction order of, 126–127, 130–132

- dependency injections in, 135, 290–291

- extending, 73

- classes, 133–134

- superclass, 137

fields in:

- abstract, 130, 132

- concrete, 128–129

for collections, 172–173

for rich interfaces, 127

implementing, 124, 269

layered, 125–126

methods of, 124–126, 135

- overriding, 127

- unimplemented, 123

- mixing into classes, 288–289

- no constructor parameters in, 132–133

- self types in, 134–135

- type parameters in, 266

- universal, 104

- vs. classes, 132

- vs. Java interfaces, 121–124, 135

- vs. structural types, 284

trampolining, 223

transform method

- of ArrayBuffer, 183

- of RewriteRule, 239

transform, transformWith methods (Future), 256

@transient annotation, 219

Traversable, TraversableOnce traits, 41

TraversableLike trait, 334

traverse method (Future), 257

trimEnd method, 36
 Try class, 31, 251
 flatMap, map methods, 255
 get method, 251
 isSuccess, isSuccess methods, 251
 try statement, 30–31
 exceptions in, 167
 trySuccess method (Promise), 259
 tuples, 47, 51–52, 287
 accessing components of, 51
 converting to maps, 52
 pattern matching for, 201–202
 zipping, 52
 type classes, 331–333
 type constraints, 269–271, 333
 type constructors, 296–299
 type errors, 150, 273
 type inference, 270
 type keyword, 280–281, 283, 287
 type parameters, 14, 97, 265–275, 287,
 328
 annotated, 217, 225
 context bounds of, 329–330
 implicit conversions for, 268
 not possible for objects, 274
 structural, 283–284
 vs. abstract types, 292
 type projections, 67, 281–283, 287
 in forSome blocks, 286
 type variables, bounds for, 266–269
 TypeAnnotation trait, 218
 types, 5–6, 279–299
 abstract, 283, 291–293, 315
 aliases for, 174, 283
 annotated, 217
 anonymous, 98
 checking, 93–94
 compound, 284–287
 converting, 6, 270
 dynamic, 150
 equality of, 270
 existential, 287
 generic, 269, 271–275
 implementing multiple traits, 269
 infix, 285–287
 invariant, 272
 naming, 296
 pattern matching by, 200–201
 primitive, 5, 36, 225
 self, 134–135, 288–295
 structural, 135, 285
 subtypes of, 270

supertypes of, 18, 102
 wrapper, 6

U

ul element (XML), 233–234
 unapply method, 147–150, 202, 206–207
 of case classes, 205
 unapplySeq method, 149–150, 202
 unary operators, 143–144
 unary_- method (BigInt), 14
 unary_op methods, 143
 @unchecked annotation, 218, 226
 @uncheckedVariance annotation, 226
 Unicode characters, 142
 uniform access principle, 58
 uniform creation principle, 173
 uniform return type principle, 182
 union method (sets), 177
 Unit class, 28, 100, 102
 value of, 18–21
 unit method (Future), 258
 universal traits, 104
 Unparsed type, 235
 until method, 175
 until statement, 166
 update method, 146–147
 updateDynamic method (Dynamic), 151–152
 URIs (Uniform Resource Identifiers),
 242
 URLs (Uniform Resource Locators)
 loading files from, 239
 reading from, 111
 redirecting input from, 115

V

val fields, 4
 declarations of, 4–5
 early definitions of, 99
 final, 99
 generated methods for, 59, 61–62, 65
 in forSome blocks, 286
 in parsers, 314
 initializing, 5, 20, 28–29
 lazy, 28–29, 99, 133, 314
 overriding, 95–96, 98
 private, 62
 scope of, 282
 specifying type of, 5
 storing functions in, 157–159
 value classes, 103–105, 208

Value method (enumerations), 75–76
 value method (Future), 250
 valueAtOneQuarter function, 160
 values
 binding to variables, 207
 naming, 200
 printing, 21
 values method (Map), 50
 van der Linden, Peter, 198
 var fields, 4
 annotated, 216
 declarations of, 4–5
 extractors in, 147
 pattern matching in, 203–204
 generated methods for, 61–62, 65
 initializing, 5
 no path elements in, 282
 overriding, 96
 private, 62
 specifying type of, 5, 266
 updating, 49
 vs. function calls, in parsers, 313
 @varargs annotation, 221
 variables
 binding to:
 node sequences, 237
 values, 207
 declaring as Java SAM interfaces, 164
 in case clauses, 199–200
 naming, 142, 200
 within loops, 24
 variance, 226
 Vector class, 174–175
 vector type (C++), 36
 view bounds, 268
 view method, 190–191
 void keyword (C++, Java), 18, 20, 102
 @volatile annotation, 219

W

wait method (AnyRef), 100
 walk method (java.nio.file.Files), 112
 walkFileTree method (Java), 112–113
 warnings, 226
 while loop, 22, 166
 whitespace
 in lexical analysis, 304
 parsing, 240, 316–317
 wildcards, 275
 for XML elements, 236
 in catch statements, 30

in imports, 8, 85–86
 in Java, 85, 286
 with keyword, 99, 124–125, 284–287
 wrapper types, 6
 write method (XML), 241
 Writer class (Java), 241

X

XHTML (Extensible Hypertext Markup Language), 235
 XhtmlParser class, 241
 XML (Extensible Markup Language), 229–243
 attributes in, 232–235, 238
 character references in, 232
 comments in, 231
 elements in, 238, 242
 entity references in, 231–232, 241
 including non-XML text into, 235
 indentation in, 242
 loading, 239
 malformed, 235
 namespaces in, 242–243
 nodes in, 230–232
 processing instructions in, 231
 saving, 233, 241–242
 self-closing tags in, 242
 transforming, 239
 XML declarations, 241
 XML literals
 braces in, 234
 defining, 230
 embedded expressions in, 233–234
 entity references in, 232
 in nested Scala code, 233
 in pattern matching, 237–238
 XPath (XML Path language), 235–236

Y

yield keyword
 as Java method, 142
 in loops, 24, 38–39

Z

zip method, 52, 181, 187–189
 zip, zipWith methods (Future), 255
 zipAll, zipWithIndex methods, 187
 of Iterable, 181
 of Iterator, 189
 zipping, 187–188

This page intentionally left blank

Expert Video Training for Experienced Programmers



From **CAY HORSTMANN**:
Java Champion, best-selling
author, and instructor

**SAVE
40%**
CODE: CHVID

- Expert-led instruction
- Lab exercises
- Detailed demonstrations
- Project files
- Online self-assessment quizzes

livelessons ◉
video instruction from technology experts



Learn Scala in Less Than 5 Hours

- Understand core language features: values and types, control structures, and functions
- Use arrays, maps, and tuples for collecting elements
- Master object-oriented programming in Scala: classes, objects, packages, inheritance, and traits
- Discover functional programming in Scala, and how to work with higher-order functions
- Leverage Scala's powerful pattern matching and case classes
- Gain a foundation for using frameworks such as Apache Spark, Play, Akka, and Spray

Learn to Put Java to Work in 10 Hours

- Learn the fundamentals of user interface programming
- Write branches and loops, and work with numbers, strings, and arrays
- Use built-in classes and build your own
- Work with inheritance and interfaces, as well as lambda expressions
- Learn what to do when your programs do the wrong thing: exception handling, logging, and debugging
- Write generic code that works for many different data types
- Test your knowledge along the way with online assessment quizzes

Save 40%—Use coupon code CHVID

informit.com/horstmann



informIT.com
the trusted technology learning source

O'REILLY®
Safari



REGISTER YOUR PRODUCT at informit.com/register Access Additional Benefits and SAVE 35% on Your Next Purchase

- Download available product updates.
- Access bonus material when applicable.
- Receive exclusive offers on new editions and related products.
(Just check the box to hear from us when setting up your account.)
- Get a coupon for 35% for your next purchase, valid for 30 days. Your code will be available in your InformIT cart. (You will also find it in the Manage Codes section of your account page.)

Registration benefits vary by product. Benefits will be listed on your account page under Registered Products.

InformIT.com—The Trusted Technology Learning Source

InformIT is the online home of information technology brands at Pearson, the world's foremost education company. At InformIT.com you can

- Shop our books, eBooks, software, and video training.
- Take advantage of our special offers and promotions (informit.com/promotions).
- Sign up for special offers and content newsletters (informit.com/newsletters).
- Read free articles and blogs by information technology experts.
- Access thousands of free chapters and video lessons.

Connect with InformIT—Visit informit.com/community

Learn about InformIT community events and programs.



informIT.com
the trusted technology learning source

Addison-Wesley • Cisco Press • IBM Press • Microsoft Press • Pearson IT Certification • Prentice Hall • Que • Sams • VMware Press